
Systems Biology Markup Language (SBML) Level 3

Proposal: Multi-component Species Features

Andrew Finney
afinney@cds.caltech.edu

March 8, 2004

Contents

1	Terms of Reference	2
2	Acknowledgements	2
3	Aims	2
4	Overview of Proposal	3
5	Tutorial on the Proposed Features	5
5.1	Terminology	5
5.2	The Definition of Chemical Entities across Compartments	5
5.3	Generalized Reactions: The Definition of Reactions across Compartments	5
5.4	Implied Species	8
5.5	Simple Multi-Component Chemical Entities	10
5.6	Multi-component Chemical Entities with explicit bonds	12
5.7	Reactions generalized to cover classes of Multi-component Chemical Entities	12
5.8	Hierarchical Species Types and Type Equivalence	24
6	Formal Definition of Proposal	28
6.1	Proposed Classes in Detail	28
6.2	Equivalence of Species Types	31
6.3	Species Equivalence	33
6.4	Semantics of Reactions	33
	References	36

1 Terms of Reference

This document describes proposed features for inclusion in Systems Biology Markup Language (SBML) Level 3. This document describes features enabling the description of large chemical entities that are composed from other chemical entities. (Entities of this type were previously classed as 'complex species'. This term is confusing (Phair, 2003) and is avoided in this document.)

This document is not a definition of SBML Level 3 or part of it. This document simply presents various features which could be incorporated into SBML Level 3 as the Systems Biology community wishes. This document is intended for detailed review by that community and to provoke alternative proposals.

This document is not the first proposal to support multi-component species (Le Novère et al., 2003) and supersedes a previous proposal by the author (Finney, 2001).

Throughout this document issues that the author believes will require further discussion have been highlighted.

For brevity the text of this document is with reference to SBML Level 2 (Finney et al., 2002) i.e. features are described in terms of changes to SBML Level 2. In addition for brevity the UML diagrams in this proposal show only new attributes and types for SBML Level 3.

All types proposed in this document will be derived from the `SBase` type.

2 Acknowledgements

This proposal has benefitted from discussions the author had with Nicolas Le Novère, Fabian Campagne, Jeremy Zucker, Robert Phair, Larry Lok, Michael Blinov and Roger Brent. In particular many of the ideas presented here are similar to those developed by the Molecular Sciences Institute and the T-10 Cell Signalling Group (Goldstein et al., 2001) at Los Alamos National Laboratories.

3 Aims

This proposal aims to support the representation of the following concepts that are not easily represented in SBML Level 2:

- the common description of biochemical entities that can then be located in different compartments
- the common description of biochemical reactions that can then be located in different compartments
- the hierarchical description of biochemical entities through the composition of other biochemical entities
- the description of biochemical entities through simple associative composition
- the description of biochemical entities through graphs of other biochemical entities where arcs represent kinds of bonding
- the description of generalized biochemical reactions that avoids the enumeration of many species states and reactions

In particular this proposal aims to enable the description of, for example, proteins which can contain many phosphorylation states, complexes of these proteins and models of signalling pathways which contain these proteins.

4 Overview of Proposal

A UML diagram for the proposed new classes is shown in figures 1 and 2. Section 5 demonstrates with examples, how instances of these classes can be assembled to achieve the aims of the proposal. Section 5 effectively defines a staged roadmap of how the features described in this proposal could be added to SBML.

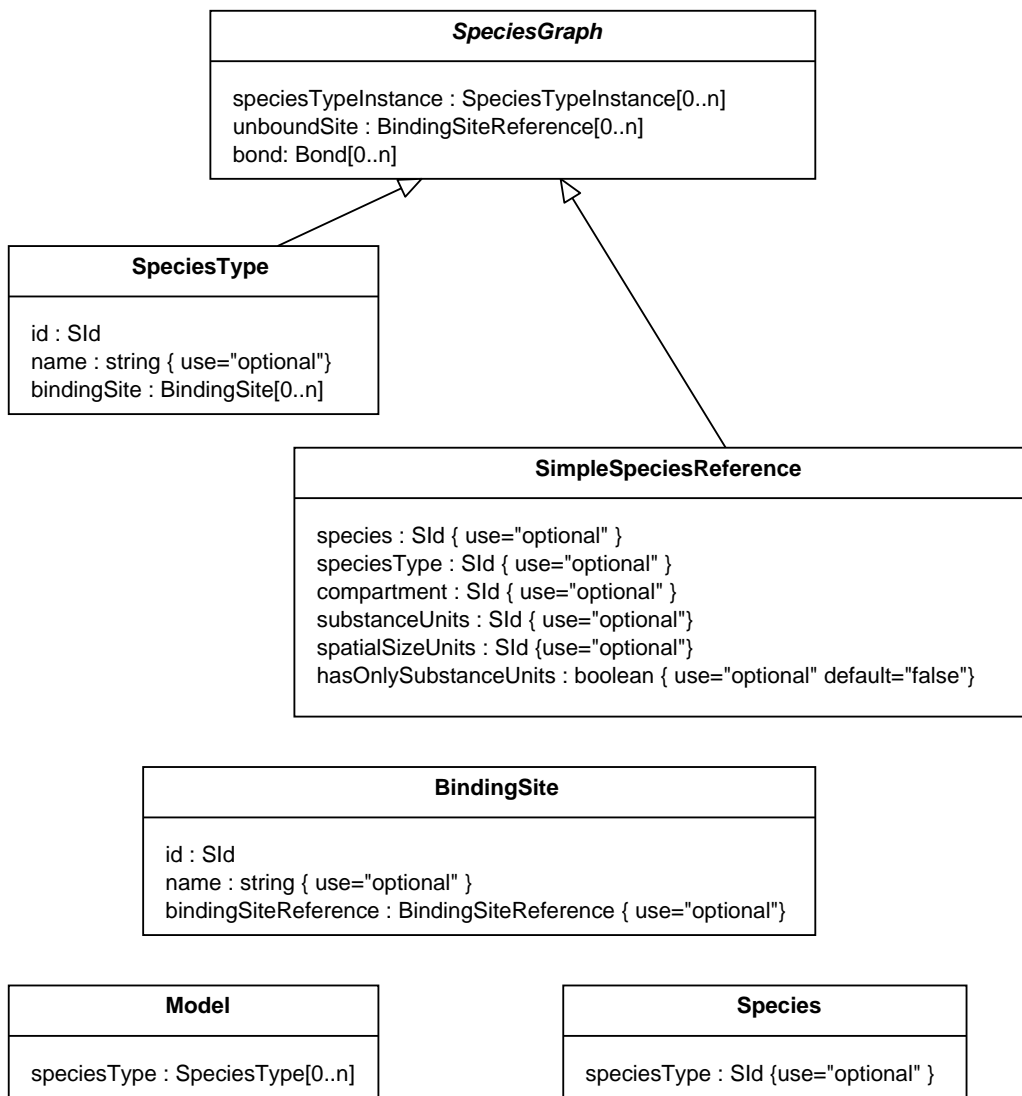


Figure 1: The types and attributes introduced into SBML by this proposal. This diagram only shows new classes and fields: all SBML Level 2 structures are assumed to be present. This diagram is continued in Figure 2

The proposal is described more formally in section 6. Section 6 can be considered as a reference section.

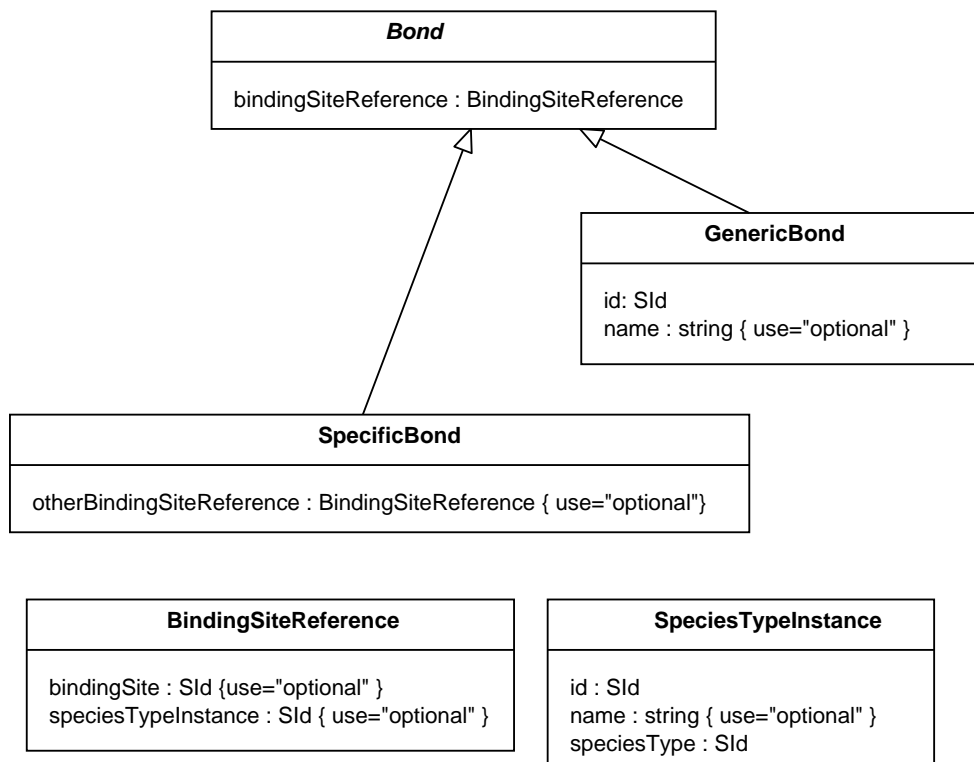


Figure 2: The types and attributes introduced into SBML by this proposal. This diagram only shows new classes and fields: all SBML Level 2 structures are assumed to be present. This diagram is a continuation of the diagram in Figure 1

5 Tutorial on the Proposed Features

5.1 Terminology

The following terminology is used in this document:

- *chemical entity* any individual chemical object e.g. a calcium ion, a phosphate, a protein, and a lipid.
- *compartment* a well stirred container of chemical entities
- *species type* a type of chemical entity, specifically a set of chemical entities with exactly the same chemical form,
- *species* a pool of chemical entities of the same species type located in a specific compartment

5.2 The Definition of Chemical Entities across Compartments

Consider a model where we have a species type which exists in more than one compartment. For example we might wish to model Aspartate in a Cytosol compartment and in the Mitochondrial Matrix. In SBML Level 2 we have to explicitly define each pool of Aspartate located in a separate compartment, using a **Species** structure as shown in Figure 3

```
<model id="malate_aspartate_shuttle1">
  <listOfCompartments>
    <compartment id="Cytosol"/>
    <compartment id="Mitochondrial_Matrix"/>
  </listOfCompartments>
  <listOfSpecies>
    <species id="Aspartate_in_Cytosol" compartment="Cytosol"/>
    <species id="Aspartate_in_Mitochondrial_Matrix" compartment="Mitochondrial_Matrix"/>
  </listOfSpecies>
</model>
```

Figure 3: *malate_aspartate_shuttle1* Model with the same type of chemical entity located in different compartments.

In SBML Level 2 there is no formal way to relate these species together. Under this proposal we can do this by representing a chemical entity type such as, Aspartate, with a **SpeciesType** structure and then refer to the **SpeciesType** from the **Species** structures. We can thus transform the *malate_aspartate_shuttle1* Model in Figure 3 to that shown in Figure 4.

This model does not introduce any new variables that are not present in *malate_aspartate_shuttle1* it simply identifies *Aspartate_in_Cytosol* and *Aspartate_in_Mitochondrial_Matrix* as being separate pools of the same chemical entity. You cannot refer to **SpeciesType** structures from MathML structures under this proposal.

The *malate_aspartate_shuttle1* example is still a valid model under this proposal. For backwards compatibility the **speciesType** attribute on **Species** is not mandatory.

It is not possible to locate a **SpeciesType** in a **Compartment** more than once i.e. it is not possible for two **Species** structures to have the same **speciesType** and **compartment** values.

5.3 Generalized Reactions: The Definition of Reactions across Compartments

Just as we might wish to give a common identify to chemical entities distributed across several compartments we might wish to have some common object describing reactions between those chemical entities that is independent of the compartments in which the reactions occur.

For example consider the representation of the transamination reaction, a reversible reaction that converts Aspartate to Oxaloacetate in both the Cytosol and Mitochondrial Matrix.

```

<model id="malate_aspartate_shuttle2">
  <listOfCompartments>
    <compartment id="Cytosol"/>
    <compartment id="Mitochondrial_Matrix"/>
  </listOfCompartments>
  <listOfSpeciesTypes>
    <speciesType id="Aspartate"/>
  </listOfSpeciesTypes>
  <listOfSpecies>
    <species
      id="Aspartate_in_Cytosol"
      speciesType="Aspartate"
      compartment="Cytosol"/>
    <species
      id="Aspartate_in_Mitochondrial_Matrix"
      speciesType="Aspartate"
      compartment="Mitochondrial_Matrix"/>
  </listOfSpecies>
</model>

```

Figure 4: *malate_aspartate_shuttle2* Model which uses a *SpeciesType* to link species of the same type of chemical entity that are located in different compartments.

We could extend *malate_aspartate_shuttle2* using *SpeciesType* structures combined with other SBML Level 2 structures the SBML Level 2 form as shown in Figure 5

Under this proposal we can replace the 2 reactions in Figure 5 with a single reaction structure as shown in Figure 6

The reaction structure represents the set of reactions which occurs in all compartments where the reactant or product are located. This reaction would only occur where the reactant is located if the reaction was not reversible. All the *SimpleSpeciesReference* structures (that is modifiers, reactants and products) refer to species in the same compartment. This means that, under this proposal, it is not possible to define a transport reaction, that is a reaction which moves chemical entities between compartments, using this simple form. However a variant form is described in section 5.3.1 which employs a similar form to transport reactions.

5.3.1 Defining the explicit location of a *SimpleSpeciesReference*

The location of a species pool can be made explicit in a *SimpleSpeciesReference* structure without referring to a *Species* structure. This can be achieved by using the proposed optional *compartment* field which refers to a *Compartment* structure to indicate the location of the given *SpeciesType*.

For example consider the transport reaction shown in Figure 7 which can be added to the model in Figure 6.

This feature could be introduced later in the SBML development road map. It is however an essential component of features introduced later.

All the *SimpleSpeciesReference* structures of a reaction should simultaneously either (a) be located (i.e. have values for the *species* or *compartment* attributes); or (b) apply to any compartment (i.e. not have values for the *species* and *compartment* attributes). *This restriction is not essential but simplifies the interpretation of the proposed format.*

5.3.2 Defining Kinetic Laws for Generalized Reactions

As defined in the examples above it is not possible to compose the kinetic law of these generalized reactions since there is no symbol that refers to either the modifiers, reactants or products or the reaction species pools. However under this proposal the *id* field of a *SimpleSpeciesReference* becomes a symbol that can be used in the *KineticLaw* of the enclosing *Reaction*.

*Here I am assuming that the *id* field on *SimpleSpeciesReference* is introduced by a new version of SBML Level 2. This *id* field is in the global symbol namespace despite, for the purposes of this proposal, only having*

```

<model id="malate_aspartate_shuttle3">
  <listOfCompartments>
    <compartment id="Cytosol"/>
    <compartment id="Mitochondrial_Matrix"/>
  </listOfCompartments>
  <listOfSpeciesTypes>
    <speciesType id="Aspartate"/>
    <speciesType id="Oxaloacetate"/>
  </listOfSpeciesTypes>
  <listOfSpecies>
    <species
      id="Aspartate_in_Cytosol"
      speciesType="Aspartate"
      compartment="Cytosol"/>
    <species
      id="Aspartate_in_Mitochondrial_Matrix"
      speciesType="Aspartate"
      compartment="Mitochondrial_Matrix"/>
    <species
      id="Oxaloacetate_in_Cytosol"
      speciesType="Oxaloacetate"
      compartment="Cytosol"/>
    <species
      id="Oxaloacetate_in_Mitochondrial_Matrix"
      speciesType="Oxaloacetate"
      compartment="Mitochondrial_Matrix"/>
  </listOfSpecies>
  <listOfReactions>
    <reaction id="Transamination_in_Cytosol" reversible="true">
      <listOfReactants>
        <speciesReference species="Aspartate_in_Cytosol"/>
      </listOfReactants>
      <listOfProducts>
        <speciesReference species="Oxaloacetate_in_Cytosol"/>
      </listOfProducts>
    </reaction>
    <reaction id="Transamination_in_Mitochondrial_Matrix" reversible="true">
      <listOfReactants>
        <speciesReference species="Aspartate_in_Mitochondrial_Matrix"/>
      </listOfReactants>
      <listOfProducts>
        <speciesReference species="Oxaloacetate_in_Mitochondrial_Matrix"/>
      </listOfProducts>
    </reaction>
  </listOfReactions>
</model>

```

Figure 5: The *malate_aspartate_shuttle3* model which has duplicate reactions for each compartment.

scope in the enclosing *Reaction*. If this is problematic then perhaps we could consider an additional attribute to declare the symbol.

As example Figure 8 shows the *Transamination* reaction, from model *malate_aspartate_shuttle4*, modified to include a rate law.

5.3.3 The Unit Attributes of SimpleSpeciesReference

To make the units of species explicit in kinetic laws under this proposal *SimpleSpeciesReference* structures have the attributes *substanceUnits*, *spatialSizeUnits* and *hasOnlySubstanceUnits*. These have the same semantics as the corresponding attributes on *Species*. When a *SimpleSpeciesReference* structure matches with a *Species* structure the unit attributes of the *SimpleSpeciesReference* structure should default to the *Species* attributes and/or be exactly equivalent to them.

EXAMPLE HERE

```

<model id="malate_aspartate_shuttle4">
  <listOfCompartments>
    <compartment id="Cytosol"/>
    <compartment id="Mitochondrial_Matrix"/>
  </listOfCompartments>
  <listOfSpeciesTypes>
    <speciesType id="Aspartate"/>
    <speciesType id="Oxaloacetate"/>
  </listOfSpeciesTypes>
  <listOfSpecies>
    <species
      id="Aspartate_in_Cytosol"
      speciesType="Aspartate"
      compartment="Cytosol"/>
    <species
      id="Aspartate_in_Mitochondrial_Matrix"
      speciesType="Aspartate"
      compartment="Mitochondrial_Matrix"/>
    <species
      id="Oxaloacetate_in_Cytosol"
      speciesType="Oxaloacetate"
      compartment="Cytosol"/>
    <species
      id="Oxaloacetate_in_Mitochondrial_Matrix"
      speciesType="Oxaloacetate"
      compartment="Mitochondrial_Matrix"/>
  </listOfSpecies>
  <listOfReactions>
    <reaction id="Transamination" reversible="true">
      <listOfReactants>
        <speciesReference speciesType="Aspartate"/>
      </listOfReactants>
      <listOfProducts>
        <speciesReference speciesType="Oxaloacetate"/>
      </listOfProducts>
    </reaction>
  </listOfReactions>
</model>

```

Figure 6: The *malate_aspartate_shuttle4* model which has a single reaction which is potentially located in all compartments.

```

<reaction id="Malate_Transport" reversible="false">
  <listOfReactants>
    <speciesReference speciesType="Malate" compartment="Cytosol"/>
  </listOfReactants>
  <listOfProducts>
    <speciesReference speciesType="Malate" compartment="Mitochondrial_Matrix"/>
  </listOfProducts>
</reaction>

```

Figure 7: The *Malate_Transport* model a transport reaction which refers to *SpeciesType* structures in specific compartments.

5.4 Implied Species

Under this proposal **Species** structures are used to indicate the initial conditions and/or attributes of species and don't represent the complete set of species. In fact this proposal does not assume that an interpreter (e.g. simulator) of models in the proposed format would use species as it's fundamental representational form, for example an interpreter may represent individual chemical entities as distinct objects. In SBML Level 2 the model's **species** list is a complete enumeration of the pools of chemical entities. Instead, in this proposal, the set of **Species** structures which have an undefined or non-zero initial amount or concentration are the *initial species* and can be used as a starting point to infer the complete set of species in the model. The complete set of species located in a given compartment can be inferred by traversing the reaction network,


```

<reaction id="Transamination" reversible="true">
  <listOfReactants>
    <speciesReference id="S1" speciesType="Aspartate"/>
  </listOfReactants>
  <listOfProducts>
    <speciesReference speciesType="Oxaloacetate"/>
  </listOfProducts>
  <kineticLaw>
    <math xmlns="http://www.w3.org/1998/MathMathML">
      <apply>
        <times/>
        <cn>1.1</cn>
        <ci>S1</ci>
      </apply>
    </math>
  </kineticLaw>
</reaction>

```

Figure 8: The Transamination reaction from Figure 6 modified to include a kinetic law.

defined by the set of **Reaction** structures, from the initial species that are located in the compartment.

So if we consider the model `malate_aspartate_shuttle4` the existence of `Oxaloacetate` in the compartment `Cytosol` can be inferred given the reaction `Transamination` and the species `Aspartate_in_Cytosol`. This means we can omit the species `Oxaloacetate_in_Cytosol` from the `malate_aspartate_shuttle4` as shown in Figure 9.

```

<model id="malate_aspartate_shuttle5">
  <listOfCompartments>
    <compartment id="Cytosol"/>
    <compartment id="Mitochondrial_Matrix"/>
  </listOfCompartments>
  <listOfSpeciesTypes>
    <speciesType id="Aspartate"/>
    <speciesType id="Oxaloacetate"/>
  </listOfSpeciesTypes>
  <listOfSpecies>
    <species
      id="Aspartate_in_Cytosol"
      speciesType="Aspartate"
      compartment="Cytosol"/>
    <species
      id="Aspartate_in_Mitochondrial_Matrix"
      speciesType="Aspartate"
      compartment="Mitochondrial_Matrix"/>
  </listOfSpecies>
  <listOfReactions>
    <reaction id="Transamination" reversible="true">
      <listOfReactants>
        <speciesReference speciesType="Aspartate"/>
      </listOfReactants>
      <listOfProducts>
        <speciesReference speciesType="Oxaloacetate"/>
      </listOfProducts>
    </reaction>
  </listOfReactions>
</model>

```

Figure 9: The `malate_aspartate_shuttle5` model with a reduced set of initial **Species** structures.

The concept of implied species could be introduced later in the SBML development road map. It is however an essential component of features introduced later.

Just as it is not possible to explicitly locate a **SpeciesType** in **Compartment** more than once implied species cannot duplicate each other or explicit **Species** structures. Any process that computes the set of implied

species must compute the equivalence of species. This is described in detail in Section 6.3.

Implied species, which don't match any **Species** structures, always have an initial concentration or substance amount of zero and are never constant nor boundary conditions. This means that constant or boundary condition pools or pools with any initial concentration must be made explicit using a **Species** structure. Two **SimpleSpeciesReference** structures that refer to the same implied species should have the same units.

Reactions of the form used in the example reaction **Malate_Transport** fit in with this scheme. If **Malate** exists in the Cytosol then we can infer the existence of a pool of **Malate** in the Mitochondrial Matrix. Given this form of reaction **Species** structures are only required to define the initial conditions of a model.

5.5 Simple Multi-Component Chemical Entities

In this proposal **SpeciesType** structures can be composed from instances of other **SpeciesType** structures. These instances are encoded as **SpeciesTypeInstance** structures. For example see the **Pheromone_Response** model, shown with XML and diagrammatic form in Figure 10.

```
<model "Pheromone_response">
  <listOfSpeciesTypes>
    <speciesType id="Ste5"/>
    <speciesType id="Ste11"/>
    <speciesType id="Ste7"/>
    <speciesType id="Fus3"/>
    <speciesType id="SteComplex">
      <listOfSpeciesTypeInstances>
        <speciesTypeInstance id="iSte5" speciesType="Ste5"/>
        <speciesTypeInstance id="iSte11" speciesType="Ste11"/>
        <speciesTypeInstance id="iSte7" speciesType="Ste7"/>
        <speciesTypeInstance id="iFus3" speciesType="Fus3"/>
      </listOfSpeciesTypeInstances>
    </speciesType>
  </listOfSpeciesTypes>
</model>
```

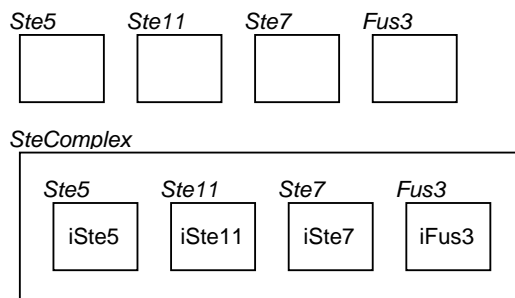


Figure 10: *The Pheromone_response model*

This indicates that **SteComplex** is a complex made up of one instance each of the proteins **Ste5**, **Ste11**, **Ste7** and **Fus3**. The individual instances are always identified to enable the components of homodimers to be separately identified.

We can also describe reactions in using this form on **SpeciesReference** structures. For example we can describe the binding of **Ste11** to **Ste5** with the **Reaction** structure shown in Figure 11.

Although the the identity of **SpeciesTypeInstance** structures is declared in each **SimpleSpeciesReference** structure these identities have scope throughout a reaction. The **SpeciesTypeInstance** **id** fields with the same value in the same reaction refer to the same chemical entity. By giving **SpeciesTypeInstance** **id** fields the same values in the reactants and products of a reaction we indicate that the entity is only modified by the reaction rather than being created or destroyed by the reaction. For example the reaction **binding_Ste5_Ste11** could be encoded as shown in Figure 12.

```

<reaction id="binding_Ste5_Ste11">
  <listOfReactants>
    <speciesReference speciesType="Ste11"/>
    <speciesReference speciesType="Ste5"/>
  </listOfReactants>
  <listOfProducts>
    <speciesReference>
      <listOfSpeciesTypeInstances>
        <speciesTypeInstance id="iSte5" speciesType="Ste5"/>
        <speciesTypeInstance id="iSte11" speciesType="Ste11"/>
      </listOfSpeciesTypeInstances>
    </speciesReference>
  </listOfProducts>
</reaction>

```




Figure 11: The *binding_Ste5_Ste11* reaction, operating in the context of the species types defined in Figure 10. This indicates that the reaction *binding_Ste5_Ste11* creates a complex consisting of *Ste5* and *Ste11* entities from unbound *Ste5* and *Ste11* entities.

```

<reaction id="binding_Ste5_Ste11_v2">
  <listOfReactants>
    <speciesReference>
      <listOfSpeciesTypeInstances>
        <speciesTypeInstance id="iSte5" speciesType="Ste5"/>
      </listOfSpeciesTypeInstances>
    </speciesReference>
    <speciesReference>
      <listOfSpeciesTypeInstances>
        <speciesTypeInstance id="iSte11" speciesType="Ste11"/>
      </listOfSpeciesTypeInstances>
    </speciesReference>
  </listOfReactants>
  <listOfProducts>
    <speciesReference>
      <listOfSpeciesTypeInstances>
        <speciesTypeInstance id="iSte5" speciesType="Ste5"/>
        <speciesTypeInstance id="iSte11" speciesType="Ste11"/>
      </listOfSpeciesTypeInstances>
    </speciesReference>
  </listOfProducts>
</reaction>

```




Figure 12: The *binding_Ste5_Ste11_v2* reaction

This encoding indicates that, for the purposes of the model, *Ste5* and *Ste11* are not modified when they bond. The distinction between *binding_Ste5_Ste11* and *binding_Ste5_Ste11_v2* is only descriptive however the latter form is used as the basis for more complex semantics later.

A model like *pheromone_response* with or without characterized reactions including kinetic laws doesn't encapsulate a model that can be simulated because it does not specify any initial species.

This proposal makes distinction between **simple** **SpeciesType** structures and **complex** **SpeciesType** structures. Simple **SpeciesType** structures do not contain any **SpeciesTypeInstance** structures whereas complex **SpeciesType** structures do.

5.6 Multi-component Chemical Entities with explicit bonds

The forms described in section 5.5 capture some but not all the relevant knowledge of chemical entities that we might wish to model. In this section I describe how chemical bond information is captured. The bond information on a **SimpleSpeciesReference** or a **SpeciesType** is a graph linking the **SpeciesTypeInstance** structures together. The description of a structure using chemical bonds requires the identification of binding sites on **SpeciesType** structures, using **BindingSite** structures, and then the enumeration of pairs of bonds on those binding sites using **SpecificBond** structures. **SpecificBond** structures consist of pairs of **BindingSiteReference** structures. We can redefine the model **pheromone_response** along these lines as shown in Figure 13 on page 14.

SpecificBond structures can be used to indicate unbound binding sites, simply by containing a single **BindingSiteReference** structure, as shown the reaction in Figure 14 on page 15.

A **SpeciesType** structure must not leave the state of a binding site undefined or ambiguous. Section 5.7 describes how a **SimpleSpeciesReference** can refer to several different **SpeciesType** structures where the class represents a range of states for one or more binding sites.

The level of decomposition of a biochemical system into chemical entities and their binding sites and bonds is not defined by this proposal. This proposal is designed to support arbitrary decomposition schemes which capture knowledge at different resolutions in the same model. The underlying chemistry represented by a given binding site state is also not defined by this proposal.

An underlying principle of this proposal is that the binding representation described in this section can be used to represent the reversible covalent modification of proteins including, for example, phosphorylation and dephosphorylation. The example model shown in Figure 15 on page 16 represents the phosphorylation of **Ste11** by **Ste20**. A diagram of this model is shown in Figure 16 on page 17.

This model deliberately does not model the involvement of ATP or ADP molecules demonstrating how the level of detail of the biological knowledge captured by the proposed standard is arbitrary. As a result not all the instances of species types in the list of reactants are present in the list of products. This is valid in this proposal: the structural details of chemical entity transformation do not have to be fully elucidated. In fact the reaction shown in Figure 17 on page 17 is valid even if it is implausible from a biochemical perspective.

SpeciesGraph structures, that is **SpeciesType** and **SimpleSpeciesReference** structures, can contain a number of disconnected components (as described previously in Section 5.5. This means that a list of **Bond** structures in a **SpeciesGraph** does not have to comprise a connected graph. In this case the **SpeciesGraph** still represents a single entity where the complete set of bonds is not specified. As an example the consider the model shown in Figure 18 on page 18. A diagram of this model is shown in Figure 19 on page 19.

5.7 Reactions generalized to cover classes of Multi-component Chemical Entities

Under this proposal the bonding concept is extended in reactions so that it is possible for a reactant, product or modifier to refer to set of closely related species that have a similar but not identical chemical structure. This is achieved the use of **GenericBond** structures within **SimpleSpeciesReference** structures.

GenericBond is a alternative **Bond** type. Reactions containing containing a **GenericBond** can potentially apply to a large set of complex species types including those not explicitly defined in the model thus reducing the number of reactions, complex species types and species that need to be enumerated in a given system. So just as the complete set of species that the reaction set operates does not need to be enumerated nor does the the complete set of species types need to be enumerated.

When a reaction is applied to the given state of one or more reactants, a **GenericBond** structure in the reaction is assigned the state of a binding site. The assignment can be to an empty entity if the match is

to an empty binding site or to a unspecified binding site on an unspecified chemical entity. The whatever is assigned to the **GenericBond** in the set reactants is transferred to the set of products.

The simple abstract example model shown in Figure 20 on page 20 uses this generalization mechanism redundantly. A diagram of this model is shown in Figure 21 on page 21. The reaction **generic** defines how entities **A** and **B** bind together without changing the state of one of the binding sites on **A**.

A more concrete example model is shown in fragments in Figures 22 to 24 on pages 21 to 23. The reactions in Figures 23 and 24 (which operate on the species types encoded in Figure 22) taken together represent the fact that the binding of **Ste11** to **Ste50** is not mutually exclusive to **Ste11** binding to **Ste5**. Figure 23 shows reaction **bind_Ste11_Ste50** which binds **Ste11** to **Ste50** and is generalized to cover all states of the **Ste11** to **Ste5** binding site. Figure 24 shows reaction **bind_Ste11_Ste5** binding **Ste11** to **Ste5** and is generalized to cover all states of the **Ste11** to **Ste50** binding site.

```

<model "pheromone_response_v2">
  <listOfSpeciesTypes>
    <speciesType id="Ste5">
      <listOfBindingSites>
        <bindingSite id="r241">
        <bindingSite id="r463">
        <bindingSite id="r744">
      </listOfBindingSites>
    </speciesType>
    <speciesType id="Ste11"/>
      <listOfBindingSites>
        <bindingSite id="site">
      </listOfBindingSites>
    </speciesType>
    <speciesType id="Ste7"/>
      <listOfBindingSites>
        <bindingSite id="site">
      </listOfBindingSites>
    </speciesType>
    <speciesType id="Fus3"/>
      <listOfBindingSites>
        <bindingSite id="site">
      </listOfBindingSites>
    </speciesType>
    <speciesType id="SteComplex">
      <listOfSpeciesTypeInstances>
        <speciesTypeInstance id="iSte5" speciesType="Ste5"/>
        <speciesTypeInstance id="iSte11" speciesType="Ste11"/>
        <speciesTypeInstance id="iSte7" speciesType="Ste7"/>
        <speciesTypeInstance id="iFus3" speciesType="Fus3"/>
      </listOfSpeciesTypeInstances>
      <listOfBonds>
        <specificBond>
          <bindingSiteReference speciesTypeInstance="iSte5" bindingSite="r241"/>
          <otherBindingSiteReference
            speciesTypeInstance="iFus3" bindingSite="site"/>
        </specificBond>
        <specificBond>
          <bindingSiteReference speciesTypeInstance="iSte5" bindingSite="r463"/>
          <otherBindingSiteReference
            speciesTypeInstance="iSte11" bindingSite="site"/>
        </specificBond>
        <specificBond>
          <bindingSiteReference speciesTypeInstance="iSte5" bindingSite="r744"/>
          <otherBindingSiteReference
            speciesTypeInstance="iSte7" bindingSite="site"/>
        </specificBond>
      </listOfBonds>
    </speciesType>
  </listOfSpeciesTypes>
</model>

```

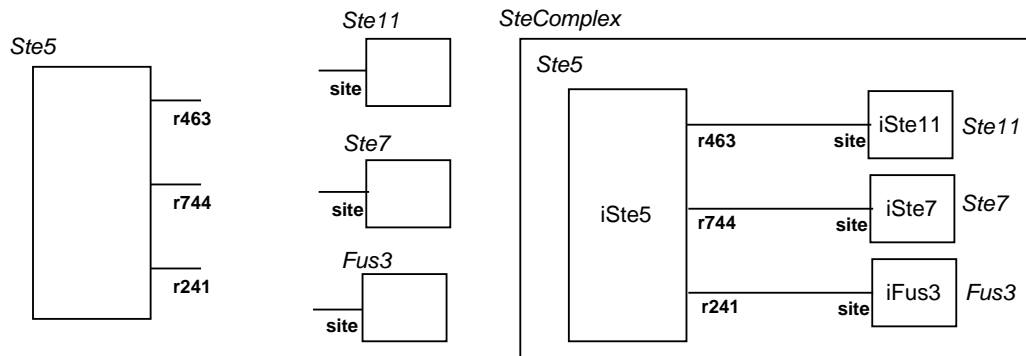


Figure 13: The *pheromone_response_v2* model.

```

<reaction id="binding_Ste5_Ste11_v3">
  <listOfReactants>
    <speciesReference>
      <listOfSpeciesTypeInstances>
        <speciesTypeInstance id="iSte5" speciesType="Ste5"/>
      </listOfSpeciesTypeInstances>
      <listOfBonds>
        <specificBond>
          <bindingSiteReference speciesTypeInstance="iSte5" bindingSite="r241"/>
        </specificBond>
        <specificBond>
          <bindingSiteReference speciesTypeInstance="iSte5" bindingSite="r463"/>
        </specificBond>
        <specificBond>
          <bindingSiteReference speciesTypeInstance="iSte5" bindingSite="r744"/>
        </specificBond>
      </listOfBonds>
    </speciesReference>
    <speciesReference>
      <listOfSpeciesTypeInstances>
        <speciesTypeInstance id="iSte11" speciesType="Ste11"/>
      </listOfSpeciesTypeInstances>
      <listOfBonds>
        <specificBond>
          <bindingSiteReference speciesTypeInstance="iSte11" bindingSite="site"/>
        </specificBond>
      </listOfBonds>
    </speciesReference>
  </listOfReactants>
  <listOfProducts>
    <speciesReference>
      <listOfSpeciesTypeInstances>
        <speciesTypeInstance id="iSte5" speciesType="Ste5"/>
        <speciesTypeInstance id="iSte11" speciesType="Ste11"/>
      </listOfSpeciesTypeInstances>
      <listOfBonds>
        <specificBond>
          <bindingSiteReference speciesTypeInstance="iSte5" bindingSite="r241"/>
        </specificBond>
        <specificBond>
          <bindingSiteReference speciesTypeInstance="iSte5" bindingSite="r463"/>
          <otherBindingSiteReference
            speciesTypeInstance="iSte11" bindingSite="site"/>
        </specificBond>
        <specificBond>
          <bindingSiteReference speciesTypeInstance="iSte5" bindingSite="r744"/>
        </specificBond>
      </listOfBonds>
    </speciesReference>
  </listOfProducts>
</reaction>

```

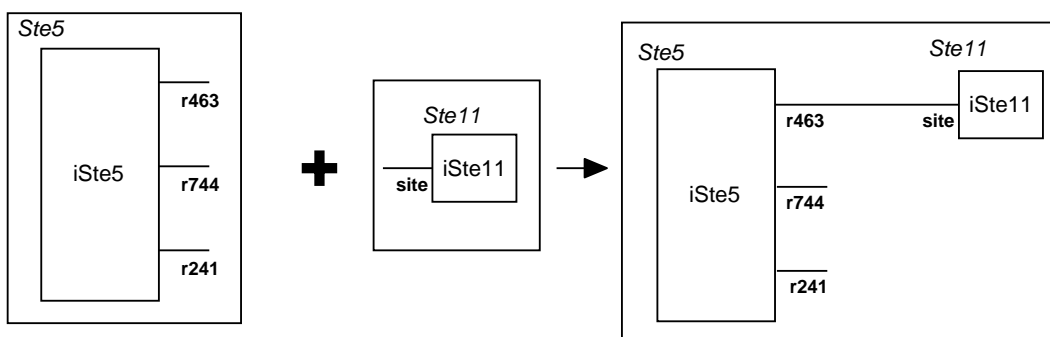


Figure 14: The *binding_Ste5_Ste11_v3* reaction

```

<model id="Phosphorylation_model">
  <listOfSpeciesTypes>
    <speciesType id="Phosphate">
      <listOfBindingSites>
        <bindingSite id="site"/>
      </listOfBindingSites>
    </speciesType>
    <speciesType id="Ste20"/>
    <speciesType id="Ste11">
      <listOfBindingSites>
        <bindingSite id="S302"/>
        <bindingSite id="T307"/>
      </listOfBindingSites>
    </speciesType>
  </listOfSpeciesTypes>
  <listOfReactions>
    <reaction id="Phosphorylation">
      <listOfReactants>
        <speciesReference>
          <listOfSpeciesTypeInstances>
            <speciesTypeInstance id="iSte11" speciesType="Ste11"/>
          </listOfSpeciesTypeInstances>
          <listOfBonds>
            <specificBond>
              <bindingSiteReference
                speciesTypeInstance="iSte11" bindingSite="S302"/>
            </specificBond>
            <specificBond>
              <bindingSiteReference
                speciesTypeInstance="iSte11" bindingSite="T307"/>
            </specificBond>
          </listOfBonds>
        </speciesReference>
      </listOfReactants>
      <listOfProducts>
        <speciesReference>
          <listOfSpeciesTypeInstances>
            <speciesTypeInstance id="iSte11" speciesType="Ste11"/>
            <speciesTypeInstance id="iPhosphate_1" speciesType="Phosphate"/>
            <speciesTypeInstance id="iPhosphate_2" speciesType="Phosphate"/>
          </listOfSpeciesTypeInstances>
          <listOfBonds>
            <specificBond>
              <bindingSiteReference speciesTypeInstance="iSte11"
                bindingSite="S302"/>
              <otherBindingSiteReference
                speciesTypeInstance="iPhosphate_1" bindingSite="site"/>
            </specificBond>
            <specificBond>
              <bindingSiteReference
                speciesTypeInstance="iSte11" bindingSite="T307"/>
              <otherBindingSiteReference
                speciesTypeInstance="iPhosphate_2" bindingSite="site"/>
            </specificBond>
          </listOfBonds>
        </speciesReference>
      </listOfProducts>
      <listOfModifiers>
        <modifierSpeciesReference>
          <listOfSpeciesTypeInstances>
            <speciesTypeInstance id="iSte20" speciesType="Ste20"/>
          </listOfSpeciesTypeInstances>
        </modifierSpeciesReference>
      </listOfModifiers>
    </reaction>
  </listOfReactions>
</model>

```

Figure 15: The *Phosphorylation_model* model, a diagram of this model is shown in Figure 16.




Figure 16: Diagram of the Phosphorylation_model model

```

<model id="demo">
  <listOfSpeciesTypes>
    <speciesType id="Ste20"/>
    <speciesType id="Ste11"/>
  </listOfSpeciesTypes>
  <listOfReactions>
    <reaction id="Implausible">
      <listOfReactants>
        <speciesReference>
          <listOfSpeciesTypeInstances>
            <speciesTypeInstance id="iSte11" speciesType="Ste11"/>
          </listOfSpeciesTypeInstances>
        </speciesReference>
      </listOfReactants>
      <listOfProducts>
        <speciesReference>
          <listOfSpeciesTypeInstances>
            <speciesTypeInstance id="iSte20" speciesType="Ste20"/>
          </listOfSpeciesTypeInstances>
        </speciesReference>
      </listOfProducts>
    </reaction>
  </listOfReactions>
</model>

```




Figure 17: The demo model which shows how a reaction operating on components can transform those components.

```

<model id="disconnected_parts">
  <listOfSpeciesTypes>
    <speciesType id="A"/>
    <speciesType id="B">
      <listOfBindingSites>
        <bindingSite id="b"/>
      </listOfBindingSites>
    </speciesType>
    <speciesType id="C">
      <listOfBindingSites>
        <bindingSite id="c"/>
      </listOfBindingSites>
    </speciesType>
    <speciesType id="D">
      <listOfSpeciesTypeInstances>
        <speciesTypeInstance id="iA" speciesType="A"/>
        <speciesTypeInstance id="iB" speciesType="B"/>
        <speciesTypeInstance id="iC" speciesType="C"/>
      </listOfSpeciesTypeInstances>
      <listOfBonds>
        <specificBond>
          <bindingSiteReference speciesTypeInstance="iB" bindingSite="b"/>
          <otherBindingSiteReference speciesTypeInstance="iC" bindingSite="c"/>
        </specificBond>
      </listOfBonds>
    </speciesType>
  </listOfSpeciesTypes>
  <listOfReactions>
    <listOfReactants>
      <speciesReference>
        <listOfSpeciesTypeInstances>
          <speciesTypeInstance id="iA" speciesType="A"/>
        </listOfSpeciesTypeInstances>
      </speciesReference>
      <speciesReference>
        <listOfSpeciesTypeInstances>
          <speciesTypeInstance id="iB" speciesType="B"/>
          <speciesTypeInstance id="iC" speciesType="C"/>
        </listOfSpeciesTypeInstances>
        <listOfBonds>
          <specificBond>
            <bindingSiteReference speciesTypeInstance="iB" bindingSite="b"/>
            <otherBindingSiteReference speciesTypeInstance="iC" bindingSite="c"/>
          </specificBond>
        </listOfBonds>
      </speciesReference>
    </listOfReactants>
    <listOfProducts>
      <speciesReference>
        <listOfSpeciesTypeInstances>
          <speciesTypeInstance id="iA" speciesType="A"/>
          <speciesTypeInstance id="iB" speciesType="B"/>
          <speciesTypeInstance id="iC" speciesType="C"/>
        </listOfSpeciesTypeInstances>
        <listOfBonds>
          <specificBond>
            <bindingSiteReference speciesTypeInstance="iB" bindingSite="b"/>
            <otherBindingSiteReference speciesTypeInstance="iC" bindingSite="c"/>
          </specificBond>
        </listOfBonds>
      </speciesReference>
    </listOfProducts>
  </listOfReactions>
</model>

```

Figure 18: *The disconnected_parts model which shows how reactions can operate on disconnected components.*

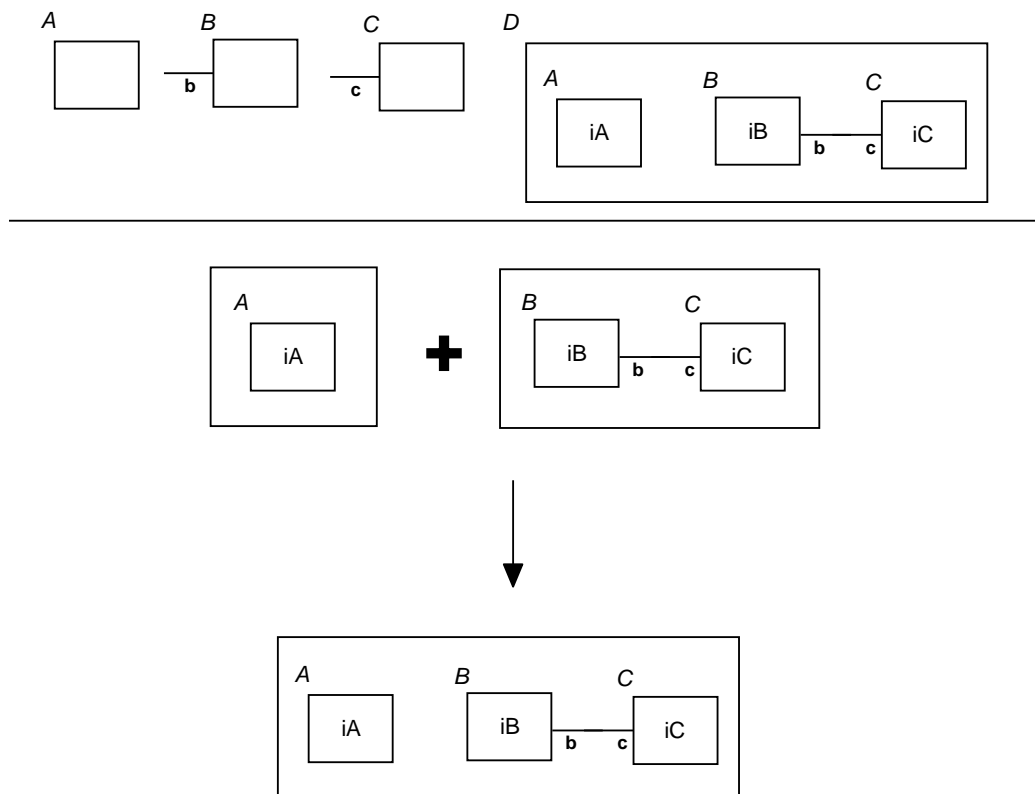


Figure 19: A diagram of the *disconnected_parts* model.

```

<model "generalized">
  <listOfSpeciesTypes>
    <speciesType id="A">
      <listOfBindingSites>
        <bindingSite id="a"/>
      </listOfBindingSite>
    </speciesType>
    <speciesType id="B">
      <listOfBindingSites>
        <bindingSite id="b1"/>
        <bindingSite id="b2"/>
      </listOfBindingSites>
    </speciesType>
  </listOfSpeciesTypes>
  <listOfReactions>
    <reaction id="generic">
      <listOfReactants>
        <speciesReference>
          <listOfSpeciesTypeInstances>
            <speciesTypeInstance id="iA" speciesType="A"/>
          </listOfSpeciesTypeInstances>
          <listOfBonds>
            <specificBond>
              <bindingSiteReference speciesTypeInstance="iA" bindingSite="a"/>
            </specificBond>
          </listOfBonds>
        </speciesReference>
        <speciesReference>
          <listOfSpeciesTypeInstances>
            <speciesTypeInstance id="iB" speciesType="B"/>
          </listOfSpeciesTypeInstances>
          <listOfBonds>
            <specificBond>
              <bindingSiteReference speciesTypeInstance="iB" bindingSite="b1"/>
            </specificBond>
            <genericBond id="X">
              <bindingSiteReference speciesTypeInstance="iB" bindingSite="b2"/>
            </genericBond>
          </listOfBonds>
        </speciesReference>
      </listOfReactants>
      <listOfProducts>
        <speciesReference>
          <listOfSpeciesTypeInstances>
            <speciesTypeInstance id="iA" speciesType="A"/>
            <speciesTypeInstance id="iB" speciesType="B"/>
          </listOfSpeciesTypeInstances>
          <listOfBonds>
            <specificBond>
              <bindingSiteReference speciesTypeInstance="iA" bindingSite="a"/>
              <bindingSiteReference speciesTypeInstance="iB" bindingSite="b1"/>
            </specificBond>
            <genericBond id="X">
              <bindingSiteReference speciesTypeInstance="iB" bindingSite="b2"/>
            </genericBond>
          </listOfBonds>
        </speciesReference>
      </listOfProducts>
    </reaction>
  </listOfReactions>
</model>

```

Figure 20: The generalized model which shows how a reaction can be applied to a set of chemical entities. A diagram of this model is shown in Figure 21.




Figure 21: A diagram of the generalized model shown in Figure 20

```

<listOfSpeciesTypes>
  <speciesType id="Ste5">
    <listOfBindingSites>
      <bindingSite id="R463_514"/>
    </listOfBindingSites>
  </speciesType>
  <speciesType id="Ste50">
    <listOfBindingSites>
      <bindingSite id="SAM"/>
    </listOfBindingSites>
  </speciesType>
  <speciesType id="Ste11">
    <listOfBindingSites>
      <bindingSite id="N_term"/>
      <bindingSite id="SAM"/>
    </listOfBindingSites>
  </speciesType>
</listOfSpeciesTypes>

```




Figure 22: The species types used in Figures 23, 24 and 25.

```

<reaction "bind_Ste11_Ste50">
  <listOfProducts>
    <speciesReference>
      <listOfSpeciesTypeInstances>
        <speciesTypeInstance id="iSte11" speciesType="Ste11"/>
      </listOfSpeciesTypeInstances>
      <listOfBonds>
        <genericBond id="X">
          <bindingSiteReference bindingSite="N_term" speciesTypeInstance="iSte11"/>
        </genericBond>
        <specificBond>
          <bindingSiteReference bindingSite="SAM" speciesTypeInstance="iSte11"/>
        </specificBond>
      </listOfBonds>
    </speciesReference>
    <speciesReference>
      <listOfSpeciesTypeInstances>
        <speciesTypeInstance id="iSte50" speciesType="Ste50"/>
      </listOfSpeciesTypeInstances>
      <listOfBonds>
        <specificBond>
          <bindingSiteReference bindingSite="SAM" speciesTypeInstance="iSte50"/>
        </specificBond>
      </listOfBonds>
    </speciesReference>
  </listOfProducts>
  <listOfReactants>
    <speciesReference>
      <listOfSpeciesTypeInstances>
        <speciesTypeInstance id="iSte11" speciesType="Ste11"/>
        <speciesTypeInstance id="iSte50" speciesType="Ste50"/>
      </listOfSpeciesTypeInstances>
      <listOfBonds>
        <genericBond id="X">
          <bindingSiteReference bindingSite="N_term" speciesTypeInstance="iSte11"/>
        </genericBond>
        <specificBond>
          <bindingSiteReference bindingSite="SAM" speciesTypeInstance="iSte11"/>
          <bindingSiteReference bindingSite="SAM" speciesTypeInstance="iSte50"/>
        </specificBond>
      </listOfBonds>
    </speciesReference>
  </listOfReactants>
</reaction>

```




Figure 23: The *bind_Ste11_Ste50* binding reaction that operate on the types defined in Figure 22

```

<reaction "bind_Ste11_Ste5">
  <listOfProducts>
    <speciesReference>
      <listOfSpeciesTypeInstances>
        <speciesTypeInstance id="iSte11" speciesType="Ste11"/>
      </listOfSpeciesTypeInstances>
      <listOfBonds>
        <genericBond id="X">
          <bindingSiteReference bindingSite="SAM" speciesTypeInstance="iSte11"/>
        </genericBond>
        <specificBond>
          <bindingSiteReference bindingSite="N_term" speciesTypeInstance="iSte11"/>
        </specificBond>
      </listOfBonds>
    </speciesReference>
    <speciesReference>
      <listOfSpeciesTypeInstances>
        <speciesTypeInstance id="iSte5" speciesType="Ste5"/>
      </listOfSpeciesTypeInstances>
      <listOfBonds>
        <specificBond>
          <bindingSiteReference bindingSite="R463_514" speciesTypeInstance="iSte5"/>
        </specificBond>
      </listOfBonds>
    </speciesReference>
  </listOfProducts>
  <listOfReactants>
    <speciesReference>
      <listOfSpeciesTypeInstances>
        <speciesTypeInstance id="iSte11" speciesType="Ste11"/>
        <speciesTypeInstance id="iSte5" speciesType="Ste5"/>
      </listOfSpeciesTypeInstances>
      <listOfBonds>
        <genericBond id="X">
          <bindingSiteReference bindingSite="SAM" speciesTypeInstance="iSte11"/>
        </genericBond>
        <specificBond>
          <bindingSiteReference bindingSite="N_term" speciesTypeInstance="iSte11"/>
          <bindingSiteReference bindingSite="R463_514" speciesTypeInstance="iSte5"/>
        </specificBond>
      </listOfBonds>
    </speciesReference>
  </listOfReactants>
</reaction>

```

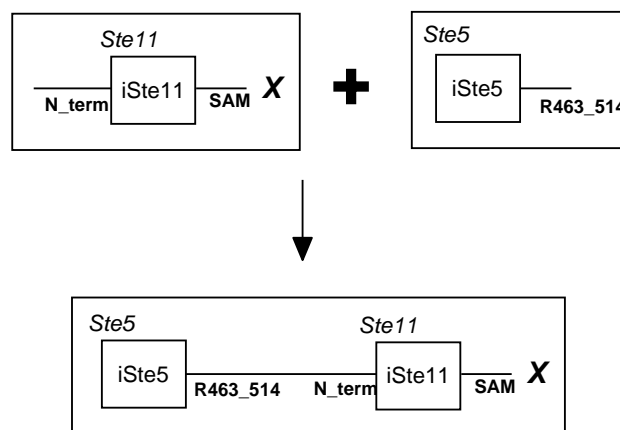


Figure 24: The `bind_Ste11_Ste5` reaction that operates on the types defined in Figure 22

5.7.1 Missing Binding Sites Infers Unchanged State

Labelling the connection point of generic bonds enables the modeler a reaction which move an component of unspecified type from one binding site to another. However in the majority of cases the binding site of such a component is not changed by the reaction i.e. the binding site state is completely irrelevant to the reaction. The encoding of these cases is simplified: if a binding site is both unchanged by a reaction *and* the reaction generalized to cover all states of that binding site then that binding site is simply omitted from the reaction entirely. This is the case for the binding sites referenced by the **GenericBond** structures in the **bind_Ste11_Ste5** reaction shown in Figure 24 on page 23 and thus we can simplify this reaction by omitting the **GenericBond** structures as shown in Figure 25 on page 25.

This simple generalization syntax can't be applied to **SpeciesType** structures. The state of all binding sites must be resolved in a **SpeciesType**.

5.8 Hierarchal Species Types and Type Equivalence

This proposal supports the hierarchial assembly of **SpeciesType** structures to an arbitrary depth. The examples referenced so far have deliberately used only structures of limited hierarchal depth. In section I will review the support for hierarchial assembly in the proposal.

SpeciesType structures can encapsulate a graph of instances of species types whilst exposing a subset of the available binding sites. The implementation of this consists of a reference, on a **BindingSite** structure, to a binding site on a chemical entity internal to the **SpeciesType**. The example model in Figure 26 on Page 26 demonstrates this feature.

This proposal contains a simple scheme for species type equivalence (described in more detail in Section 6.2). This equivalence is used for resolving whether for example the products of different reactions refer to the same species. In this scheme a distinction is made between species types enclosing one or more other species type instances and those do not. I'll call those types which do not contain species type instances *simple* all other types are *complex*. Simple species types are never equivalent however complex species types can be. To evaluate the equivalence of *complex* types we first normalize them into an equivalent form where all species type instances are *simple* species types. This normalization process simply removes the intermediate levels in the hierarchy. Species type equivalence then considers a normalized type as a graph formed by the species type instances, which are graph nodes, and bonds, which are graph arcs. Two species types are equivalent if their graphs are equivalent. The identity of species types instances normalized complex species types are not relevant for the purposes of comparing these graphs. The binding sites of normalized complex species types are also not relevant.

The **Complex2 SpeciesType** in Figure 27 on Page 27 is equivalent to the **Complex SpeciesType** in Figure 26 on Page 26.


```

<reaction "bind_Ste11_Ste5">
  <listOfProducts>
    <speciesReference>
      <listOfSpeciesTypeInstances>
        <speciesTypeInstance id="iSte11" speciesType="Ste11"/>
      </listOfSpeciesTypeInstances>
      <listOfBonds>
        <specificBond>
          <bindingSiteReference bindingSite="N_term" speciesTypeInstance="iSte11"/>
        </specificBond>
      </listOfBonds>
    </speciesReference>
    <speciesReference>
      <listOfSpeciesTypeInstances>
        <speciesTypeInstance id="iSte5" speciesType="Ste5"/>
      </listOfSpeciesTypeInstances>
      <listOfBonds>
        <specificBond>
          <bindingSiteReference bindingSite="R463_514" speciesTypeInstance="iSte5"/>
        </specificBond>
      </listOfBonds>
    </speciesReference>
  </listOfProducts>
  <listOfReactants>
    <speciesReference>
      <listOfSpeciesTypeInstances>
        <speciesTypeInstance id="iSte11" speciesType="Ste11"/>
        <speciesTypeInstance id="iSte5" speciesType="Ste5"/>
      </listOfSpeciesTypeInstances>
      <listOfBonds>
        <specificBond>
          <bindingSiteReference bindingSite="N_term" speciesTypeInstance="iSte11"/>
          <bindingSiteReference bindingSite="R463_514" speciesTypeInstance="iSte5"/>
        </specificBond>
      </listOfBonds>
    </speciesReference>
  </listOfReactants>
</reaction>

```

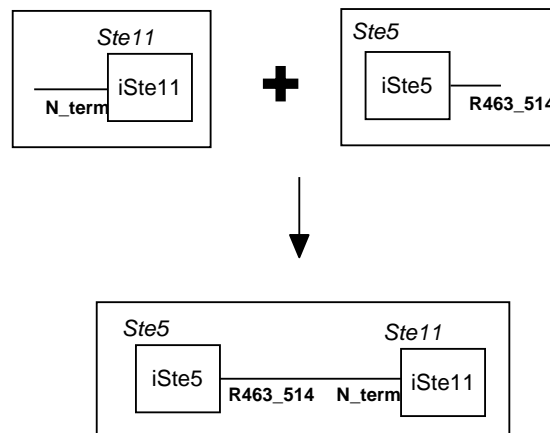


Figure 25: The `bind_Ste11_Ste5.v2` reaction that operates on the types defined in Figure 22 and is a simplification of the `bind_Ste11_Ste5` reaction shown in Figure 24. The SAM binding site is not changed by this reaction and the reaction is generalized to cover all states of that binding site. As a result the SAM binding site can and has been omitted from the model.

```

<model id="Hierarchical">
  <listOfSpeciesTypes>
    <speciesType id="A">
      <listOfBindingSites>
        <bindingSite id="site"/>
      </listOfBindingSites>
    </speciesType>
    <speciesType id="B">
      <listOfBindingSites>
        <bindingSite id="x"/>
        <bindingSite id="y"/>
      </listOfBindingSites>
    </speciesType>
    <speciesType id="C">
      <listOfSpeciesTypeInstances>
        <speciesTypeInstance id="iA" speciesType="A"/>
        <speciesTypeInstance id="iB" speciesType="B"/>
      </listOfSpeciesTypeInstances>
      <listOfBonds>
        <specificBond>
          <bindingSiteReference speciesTypeInstance="iA" bindingSite="site"/>
          <otherBindingSiteReference speciesTypeInstance="iB" bindingSite="x"/>
        </specificBond>
      </listOfBonds>
      <listOfBindingSites>
        <bindingSite id="p">
          <bindingSiteReference speciesTypeInstance="iB" bindingSite="y"/>
        </bindingSite>
      </listOfBindingSites>
    </speciesType>
    <speciesType id="Complex">
      <listOfSpeciesTypeInstances>
        <speciesTypeInstance id="iA" speciesType="A"/>
        <speciesTypeInstance id="iC" speciesType="C"/>
      </listOfSpeciesTypeInstances>
      <listOfBonds>
        <specificBond>
          <bindingSiteReference speciesTypeInstance="iA" bindingSite="site"/>
          <otherBindingSiteReference speciesTypeInstance="iC" bindingSite="p"/>
        </specificBond>
      </listOfBonds>
    </speciesType>
  </listOfSpeciesTypes>
</model>

```




Figure 26: The Hierarchical model which demonstrates the graph encapsulation facilities of the proposal. The BindingSite p on SpeciesType C exposes the uncommitted BindingSite y on SpeciesTypeInstance iB.

```

<speciesType id="Complex2">
  <listOfSpeciesTypeInstances>
    <speciesTypeInstance id="iA1" speciesType="A"/>
    <speciesTypeInstance id="iA2" speciesType="A"/>
    <speciesTypeInstance id="iB" speciesType="C"/>
  </listOfSpeciesTypeInstances>
  <listOfBonds>
    <specificBond>
      <bindingSiteReference speciesTypeInstance="iA1" bindingSite="site"/>
      <otherBindingSiteReference speciesTypeInstance="iB" bindingSite="y"/>
    </specificBond>
    <specificBond>
      <bindingSiteReference speciesTypeInstance="iA2" bindingSite="site"/>
      <otherBindingSiteReference speciesTypeInstance="iB" bindingSite="x"/>
    </specificBond>
  </listOfBonds>
</speciesType>

```




Figure 27: *The Complex2 SpeciesType that is equivalent to the Complex type in Figure 26 on Page 26*

6 Formal Definition of Proposal

6.1 Proposed Classes in Detail

This section describes in detail each class in the proposal as shown in figure 1. As in the diagram only new or extended classes are described in this section. All level 2 structures and basic semantics are assumed to be part of the proposal. For each new or extended class the definition of the class and its fields are described.

`id` fields on structures, described here, enclosed within a **SpeciesType** structure are unique to that structure only. `id` fields on structures, described here, enclosed within a **Reaction** structure are unique to that reaction only (apart from the `id` fields of **SimpleSpeciesReferences** which are unique to a whole model).

The definition of **species** and **speciesType** may not be as one would expect in conventional biochemical terminology because **species** retains its definition from SBML Level 1 and 2.

6.1.1 Bond

The abstract base class **Bond** represents one or more chemical bonds or forces holding two chemical entities together enabling them to form a larger chemical entity. A **bond** can represent covalent and non-covalent bonds. The existence of a **bond** can imply some modification of the chemical entities. For example a phosphorylated protein can be represented in a model as a **bond** between a protein and a phosphate group. In such a model the loss of the hydrogen atom bound to the phosphorylation site may or may not be represented explicitly and does not affect the identity of the protein.

A **Bond** structure consists of one **BindingSiteReference** field, **bindingSite**, which indicates where the bond has effect.

The linkage of chemical entities together to form a larger chemical entity can be represented without using **bonds**. See section 5.5.

6.1.2 BindingSite

A logical site where a bond may form on the chemical entity represented by the enclosing **SpeciesType** structure. A **BindingSite** may represent a set of physical binding sites which are treated as a single entity for the purposes of the model. A **BindingSite** structure consists of

- `id`, a mandatory `SId` field to identify the site
- `name`, an optional string field (see SBML Level 2)
- **bindingSiteReference**, an optional field containing a **BindingSiteReference** structure. When this attribute is present the binding site is exposing an internal binding site on a **SpeciesTypeInstance**. **bindingSiteReference** refers to a binding site internal to the enclosing **SpeciesType** which is not referenced by another **bindingSiteReference**.

A **bindingSiteReference** value must be present if the enclosing **SpeciesType** contains one or more **SpeciesTypeInstance** structures. This means that a **SpeciesType** cannot 'introduce' a new binding site that doesn't exist on the chemical entities that make up the **SpeciesType**. This restriction is designed to ensure that evaluating type equivalence is straightforward.

6.1.3 BindingSiteReference

A reference to an instance of a **BindingSite** in a given **SpeciesGraph**. A **BindingSiteReference** structure consists of

- **speciesTypeInstance**, a `SId` field, which refers to a **SpeciesTypeInstance** within the enclosing **SpeciesGraph**; and
- **bindingSite**, a `SId` field, which refers to a **binding site** on that instance. The **bindingSite** must be declared on the **SpeciesType** referenced by the **speciesTypeInstance**.

The combination of attribute values of a given **BindingSiteReference** structure cannot occur in any other **BindingSiteReference** structure within the same **SpeciesGraph** structure.

6.1.4 GenericBond

Represents a chemical bond between a specific binding site and an unspecified chemical entity which is unchanged by the reaction. **GenericBond** is a subtype of **Bond** and in turn **Sbase**.

On **GenericBond** the **id** field, inherited from **SBase**, becomes mandatory and identifies a connection to the unspecified chemical entity. The **bindingSiteReference** field represents the binding site that the entity is connected to. **GenericBond** structures can only occur within **Reactions**. (Specifically they can only occur in the **bond** array/list field in **SimpleSpeciesReference**.) The values of **GenericBond** **id** fields are specific to one reaction. All **GenericBond** **id** fields with the same value in the same reaction refer to the same chemical entity. **GenericBond** **id** fields in different reactions refer to different chemical entities.

6.1.5 Model

See SBML Level 2 for the existing definition of **model**. This proposal adds a **speciesType** field which consists of a list of **SpeciesType** structures to **Model** structures.

This proposal changes the definition of the **species** list on **model**. This list does not necessarily comprise the complete set of pools of chemical entities in the model. The set of **Species** structures in this list which have an undefined or non-zero initial amount or concentration, *initial species*, are used to infer the complete set of pools of chemical entities in the model. The complete set of species located in a given compartment can be inferred by traversing the reaction network, defined by the set of **Reaction**, from the initial species that are located in the compartment.

6.1.6 Reaction

A reaction is either located inside a specific compartment, across more than 2 or more compartments or potentially in all compartments. This final case is indicated by the absence of **compartment** and **species** fields on all enclosed **SimpleSpeciesReference** structures. In this case the **SimpleSpeciesReference** structures match with **Species** in the same **Compartment** across all compartments in the model. See section 5.3

A **SpeciesTypeInstance** **id** field assignment cannot appear more than once within the set of reactant **SpeciesReference** structures on a given reaction. A **SpeciesTypeInstance** **id** field assignment cannot appear more than once within the set of product **SpeciesReference** structures on a given reaction.

A **BindingSiteReference** **bindingSite** field assignment must appear in both of reactant and product **SpeciesReference** sets on a given reaction. This means that a binding site state determined in the reactants can not be 'lost' from the set of products.

6.1.7 SimpleSpeciesReference

See SBML Level 2 for the existing definition of **SimpleSpeciesReference**. **SimpleSpeciesReference** is the base class for: (a) **SpeciesReference** the type used to represent the reactants and products of a reaction and (b) **ModifierSpeciesReference** the type used to represent the modifiers of a reaction. **SimpleSpeciesReference** structures can only occur within a reaction.

In this proposal a **SimpleSpeciesReference** can refer to a set of **Species** both those that are explicitly defined and those that are created through generalized reactions (see section 5.7). (A **SimpleSpeciesReference** on its own does not imply the existence of a **Species**.)

In this proposal a **SimpleSpeciesReference** becomes a subtype of **SpeciesGraph**.

SimpleSpeciesReference has the following fields:

- **id**, this optional **SId** field is not introduced by this proposal but instead it has a new role: the value of this field can be used as a symbol, enclosed in **MathMLci** elements, within the **KineticLaw** structure of the enclosing **Reaction** structure.

- **substanceUnits**, **spatialSizeUnits** and **hasOnlySubstanceUnits**, these optional fields have the same semantics as the corresponding attributes on **Species**. These attributes default to the values of matching **Species** structures before following the **Species** semantics. **SimpleSpeciesReference** structures that match with a **Species** structure should default to the **Species** attributes and/or be exactly equivalent to them. Two **SimpleSpeciesReference** structures that refer to the same inferred species should have the same units.
- **species**, this **SIId** field is present in Level 2 however we now make this field optional. This field refers to a **Species** that is involved in the reaction. If this field is present then the fields inherited from **SpeciesGraph** as well as the **compartment**, **bond** and **speciesType** fields are not available.
- **speciesType**, this **SIId** field refers to a **SpeciesType** that is involved in the reaction. If this field is present then the fields inherited from **SpeciesGraph** as well as the **Species** and **bond** fields are not available. If the **compartment** field is present then the **SimpleSpeciesReference** refers to the **Species** of the given **SpeciesType** located in the given **compartment**; otherwise the **SimpleSpeciesReference** refers to a set of **Species** of the given **SpeciesType**.
- **compartment**, this **SIId** field refers to a **Compartment** where the matching **Species** are located. If this field is not present a given **SimpleSpeciesReference** structure then it must not be present on any other **SimpleSpeciesReference** structure in the same enclosing **reaction**.

*We might want to consider restricting stoichiometry if **SpeciesTypeInstance** structures are used.*

6.1.8 Species

See SBML Level 2 for the existing definition of **Species**. In this proposal a **Species** structure represents a pool of a given chemical entity located in a specific compartment. This proposal introduces one optional **SIId** field, **speciesType** which refers to the **SpeciesType** (chemical entity) to be located in the **Compartment** referenced by the **compartment** field. There can only one **Species** structure in a model with a given pair of values for the **speciesType** and **compartment** attributes i.e. a given **SpeciesType** cannot be located in the same **Compartment** more than once.

revise this: When the **speciesType** field is not present then the **Species** structure is equivalent to a **Species** structure which does contain a **speciesType** field. This field would refer to a **SpeciesType** that is not referenced anywhere else in the model. In short a **Species** structure without a **speciesType** field has a 'hidden' **SpeciesType** associated with it.

6.1.9 SpeciesGraph

SpeciesGraph is an abstract base class. A **SpeciesGraph** structure represents a chemical entity or a set of chemical entities of a specific common form. The form of these entities is defined as a graph where the nodes are **SpeciesTypeInstances** and the arcs are **Bond** structures. The graph can be disconnected indicating that the detail of how parts of the chemical entities are associated are not relevant to the model (see 5.5).

A **SpeciesGraph** structure is composed of the following fields:

- **speciesTypeInstance**, this is an optional list of **SpeciesTypeInstance** structures that form the **SpeciesGraph**. If this list is not present then the **SpeciesGraph** simply represents an chemical entity for which the detail of its composition is not relevant to the model.
- **unboundSite**, this is an optional list of **BindingSiteReference** structures which refer to the binding sites on the **SpeciesTypeInstance** structures in the **SpeciesGraph** which are unbound (not part of a **Bond**). In model an unbound binding site may have some implied chemical structure which is not made explicit i.e. the binding site may not be physically unoccupied and the entity occupying the site is simply not modelled.
- **bond**, an optional list of **Bond** structures that can include both **SpecificBond** and **GenericBond** structures. This list links the chemical entities enumerated in the **speciesTypeInstance** field.

6.1.10 SpeciesType

The class **SpeciesType** represents a specific chemical entity. **SpeciesType** is derived from **SpeciesGraph**, and has the following fields:

- **id** a mandatory **SIId** field that identifies the **SpeciesType**
- **name**, an optional string field (see SBML Level 2)
- **bindingSite**, an optional **BindingSite** list, which contains the set of binding sites that are located on the **SpeciesType**.

A simple example of the use **SpeciesType** structures is given in section 5.2.

Consider all the **BindingSite** structures of the **SpeciesType** structures referenced by the species type instances in a given **SpeciesType** structure. Each of these **BindingSite** structures should be referenced exactly once by a **Bond** structure enclosed in the **SpeciesType** structure. This means that the status of a **BindingSite** can't be left undefined or ambiguous by a **SpeciesType**.

6.1.11 SpeciesTypeInstance

A **SpeciesTypeInstance** structure represents the occurrence of a chemical entity of a given **SpeciesType** within a **SpeciesGraph**. A **SpeciesTypeInstance** structure has the following fields:

- **id** a mandatory **SIId** field that identifies the **SpeciesTypeInstance**. This field is unique to the enclosing **SpeciesGraph** structure and the enclosing **Reaction** structure if it exists.
- **name**, an optional string field (see SBML Level 2)
- **speciesType**, a mandatory **SIId** field which refers to the **SpeciesType** that the **SpeciesTypeInstance** is an instance of.

6.1.12 SpecificBond

A subtype of **Bond**. **SpecificBond** represents either (a) one or more chemical bonds between two explicitly identified binding sites or (b) an unoccupied binding site. A **bond** structure consists of two structures **bindingSite** and **otherBindingSiteReference** which are **bindingSiteReference** structures.

bindingSite is inherited from **Bond**. **otherBindingSiteReference** is optional. The **SpecificBond** structure represents the state in which **bindingSite** is unbound if **otherBindingSiteReference** is not present. If **otherBindingSiteReference** is present the **bindingSite** and **otherBindingSiteReference** fields represent the 2 binding sites that are linked by a bond. Neither binding site has privileged semantics. See section 5.6 for examples.

6.2 Equivalence of Species Types

This section defines the equivalence of **SpeciesType** structures. This is defined as a process with 3 stages described by the following sections in order. First the **SpeciesType** structures are transformed into a standard *complex* form of **SpeciesGraph**, as described in Section 6.2.1. Second these **SpeciesGraph** are normalized (flattened), as described in Section 6.2.2. Finally if these normalized structures can be matched, as described in Section 6.2.3, then the **SpeciesType** structures are equivalent.

6.2.1 Normalization of Species Types

This section describes how a **SpeciesType** is normalized.

A **SpeciesGraph** is *simple* if it contains no **SpeciesTypeInstance** structures. A **SpeciesGraph** is *complex* if it contains one or more **SpeciesTypeInstance** structures.

A normalized **SpeciesType** is always complex. A simple **SpeciesType** structure is normalized by creating a **SpeciesGraph** that contains one **SpeciesTypeInstance** structure which refers to the simple **SpeciesType**.

All binding sites of the simple **SpeciesType** are referenced as unbound in a set of **SpecificBonds** within the new normalized **SpeciesGraph**.

A complex **SpeciesType** with one or more **BindingSite** structures is transformed to remove the those structures. Each **BindingSite** structure is replaced by a **SpecificBond** structure containing just the **BindingSiteReference** structure that was contained in the **BindingSite** structure. This means that the original binding sites are made unbound by the normalization process (instead of being available for binding in a different context).

Once the rules have been applied the **SpeciesType** is normalized as described in Section 6.2.2.

6.2.2 Normalization of Species Graphs

This section describes how the form of species graphs is normalized for the matching process. This normalization process simply reduces the given **Species Graph** to a single hierarchical level.

A **SpeciesGraph** structure (i.e. either a **SpeciesType** or **SimpleSpecies** structure) is normalized if the set of **SpeciesType** structures, referred to by the set of **SpeciesTypeInstance** structures, are simple. The normalization process consists of iteratively replacing any **SpeciesTypeInstance** structures that refer to complex **SpeciesType** structures with a set of new **SpeciesTypeInstance** structures and **SpecificBond** structures that are copies of that occurring within the complex **SpeciesType** structure. These new structures are linked into the same 'outer' **Bond** structures as the initial complex **SpeciesType** structure.

6.2.3 Normalized Species Graph Equivalence

When evaluating the equivalence of **SpeciesGraph** structures the following aspects are not directly relevant:

- id on **SpeciesTypeInstance**
- the order of structures within lists **SpecificBond** structures representing unbound binding sites

The id on **SpeciesTypeInstance** is only used to provide linkage within a single graph.

In this section we consider a normalized **SpeciesGraph** to be a formal graph and equivalence to be a graph matching operation. A formal representation of graph is

Definition (A Simple Graph) A *simple graph* G is a tuple $G = (G_V, G_E, L, I, J, s, t, \ell, i, j)$ consisting of

- a finite set of nodes (or “vertices”) G_V and a finite set of arcs (or “edges”) G_E where $G_V \cap G_E = \emptyset$,
- two total mappings $s, t : G_E \rightarrow G_V$ (“source and target”),
- a set of node labels L ,
- a set of arc source labels I ,
- a set of arc target labels J ,
- a total mapping $\ell : G_V \rightarrow L$ (“node labelling”)
- a total mapping $i : G_E \rightarrow I$ (“arc source labelling”)
- a total mapping $j : G_E \rightarrow J$ (“arc target labelling”)

The nodes and arcs of a graph are also collectively called the “objects” of the graph (or “graph objects”). Note that in this definition that we are labelling the source and target 'ends' of each arc.

modified from (Rudolf, 1998)

In this formulation of the formal graphs the graph nodes, G_V , are the **SpeciesTypeInstance** structures and the arcs, G_E , are the **SpecificBond** structures. The source and target of an arc is determined by the

speciesTypeInstance field of the **BindingSiteReference** structures enclosed in the given **SpeciesBond**. The arc direction (the distinction between the source and target) is determined by a consistent ordering over **speciesType** attributes of the **speciesTypeInstance** structures.

The label, $\ell(v)$ of a node, v is its **speciesType** attribute. The arc source label, $i(e)$ of an arc, e , is formed from the **bindingSite** attribute of the given source **BindingSiteReference** structure. The arc target label, $j(e)$ of an arc, e , is formed from the **bindingSite** attribute of the given target **BindingSiteReference** structure.

Two **SpeciesGraph** structures are equivalent if their normalized forms as graphs match in both directions. Formally a match of one graph into another is given by a graph morphism, which is a mapping of one graph's object sets into the other's, with some restrictions to preserve the graph's structure and its typing information:

Definition (Graph Morphism) A *graph morphism* $m : L \rightarrow G$ between two simple graphs

- $L = (L_V, L_E, L_L, I_L, J_L, s_L, t_L, \ell_L, i_L, j_L)$ and
- $G = (G_V, G_E, L_G, I_G, J_G, s_G, t_G, \ell_G, i_G, j_G)$

is a pair of total mappings $m = (m_V : L_V \rightarrow G_V, m_E : L_E \rightarrow G_E)$, where the following restrictions apply:

1. $\forall e \in L_E :$
 - $m_V(s_L(e)) = s_G(m_E(e))$
 - $m_V(t_L(e)) = t_G(m_E(e))$
2. $\forall v \in L_V : \ell_L(v) = \ell_G(m_V(v))$
3. $\forall e \in L_E : i_L(e) = i_G(m_E(e))$
4. $\forall e \in L_E : j_L(e) = j_G(m_E(e))$

modified from (Rudolf, 1998)

6.2.4 Implications of SpeciesType Equivalence

Two or more equivalent **SpeciesType** structures can co-exist in the same model (the **id** attributes must have different values). Section 6.3 describes how **Species** equivalence is derived from **SpeciesType** equivalence.

6.3 Species Equivalence

Two **Species** are equivalent if they are located in the same compartment and their associated **SpeciesType** structures are equivalent. A model must not contain equivalent **Species**. The species implied by the reactions in a model, described in Section 6.4, are inherently not equivalent.

6.4 Semantics of Reactions

6.4.1 Simple Framework for Operational Semantics

For the definition of the semantics of reactions we will consider a model with a simulator to be form of AI *production system*.

The major elements of an AI production system are a *global database*, a set of *production rules*, and a *control system*. [...] Depending on the application, this database may be as simple as a small matrix of numbers or as complex as a large, relational, index file structure. (The reader should not confuse the phrase, “global database,” as it is used [here], with the databases of database systems.)

The production rules operate on the global database. Each rule has a *precondition* that is either satisfied or not by the global database. If the precondition is satisfied, the rule can be *applied*. Application of the rule changes the database. The control system chooses which applicable rule should be applied and ceases computation when a termination condition on the global database is satisfied. (Nilsson, 1982)

In the context of this proposal reactions are considered to be rules. The global database is formed by the modelled species which are pools of chemical entities. In SBML Level 2 the set of **Species** structures defines the complete set of pools of chemical entities that are created by, destroyed by and affect the set of reactions. In this proposal the set of **Species** structures only describes the initial conditions of the model. The reactions themselves define the species or set of pools of entities that are relevant to the model. It is expected that any analysis of a model will require the computation of the set of species. This proposal does not depend on any particular representation scheme for species within a given software analysis system.

A reaction matches the set of reactants, its precondition, to the set of species, the global database. The effect of the reaction is to remove reactant entities and create product entities within the set of species. The existence of a modifier entity is not a pre-condition of a reaction however the number of chemical entities in the modifier's pool will affect the reaction according to the reaction's rate law.

For the purposes of this definition the control system is idealized. Real software systems will in practice create approximations of this control system. The ideal control system simply applies all matching reactions concurrently at the rate defined by the reaction's kinetic laws. Unlike a AI production system this proposal does not define any termination conditions.

6.4.2 Matching of Species

This section describes how reactions are matched to the global database of species.

Unification of Species to Reactants and Modifiers

The matching of reactant and modifiers, **SimpleSpeciesReference** structures, to species is similar to equivalence between species. The cases we will consider with respect to the **SimpleSpeciesReference** are:

- *species attribute is present*, in this case the match is with the indicated **Species**.
- *speciesType and compartment attributes are present*, in this case the match is with the species located in the given **Compartment** which has a **SpeciesType** equivalent to the given **SpeciesType**.
- *speciesType is present but the compartment attribute is not present*, in this case the match is with any species which has a **SpeciesType** equivalent to the given **SpeciesType**. The **Compartment** of the species must be the same as all other species involved in the reaction.
- *all of the remaining cases* are described below.

In the remaining cases the matching between **SimpleSpeciesReference** and species is considered to be a variant of **SpeciesType** equivalence. If the **compartment** attribute is not present on the **SimpleSpeciesReference** then the match is with any species matching the **SpeciesType** located in the same **Compartment** as all the other species involved in the reaction. If the **compartment** attribute is present then the match is with any species matching the **SpeciesType** located in the given **Compartment**.

To enable the definition of a match between a **SimpleSpeciesReference** and a **SpeciesType** we first require a definition of a *generic graph* and definition of **SimpleSpeciesReference** in terms of a *generic graph*. (The use *generic graph* here is to capture the semantics of **GenericBond** structures which may be contained within a **SimpleSpeciesReference**.)

Definition (A Generic Graph) A *generic graph* G is a tuple $G = (G_V, G_E, G_X, L, I, J, s, t, u, \ell, i, j)$ consisting of

- a finite set of nodes (or "vertices") G_V , a finite set of arcs (or "edges") G_E and a finite set of generic nodes G_X where

1. $G_V \cap G_E = \emptyset$
 2. $G_V \cap G_X = \emptyset$
 3. $G_E \cap G_X = \emptyset$
- two total mappings $s, t : G_E \rightarrow G_V$ (“source and target”),
 - a total mapping $u : G_X \rightarrow G_V$ (“generic connection”),
 - a set of arc source labels I ,
 - a set of arc target labels J ,
 - a total mapping $\ell : G_V \cup G_X \rightarrow L$ (“node labelling”)
 - a total mapping $i : G_E \rightarrow I$ (“arc source labelling”)
 - a total mapping $j : G_E \rightarrow J$ (“arc target labelling”)

The nodes, generic nodes and arcs of a graph are also collectively called the “objects” of the generic graph (or “generic graph objects”).

The formulation of a **SimpleSpeciesReference** as a generic graph similar to that of a **SpeciesType** to a simple graph. The graph nodes, G_V , are the **SpeciesTypeInstance** structures, arcs, G_E , are the **SpecificBond** structures and generic nodes, G_X , are the **GenericBond** structures. The source, s and target, t , of an arc are determined by the **speciesTypeInstance** field of the **BindingSiteReference** structures enclosed in the given **SpeciesBond**. The arcs are directed where arc direction should be determined using an consistent ordering over the set of simple species types. The generic connections of a generic node, u are determined by the **speciesTypeInstance** field of the **BindingSiteReference** structure enclosed in the given **GenericBond**. For our purposes the label or type of a node is its **speciesType** attribute and the label or type of an arc is formed from the pair of **bindingSite** attributes of the given **SpecificBond** structures. The arc direction determines the ordering of each arc’s **bindingSite** attributes in the arc’s label. The label of generic node is the **bindingSite** attribute of the given **GenericBond** structure.

A match of generic graph (from a **SimpleSpeciesReference**) into a simple graph (from a **SpeciesType** as described in Section 6.2.3) is given by a graph morphism, which is a mapping of the generic graph’s object sets into the simple graph’s, with some restrictions to preserve the generic graph’s structure and it’s typing information:

Definition (Generic Graph Morphism) A *generic graph morphism* $m : L \rightarrow G$ between

- a generic graph $L = (L_V, L_E, L_X, L_L, I_L, J_L, s_L, t_L, u_L, \ell_L, i_L, j_L)$ and
- a simple graph $G = (G_V, G_E, L_G, I_G, J_G, s_G, t_G, \ell_G, i_G, j_G)$

is a pair of total mappings $m = (m_V : L_V \rightarrow G_V, m_E : L_E \cup L_X \rightarrow G_E)$ where the following restrictions apply:

1. $\forall x \in L_X : m_V(u(x)) = s_G(m_E(x)) \vee m_V(u(x)) = t_G(m_E(x))$
2. $\forall e \in L_E :$
 - $m_V(s_L(e)) = s_G(m_E(e))$
 - $m_V(t_L(e)) = t_G(m_E(e))$
3. $\forall v \in L_V : \ell_L(v) = \ell_G(m_V(v))$
4. $\forall e \in L_E : i_L(e) = i_G(m_E(e))$
5. $\forall e \in L_E : j_L(e) = j_G(m_E(e))$

$$6. \forall x \in L_X : \ell_L(x) = \begin{cases} i_G(m_E(x)) & \text{if } m_V(u(x)) = s_G(m_E(x)) \\ j_G(m_E(x)) & \text{if } m_V(u(x)) = t_G(m_E(x)) \end{cases}$$

The mapping m_E from the set of generic bonds, L_X to specific bonds, G_E represents a set of assignments which are used to determine the effect of a reaction. This is described further in Section 6.4.4.

6.4.3 Matching Reactions

A **Reaction** which does not contain any **GenericBond** structures represents a single *reaction instance*. A **Reaction** that does contain **GenericBond** structures represents a set of reaction instances where each instance in this set is a concrete reaction with generic graph mappings for all the reactant and modifier **SimpleSpeciesReference** structures. There is a separate reaction instance for all possible combinations of these graph mappings. Each reaction instance exists concurrently with the other instance and operates at the rate determined by the reaction's kinetic law.

*This means that some care must be taken when formulating the kinetic laws of reactions containing **GenericBond** structures. Typically to obtain the correct results kinetic laws should be defined so that the rate of the reaction is directly proportional to the amount of the species matched using **GenericBond** structures.*

6.4.4 The Effect of Reactions

A reaction reduces the amount of the reactant species and increases the amount of the product species. A reaction does not affect the amount of a Modifier species (unless its a reactant or product). If the species graph for a product which doesn't contain **GenericBond** structures it can be matched to a species as for reactants.

In a reaction instance any product **GenericBond** structures are eliminated by viewing the effect of a reaction as graph transformation process. The resulting product species graphs can then be matched to species. In this transformation process we take each set of matching reactant species graphs. All the **SpeciesTypeInstance** and **SpecificBond** structures in these graphs that are mapped from equivalent structures in the reactant **SpeciesReference** structures are discarded. The **SpecificBond** structures that are mapped from **GenericBond** structures modified. The **BindingSiteReference** that matches that contained in the **GenericBond** structure is discarded. A mapping, the *generic bond assignment*, is created from the **GenericBond** structures to the remaining **BindingSiteReference** structures. The result is that subgraphs of the matching reactant species graphs remain which are composed of only those components not made explicit by the corresponding reactant **SpeciesReference** structures.

The product graphs are created starting from the species graphs in the **SpeciesReference** structures. The **GenericBond** structures are replaced by **SpecificBond** structures where one **BindingSiteReference** is taken from the product **GenericBond** and the other **BindingSiteReference** is mapped from the reactant **GenericBond** with the same *id* value via the *generic bond assignment*. These structures are then combined with the subgraphs that are left from the reactant species to form the set of product species graphs. These can then be matched with the global database of species.

6.4.5 Matching more than one SimpleSpeciesReference to a single species

The above description is written from the perspective that the modifiers, reactants and products of a reaction are always separate species. However in practice the modeler may wish to define reactions where the relation between reactant and products varies across the set of reaction instances. As an example consider

6.4.6 Reversible reactions

References

Finney, A. (2001). Internal discussion document possible extension to the Systems Biology Markup Language Complex Species and Species Graphs. Available via the World Wide Web at <http://www.cds.caltech.edu/~{afinney/CplxSpecies.pdf>.

- Finney, A., Hucka, M., and Bolouri, H. (2002). Systems Biology Markup Language (SBML) Level 2: Structures and facilities for model definitions. Available via the World Wide Web at <http://www.sbml.org/>.
- Goldstein, B., Faeder, J. R., Hlavacek, W. S., Blinov, M. L., Redondoc, A., and Wofsy, C. (2001). Modeling the early signaling events mediated by Fc ϵ RI. *Molecular Immunology*, 38:1213–1219.
- Le Novère, N., Shimizu, T. S., and Finney, A. (2003). Systems Biology Markup Language (SBML) Level 3 Proposal: Multistate Features. Available via the World Wide Web at <http://sbml.org/multistates.pdf>.
- Nilsson, N. J. (1982). *Principles of Artificial Intelligence*. Springer-Verlag.
- Phair, R. (2003). Posting to complex species sbml wiki. Available via the World Wide Web at <http://sbml.org/wiki/Complexes>.
- Rudolf, M. (1998). Utilizing constraint satification techniques for efficient graph pattern matching. In *6th International Workshop on Theory and Application of Graph Transformations*.