
Systems Biology Markup Language (SBML) Level 3 Core

Michael Hucka (Chair)	<i>California Institute of Technology, US</i>
Frank Bergmann	<i>University of Washington, US</i>
Stefan Hoops	<i>Virginia Bioinformatics Institute, US</i>
Sarah M. Keating	<i>California Institute of Technology, US</i>
Sven Sahle	<i>University of Heidelberg, DE</i>
Darren J. Wilkinson	<i>Newcastle University, GB</i>

sbml-editors@sbml.org

SBML Level 3 Version 1 Core

Release 1

PUBLIC REVIEW DRAFT—August 22, 2009

Corrections and other changes to this SBML language specification may appear over time.
Notifications of new releases are broadcast on the mailing list sbml-announce@sbml.org

The latest release of the SBML Level 3 Version 1 Core specification is available at
<http://sbml.org/specifications/sbml-level-3/version-1/core>

This release of the specification is available at
<http://sbml.org/specifications/sbml-level-3/version-1/core/release-1/>

The list of known issues in all releases of SBML Level 3 Version 1 Core is available at
<http://sbml.org/specifications/sbml-level-3/version-1/core/errata/>

XML Schemas are available at
<http://sbml.org/xml-schemas/>

Contents

1	Introduction	3
1.1	Developments, discussions, and notifications of updates	3
1.2	SBML Levels, Versions, and Releases	3
1.3	SBML Level 3 Packages	4
1.4	Document conventions	4
2	Overview of SBML	8
3	Preliminary definitions and principles	10
3.1	Primitive data types	10
3.2	Type <i>SBase</i>	13
3.3	The id and name attributes on SBML components	17
3.4	Mathematical formulas in SBML Level 3	18
4	SBML components	29
4.1	The SBML container	29
4.2	Model	30
4.3	Function definitions	34
4.4	Unit definitions	36
4.5	Compartments	41
4.6	Species	43
4.7	Parameters	48
4.8	Initial assignments	50
4.9	Rules	53
4.10	Constraints	57
4.11	Reactions	59
4.12	Events	71
5	The Systems Biology Ontology and the sboTerm attribute	78
5.1	Principles	78
5.2	Using SBO and sboTerm	79
5.3	Relationships to the SBML annotation element	84
5.4	Discussion	85
6	A standard format for the annotation element	86
6.1	Motivation	86
6.2	XML namespaces in the standard annotation	86
6.3	General syntax for the standard annotation	87
6.4	Use of URIs	88
6.5	Relation elements	88
6.6	History	90
6.7	Examples	91
7	Example models expressed in XML using SBML	97
7.1	A simple example application of SBML	97
7.2	Example of a discrete version of a simple dimerization reaction	99
7.3	Example involving assignment rules	102
7.4	Example involving algebraic rules	104
7.5	Example with combinations of boundaryCondition and constant values on Species with RateRule objects	106
7.6	Example of translation from a multi-compartmental model to ODEs	108
7.7	Example involving function definitions	111
7.8	Example involving delay functions	112
7.9	Example involving events	113
7.10	Example involving two-dimensional compartments	115
7.11	Example of a reaction located at a membrane	119
8	Recommended practices	122
8.1	Recommended practices concerning common SBML attributes and objects	122
8.2	Recommended practices concerning specific SBML components	124
A	XML Schema for SBML	130
B	XML Schema for MathML subset	131
C	Validation and Consistency checks for SBML	132
D	A method for assessing whether an SBML model is overdetermined	133
E	Mathematical consequences of the fast attribute on Reaction	136
F	Processing and validating SBase notes and Constraint message content	138
	Acknowledgments	139
	References	140

1 Introduction

This document defines Version 1 of the **S**ystems **B**iology **M**arkup **L**anguage (SBML) Level 3 Core, an electronic model representation format for systems biology. SBML is oriented towards describing biological processes of the sort common in research on a number of topics, including metabolic pathways, cell signaling pathways, and many others. SBML is defined neutrally with respect to programming languages and software encoding; however, it is oriented primarily towards allowing models to be encoded using XML, the eXtensible Markup Language (Bray et al., 2000). This document contains many examples of SBML models written in XML, as well as the text of an XML Schema (Biron and Malhotra, 2000; Fallside, 2000; Thompson et al., 2000) defining the basic syntax of SBML Level 3 Version 1 Core. The Schema and other materials and software are available for downloading from the SBML project web site, <http://sbml.org/>.

The SBML project is not an attempt to define a universal language for representing quantitative models. The rapidly evolving views of biological function, coupled with the vigorous rates at which new computational techniques and individual tools are being developed today, are incompatible with a one-size-fits-all idea of a universal language. A more realistic alternative is to acknowledge the diversity of approaches and methods being explored by different software tool developers, and seek a common intermediate format—a *lingua franca*—enabling communication of the most essential aspects of the models.

The definition of the model description language presented here does not specify *how* programs should communicate or read/write SBML. We assume that for a simulation program to communicate a model encoded in SBML, the program will have to translate its internal data structures to and from SBML, use a suitable transmission medium and protocol, etc., but these issues are outside of the scope of this document.

1.1 Developments, discussions, and notifications of updates

SBML has been, and continues to be, developed in collaboration with an international community of researchers and software developers. As in many projects, the primary mode of interaction between members is electronic mail. Discussions about SBML take place on the mailing list sbml-discuss@caltech.edu. The mailing list archives and a web browser-based interface to the list are available at <http://sbml.org/Forums/>.

A low-volume, broadcast-only mailing list is available where notifications of revisions to the SBML specification, notices of votes on SBML technical issues, and other critical matters are announced. This list is sbml-announce@caltech.edu and anyone may subscribe to it freely. This list will never be used for advertising and its membership list will never be disclosed. *It is vitally important that all users of SBML stay informed about new releases and other developments by subscribing to this list*, even if they do not wish to participate in discussions on the sbml-discuss@caltech.edu list. Please visit <http://sbml.org/>, for information about how to subscribe to sbml-announce@caltech.edu as well as for access to the list archives.

In the [Acknowledgments](#) section, we attempt to acknowledge as many contributors to SBML's development as we can, but as SBML evolves, it becomes increasingly difficult to detail the individual contributions on a project that has truly become an international community effort.

1.2 SBML Levels, Versions, and Releases

Major editions of SBML are termed *levels* and represent substantial changes to the composition and structure of the language. The edition of SBML defined in this document, SBML Level 3, represents an evolution of the language resulting from the practical experiences of users and developers working with SBML since its introduction in the year 2001 (Hucka et al., 2001, 2003). All of the constructs of Level 1 can be mapped to Level 2; likewise, all of the constructs from Level 2 can be mapped to Level 3 (when Level 3 is considered in terms of the Core and Level 3 packages; see next section). In addition, a subset of Level 3 constructs can be mapped to Level 2, and a subset of Level 2 constructs can be mapped to Level 1. However, the levels remain distinct; a valid SBML Level 1 document is not a valid SBML Level 2 document, and so on.

In practice, once a new Level of SBML is defined, no further development is undertaken on lower Levels. An exception is made for the correction of problems and other issues that may be identified in the specifications of lower Levels; such corrections are handled as described below.

Minor revisions of SBML are termed *versions* and constitute changes within a Level to correct, adjust, and refine language features. The present document defines Level 3 Version 1 Core. A separate document provides information about the changes between SBML Level 3 and SBML Level 2.

Specification documents inevitably require minor editorial changes as its users discover errors and ambiguities. As a practical reality, these discoveries occur over time. In the context of SBML, such problems are formally announced publicly as *errata* in a given specification document. Borrowing concepts from the World Wide Web Consortium (Jacobs, 2004), we define SBML errata as changes of the following types: (a) formatting changes that do not result in changes to textual content; (b) corrections that do not affect conformance of software implementing support for a given combination of SBML Level and Version; and (c) corrections that *may* affect such software conformance, but add no new language features. A change that affects conformance is one that either turns conforming data, processors, or other conforming software into non-conforming software, or turns non-conforming software into conforming software, or clears up an ambiguity or insufficiently documented part of the specification in such a way that software whose conformance was once unclear now becomes clearly conforming or non-conforming (Jacobs, 2004). In short, errata do not change the fundamental semantics or syntax of SBML; they clarify and disambiguate the specification and correct errors. (New syntax and semantics are only introduced in SBML Versions and Levels.) An electronic tracking system for reporting and monitoring such issues is available at <http://sbml.org/issue-tracker>.

SBML errata result in new *Releases* of the specification. Each release is numbered, with the first release of the specification being number 1. Subsequent releases of an SBML specification document contain a section for the accumulated issues corrected since the first release. The errata acknowledged in SBML Level 3 Version 1 Core since the publication of Release 1 are also publicly listed at <http://sbml.org/specifications/sbml-level-3/version-1/core/errata/>. Announcements of errata, updates to the SBML specification and other major changes are made on the sbml-announce@caltech.edu mailing list.

1.3 SBML Level 3 Packages

SBML Level 3 is being developed as a modular language, with a central core comprising a self-sufficient model definition language, and extension packages layered on top of this core to provide additional, optional sets of features. This document defines SBML Level 3 Core. The definition is based largely on SBML Level 2, with some modifications to address sources of problems found by experience with Level 2, and some simplifications to remove Level 2 constructs that are expected to be supported more thoroughly through SBML Level 3 packages. Section 4.1.2 describes the mechanism by which models defined in SBML Level 3 can declare which packages they use.

The specifications for packages available for SBML Level 3 is maintained separately on the SBML website at <http://sbml.org/Documents/Specifications>. A list of packages is not provided in *this* specification document (i.e., for Level 3 Core) because the development of packages for Level 3 proceeds independently, and new ones may be introduced over time after Level 3 Core is published. The SBML website provides information ongoing activities in this area, as well as about the process whereby individuals and groups may propose new packages.

1.4 Document conventions

In this section, we describe conventions used in this document to communicate information more effectively.

1.4.1 Color conventions

Through this document, we use blue text to indicate a hyperlink from one point in this document to another. Clicking your computer's pointing device on blue-colored text will cause a jump to the section, figure, table or page to which the link refers. (Of course, this capability is only available when using electronic media that support hyperlinking, such as PDF and HTML.)

1.4.2 Typographical conventions for names

We use the following typographical conventions to distinguish objects and data types from other entities:

AbstractClass: Abstract classes are classes that are never instantiated directly, but rather serve as parents of other classes. Their names begin with a capital letter and they are printed in a slanted, bold, sans-serif typeface. In electronic document formats, the class names are also hyperlinked to their definitions in the specification. For example, in the PDF and HTML versions of this document, clicking on the word [SBase](#) will send the reader to the section containing the definition of this class.

Class: Names of ordinary (concrete) classes begin with a capital letter and are printed in an upright, bold, sans-serif typeface. In electronic document formats, the class names are also hyperlinked to their definitions in the specification. For example, in the PDF and HTML versions of this document, clicking on the word [Species](#) will send the reader to the section containing the definition of this class.

Something, otherThing: Attributes of classes, data type names, literal XML, and generally all tokens *other* than SBML UML class names, are printed in an upright typewriter typeface. Primitive types defined by SBML begin with a capital letter, but unfortunately, XML Schema 1.0 does not follow any convention and primitive XML types may either start with a capital letter (e.g., `ID`) or not (e.g., `double`).

1.4.3 UML notation

Previous specifications of SBML used a notation that was at one time (in the days of SBML Level 1) fairly close to UML, the Unified Modeling Language ([Eriksson and Penker, 1998](#); [Oestereich, 1999](#)), though many details were omitted from the UML diagrams themselves. Over the years, the notation used in successive specifications of SBML grew increasingly less UML-like. Beginning with SBML Level 2 Version 3, we have completely overhauled the specification's use of UML and once again define the XML syntax of SBML using, as much as possible, proper and complete UML 1.0. We then systematically map this UML notation to XML, using XML Schema 1.0 ([Biron and Malhotra, 2000](#); [Fallside, 2000](#); [Thompson et al., 2000](#)) to express the overall syntax of SBML. In the rest of this section, we summarize the UML notation used in this document and explain the few embellishments needed to support transformation to XML form. A complete Schema for SBML is given in Appendix A.

We see three main advantages to using UML as a basis for defining SBML data objects. First, compared to using other notations or a programming language, the UML visual representations are generally easier to grasp by readers who are not computer scientists. Second, the notation is implementation-neutral: the objects can be encoded in any concrete implementation language—not just XML, but C, Java and other languages as well. Third, UML is a de facto industry standard that is documented in many resources. Readers are therefore more likely to be familiar with it than other notations.

Object class definitions

Object classes in UML diagrams are drawn as simple tripartite boxes, as shown in Figure 1 (left). UML allows for operations as well as data attributes to be defined, but SBML only uses data attributes, so all SBML class diagrams use only the top two portions of a UML class box (Figure 1, right).




Figure 1: (Left) The general form of a UML class diagram. (Right) Example of a class diagram of the sort seen in SBML. SBML classes never use operators, so SBML class diagrams only show the top two parts.

As mentioned above, the names of ordinary (concrete) classes begin with a capital letter and are printed in an upright, bold, sans-serif typeface. The names of attributes begin with a lower-case letter and generally use a mixed case (sometimes called “camel case”) style when the name consists of multiple words. Attributes and their data types appear in the part below the class name, with one attribute defined per line. The colon

character on each line separates the name of the attribute (on the left) from the type of data that it stores (on the right). The subset of data types permitted for SBML attributes is given in Section 3.1.

In the right-hand diagram of Figure 1, the symbols **attribute** and **anotherAttribute** represent attributes of the object class **ExampleClass**. The data type of **attribute** is **int**, and the data type of **anotherAttribute** is **double**. In the scheme used by SBML for translating UML to XML, object attributes map directly to XML attributes. Thus, in XML, **ExampleClass** would yield an element of the form `<element attribute="42" anotherAttribute="10.0">`.

Notice that the element name is not `<ExampleClass ...>`. Somewhat paradoxically, the name of the element is *not* the name of the UML class defining its structure. The reason for this may be subtle at first, but quickly becomes obvious: object classes define the form of an object's *content*, but a class definition by itself does not define the *label* or symbol referring to an instance of that content. It is this label that becomes the name of the XML element. In XML, this symbol is most naturally equated with an element name. This point will hopefully become more clear with additional examples below.

Subelements

We use UML composite aggregation to indicate a class object can have other class objects as parts. Such containment hierarchies map directly to element-subelement relationships in XML. Figure 2 gives an example.




Figure 2: Example illustrating composite aggregation: the definition of one class of objects employing another class of objects in a part-whole relationship. In this particular example, an instance of a **Whole** class object must contain exactly one instance of a **Part** class object, and the symbol referring to the **Part** class object is **inside**. In XML, this symbol becomes the name of a subelement and the content of the subelement follows the definition of **Part**.

The line with the black diamond indicates composite aggregation, with the diamond located on the “container” side and the other end located at the object class being contained. The label on the line is the symbol used to refer to instances of the contained object, which in XML, maps directly to the name of an XML element. The class pointed to by the aggregation relationship (**Part** in Figure 2) defines the *contents* of that element. Thus, if we are told that some element named **barney** is of class **Whole**, the following is an example XML fragment consistent with the class definition of Figure 2:

```

<barney A="110" B="some string">
  <inside C="444.4">
</barney>
  
```

Sometimes numbers are placed above the line near the “contained” side of an aggregation to indicate how many instances can be contained. The common cases in SBML are the following: `[0..*]` to signify a list containing zero or more; `[1..*]` to signify a list containing at least one; and `[0..1]` to signify exactly zero or one. The absence of a numerical label means “exactly 1”. This notation appears throughout this specification document.

Inheritance

Classes can inherit properties from other classes. Since SBML only uses data attributes and not operations, inheritance in SBML simply involves data attributes from a parent class being inherited by child classes. Inheritance is indicated by a line between two classes, with an open triangle next to the parent class; Figure 3 illustrates this. In this example, the instances of object class **Child** would have not only attributes C and D, but also attributes A and B. All of these attributes would be required (not optional) on instances of class **Child** because they are mandatory on both **Parent** and **Child**.




Figure 3: Inheritance.

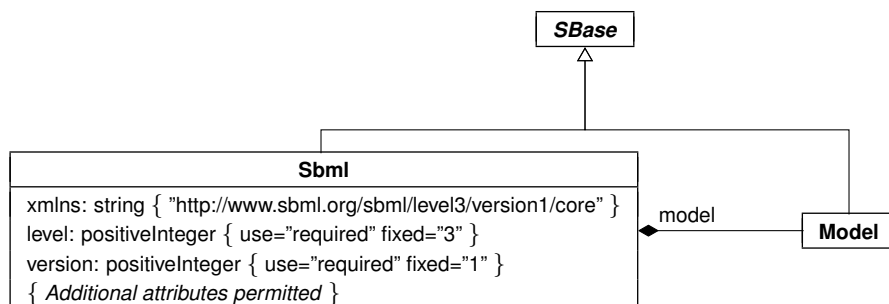


Figure 4: A more complex example definition drawing on the concepts introduced so far in this section. Both **SBml** and **Model** are derived from **SBase**; further, **SBml** contains a single **Model** object named `model1`. Note the constraints on the values of the attributes in **SBml**; they are enclosed in braces and written in XML Schema language. The particular constraints here state that both the `level` and `version` attributes must be present, and that the values are fixed as indicated. In addition, other attributes are permitted (for example, such as those added by Level 3 packages).


Additional notations for XML purposes

Not everything is easily expressed in plain UML. For example, it is often necessary to indicate some constraints placed on the values of an attribute. In computer programming uses of UML, such constraints are often expressed using Object Constraint Language (OCL), but since we are most interested in the XML rendition of SBML, in this specification we use XML Schema 1.0 (when possible) as the language for expressing value constraints. Constraints on the values of attributes are written as expressions surrounded by braces (`{ }`) after the data type declaration, as in the example of Figure 4.

In other situations, when something cannot be concisely expressed using a few words of XML Schema, we write constraints using English language descriptions surrounded by braces (`{ }`). To help distinguish these from literal XML Schema, we set the English text in a slanted typeface. The text accompanying all SBML component definitions provides explanations of the constraints and any other conditions applicable to the use of the components.

2 Overview of SBML

The following is an example of a simple network of biochemical reactions that can be represented in SBML:



In this particular set of chemical equations above, the symbols in square brackets (e.g., “[*S*₁]”) represent concentrations of molecular species, the arrows represent reactions, and the formulas above the arrows represent the rates at which the reactions take place. (And while this example uses concentrations, it could equally have used other measures such as molecular counts.) Broken down into its constituents, this model contains a number of components: reactant species, product species, reactions, reaction rates, and parameters in the rate expressions. To analyze or simulate this network, additional components must be made explicit, including compartments for the species, and units on the various quantities.

SBML allows models of arbitrary complexity to be represented. Each type of component in a model is described using a specific type of data object that organizes the relevant information. The top level of an SBML model definition consists of lists of these components, with every list being optional:

<i>beginning of model definition</i>	
<i>list of function definitions (optional)</i>	(Section 4.3)
<i>list of unit definitions (optional)</i>	(Section 4.4)
<i>list of compartments (optional)</i>	(Section 4.5)
<i>list of species (optional)</i>	(Section 4.6)
<i>list of parameters (optional)</i>	(Section 4.7)
<i>list of initial assignments (optional)</i>	(Section 4.8)
<i>list of rules (optional)</i>	(Section 4.9)
<i>list of constraints (optional)</i>	(Section 4.10)
<i>list of reactions (optional)</i>	(Section 4.11)
<i>list of events (optional)</i>	(Section 4.12)
<i>end of model definition</i>	

The meaning of each component is as follows:

Function definition: A named mathematical function that may be used throughout the rest of a model.

Unit definition: A named definition of a new unit of measurement. Named units can be used in the expression of quantities in a model.

Compartment: A well-stirred container of finite size where species may be located. Compartments may or may not represent actual physical structures.

Species: A pool of entities of the same kind located in a compartment and participating in reactions (processes). In biochemical network models, common examples of species include ions, proteins and other molecules; however, in practice, an SBML species can be any kind of entity that makes sense in the context of a given model.

Parameter: A quantity with a symbolic name. In SBML, the term *parameter* is used in a generic sense to refer to named quantities regardless of whether they are constants or variables in a model. SBML Level 3 provides the ability to define parameters that are global to a model as well as parameters that are local to a single reaction.

Initial Assignment: A mathematical expression used to determine the initial conditions of a model. This type of object can only be used to define how the value of a variable can be calculated from other values and variables at the start of simulated time.

1 *Rule:* A mathematical expression added to the set of equations constructed based on the reactions defined
2 in a model. Rules can be used to define how a variable's value can be calculated from other variables,
3 or used to define the rate of change of a variable. The set of rules in a model can be used with the
4 reaction rate equations to determine the behavior of the model with respect to time. Rules constrains
5 the model for the entire duration of simulated time.

6 *Constraint:* A means of detecting out-of-bounds conditions during a dynamical simulation and optionally
7 issuing diagnostic messages. Constraints are defined by an arbitrary mathematical expression comput-
8 ing a true/false value from model variables, parameters and constants. An SBML constraint applies at
9 all instants of simulated time; however, the set of constraints in model should not be used to *determine*
10 the behavior of the model with respect to time.

11 *Reaction:* A statement describing some transformation, transport or binding process that can change the
12 amount of one or more species. For example, a reaction may describe how certain entities (reactants) are
13 transformed into certain other entities (products). Reactions have associated kinetic rate expressions
14 describing how quickly they take place.

15 *Event:* A statement describing an instantaneous, discontinuous change in one or more variables of any type
16 (species, compartment, parameter, etc.) when a triggering condition is satisfied.

17 A software package can read an SBML model description and translate it into its own internal format for
18 model analysis. For example, a package might provide the ability to simulate the model by constructing
19 differential equations representing the network and then perform numerical time integration on the equations
20 to explore the model's dynamic behavior. By supporting SBML as an input and output format, different
21 software tools can all operate on an identical external representation of a model, removing opportunities for
22 errors in translation and assuring a common starting point for analyses and simulations.

3 Preliminary definitions and principles

This section covers certain concepts and constructs that are used repeatedly in the rest of SBML Level 3.

3.1 Primitive data types

Most primitive types in SBML are taken from the data types defined in XML Schema 1.0 (Biron and Malhotra, 2000; Fallside, 2000; Thompson et al., 2000). A few other primitive types are defined by SBML itself. What follows is a summary of the XML Schema types and the definitions of the SBML-specific types. Note that while we have tried to provide accurate and complete summaries of the XML Schema types, the following should not be taken to be normative definitions of these types. Readers should consult the XML Schema 1.0 specification for the normative definitions of the XML types used by SBML.

3.1.1 Type string

The XML Schema 1.0 type **string** is used to represent finite-length strings of characters. The characters permitted to appear in XML Schema **string** include all Unicode characters (Unicode Consortium, 1996) except for two delimiter characters, 0xFFFE and 0xFFFF (Biron and Malhotra, 2000). In addition, the following quoting rules specified by XML for character data (Bray et al., 2000) must be obeyed:

- The ampersand (&) character must be escaped using the entity `&`.
- The apostrophe (') and quotation mark (") characters must be escaped using the entities `'` and `"`, respectively, when those characters are used to delimit a string attribute value.

Other XML built-in character or entity references, e.g., `<` and `&x1A;`, are permitted in strings.

3.1.2 Type boolean

The XML Schema 1.0 type **boolean** is used as the data type for SBML object attributes that represent binary true/false values. XML Schema 1.0 defines the possible literal values of **boolean** as the following: “true”, “false”, “1”, and “0”. The value “1” maps to “true” and the value “0” maps to “false”.

Note that there is a discrepancy between the value spaces of type **boolean** as defined by XML Schema 1.0 and MathML: the latter uses only “true” and “false” to represent boolean values and “0” and “1” are interpreted as numbers. Software tools should take care to not to use “0” and “1” as boolean values in MathML expressions. See further discussion in Section 3.4.4.

3.1.3 Type int

The XML Schema 1.0 type **int** is used to represent decimal integer numbers in SBML. The literal representation of an **int** is a finite-length sequence of decimal digit characters with an optional leading sign (“+” or “-”). If the sign is omitted, “+” is assumed. The value space of **int** is the same as a standard 32-bit signed integer in programming languages such as C, i.e., 2147483647 to -2147483648.

3.1.4 Type positiveInteger

The XML Schema 1.0 type **positiveInteger** is used to represent nonzero, nonnegative, decimal integers: i.e., 1, 2, 3, ... The literal representation of an integer is a finite-length sequence of decimal digit characters, optionally preceded by a positive sign (“+”). There is no restriction on the absolute size of **positiveInteger** values in XML Schema; however, the only situations where this type is used in SBML involve very low-numbered integers. Consequently, applications may safely treat **positiveInteger** as unsigned 32-bit integers.

3.1.5 Type double

The XML Schema 1.0 type **double** is the data type of floating point numerical quantities in SBML. It is restricted to IEEE double-precision 64-bit floating point type IEEE 754-1985. The value space of **double** consists of (a) the numerical values $m \cdot 2^x$, where m is an integer whose absolute value is less than 2^{53} ,

and x is an integer between -1075 and 970, inclusive, (b) the special value positive infinity (INF), (c) the special value negative infinity (-INF), and (d) the special value not-a-number (NaN). The order relation on the values is the following: $x < y$ if and only if $y - x$ is positive for values of x and y in the value space of **double**. Positive infinity is greater than all other values other than NaN. NaN is equal to itself but is neither greater nor less than any other value in the value space. (Software implementors should consult the XML Schema 1.0 definition of **double** for additional details about equality and relationships to IEEE 754-1985.)

The general form of **double** numbers is “ $x\text{e}y$ ”, where x is a decimal number (the mantissa), “e” is a separator character, and y is an exponent; the meaning of this is “ x multiplied by 10 raised to the power of y ”, i.e., $x \cdot 10^y$. More precisely, a **double** value consists of a mantissa with an optional leading sign (“+” or “-”), optionally followed by the character E or e followed by an integer (the exponent). The mantissa must be a decimal number: an integer optionally followed by a period (.) optionally followed by another integer. If the leading sign is omitted, “+” is assumed. An omitted E or e and exponent means that a value of 0 is assumed for the exponent. If the E or e is present, it must be followed by an integer or an error results. The integer acting as an exponent must consist of a decimal number optionally preceded by a leading sign (“+” or “-”). If the sign is omitted, “+” is assumed. The following are examples of legal literal **double** values:

-1E4, +4, 234.234e3, 6.02E-23, 0.3e+11, 2, 0, -0, INF, -INF, NaN

As described in Section 3.4, SBML uses a subset of the MathML 2.0 standard (W3C, 2000b) for expressing mathematical formulas in XML. This is done by stipulating that the MathML language be used whenever a mathematical formula must be written into an SBML model. Doing this, however, requires facing two problems: first, the syntax of numbers in scientific notation (“e-notation”) is different in MathML from that just described for **double**, and second, the value space of integers and floating-point numbers in MathML is not defined in the same way as in XML Schema 1.0. We elaborate on these issues in Section 3.4.2; here we summarize the solution taken in SBML. First, within MathML, the mantissa and exponent of numbers in “e-notation” format must be separated by one <sep/> element. This leads to numbers of the form <cn type=“e-notation”> 2 <sep/> -5 </cn>. Second, SBML stipulates that the representation of numbers in MathML expressions obey the same restrictions on values as defined for types **double** and **int** (Section 3.1.3).

3.1.6 Type ID

The XML Schema 1.0 type ID is identical to the XML 1.0 type ID. The literal representation of this type consists of strings of characters restricted as summarized in Figure 5.

```
NameChar ::= letter | digit | '.' | '-' | '_' | ':' | CombiningChar | Extender
ID        ::= ( letter | '-' | ':' ) NameChar*
```

Figure 5: Type ID expressed in the variant of BNF used by the XML 1.0 specification (Bray et al., 2004). The characters (and) are used for grouping, the character * indicates “zero or more times”, and the character | indicates “or”. The production letter consists of the basic upper and lower case alphabetic characters of the Latin alphabet along with a large number of related characters defined by Unicode 2.0; similarly, the production digit consists of the numerals 0..9 along with related Unicode 2.0 characters. The CombiningChar production is a list of characters that add such things as accents to the preceding character. (For example, the Unicode character #x030A when combined with ‘a’ produces ‘â’.) The Extender production is a list of characters that extend the shape of the preceding character. Please consult the XML 1.0 specification (Bray et al., 2004) for the complete definitions of letter, digit, CombiningChar, and Extender.

In SBML, type ID is the data type of the **metaid** attribute on **SBase**, described in Section 3.2. An important aspect of ID is the XML requirement that a given value of ID must be unique throughout an XML document. All data values of type ID are considered to reside in a single common global namespace spanning the entire XML document, regardless of the attribute where type ID is used and regardless of the level of nesting of the objects (or XML elements).

3.1.7 Type SId

The type SId is the type of the **id** attribute found on the majority of SBML components. SId is a data type derived from the basic XML type **string**, but with restrictions about the characters permitted and the sequences in which those characters may appear. The definition is shown in Figure 6 on the following page.

```

1         letter ::= 'a'..'z','A'..'Z'
2         digit  ::= '0'..'9'
3         idChar ::= letter | digit | '_'
4         SId    ::= ( letter | '_' ) idChar*

```

Figure 6: The definition of the type `SId`. (Please see the caption of Figure 5 for an explanation of the notation.)

The equality of `SId` values is determined by an exact character sequence match; i.e., comparisons of these identifiers must be performed in a case-sensitive manner. This applies to all uses of `SId`.

The `SId` is purposefully not derived from the XML `ID` type (Section 3.1.6). Using `ID` would force all SBML identifiers to exist in a single global namespace, affecting not only `Reaction` local parameter definitions but also SBML packages for (e.g.) hierarchical model composition. Further, the use of `ID` for SBML identifiers would have limited utility because MathML 2.0 `ci` elements are not of the type `IDREF` (see Section 3.4). Since the `IDREF`/`ID` linkage cannot be exploited in MathML constructs, the utility of XML's `ID` type is greatly reduced. Finally, unlike `ID`, `SId` does not include Unicode character codes; the identifiers are plain text.

3.1.8 Type `SIdRef`

Type `SIdRef` is used for all attributes that refer to identifiers of type `SId` in a model. This type is derived from `SId`, but with the restriction that the value of an attribute having type `SIdRef` must equal the value of *some* `SId` attribute in the model where it appears. In other words, a `SIdRef` value must be an existing identifier in a model.

As with `SId`, the equality of `SIdRef` values is determined by exact character sequence match; i.e., comparisons of these identifiers must be performed in a case-sensitive manner.

3.1.9 Type `UnitSId`

The type `UnitSId` is derived from `SId` (Section 3.1.7) and has identical syntax. The `UnitSId` type is used as the data type for the identifiers of units (Section 4.4.1) in SBML objects. The purpose of having a separate type for such identifiers is to enable the space of possible unit identifier values to be separated from the space of all other identifier values in SBML. The equality of `UnitSId` values is determined by an exact character sequence match; i.e., comparisons of these identifiers must be performed in a case-sensitive manner.

A number of reserved symbols are defined in the space of values of `UnitSId`. These reserved symbols are the list of base unit names defined in Table 1 on page 38.

3.1.10 Type `UnitSIdRef`

Type `UnitSIdRef` is used for all attributes that refer to identifiers of type `UnitSId`, which are the identifiers of units (Section 4.4.1) in SBML objects. This type is derived from `UnitSId`, but with the restriction that the value of an attribute having type `UnitSIdRef` must match either the value of a `UnitSId` attribute in the model, or one of the predefined units in Table 1 on page 38. In other words, the value of a `UnitSIdRef` attribute must be an existing unit identifier in the model or in SBML.

As with `UnitSId`, the equality of `UnitSIdRef` values is determined by exact character sequence match; i.e., comparisons of these identifiers must be performed in a case-sensitive manner.

3.1.11 Type `SBOTerm`

The type `SBOTerm` is used as the data type of the attribute `sboTerm` on `SBase`. The type consists of strings of characters matching the restricted pattern described in Figure 7.

```

39         digit    ::= '0'..'9'
40         SBOTerm  ::= 'SBO:' digit digit digit digit digit digit digit

```

Figure 7: The definition of `SBOTerm`. (Please see the caption of Figure 5 for an explanation of the notation.)

Examples of valid string values of type **SBOTerm** are “SB0:0000014” and “SB0:0003204”. These values are meant to be the identifiers of terms from an ontology whose vocabulary describes entities and processes in computational models. Section 5 provides more information about the ontology and principles for the use of these terms in SBML models.

3.2 Type **SBase**

Nearly every object composing an SBML Level 3 model definition has a specific data type that is derived directly or indirectly from a single abstract type called **SBase**. In addition to serving as the parent class for most other classes of objects in SBML, this base type is designed to allow a modeler or a software package to attach arbitrary information to each major element or list in an SBML model. The definition of **SBase** is presented in Figure 8.

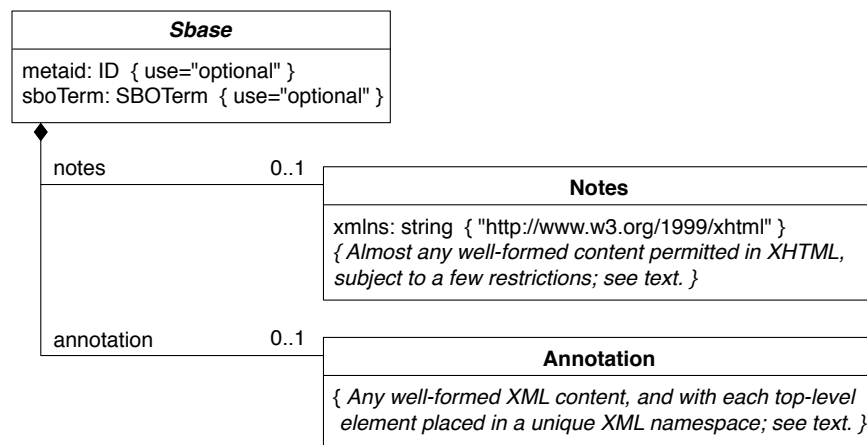


Figure 8: The definition of **SBase**. Please refer to Section 1.4 for a summary of the UML notation used here.

SBase contains two attributes and two subelements, all of which are optional: **metaid**, **sboterm**, **notes** and **annotation**. These are discussed separately in the following subsections.

3.2.1 The **metaid** attribute

The **metaid** attribute is present for supporting metadata annotations using RDF (Resource Description Format; Lassila and Swick, 1999). It has a data type of XML ID (the XML identifier type; see Section 3.1.6), which means each **metaid** value must be globally unique within an SBML file. The **metaid** value serves to identify a model component for purposes such as referencing that component from metadata placed within **annotation** elements (see Section 3.2.4). Such metadata can use RDF **description** elements, in which an RDF attribute called “**rdf:about**” points to the **metaid** identifier of an object defined in the SBML model. This topic is discussed in greater detail in Section 6.

3.2.2 The **sboterm** attribute

The attribute called **sboterm** is provided on **SBase** to support the use of the Systems Biology Ontology (SBO; see Section 5). When a value is given to this attribute, it must conform to the data type **SBOTerm** (Sections 3.1.11). SBO terms are a type of optional annotation, and each different class of SBML object derived from **SBase** imposes its own requirements about the values permitted for **sboterm**. Specific details on the permitted values are provided with the definitions of SBML classes throughout this specification document, and a broader discussion is provided in Section 5.

3.2.3 The notes element

The element **notes** in *SBase* is a container for XHTML 1.0 (Pemberton et al., 2002) content. It is intended to serve as a place for storing optional information intended to be seen by humans. An example use of the **notes** element would be to contain formatted user comments about the model element in which the **notes** element is enclosed. Every object derived directly or indirectly from type *SBase* can have a separate value for **notes**, allowing users considerable freedom when adding comments to their models.

XML namespace requirements for notes

The XML content of **notes** elements must declare the use of the XHTML XML namespace. This can be done in multiple ways. One way is to place a namespace declaration for the appropriate namespace URI (which is <http://www.w3.org/1999/xhtml>) on the top-level *Sbml* object (see Section 4.1) and then reference the namespace in the **notes** content using a prefix. The following example illustrates this approach:

```
<sbml xmlns="http://www.sbml.org/sbml/level3/version1/core" level="3" version="1"
  xmlns:xhtml="http://www.w3.org/1999/xhtml">
  ...
  <notes>
    <xhtml:body>
      <xhtml:center><xhtml:h2>A Simple Mitotic Oscillator</xhtml:h2></xhtml:center>
      <xhtml:p>A minimal cascade model for the mitotic oscillator
        involving cyclin and cdc2 kinase</xhtml:p>
    </xhtml:body>
  </notes>
  ...
```

Another approach is to declare the XHTML namespace within the **notes** content itself, as in the following example:

```
...
<notes>
  <body xmlns="http://www.w3.org/1999/xhtml">

    <center><h2>A Simple Mitotic Oscillator</h2></center>

    <p>A minimal cascade model for the mitotic oscillator
      involving cyclin and cdc2 kinase</p>

  </body>
</notes>
...
```

The `xmlns="http://www.w3.org/1999/xhtml"` declaration on **body** as shown above changes the default XML namespace within it, such that all of its content is by default in the XHTML namespace. This is a particularly convenient approach because it obviates the need to prefix every element with a namespace prefix (i.e., “xhtml:” in the earlier example). Other approaches are also possible.

The content of notes

SBML does not require the content of **notes** to be any particular XHTML element; the content can be almost any well-formed XHTML content. There are only two simple restrictions. The first restriction comes from the requirements of XML: the **notes** element must not contain an XML declaration nor a DOCTYPE declaration. That is, **notes** must *not* contain

```
<?xml version="1.0" encoding="UTF-8"?>
```

nor (where the following is only one specific example of a DOCTYPE declaration)

```
<!DOCTYPE html PUBLIC "-//W3C//DTD XHTML 1.0 Strict//EN"
  "http://www.w3.org/TR/xhtml1/DTD/xhtml1-strict.dtd">
```

The second restriction is intended to balance freedom of content with the complexity of implementing software that can interpret the content. The content of **notes** in SBML can consist only of the following possibilities:

1. A complete XHTML document (minus the XML and DOCTYPE declarations, of course), that is, XHTML content beginning with the `html` tag. The following is an example skeleton:

```
<notes>
  <html xmlns="http://www.w3.org/1999/xhtml">
    ...
  </html>
</notes>
```

2. The `body` element from an XHTML document. The following is an example skeleton:

```
<notes>
  <body xmlns="http://www.w3.org/1999/xhtml">
    ...
  </body>
</notes>
```

3. Any XHTML content that would be permitted within a `body` element. Each such XHTML element must be placed in the XML namespace for XHTML, which can be done in a variety of standard XML ways. The following is an example fragment:

```
<notes>
  <p xmlns="http://www.w3.org/1999/xhtml">
    ...
  </p>
  <p xmlns="http://www.w3.org/1999/xhtml">
    ...
  </p>
</notes>
```

Another way to summarize the restrictions above is simply to say that the content of an SBML `notes` element can be only be a complete `html` element, a `body` element, or whatever is permitted inside a `body` element. In practice, this does not limit in any meaningful way what can be placed inside a `notes` element; for example, if an application or modeler wants to put a complete XHTML page, including a `head` element, it can be done by putting in everything starting with the `html` container. However, the restrictions above do make it somewhat simpler to write software that can read and write the `notes` content. Appendix F describes one possible approach to doing just that.

3.2.4 The annotation element

Whereas the `notes` element described above is a container for content to be shown directly to humans, the `annotation` element is a container for optional software-generated content *not* meant to be shown to humans. Every object derived from *SBase* can have its own value for `annotation`. The element's content type is XML type `any`, allowing essentially arbitrary well-formed XML data content. SBML places only a few restrictions on the organization of the content; these are intended to help software tools read and write the data as well as help reduce conflicts between annotations added by different tools.

The use of XML namespaces in annotation

At the outset, software developers should keep in mind that multiple software tools may attempt to read and write annotation content. To reduce the potential for collisions between annotations written by different applications, SBML Level 3 Version 1 Core stipulates that tools must use XML namespaces (Bray et al., 1999) to specify the intended vocabulary of every annotation. The application's developers must choose a URI (*Universal Resource Identifier*; Harold and Means 2001; W3C 2000a) reference that uniquely identifies the vocabulary the application will use, and a prefix string for the annotations. Here is an example. Suppose an application uses the URI `http://www.mysim.org/ns` and the prefix `mysim` when writing annotations related to molecules. The content of an annotation might look like the following:

```
<annotation>
  <mysim:molecule xmlns:mysim="http://www.mysim.org/ns"
    mysim:weight="18.02" mysim:atomCount="3"/>
</annotation>
```


In this particularly simple example, the content consists of a single XML element (**molecule**) with two attributes (**weight**, **atomCount**), all of which are prefixed by the string **mysim**. (Presumably this particular content would have meaning to the hypothetical application in question.) The content in this particular example is small, but it should be clear that there could easily have been an arbitrarily large amount of data placed inside the **mysim:molecule** element.

The key point of the example above is that application-specific annotation data are entirely contained inside a single *top-level element* within the SBML **annotation** container. SBML Level 3 Version 1 places the following restrictions on annotations:

- Within a given SBML **annotation** element, there can only be one top-level element using a given namespace. An annotation element can contain multiple top-level elements but each must be in a different namespace.
- The ordering of top-level elements within a given **annotation** element is *not* significant. An application should not expect that its annotation content appears first in the **annotation** element, nor in any other particular location. Moreover, the ordering of top-level annotation elements may be changed by different applications as they read and write the same SBML file.

The use of XML namespaces in this manner is intended to improve the ability of multiple applications to place annotations on SBML model elements with reduced risks of interference or name collisions. Annotations stored by different simulation packages can therefore coexist in the same model definition. The rules governing the content of **annotation** elements are designed to enable applications to easily add, change, and remove their annotations from SBML elements while simultaneously preserving annotations inserted by other applications when mapping SBML from input to output.

As a further simplification for developers of software and to improve software interoperability, applications are only required to preserve other annotations (i.e., annotations they do not recognize) when those annotations are self-contained entirely within **annotation**, complete with namespace declarations. The following is an example:

```
<annotation>
  <topLevelElement xmlns="URI">
    ... content in the namespace identified by "URI"...
  </topLevelElement>
</annotation>
```

Some more examples hopefully will make these points more clear. The following example is invalid because it contains two top-level elements using the same XML namespace. Note that it does not matter that these are two different top-level elements (**molecule** and **atom**); what matters for SBML is that these separate elements are both in the same namespace rather than having been collected and placed inside one overall container element for that namespace:

```
<annotation>
  <mysim:molecule xmlns:mysim="http://www.mysim.org/ns"
    mysim:weight="18.02" mysim:atomCount="3"/>
  <mysim:atom xmlns:mysim="http://www.mysim.org/ns"
    mysim:weight="18.02" mysim:atomCount="3"/>
</annotation>
```

On the other hand, the following example is valid:

```
<annotation>
  <mysim:molecule xmlns:mysim="http://www.mysim.org/ns"
    mysim:weight="18.02" mysim:atomCount="3">
    <struct:bonds xmlns:size="http://www.struct.org/ns"
      struct:number="2" struct:type="ionic" />
  </mysim:molecule>
  <othersim:icon xmlns:othersim="http://www.othersim.com/">
    WS2002
  </othersim:icon>
</annotation>
```

For completeness, we note that annotations legally can be empty (but such annotations have no meaning):

```
<annotation />
```

It is worth keeping in mind that although XML namespace names must be URIs, they are (like all XML namespace names) *not required* to be directly usable in the sense of identifying an actual, retrieval document or resource on the Internet (Bray et al., 1999). URIs such as <http://www.mysim.org/> may appear as though they are (e.g.,) Internet addresses, but there are not the same thing. This style of URI strings, using a domain name and other parts, is only a simple and commonly-used way of creating a unique name string.

Finally, note that the namespaces being referred to here are XML namespaces specifically in the context of the **annotation** element on **SBase**. The namespace issue here is unrelated to the namespaces discussed in Section 3.3.1 in the context of component identifiers in SBML.

Content of annotations and implications for software tools

The **annotation** element in the definition of **SBase** exists in order that software developers may attach optional application-specific data to the elements in an SBML model. However, it is important that this facility not be misused. In particular, it is *critical* that data essential to a model definition or that can be encoded in existing SBML elements is *not* stored in **annotation**. Parameter values, functional dependencies between model elements, etc., should not be recorded as annotations. It is crucial to keep in mind the fact that data placed in annotations can be freely ignored by software applications. If such data affect the interpretation of a model, then software interoperability is greatly impeded. Recommendations regarding the use of any sort of annotation are given in Section 8.1.4.

3.3 The id and name attributes on SBML components

As will become apparent below, most objects in SBML include two common attributes: **id** and **name**. These attributes are not defined on **SBase** (as explained in Section 3.3.3 below), but where they do appear, the common rules of usage described below apply.

3.3.1 The id attribute and identifier scoping

The **id** attribute is mandatory on most (but not all) objects in SBML. It is used to identify a component within the model. Other SBML objects can refer to the component using this identifier. The data type of **id** is always either **SId** (Section 3.1.7) or **UnitSId** (Section 3.1.9), depending on the object in question.

A model can contain a large number of components representing different parts. This leads to a problem in deciding the scope of an identifier: in what contexts does a given identifier *X* represent the same thing? The approaches used in existing simulation packages tend to fall into two categories which we may call global and local. The *global* approach places all identifiers into a single global space of identifiers, so that an identifier *X* represents the same thing wherever it appears in a given model definition. The *local* approach places symbols in separate identifier namespaces, depending on the context, where the context may be, for example, individual reaction rate expressions. The latter approach means that a model may use the same identifier *X* in different rate expressions and have each instance represent a different quantity.

The scoping rules in SBML Level 3 are intended as a compromise to help support both scenarios:

- The identifier (i.e., the value of the attribute **id**) of every **FunctionDefinition**, **Compartment**, **Species**, **Parameter**, **Reaction**, **SpeciesReference**, **ModifierSpeciesReference**, **Event**, and **Model**, must be unique across the set of all such identifiers in the model. This means, for example, that a reaction and a species definition cannot both have the same identifier.
- The identifier of every **UnitDefinition** must be unique across the set of all such identifiers in the model. However, unit identifiers live in a separate space of identifiers from other identifiers in the model, by virtue of the fact that the data type of unit identifiers is **UnitSId** (Section 3.1.9) and not **SId**.
- Each **Reaction** instance (see Section 4.11) establishes a separate private local space for local parameters represented by objects of class **LocalParameter**. Within the definition of that reaction, local parameter

identifiers override (shadow) identical identifiers (whether those identifiers refer to parameters, species or compartments) outside of that reaction. Of course, the corollary of this is that local parameters inside a [Reaction](#) object instance are not visible to other objects outside of that reaction.

3.3.2 The name attribute

In contrast to the `id` attribute, the `name` attribute is optional and is not intended to be used for cross-referencing purposes within a model. Its purpose instead is to provide a human-readable label for the component. The data type of `name` is the type `string` defined in XML Schema ([Biron and Malhotra, 2000](#); [Thompson et al., 2000](#)) and discussed further in Section 3.1. SBML imposes no restrictions as to the content of `name` attributes beyond those restrictions defined by the `string` type in XML Schema. In addition, there are no restrictions on the uniqueness of `name` values in a model (unlike the restrictions on `id` values discussed in Section 3.3.1).

3.3.3 Why id and name are not defined on SBase

Although many SBML components feature `id` and `name`, these attributes are purposefully not defined on [SBase](#). There are several reasons for this.

- The presence of an SBML identifier attribute (`id`) necessarily requires specifying scoping rules for the corresponding identifiers. However, the [SBase](#) abstract type is used as the basis for defining components whose scoping rules are in some cases different from each other. (See Section 3.3.1 for more details). If [SBase](#) were to have an `id` attribute, then the specification of [SBase](#) would need a default scoping rule and this would then have to be overloaded on derived classes that needed different scoping. This would make the SBML specification even more complex.
- Identifiers are optional on some SBML components and required on most others. If `id` were defined as optional on [SBase](#), most component classes would separately have to redefine `id` as being mandatory—hardly an improvement over the current arrangement. Conversely, if `id` were defined as mandatory on [SBase](#), it would prevent it from being optional on components where it *is* currently optional.
- The [SBase](#) abstract type is used as the base type for certain objects such as [Sbml](#), [AssignmentRule](#), etc., which do not have identifiers because these components do not need to be referenced by other components. If [SBase](#) had a mandatory `id` attribute, *all* objects of these other types in a model would then need to be assigned unique identifiers. Similarly, because [SBase](#) is the base type of the `listOf_____` lists, putting `id` on [SBase](#) would require all of these lists in a model to be given identifiers. This would be a needless burden on software developers, tools, and SBML users, requiring them to generate and store additional identifiers for objects that never need them.
- [SBase](#) does not have a `name` simply because such an attribute is always paired with an `id`. Without `id` on [SBase](#), it does not make sense to have `name`.

3.4 Mathematical formulas in SBML Level 3

Mathematical expressions in SBML Level 3 are represented using MathML 2.0 ([W3C, 2000b](#)). MathML is an international standard for encoding mathematical expressions using XML. There are two principal facets of MathML, one for encoding content (i.e., the semantic interpretation of a mathematical expression), and another for encoding presentation or display characteristics. SBML only makes direct use of a subset of the content portion of MathML. However, it is not possible to produce a completely smooth and conflict-free interface between MathML and other standards used by SBML (in particular, XML Schema). Two specific issues and their resolutions are discussed in Sections 3.4.2.

The XML namespace URI for all MathML elements is <http://www.w3.org/1998/Math/MathML>. Everywhere MathML content is allowed in SBML, the MathML elements must be properly placed within the MathML 2.0 namespace. In XML, this can be accomplished in a number of ways, and the examples throughout this specification illustrate the use of this namespace and MathML in SBML. Please refer to the W3C document by [Bray et al. \(1999\)](#) for more technical information about using XML namespaces.

3.4.1 Subset of MathML used in SBML Level 3

The subset of MathML elements used in SBML is listed below:

- *token*: `cn`, `ci`, `csymbol`, `sep`
- *general*: `apply`, `piecewise`, `piece`, `otherwise`, `lambda` (the last is restricted to use in [FunctionDefinition](#))
- *relational operators*: `eq`, `neq`, `gt`, `lt`, `geq`, `leq`
- *arithmetic operators*: `plus`, `minus`, `times`, `divide`, `power`, `root`, `abs`, `exp`, `ln`, `log`, `floor`, `ceiling`, `factorial`
- *logical operators*: `and`, `or`, `xor`, `not`
- *qualifiers*: `degree`, `bvar`, `logbase`
- *trigonometric operators*: `sin`, `cos`, `tan`, `sec`, `csc`, `cot`, `sinh`, `cosh`, `tanh`, `sech`, `csch`, `coth`, `arcsin`, `arccos`, `arctan`, `arcsec`, `arccsc`, `arccot`, `arcsinh`, `arccosh`, `arctanh`, `arcsech`, `arccsch`, `arccoth`
- *constants*: `true`, `false`, `notanumber`, `pi`, `infinity`, `exponentiale`
- *annotation*: `semantics`, `annotation`, `annotation-xml`

The inclusion of logical operators, relational operators, `piecewise`, `piece`, and `otherwise` elements facilitates the encoding of discontinuous expressions.

As defined by MathML 2.0, the semantic interpretation of the mathematical functions listed above follows the definitions of the functions laid out by [Abramowitz and Stegun \(1977\)](#) and [Zwillinger \(1996\)](#). Readers are directed to these sources and the MathML specification for information about such things as which principal values of the inverse trigonometric functions to use.

Software authors should take particular note of the MathML semantics of the N-ary operators `plus`, `times`, `and`, `or` and `xor`, when they are used with different numbers of arguments. The MathML specification ([W3C, 2000b](#)) appendix C.2.3 describes the semantics for these operators with zero, one, and more arguments.

The following are the only attributes permitted on MathML elements in SBML (in addition to the `xmlns` attribute on `math` elements):

- `style`, `class` and `id` on any element;
- `encoding` on `csymbol`, `annotation` and `annotation-xml` elements;
- `definitionURL` on `csymbol` and `semantics` elements; and
- `type` and `sbml:units` (see Section [3.4.2](#)) on `cn` elements.

Missing values for the MathML attributes are to be treated in the same way as defined by MathML 2.0. These restrictions on attributes are designed to confine the MathML elements to their default semantics and to avoid conflicts in the interpretation of the type of token elements.

3.4.2 Numbers and `cn` elements

In MathML, literal numbers are written as the content portion of a particular element called `cn`. This element takes an optional attribute, `type`, used to indicate the *type* of the number (such as whether it is meant to be an integer or a floating-point quantity). Here is an example of its use:

```
<math xmlns="http://www.w3.org/1998/Math/MathML">
  <apply>
    <times/> <cn type="integer"> 42 </cn> <cn type="real"> 3.3 </cn>
  </apply>
</math>
```

The content of a `cn` element must be a number. The number can be preceded and succeeded by whitespace (see Section [3.4.5](#)). The following are the only permissible values for the `type` attribute on MathML `cn` elements: “`e-notation`”, “`real`”, “`integer`”, and “`rational`”. The value of the `type` attribute defaults to “`real`” if it is not specified on a given `cn` element.

Value space restrictions on `cn` content

SBML imposes certain restrictions on the value space of numbers allowed in MathML expressions. According to the MathML 2.0 specification, the values of the content of `cn` elements do not necessarily have to conform to any specific floating point or integer representations designed for CPU implementation. For example, in strict MathML, the value of a `cn` element could exceed the maximum value that can be stored in a IEEE 64 bit floating point number (IEEE 754). This is different from the XML Schema type `double` that is used in the definition of floating point attributes of objects in SBML; the XML Schema `double` is restricted to IEEE double-precision 64-bit floating point type IEEE 754-1985. To avoid an inconsistency that would result between numbers elsewhere in SBML and numbers in MathML expressions, SBML Level 3 Version 1 Core imposes the following restriction on MathML content appearing in SBML:

- Integer values (i.e., the values of `cn` elements having `type="integer"` and both values in `cn` elements having `type="rational"`) must conform to the `int` type used elsewhere in SBML (Section 3.1.3)
- Floating-point values (i.e., the content of `cn` elements having `type="real"` or `type="e-notation"`) must conform to the `double` type used elsewhere in SBML (Section 3.1.5)

Syntactic differences in the representation of numbers in scientific notation

It is important to note that MathML uses a style of scientific notation that differs from what is defined in XML Schema, and consequently what is used in SBML attribute values. The MathML 2.0 type “`e-notation`” (as well as the type “`rational`”) requires the mantissa and exponent to be separated by one `<sep/>` element. The mantissa must be a real number and the exponent part must be a signed integer. This leads to expressions such as

```
<cn type="e-notation"> 2 <sep/> -5 </cn>
```

for the number $2 \cdot 10^{-5}$. It is especially important to note that the expression

```
<cn type="e-notation"> 2e-5 </cn>
```

is *not valid* in MathML 2.0 and therefore cannot be used in MathML content in SBML. However, elsewhere in SBML, when an attribute value is declared to have the data type `double` (a type taken from XML Schema), the compact notation “`2e-5`” is in fact allowed. In other words, within MathML expressions contained in SBML (and *only* within such MathML expressions), numbers in scientific notation must take the form `<cn type="e-notation"> 2 <sep/> -5 </cn>`, and everywhere else they must take the form “`2e-5`”.

This is a regrettable difference between two standards that SBML relies upon, but it is not feasible to redefine these types within SBML because the result would be incompatible with parser libraries written to conform with the MathML and XML Schema standards. It is also not possible to use XML Schema to define a data type for SBML attribute values permitting the use of the `<sep/>` notation, because XML attribute values cannot contain XML elements—that is, `<sep/>` cannot appear in an XML attribute value.

Units associated with numbers in MathML `cn` expressions

What units should be attributed to numbers appearing inside MathML `cn` elements? One answer is to assume that the units should be “whatever units appropriate in the context where the number appears”. This implies that units can always be assigned unambiguously to any number by inspecting the expression in which it appears, and this turns out to be false. Another answer is that numbers should be considered “dimensionless”. Many people argue that this is the correct interpretation, but even if it is, there is an overriding practical reason why it cannot be adopted for SBML’s domain of application: when numbers appear in expressions in SBML, they are *rarely intended* by the modeler to have the unit “`dimensionless`” even if the unit is not declared—instead, the numbers are *supposed* to have specific units, but the units are usually undeclared. (Being “dimensionless” is not the same as having *undeclared* units!) If SBML defined numbers as being *by default* dimensionless, it would result in many models being technically incorrect without the modeler being aware of it unless their software tools performed dimensional analysis. Many software tools do not perform unit analysis, and so potential errors due to inconsistent units in a model would not

be detected until other researchers and database curators attempted to use the model in software packages that *did* check units. We believe the negative impact on interoperability would be too high.

SBML borrows an idea from CellML (Hedley et al., 2001), another model definition language with similar goals as SBML, and allows an additional attribute to appear on MathML `cn` elements; the value of this attribute can be used to indicate the units associated with the number in the content of the `cn` element. The attribute is named `units` but, because it appears inside MathML element (which is in the XML namespace for MathML and not the namespace for SBML), it must always be prefixed with an XML namespace prefix for the SBML Level 3 Version 1 Core namespace. The value of the attribute must have the data type `UnitSIdRef` (Section 3.1.10) and can be the identifier of a `UnitDefinition` object in the model or a base unit listed in Table 1 on page 38. The following example illustrates how this attribute can be used to define a number with value “10” and unit of measurement “second”:

```
<math xmlns="http://www.w3.org/1998/Math/MathML"
      xmlns:sbml="http://www.sbml.org/sbml/level3/version1/core">
  <cn type="integer" sbml:units="second"> 10 </cn>
</math>
```

In this example, we chose to use the string “`sbml`” as the XML namespace prefix for the SBML Level 3 Version 1 Core namespace, which leads to the use of `sbml:units` as the attribute on the `cn` element. We could have used another prefix string besides “`sbml`”, and the definition of the prefix could also have appeared on a higher-level element in the model. Section 4.1 provides more information about the XML namespace for SBML Level 3 Version 1 Core.

An alternative approach to specifying units is to avoid using `cn` elements altogether, and always use `ci` elements to reference `Parameter` objects having both value and units defined. In the example above, we could have avoided putting the literal number “10” inside the mathematical expression, and instead, defined a parameter in the model, given it the value “10” and unit “second”, and finally, referred to that parameter in the `math` content above. The approach of using named parameters provides additional power and advantages over simply using `sbml:units` attributes on `cn` elements; for example, `Parameter` allows the association of terms from the Systems Biology Ontology (SBO; Section 5) as well as MIRIAM annotations (Section 3.2.4).

In summary, a literal number within MathML content without an SBML `units` attribute has no declared units. Either of the two approaches described above (i.e., avoiding `cn` in favor of `ci` elements and `Parameter` objects, or using `cn` with a `sbml:units` attribute) leads to mathematical formulas whose units can be fully determined, enabling software tools to perform dimensional analysis and, potentially, detect and report problems with the model. Conversely, in the absence of an SBML `units` attribute on a MathML `cn` element, no units are associated with the number within the `cn` element; in other words, *the units are undeclared*. If the example above lacked the attribute `sbml:units`, the value “10” would have no declared units.

Finally, although SBML provides ways of associating units with numbers and entities, SBML does not stipulate that implicit unit conversions be performed. Section 3.4.11 explores this topic in more detail.

3.4.3 Use of `ci` elements in MathML expressions in SBML

The content of a `ci` element must be an SBML identifier that is declared elsewhere in the model. The identifier can be preceded and succeeded by whitespace within the `ci`. The set of possible identifiers that can appear in a `ci` element depends on the containing element in which the `ci` is used:

- If a `ci` element appears in the `math` body of a `FunctionDefinition` object (Section 4.3), the referenced identifier must be either (i) one of the declared arguments to that function, or (ii) the identifier of another `FunctionDefinition` object in the model.
- Otherwise, the referenced identifier must be that of a `Compartment`, `FunctionDefinition`, `Parameter`, `Reaction`, `Species` or `SpeciesReference` object defined in the model. The following are the only possible interpretations of using such an identifier in SBML:
 - *Compartment identifier*: When a `Compartment` identifier occurs in a `ci` element, it represents the size of the compartment. The units of measurement associated with the size of the compartment are those given by the `Compartment` instance’s `units` attribute value; see Section 4.5.4.

- *Function identifier*: When a **FunctionDefinition** identifier occurs in a **ci** element, it represents a call to that function. Function references in MathML occur in the context of using MathML’s **apply** and often involve supplying arguments to the function; see Section 4.3. The units associated with the value returned by the function call are the overall units of the mathematical expression contained in the function definition.
- *Parameter identifier*: When a **Parameter** identifier occurs in a **ci** element, it represents the numerical value assigned to that parameter. The units associated with the parameter’s value are those given by the **Parameter** instance’s **units** attribute; see Section 4.7.3.
- *Reaction identifier*: When a **Reaction** identifier occurs in a **ci** element, it represents the rate of that reaction as defined by the **math** expression in the **KineticLaw** object within the **Reaction**. The units associated with that rate are *extent units/time units*, where the *extent units* are given by the **Model** attribute **extentUnits** (Section 4.2.6) and *time units* are given by the **Model** attribute **timeUnits** (Section 4.2.4). If a **Reaction** instance has no **KineticLaw**, using its reaction identifier this way is an error.
- *Species identifier*: When a **Species** identifier occurs in a **ci** element, it represents the quantity of that species in units of either *amount of substance* or *concentration*, depending on the species’ definition; see Section 4.6.5.
- *Species reference identifier*: When a **SpeciesReference** identifier occurs in a **ci** element, it represents the stoichiometry of the reactant or product in the reaction where the **SpeciesReference** object is defined. More precisely, the value of the **SpeciesReference** identifier is equal to the value of the **stoichiometry** attribute of that **SpeciesReference** object instance; see Section 4.11. The unit associated with the value is always **dimensionless**.

The content of **ci** elements in MathML formulas outside of a **KineticLaw** or **FunctionDefinition** must always refer to objects declared in the top level global namespace; i.e., SBML uses “early binding” semantics. Inside of **KineticLaw**, **ci** elements can additionally refer to identifiers of **LocalParameter** objects defined within that **KineticLaw** instance; see Section 4.11.6 for more information.

3.4.4 Interpretation of boolean values

As noted already in Section 3.1.2, there is another unfortunate difference between the XML Schema 1.0 and MathML 2.0 standards that impacts mathematical expressions in SBML: in XML Schema, the value space of type **boolean** includes “**true**”, “**false**”, “**1**”, and “**0**”, whereas in MathML, only “**true**” and “**false**” count as boolean values.

The impact of this difference thankfully is minimal because the XML Schema definition is only used for attribute values on SBML objects, and those values turn out never to be accessible from MathML content in SBML—values of boolean attributes on SBML objects can never enter into MathML expressions. Nevertheless, software authors and users should be aware of the difference and in particular that “**0**” and “**1**” are interpreted as numerical quantities in mathematical expressions. There is no automatic conversion of “**0**” or “**1**” to boolean values in contexts where booleans are expected. This allows stricter type checking and unit verification during the validation of mathematical expressions.

3.4.5 Handling of whitespace

MathML 2.0 defines “whitespace” in the same way as XML does, i.e., the space character (Unicode hexadecimal code 0020), horizontal tab (code 0009), newline or line feed (code 000A), and carriage return (code 000D). In MathML, the content of elements such as **cn** and **ci** can be surrounded by whitespace characters. Prior to using the content, this whitespace is “trimmed” from both ends: all whitespace at the beginning and end of the content is removed (Ausb Brooks et al., 2003). For example, in **<cn> 42 </cn>**, the amount of white space on either side of the “42” inside the **<cn> ... </cn>** container does not matter. Prior to interpreting the content, the whitespace is removed altogether.

3.4.6 Use of `csymbol` elements in MathML expressions in SBML

SBML Level 3 uses the MathML `csymbol` element to denote certain built-in mathematical entities without introducing reserved names into the component identifier namespace. The `encoding` attribute of `csymbol` must be set to “text”. The `definitionURL` should be set to one of the following URIs defined by SBML:

- <http://www.sbml.org/sbml/symbols/time>. This represents the current simulation time. See Section 3.4.7 for more information. The unit of measurement associated with time is determined by the value of the `timeUnits` attribute on `Model`.
- <http://www.sbml.org/sbml/symbols/delay>. This represents a delay function. The delay function has the form $\text{delay}(x, d)$, taking two MathML expressions as arguments. The function’s value is the value of argument x , but taken at a time d before the current time. There are no restrictions on the form of x . Since the parameter d represents a time value, the units of d are expected to match the time units of the model; those are in turn specified by the value of the attribute `timeUnits` on `Model`. The value of the d parameter, when evaluated, must be numerical (i.e., a number in MathML real, integer, or “e-notation” format) and be greater than or equal to 0. The units of the return value of the delay functions are those of the parameter x . See Section 3.4.7 below for additional considerations surrounding the use of this `csymbol`.
- <http://www.sbml.org/sbml/symbols/avogadro>. This represents the numerical value of Avogadro’s constant. The value of Avogadro’s constant is determined experimentally; for the purposes of SBML Level 3 Version 1, the numerical value is taken to be the one recommended by the 2006 edition of CODATA (Mohr et al., 2008), but the unit is **dimensionless**. In other words, the value of this `csymbol` is equivalent to the following:

$$(6.02214179 \cdot 10^{23}) \cdot \text{dimensionless}$$

If the value of Avogadro’s constant is revised by international standards-setting organizations in the future, a future Version of the SBML Level 3 specification may stipulate a new value to be used for this `csymbol` constant. However, all software reading models expressed in *this* Version of SBML Level 3 should *always* use the value of Avogadro’s constant given above. (In other words, changes will apply only beginning with a possible new Version of SBML Level 3 and not this existing version.)

The following examples demonstrate these concepts. The XML fragment below encodes the formula $x + t$, where t stands for time.

```
<math xmlns="http://www.w3.org/1998/Math/MathML">
  <apply>
    <plus/>
    <ci> x </ci>
    <csymbol encoding="text" definitionURL="http://www.sbml.org/sbml/symbols/time">
      t
    </csymbol>
  </apply>
</math>
```

In the fragment above, the use of the token `t` is mostly a convenience for human readers—the string inside the `csymbol` could have been almost anything, because it is essentially ignored by MathML parsers and SBML. Some MathML and SBML processors will take note of the token and use it when presenting the mathematical formula to users, but the token used has no impact on the interpretation of the model and it does *not* enter into the SBML component identifier namespace. In other words, the SBML model cannot refer to `t` in the example above. The content of the `csymbol` element is for rendering purposes only and can be ignored by the parser.

As a further example, the following XML fragment encodes the equation $k + \text{delay}(x, 0.1)$ or, alternatively, $k_t + x_{t-0.1}$:

```

1      <math xmlns="http://www.w3.org/1998/Math/MathML"
2          xmlns:sbml="http://www.sbml.org/sbml/level3/version1/core">
3          <apply>
4              <plus/>
5              <ci> k </ci>
6              <apply>
7                  <csymbol encoding="text" definitionURL="http://www.sbml.org/sbml/symbols/delay">
8                      delay
9                  </csymbol>
10                 <ci> x </ci>
11                 <cn sbml:units="second"> 0.1 </cn>
12             </apply>
13         </apply>
14     </math>

```

Finally, the use of Avogadro’s number is illustrated in the following XML fragment:

```

16     <math xmlns="http://www.w3.org/1998/Math/MathML">
17         <apply>
18             <times/>
19             <apply>
20                 <csymbol encoding="text" definitionURL="http://www.sbml.org/sbml/symbols/avogadro">
21                     Avogadro's number
22                 </csymbol>
23                 <ci> x </ci>
24             </apply>
25         </apply>
26     </math>

```

3.4.7 Simulation time

The principal use of SBML is to represent quantitative dynamical models whose behaviors manifest themselves over time. In defining an SBML model using constructs such as reactions, time is most often implicit and does not need to be referred to in the mathematical expressions themselves. However, sometimes an explicit time dependency needs to be stated, and for this purpose, the *time* **csymbol** (described above in Section 3.4.6) may be used. This *time* symbol refers to “instantaneous current time” in a simulation, frequently given the literal name *t* in one’s equations.

An assumption in SBML is that “start time” or “initial time” in a simulation is zero, that is, if t_0 is the initial time in the system, $t_0 = 0$. This corresponds to the most common scenario. Initial conditions in SBML take effect at time $t = 0$. There is no mechanism in SBML for setting the initial time to a value other than 0. To refer to a different time in a model, one approach is to define a **Parameter** for a new time variable and use an **AssignmentRule** in which the assignment expression subtracts a value from the **csymbol** *time*. For example, if the desired offset is 2 time units, the MathML expression would be

```

40     <math xmlns="http://www.w3.org/1998/Math/MathML"
41         xmlns:sbml="http://www.sbml.org/sbml/level3/version1/core">
42         <apply>
43             <minus/>
44             <csymbol encoding="text" definitionURL="http://www.sbml.org/sbml/symbols/time"/>
45                 <cn sbml:units="second"> 2 </cn>
46         </apply>
47     </math>

```

SBML’s assignment rules (Section 4.9.3) can be used to express mathematical statements that hold true at all moments, so using an assignment rule with the expression above will result in the value being equal to $t - 2$ at every point in time. A parameter assigned this value could then be used elsewhere in the model.

3.4.8 Initial conditions and special considerations

The identifiers of **Species**, **SpeciesReference**, **Compartment**, **Parameter**, and **Reaction** object instances in a given SBML model refer to the main variables in a model. Depending on certain attributes of these objects (e.g., the attribute **constant** on species, species references, compartments and parameters—this and other conditions are explained in the relevant sections elsewhere in this document), some of the variables may have

constant values throughout a simulation, and others' values may change. These changes in values over time are determined by the system of equations constructed from the model's reactions, initial assignments, rules, and events.

As described in Section 3.4.7, an SBML model's simulation is assumed to begin at $t = 0$. The availability of the `delay` **csymbol** (Section 3.4.6) introduces the possibility that at $t \geq 0$, mathematical expressions in a model may draw on values of model components from time *prior* to $t = 0$. A simulator may therefore need to compute the values of variables at time points $t_i \leq 0$ to allow the calculation of values required for the evaluation of delay expressions in the model for $t \geq 0$. If there are no delays in the model, then $t_i = 0$.

The following is how the definitions of the model should be applied:

1. At time t_i :

- Every **Species**, **SpeciesReference**, **Compartment**, and **Parameter** whose definition includes an initial value is assigned that value. If an element has `constant="false"`, its value may be changed by other constructs or reactions in a model according to the steps below; if `constant="true"`, only an **InitialAssignment** can override the value.
- All **InitialAssignment** definitions take effect at t_i and continue to have effect up to and including $t = 0$, overriding any initial values on **Species**, **SpeciesReference**, **Compartment**, and **Parameter**. Since **InitialAssignments** contain mathematical formulas, different values may be computed at each time step t in $t_i \leq t \leq 0$.

2. For time $t \geq t_i$:

- **AssignmentRule** and **AlgebraicRule** definitions are in effect from this point in time forward and may influence the values of **Species** quantity, **SpeciesReference** stoichiometry, **Compartment** size, and **Parameter** values. (Note there cannot be both an **AssignmentRule** and an **InitialAssignment** for the same identifier; see Section 4.9.)

3. At time $t = 0$:

- The system of equations constructed by combining **AssignmentRule** equations, **AlgebraicRule** equations, **RateRule** equations, and the equations constructed from the **Reaction** definitions in the model, are used to obtain consistent initial conditions for numerical solver algorithms. (Note that there cannot be both an **AssignmentRule** and a **RateRule** for the same identifier, or both an **AssignmentRule** and an **InitialAssignment** for the same identifier; see Section 4.9.3.)
- **Constraint** definitions begin to take effect (and a constraint violation may result; see Section 4.10).

4. For time $t > 0$:

- **RateRule** definitions can begin to take effect.
- **Event** definitions can begin to take effect. (Note that an **Event** cannot be defined to change the value of a variable that is also the subject of an **AssignmentRule**; see Section 4.12.)
- System simulation proceeds.

To reiterate: in modeling situations that do not involve the use of the `delay` **csymbol**, then t_i becomes $t_i = 0$, but this does not alter the steps numbers 1–4 above.

3.4.9 MathML expression data types

MathML operators in SBML each return results in one of two possible types: boolean and numerical. By *numerical* type, we mean either (1) a number in MathML real, integer, rational, or “e-notation” format; or (2) the `csymbol` for delay or the `csymbol` for time described in Section 3.4.6. The following guidelines summarize the different possible cases.

The relational operators (`eq`, `neq`, `gt`, `lt`, `geq`, `leq`), the logical operators (`and`, `or`, `xor`, `not`), and the boolean constants (`false`, `true`) always return boolean values. As noted in Section 3.4.4, the numbers `0` and `1` do not count as boolean values in MathML contexts in SBML.

The type of an operator referring to a **FunctionDefinition** is determined by the type of the top-level operator of the expression in the **math** element of the **FunctionDefinition** instance, and can be boolean or numerical.

All other operators, values and symbols return numerical results.

The roots of the expression trees used in the following contexts must yield boolean values:

- the arguments of the MathML logical operators (**and**, **or**, **xor**, **not**);
- the second argument of a MathML **piece** operator;
- the **trigger** element of an SBML **Event**; and
- the **math** element of an SBML **Constraint**.

The roots of the expression trees used in the following contexts can optionally yield boolean values:

- the arguments to the **eq** and **neq** operators;
- the first arguments of MathML **piece** and **otherwise** operators; and
- the top level expression of a function definition.

The roots of expression trees in other contexts must yield numerical values.

The type of expressions should be used consistently. The set of expressions that make up the first arguments of the **piece** and **otherwise** operators within the same **piecewise** operator should all return values of the same type. The arguments of the **eq** and **neq** operators should return the same type.

3.4.10 Consistency of units in mathematical expressions and treatment of unspecified units

Strictly speaking, physical validity of mathematical formulas requires not only that physical quantities added to or equated with each other have the same fundamental dimensions and units of measurement; it also requires that the application of operators and functions to quantities produces sensible results. Yet, in real-life models today, these conditions are often and sometimes legitimately disobeyed.

In a public vote held in late 2007, the SBML community decided to revoke the requirement (present up through Level 2 Version 3) for strict unit consistency in SBML. As a result, SBML Level 3 follows this decision; the units on quantities and the results of mathematical formulas in a model *should* be consistent, but it is not a strict error if they are not. The following are thus formulated as recommendations that *should* be followed except in special circumstances.

Recommendations for unit consistency of mathematical expressions

The consistency of units is defined in terms of dimensional analysis applied recursively to every operator and function and every argument to them. The following conditions should hold true in a model (and software developers may wish to consider having their software warn users if one or more of the following conditions is not true):

1. All arguments to the following operators should have the same units (regardless of what those units happen to be): **plus**, **minus**, **eq**, **neq**, **gt**, **lt**, **geq**, **leq**.
2. The units of each argument in a call to a **FunctionDefinition** should match the units expected by the **lambda** expression within the **math** expression of that **FunctionDefinition** instance.
3. All of the possible return values from **piece** and **otherwise** subelements of a **piecewise** expression should have the same units, regardless of what those units are. (Without this guideline, the **piecewise** expression would return values having different units depending on which case evaluated to true.)
4. For the *delay* **csymbol** (Section 3.4.6) function, which has the form *delay*(*x*,*d*), the second argument *d* should match the model's unit of *time* (as determined by the value of the “**timeUnits**” attribute on **Model**).
5. The units of the value returned by the *delay* **csymbol** (Section 3.4.6) function, should match the units of the first argument *x*.

6. The units of each argument to the following operators should be “**dimensionless**”: **exp**, **ln**, **log**, **factorial**, **sin**, **cos**, **tan**, **sec**, **csc**, **cot**, **sinh**, **cosh**, **tanh**, **sech**, **csch**, **coth**, **arcsin**, **arccos**, **arctan**, **arcsec**, **arccsc**, **arccot**, **arcsinh**, **arccosh**, **arctanh**, **arcsech**, **arccsch**, **arccoth**.
7. The two arguments to **power**, which are of the form $power(a, b)$ with the meaning a^b , should be as follows: (1) if the second argument is an integer, then the first argument can have any units; (2) if the second argument b is a rational number n/m , it should be possible to derive the m -th root of $(a\{\text{units}\})^n$, where $\{\text{units}\}$ signifies the units associated with a ; otherwise, (3) the units of the first argument should be “**dimensionless**”. The second argument (b) should always have units of “**dimensionless**”.
8. The two arguments to **root**, which are of the form $root(n, a)$ with the meaning $\sqrt[n]{a}$ and where the degree n is optional (defaulting to “2”), should be as follows: (1) if the optional degree qualifier n is an integer, then it should be possible to derive the n -th root of a ; (2) if the optional degree qualifier n is a rational n/m then it should be possible to derive the n -th root of $(a\{\text{units}\})^m$, where $\{\text{units}\}$ signifies the units associated with a ; otherwise, (3) the units of a should be “**dimensionless**”.
9. Where the units of literal numbers have not been specified directly in SBML, it is possible for the units of a **FunctionDefinition** object’s return value to be effectively different in different contexts where it is called (see below). If a **FunctionDefinition**’s mathematical formula contains literal constants (i.e., numbers within MathML **cn** elements with no **sbml:units** attribute), the units of the constants should be identical in all contexts the function is called.

The units of other operators such as **abs**, **floor**, and **ceiling**, can be anything.

Item number 9 above, regarding **FunctionDefinition**, merits additional elaboration. An example may help illustrate the problem. Suppose the formula $x + 5$ is defined as a function, where x is an argument and the literal number 5 has unspecified units. If this function is called with an argument in moles, the only possible consistent unit for the return value is mole. If in another context in the same model, the function is called with an argument in seconds, the function return value can only be treated as being in seconds. Now suppose that a modeler decides to change all uses of seconds to milliseconds in the model. To make the function definition return the same quantity in terms of seconds, the 5 in the formula would need to be changed, but doing so would change the result of the function everywhere it is called—with the wrong consequences in the context where moles were intended. This illustrates the subtle danger of using numbers with unspecified units in function definitions. There are at least two approaches for avoiding this: (1) define separate functions for each case where the units of the constants are supposed to be different, optionally explicitly defining the units of the literal number; or (2) declare the necessary constants as **Parameter** objects in the model (with declared units!) and pass those parameters as arguments to the function, avoiding the use of literal numbers in the function’s formula.

Treatment of unspecified units

If an expression contains literal numbers and/or SBML components without declared units, the consistency or inconsistency of units may be impossible to determine. In the absence of a verifiable *inconsistency*, an expression in SBML is accepted as-is; the writer of the model is assumed to have written what they intended. However, this is *not* equivalent to assuming the expression *does* have consistent units. The lack of declared units on quantities in an SBML model does not render the model invalid insofar as the SBML specification is concerned, but it reduces the types of consistency checks and useful operations (such as conversions and translations) that software systems can perform.

In some cases, it may be possible to determine that expressions containing unspecified units are inconsistent regardless of what units would be attributed to the unspecified quantities. For example, the expression

$$\frac{dX}{dt} = \frac{[Y] \cdot [Z]^n}{[Z]^m + 1} \cdot V$$

with X , Y and Z in units of substance, V in units of volume, and $m \neq n$, cannot ever be consistent, no matter what units the literal 1 takes on. (This also illustrates the need not to stop verifying the units of an

expression immediately upon encountering an unspecified quantity—the rest of the expression may still be profitably evaluated and checked for inconsistency.)

3.4.11 *SBML does not define implicit unit conversions*

Implicit unit conversions do not exist in SBML. Consider the following example. Suppose that in some model, a species S_1 has been declared as having a mass of 1 kg, and a second species S_2 has been declared as having a mass of 500 g. What should be the result of evaluating an expression such as $S_1 > S_2$? If the numbers alone are considered,

$$1 > 500$$

would evaluate to “**false**”, but if the units were implicitly converted by the software tool interpreting the model,

$$1 \text{ kg} > 500 \text{ g}$$

would evaluate to “**true**”. This is a trivial example, but the problem for SBML is that implicit unit conversions of this kind can lead to controversial situations where even humans do not agree on the answer. Consequently, SBML only requires that mathematical expressions be evaluated numerically. It is up to the model writer to ensure that the units on both sides of an expression match, by inserting explicit unit conversion factors if necessary.

4 SBML components

In this section, we define each of the major components of SBML. We use the UML notation described in Section 1.4.3 for defining classes of objects. We also illustrate the use of SBML components by giving partial model definitions in XML. Section 7 provides many complete example models encoded in SBML.

4.1 The SBML container

All well-formed SBML documents must begin with an *XML declaration*, which specifies both the version of XML assumed and the document character encoding. The declaration begins with the characters `<?xml` followed by the XML **version** and **encoding** attributes. SBML Level 3 uses XML version 1.0 and requires a document encoding of UTF-8. Following this declaration, the outermost portion of a model expressed in Level 3 consists of an object of class **Sbml**, defined in Figure 9. This class contains three required attributes (**level**, **version** and **xmlns**), and a required **model** element.

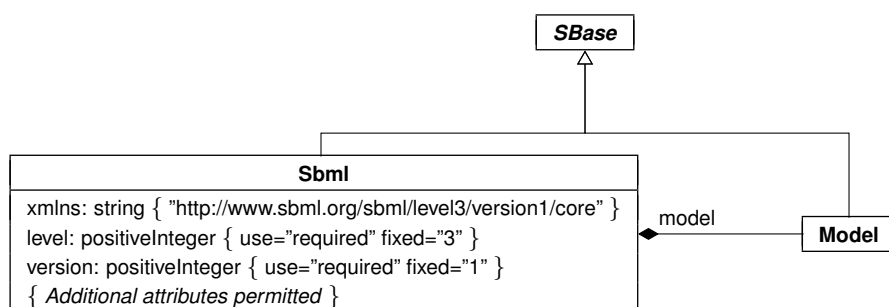


Figure 9: The definition of **Sbml** for SBML Level 3 Version 1 Core. The class **Model** is defined in Section 4.2. Note that **Sbml** and **Model** are subclasses of **SBase**, and therefore inherit the attributes of that abstract class.

The following is an abbreviated example of **Sbml** translated into XML form for an SBML Level 3 Version 1 Core document (and here, ellipses are used to indicate content elided from this example):

```
<?xml version="1.0" encoding="UTF-8"?>
<sbml xmlns="http://www.sbml.org/sbml/level3/version1/core" level="3" version="1">
  ...
  <model ...>
    ...
  </model>
</sbml>
```

The attribute **xmlns** declares the XML namespace used within the **sbml** element. The URI for SBML Level 3 Version 1 Core is <http://www.sbml.org/sbml/level3/version1/core>. All SBML Level 3 Version 1 Core elements and attributes must be placed in this namespace either by assigning the default namespace as shown in the example above, or using a tag prefix on every element. The **sbml** element may contain additional attributes, in particular, attributes to support the inclusion of SBML Level 3 packages; see Section 4.1.2. For purposes of checking conformance to the SBML Level 3 *Core* specification, only the elements and attributes in the SBML Level 3 Core XML namespace are considered.

4.1.1 The model element

The actual model contained within an SBML document is defined by an instance of the **Model** class element. The structure of this object and its use are described in Section 4.2. Every SBML document must contain one model definition. (As a result of extension packages defined in SBML Level 3, it is possible that a model is composed of multiple submodels; however, there must still be *one* top-level model defining the structure of the overall composition.)

4.1.2 Package declarations

SBML Level 3 is modular, in the sense of having a defined core set of features and optional packages adding features on top of the core. This modular approach means that models can declare which feature-sets they use, and likewise, software tools can declare which packages they support. The mechanism for models to declare which packages they use involves two parts: a standard XML namespace declaration, and an attribute that every package must declare in this namespace.

1. Every SBML Level 3 package is identified uniquely by an XML namespace URI. The use of a given SBML Level 3 package must be declared by a model using the standard XML namespace declaration approach. The declaration is made using the character sequence “**xmlns:**” (without the quotes), followed by additional characters providing a prefix by which elements and attributes in that namespace are known in the rest of the SBML document, and finally followed by the namespace URI as a value. The following is an example of namespace declarations for a package nicknamed “**multi**” and another package nicknamed “**layout**” (and here, ellipses are used to indicate content elided from this example):

```
<sbml xmlns="http://www.sbml.org/sbml/level3/version1/core" level="3" version="1"
      xmlns:multi="http://www.sbml.org/sbml/level3/version1/multi/version1"
      xmlns:layout="http://www.sbml.org/sbml/level3/version1/layout/version1" ...>
  ...
</sbml>
```

There are no restrictions on the prefixes used for XML namespaces referring to SBML Level 3 packages beyond those imposed by the relevant specifications of XML 1.0 and XML namespaces. (In other words, the prefix strings “**multi**” and “**layout**” in the example above are arbitrarily chosen, and could have been something else.)

2. SBML Level 3 requires that every package defines the addition of at least one attribute named **required**. The attribute, being in the namespace of the Level 3 package in question, must be referenced by the XML namespace prefix described in point number 1 above. The value of the **required** attribute indicates whether understanding the package is required for complete mathematical interpretation of a model, or whether the package is optional. A value of **required**=“**true**” indicates that interpreting the package is required. The following is an example:

```
<sbml xmlns="http://www.sbml.org/sbml/level3/version1/core" level="3" version="1"
      xmlns:multi="http://www.sbml.org/sbml/level3/version1/multi/version1"
      xmlns:layout="http://www.sbml.org/sbml/level3/version1/layout/version1"
      multi:required="true"
      layout:required="false" ... >
  ...
</sbml>
```

If a package is declared optional, it means that the elements and attributes added by the package can be ignored without any loss of meaning of the model’s full mathematics. “Ignoring” a package can be accomplished in multiple ways; a reader could either skip those elements or attributes altogether during parsing, or read them but not interpret them, or do something similar.

The XML namespace declaration for an SBML Level 3 package is an indication that a model makes use of features defined by that package, while the **required** attribute indicates whether the features may be ignored without compromising the mathematical meaning of the model. Both are necessary for a complete reference to an SBML Level 3 package. (On the other hand, no declaration is necessary for the Level 3 Core package, since it is the base package and support for it is required in any case.)

4.2 Model

The definition of **Model** is shown in Figure 10 on the next page. Only one instance of a **Model** object is allowed per instance of an SBML Level 3 Version 1 Core document or data stream, and it must be located inside the `<sbml> ... </sbml>` element as described in Section 4.1.

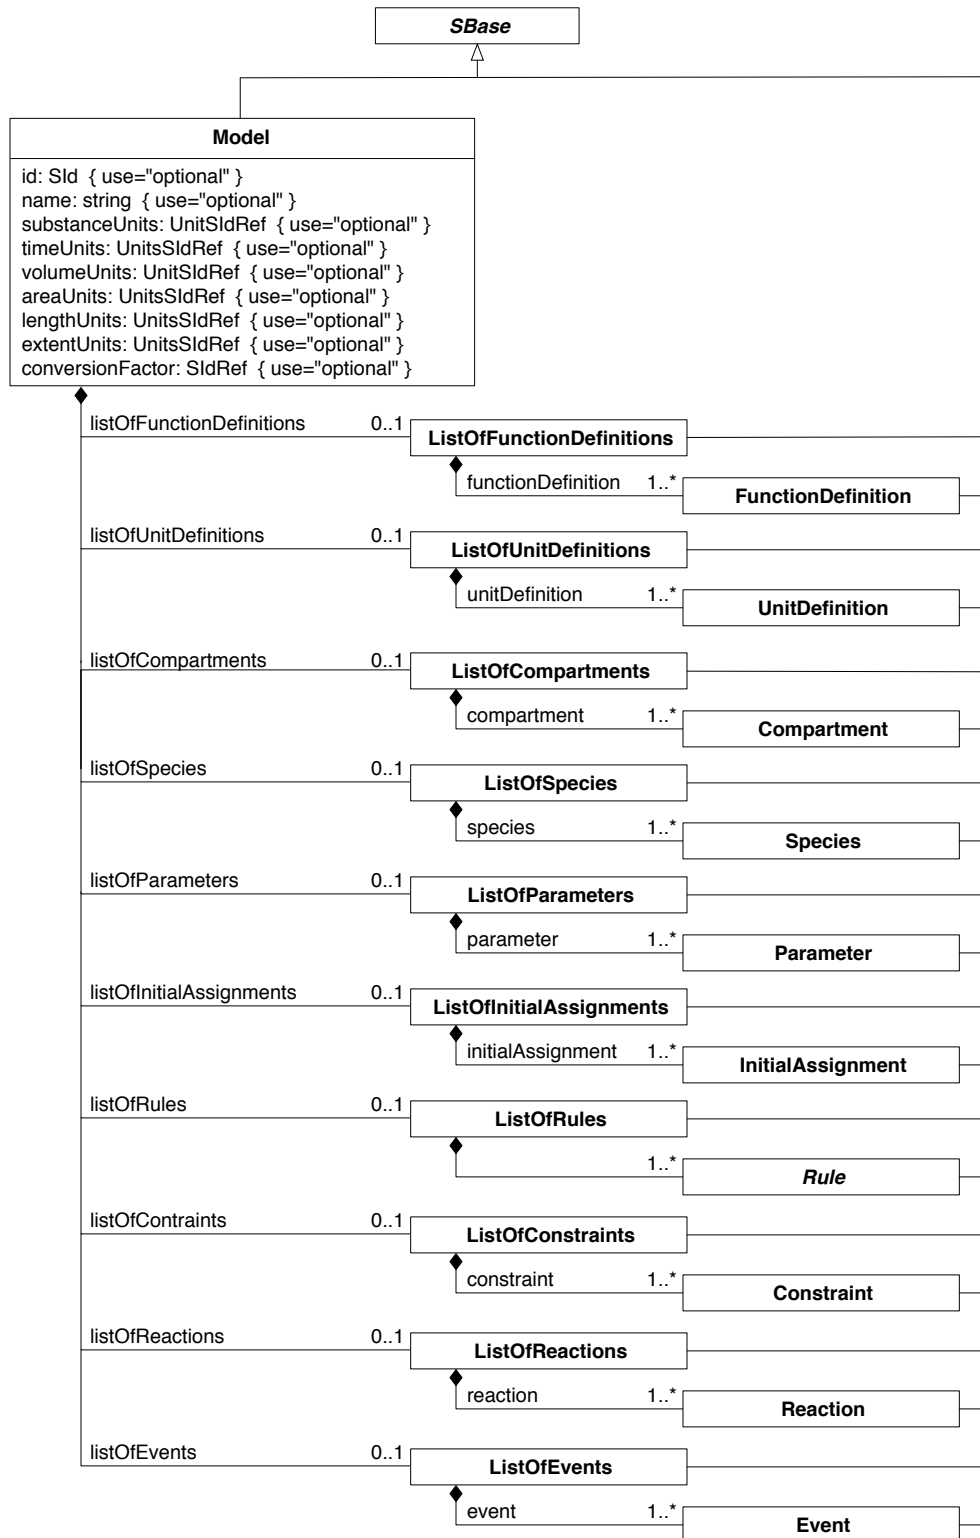


Figure 10: The definition of *Model* and the many helper classes *ListOfFunctionDefinitions*, *ListOfUnitDefinitions*, *ListOfCompartments*, *ListOfSpecies*, *ListOfParameters*, *ListOfInitialAssignments*, *ListOfRules*, *ListOfConstraints*, *ListOfReactions*, and *ListOfEvents*.

Model serves as a container for components of classes **FunctionDefinition**, **UnitDefinition**, **Compartment**, **Species**, **Parameter**, **InitialAssignment**, **Rule**, **Constraint**, **Reaction** and **Event**. Instances of the classes are placed inside instances of classes **ListOfFunctionDefinitions**, **ListOfUnitDefinitions**, **ListOfCompartments**, **ListOfSpecies**, **ListOfParameters**, **ListOfInitialAssignments**, **ListOfRules**, **ListOfConstraints**, **ListOfReactions**, and **ListOfEvents**. The “list” classes are defined in Figure 10. All of the lists are optional, but if a given list container is present within the model, the list must not be empty; that is, it must have length one or more. The resulting XML data object for a full model containing every possible list would have the following form:

```
<?xml version="1.0" encoding="UTF-8"?>
<sbml xmlns="http://www.sbml.org/sbml/level3/version1/core" level="3" version="1">
  <model id="My_Model">
    <listOfFunctionDefinitions>
      one or more <functionDefinition> ... </functionDefinition> elements } optional
    </listOfFunctionDefinitions>
    <listOfUnitDefinitions>
      one or more <unitDefinition> ... </unitDefinition> elements } optional
    </listOfUnitDefinitions>
    <listOfCompartments>
      one or more <compartment> ... </compartment> elements } optional
    </listOfCompartments>
    <listOfSpecies>
      one or more <species> ... </species> elements } optional
    </listOfSpecies>
    <listOfParameters>
      one or more <parameter> ... </parameter> elements } optional
    </listOfParameters>
    <listOfInitialAssignments>
      one or more <initialAssignment> ... </initialAssignment> elements } optional
    </listOfInitialAssignments>
    <listOfRules>
      one or more elements of subclasses of Rule } optional
    </listOfRules>
    <listOfConstraints>
      one or more <constraint> ... </constraint> elements } optional
    </listOfConstraints>
    <listOfReactions>
      one or more <reaction> ... </reaction> elements } optional
    </listOfReactions>
    <listOfEvents>
      one or more <event> ... </event> elements } optional
    </listOfEvents>
  </model>
</sbml>
```

Although the lists are optional, there are dependencies between SBML components, such that defining some components requires defining others. For example, defining a species requires defining a compartment, and defining a reaction requires defining a species. Such dependencies are explained throughout this document.

4.2.1 The id and name attributes

The **Model** object has an optional attribute, **id**, used to give the model an identifier. The value of **id** must conform to the syntax permitted by the **SId** data type described in Section 3.1.7. **Model** also has an optional **name** attribute, of type **string**. The **name** and **id** attributes must be used as described in Section 3.3.

4.2.2 The sboTerm attribute

Model inherits an optional **sboTerm** attribute of type **SBOTerm** from its parent class **SBase** (see Sections 3.1.11 and 5). When a value is given to this attribute in a **Model** instance, it should be an SBO identifier belonging to the branch for type **Model** indicated in Table 5. The term chosen should be the most precise (narrow) one that captures the overall process or phenomenon represented by the overall SBML model.

As discussed in Section 5, SBO labels are optional information on a model. Applications are free to ignore **sboTerm** values. A model must be interpretable without the benefit of SBO labels.

4.2.3 The substanceUnits attribute

The `substanceUnits` attribute specifies the units of measurement associated with substance quantities of `Species` objects that do not themselves specify the units. The attribute's value must be of type `UnitSIdRef` (Section 3.1.10). A list of recommended units is given in Section 8.2.1.

If a given `Species` object definition does not specify units of substance via the `substanceUnits` attribute on `Species` (described in Section 4.6), then the species quantity inherits the units specified by the `Model` `substanceUnits` attribute. If the `Model` does not define a value for this attribute, then there are no units to inherit, and all species that do not specify individual `substanceUnits` attribute values then have *no* declared units for their quantities. Section 4.6.4 provides more information about the units of species quantities.

Note that the overall units associated with the identifier of a species (when that identifier appears in a model's mathematical expressions) are *not solely determined* by setting `substanceUnits` on `Model` or `Species`. Sections 4.6.5 and 4.6.8 explain this point in more detail.

4.2.4 The timeUnits attribute

The `timeUnits` attribute specifies the units in which time is measured in the model. The value of this attribute must be of type `UnitSIdRef` (Section 3.1.10). A list of recommended units is given in Section 8.2.1.

This attribute on `Model` is the *only* way to specify the units of time for a model. It is a global attribute; time is measured in the same units everywhere in the model. This is particularly relevant to `Reaction` and `RateRule` objects in a model: all `Reaction` and `RateRule` objects in SBML define per-time values, and the time units are those given by the `timeUnits` attribute on the `Model` object instance. If the `Model` `timeUnits` attribute has no value, it means that the units of time are not defined for the model's reactions and rate rules. Leaving the units unspecified in an SBML model does not result in an invalid model; however, as a matter of best practice, we strongly recommend that all models specify the units of measurement for time.

4.2.5 The volumeUnits, areaUnits and lengthUnits attributes

The attributes `volumeUnits`, `areaUnits` and `lengthUnits` together are used to set the units of measurement for the sizes of `Compartment` objects in the model when those objects do not otherwise specify the units. The three attributes correspond to the different possible integral dimensions of compartments; that is, `volumeUnits` is applicable to compartments having `spatialDimensions`="3", `areaUnits` is applicable to compartments having `spatialDimensions`="2", and `lengthUnits` to compartments having `spatialDimensions`="1". The values of these attributes must be of type `UnitSIdRef` (Section 3.1.10). A list of recommended units is given in Section 8.2.1. The attributes are not applicable to compartments whose `spatialDimensions` attribute value is not one of "1", "2" or "3".

If a given `Compartment` object instance does not provide a value for its `units` attribute, then that compartment's size units are inherited from the units specified by the `Model` `volumeUnits`, `areaUnits` or `lengthUnits` attribute, as appropriate based on the `Compartment` object's `spatialDimensions` attribute value. If the `Model` object does not define the relevant units attribute, then there are no units to inherit, and all compartments that do not set a value for the `Compartment` `units` attribute then have *no* declared units for the compartment size. Section 4.5.4 provides more information about the units of compartment sizes.

The use of three separate attributes is a carry-over from SBML Level 2. Note that it is entirely possible for a model to define a value for two or more of the attributes `volumeUnits`, `areaUnits` and `lengthUnits` simultaneously, because SBML models may contain compartments with different numbers of dimensions.

4.2.6 The extentUnits attribute

Reactions are processes that occur over time. These processes involve events of some sort, where a single "reaction event" is one in which some set of entities (known as reactants, products and modifiers in SBML) interact, once. The *extent* of a reaction is a measure of how many times the reaction has occurred, while the time derivative of the extent gives the instantaneous rate at which the reaction is occurring. Thus, what is colloquially referred to as the "rate of the reaction" is in fact equal to the rate of change of reaction extent.

The combination of **extentUnits** and **timeUnits** defines the units of kinetic laws in SBML; this establishes how the numerical value of each **KineticLaw**'s mathematical formula (Section 4.11.6) is meant to be interpreted in a model. The units of the kinetic laws are taken to be **extentUnits** divided by **timeUnits**. A list of recommended units is given in Section 8.2.1.

Note that this embodies an important principle in SBML models: *all reactions in an SBML model must have the same units* for the rate of change of extent. In other words, the units of all reaction rates in the model *must be the same*. There is only one global value for **extentUnits** and one global value for **timeUnits**.

4.2.7 The conversionFactor attribute

The attribute **conversionFactor** defines a global value inherited by all **Species** object instances that do not define separate values for their **conversionFactor** attribute. The value of this attribute must be of type **SIdRef** (Section 3.1.8) and refer to a **Parameter** object instance defined in the model. The **Parameter** object in question must be a constant; i.e., have its **constant** attribute set to “true”.

If a given **Species** object definition does not specify a conversion factor via the **conversionFactor** attribute on **Species** (described in Section 4.6), then the species quantity inherits the conversion factor specified by the **Model** **conversionFactor** attribute. If the **Model** does not define a value for this attribute, then there is no conversion factor to inherit. Section 4.11.7 describes how to interpret the effects of reactions on species in that situation. More information about conversion factors in SBML is provided in Sections 4.6 and 4.11.

4.2.8 The ListOf container classes

The various **ListOf**_____ classes defined in Figure 10 are merely containers used for organizing the main components of an SBML document. All are derived from the abstract class **SBase** (Section 3.2), and inherit **SBase**'s various attributes and subelements such as **metaid** and **annotation**. The **ListOf**_____ classes do not add any attributes of their own.

There are several motivations for grouping SBML components within XML elements named after **listOfClassName**s rather than placing all the components directly at the top level. First, the fact that the container classes are derived from **SBase** means that software tools can add information about the lists themselves into each list container's **annotation**, a feature that a number of today's software tools exploit. Second, we believe the grouping leads to a more modular structure that is helpful when working with elements from multiple SBML Level 3 packages. Third, we believe that it makes visual reading of models in XML easier, for situations when humans must inspect and edit SBML directly.

4.3 Function definitions

The **FunctionDefinition** object associates an identifier with a function definition. This identifier can then be used as the function called in subsequent MathML **apply** elements. **FunctionDefinition** is shown in Figure 11.




Figure 11: The definition of class **FunctionDefinition**. The contents of the **Lambda** class is a single MathML lambda expression (or a lambda surrounded by a semantics element). A function definition must contain exactly one math element defined by the **Lambda** class; note also that **Lambda** is not derived from **SBase**, which means that the attributes defined on **SBase** are not available on the math element. A sequence of one or more instances of **FunctionDefinition** objects can be located in an instance of **ListOfFunctionDefinitions** in **Model**, as shown in Figure 10.

Function definitions in SBML (also informally known as “user-defined functions”) have purposefully limited capabilities. As is made more clear below, a function cannot reference parameters or other model quantities

outside of itself; values must be passed as parameters to the function. Moreover, recursive and mutually-recursive functions are not permitted. The purpose of these limitations is to balance power against complexity of implementation. With the restrictions as they are, function definitions could be implemented as textual substitutions—they are simply macros. Software implementations therefore do not need the full function-definition machinery typically associated with programming languages.

4.3.1 The `id` and `name` attributes

The `id` and `name` attributes have types `SId` and `string`, respectively, and operate in the manner described in Section 3.3. MathML `ci` elements in an SBML model can refer to the function defined by a [FunctionDefinition](#) using the value of its `id` attribute.

4.3.2 The `math` element

The `math` element is a container for MathML content that defines the function. The content of this element can only be a MathML `lambda` element or a MathML `semantics` element containing a `lambda` element. [FunctionDefinition](#) is the only place in SBML where a `lambda` element can be used. The `lambda` element must begin with zero or more `bvar` elements, followed by any other of the elements in the MathML subset listed in Section 3.4.1 *except* `lambda` (i.e., a `lambda` element cannot contain another `lambda` element).

A further restriction on the content of `math` is that it cannot contain references to variables other than the variables declared to the `lambda` itself. That is, the contents of MathML `ci` elements inside the body of the `lambda` can only be the variables declared by its `bvar` elements, or the identifiers of other [FunctionDefinitions](#) defined in the same model. This restriction also applies to the `csymbol` for *time*. Functions must be written so that all variables or parameters used in the MathML content are passed to them via their function parameters.

4.3.3 The `sboTerm` attribute

[FunctionDefinition](#) inherits an optional `sboTerm` attribute of type `SBOTerm` from its parent class [SBase](#) (see Sections 3.1.11 and 5). When a value is given to this attribute in a [FunctionDefinition](#) instance, it should be an SBO identifier belonging to the branch for type [FunctionDefinition](#) indicated in Table 5. The relationship is of the form “the function definition *is a* X”, where X is the SBO term. The term chosen should be the most precise (narrow) one that captures the role of the function in the model.

As discussed in Section 5, SBO labels are optional information on a model. Applications are free to ignore `sboTerm` values. A model must be interpretable without the benefit of SBO labels.

4.3.4 Calling user-defined functions

Within MathML expressions in an SBML model, all calls to a function defined by a [FunctionDefinition](#) must use the same number of arguments as specified in the function’s definition. The number of arguments is equal to the number of `bvar` elements inside the `lambda` element of the function definition.

Note that [FunctionDefinition](#) does not have a separate attribute for defining the units of the value returned by the function. The units associated with the function’s return value, when the function is called from within MathML expressions elsewhere in SBML, are simply the overall units of the expression in [FunctionDefinition](#)’s `math` when applied to the arguments supplied in the call to the function.

4.3.5 Examples

The following abbreviated SBML example shows a [FunctionDefinition](#) object instance defining `pow3` as the identifier of a function computing the mathematical expression x^3 , and after that, the invocation of that function in the mathematical formula of a rate law. Note how the invocation of the function uses its identifier.

```
<model ...>
  ...
  <listOfFunctionDefinitions>
    <functionDefinition id="pow3">
```

```

1      <math xmlns="http://www.w3.org/1998/Math/MathML"
2          xmlns:sbml="http://www.sbml.org/sbml/level3/version1/core">
3          <lambda>
4              <bvar><ci> x </ci></bvar>
5              <apply>
6                  <power/>
7                  <ci> x </ci>
8                  <cn sbml:units="dimensionless"> 3 </cn>
9              </apply>
10         </lambda>
11     </math>
12 </functionDefinition>
13 </listOfFunctionDefinitions>
14 ...
15 <listOfReactions>
16     <reaction id="reaction_1" reversible="true" fast="false">
17         ...
18         <kineticLaw>
19             <math xmlns="http://www.w3.org/1998/Math/MathML">
20                 <apply>
21                     <ci> pow3 </ci>
22                     <ci> S1 </ci>
23                 </apply>
24             </math>
25         </kineticLaw>
26         ...
27     </reaction>
28 </listOfReactions>
29 ...
30 </model>

```

4.4 Unit definitions

Units of measurement may be supplied in a number of contexts in an SBML model and associated with a variety of components. The overall units of any mathematical formula appearing in SBML are those that arise naturally from the components and mathematical expressions comprising the formula, or in other words, the units obtained by doing dimensional analysis on the formula.

Rather than requiring a complete unit definition on every object, SBML provides a facility for defining units that can be referenced throughout a model. The SBML unit definition facility uses two classes of objects, **UnitDefinition** and **Unit**. Their definitions are shown in Figure 12 and explained in more detail in Sections 4.4.1 and 4.4.2 below.

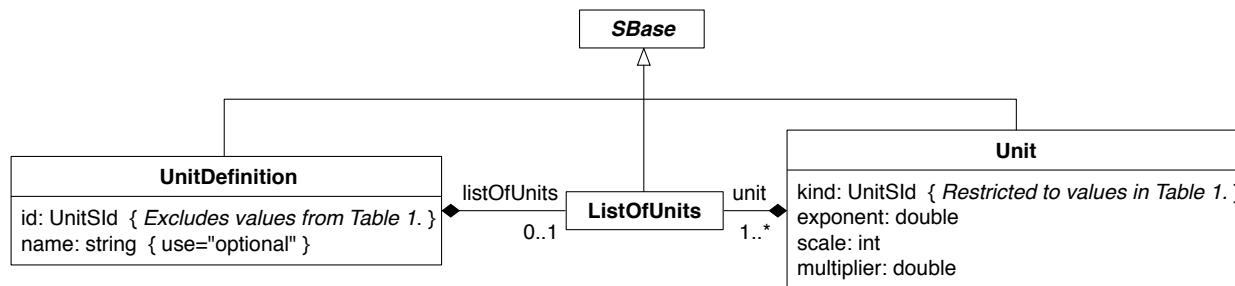


Figure 12: The definition of classes **UnitDefinition** and **Unit**. A sequence of one or more instances of **UnitDefinition** can be located in an instance of **ListOfUnitDefinitions** in **Model** (Figure 10). **ListOfUnits** has no attributes (beyond those it inherits from class **SBBase**); it merely acts as a container for one or more instances of **Unit** objects. Note that the only permitted values of kind on **Unit** are the reserved words in Table 1 on page 38, but these symbols are excluded from the permitted values of **UnitDefinition**'s id because SBML's unit system does not allow redefining the base units.

The approach to defining units in SBML is compositional; for example, meter second^{-2} is constructed by combining a **Unit** object representing *meter* with another **Unit** object representing second^{-2} . The combina-

tion is wrapped inside a **UnitDefinition**, which provides for assigning an identifier and optional name to the combination. The identifier can then be referenced from elsewhere in a model.

The vast majority of modeling situations requiring new SBML unit definitions involve simple multiplicative combinations of base units and factors. An example of this might be “moles per litre per second”. What distinguishes these sorts of simpler unit definitions from more complex ones is that they may be expressed without the use of an additive offset from a zero point. The use of offsets complicates all unit definition systems, yet in the domain of SBML the real-life cases requiring offsets are few (and in fact, to the best of our knowledge, only involve temperature). Consequently, the SBML unit system has been consciously designed in a way that attempts to simplify implementation of unit support for the most common cases in systems biology, at the cost of requiring units with offsets to be handled explicitly by the modeler. Section 8.2.1 discusses suggested approaches for handling situations requiring units with offsets.

4.4.1 UnitDefinition

A unit definition in SBML consists of an instance of a **UnitDefinition** object, shown in Figure 12.

The id and name attributes

The required attribute **id** and optional attribute **name** have data types **UnitSIdRef** (Section 3.1.10) and **string**, respectively. The **id** attribute is used to give the defined unit a unique identifier by which other parts of the model may refer to the unit. The **name** attribute is intended to be used for giving the unit definition an optional human-readable name; see Section 3.3.2 for more guidelines about the use of names.

There is one important restriction about the use of unit definition **id** values: the **id** of a **UnitDefinition** must *not* contain a value from Table 1, the list of reserved base unit names. This constraint simply prevents the redefinition of base units.

The list of Units

A **UnitDefinition** object may contain a **ListOfUnits** container which must contain one or more **Unit** objects. Section 4.4.2 explains the meaning and use of **Unit**.

Example

The following skeleton of a unit definition illustrates an example use of **UnitDefinition**:

```
<model ...>
  <listOfUnitDefinitions>
    <unitDefinition id="unit1">
      <listOfUnits>
        ...
      </listOfUnits>
    </unitDefinition>
    <unitDefinition id="unit2">
      <listOfUnits>
        ...
      </listOfUnits>
    </unitDefinition>
  </listOfUnitDefinitions>
  ...
</model>
```

4.4.2 Unit

A **Unit** object represents a (possibly transformed) reference to a base unit chosen from the list in Table 1 on the next page. The attribute **kind** indicates the chosen base unit, whereas the attributes **exponent**, **scale**, and **multiplier** define how the base unit is being transformed. These various attributes are described in detail below.

The kind attribute

The **Unit** attribute **kind** specifies a base unit to serve as the starting point for a new unit definition. The value of the attribute must be taken from the list of reserved words given in Table 1. These reserved symbols are defined in the value space of the data type **UnitSid** (Section 3.1.9).

ampere	farad	joule	lux	radian	volt
avogadro	gram	katal	metre	second	watt
becquerel	gray	kelvin	mole	siemens	weber
candela	henry	kilogram	newton	sievert	
coulomb	hertz	litre	ohm	steradian	
dimensionless	item	lumen	pascal	tesla	

Table 1: Base units defined in SBML. These symbols are predefined values of type **UnitSid** (Section 3.1.9). All are names of base or derived SI units (Bureau International des Poids et Mesures, 2006), except for “avogadro”, “dimensionless” and “item”, which are SBML additions important for handling certain common situations. The unit “avogadro” is the unit “dimensionless” multiplied with Avogadro’s number, the unit “dimensionless” is intended for cases where a quantity is a ratio whose units cancel out, and “item” is used for expressing such things as “N items” (e.g., “100 molecules”). Also, note that the gram and litre are not strictly part of SI; however, they are frequently used in SBML’s areas of application and therefore are included as predefined unit identifiers. (The standard SI unit of mass is in fact the kilogram, and volume is defined in terms of cubic metres.) Comparisons of these values, like all values of type **UnitSid**, must be performed in a case-sensitive manner.

Note that the set of acceptable values for the attribute **kind** does *not* include units defined by **UnitDefinition** objects. This means that the units definition system in SBML is not hierarchical: user-defined units cannot be built on top of other user-defined units, only on top of base units.

The presence of **avogadro** in Table 1 warrants an explanation. The Bureau International des Poids et Mesures specifically states, “The mole is the amount of substance of a system which contains as many elementary entities as there are atoms in 0.012 kilogram of carbon 12”—in other words, the SI unit **mole** is technically *a unit of measure for substance amount*. Although people sometimes use “mole” loosely to refer to other things (e.g., “a mole of *X*” to mean a number of *X* equal to Avogadro’s number, $6.022 \cdot 10^{23}$), such usage is not strictly correct. We believe it is even less correct in the context of reactions: although in SBML they are called “reactions”, there is nothing preventing the SBML **Reaction** construct from being used to represent other kinds of processes not involving substances. Consequently, we avoid using “mole” loosely where substances may not be involved, and instead use “Avogadro’s number of *X*”. In order to make it easier for models to define units in these terms, the SBML unit system includes the pseudo-unit “avogadro”, whose definition is identical to the definition of the *avogadro* **csymbol** described in Section 3.4.6. The numerical value is taken to be the one recommended by CODATA (Mohr et al., 2008), but the unit is **dimensionless**. In other words, it is defined as

$$(6.02214179 \cdot 10^{23}) \cdot \text{dimensionless}$$

where the dot (·) indicates simple scalar multiplication. If the value of Avogadro’s constant is revised by international standards-setting organizations in the future, a future Version of the SBML Level 3 specification may stipulate a new value to be used for **avogadro**. However, all software reading models expressed in *this* Version of SBML Level 3 should *always* use the value of Avogadro’s constant given above.

For purposes of comparing units, the following relationship should be assumed:

$$\text{mole} = \text{avogadro} \cdot \text{item} \quad (1)$$

The exponent, scale and multiplier attributes

The three other attributes **exponent**, **scale** and **multiplier** work together to permit the use of **Unit** for expressing new units in terms of the base units listed in Table 1. The formula underlying this definition is the following:

$$\{u_{\text{new}}\} = (\text{multiplier} \cdot 10^{\text{scale}} \cdot \{u_{\text{kind}}\})^{\text{exponent}} \quad (2)$$

This formula defines a new unit, $\{u_{\text{new}}\}$, in terms of another unit, $\{u_{\text{kind}}\}$. The unit $\{u_{\text{kind}}\}$ is one of the units listed in Table 1; in a given **Unit** object, it is chosen by setting the **kind** attribute. Each of the other components on the right-hand side of Equation 2 correspond to the remaining attributes in a **Unit** object instance, and their meanings are as follows:

- The **multiplier** attribute can be used to multiply the **kind** unit by a real-numbered factor. This enables the definition of units that are not power-of-ten multiples of SI units. For instance, a **multiplier** of 0.3048 could be used to define “foot” as a measure of length in terms of a “metre”. A value of **multiplier** must always be provided in a **Unit** object instance, but the value can be “1”.
- The **scale** attribute can be used to set the exponent for a power-of-ten multiplier that rescales the unit. For example, a unit having a **kind** value of “gram” and a **scale** value of “-3” signifies $10^{-3} \cdot \text{gram}$, or milligrams. A value of **scale** must always be provided in a **Unit** object instance; in those cases where a unit does not need to be scaled by a power of ten, the value of **scale** can be set to “0” (zero), because $10^0 = 1$.
- The **exponent** attribute can be used to specify an overall exponent on the unit definition. This provides a way to define units such as “cubic meter” in terms of the base unit “metre” (for which an **exponent** value of “3” would be appropriate). A value of **exponent** must always be provided in a **Unit** object instance.

4.4.3 Semantics and use of **Unit** and **UnitDefinition**

A single **Unit** object instance takes one of the base units from Table 1 and specifies how it should be transformed. A **UnitDefinition** object instance combines one or more **Unit** objects to define a new, composed unit, $\{u\}$. The new unit $\{u\}$ created by a **UnitDefinition** is defined as the product of all the **Unit** objects contained in the **ListOfUnits** within the **UnitDefinition** object instance. More formally,

$$\{u\} = \{u_1\} \cdot \{u_2\} \cdot \dots \cdot \{u_n\} \quad (3)$$

where the $\{u_i\}$ ’s are individual **Unit** definitions as defined by Equation 2. Now, let the value of the **multiplier** attribute of a given unit $\{u_i\}$ be represented by the variable m_i . Similarly, let the value of the **scale** attribute be represented by s_i , and the value of the **exponent** attribute be represented by x_i . Equation 3 can be rewritten in expanded form as

$$\begin{aligned} \{u\} &= (m_1 \cdot 10^{s_1} \cdot \{u_{b_1}\})^{x_1} \cdot (m_2 \cdot 10^{s_2} \cdot \{u_{b_2}\})^{x_2} \cdot \dots \cdot (m_n \cdot 10^{s_n} \cdot \{u_{b_n}\})^{x_n} \\ &= m_1^{x_1} \cdot m_2^{x_2} \cdot \dots \cdot m_n^{x_n} \cdot 10^{(s_1 x_1 + s_2 x_2 + \dots + s_n x_n)} \cdot \{u_{b_1}\}^{x_1} \cdot \{u_{b_2}\}^{x_2} \cdot \dots \cdot \{u_{b_n}\}^{x_n} \\ &= m \cdot 10^s \cdot \{u_{b_1}\}^{x_1} \cdot \{u_{b_2}\}^{x_2} \cdot \dots \cdot \{u_{b_n}\}^{x_n} \end{aligned} \quad (4)$$

where the terms m and s in the last line (Equation 4) are defined as

$$\begin{aligned} m &= m_1^{x_1} \cdot m_2^{x_2} \cdot \dots \cdot m_n^{x_n} \\ s &= s_1 x_1 + s_2 x_2 + \dots + s_n x_n \end{aligned}$$

Equation 4 expresses how a **UnitDefinition** object instance combines multiple **Unit** object instances to produce a new unit definition in an SBML model.

Examples

As a concrete example to illustrate the definitions above, let us define a unit for “foot” in terms of the base unit “metre”. This requires using **multiplier**=“0.3048”, **scale**=“0”, and **exponent**=“1”:

$$\text{foot} = 0.3048 \cdot 10^0 \cdot \text{metre}$$

The following fragment of SBML illustrates how this could be represented in XML:

```

1      <listOfUnitDefinitions>
2          <unitDefinition id="foot">
3              <listOfUnits>
4                  <unit kind="metre" multiplier="0.3048" scale="0" exponent="1"/>
5              </listOfUnits>
6          </unitDefinition>
7      </listOfUnitDefinitions>

```

To give another example, the following illustrates the definition of an abbreviation “**mmols**” for the units *millimoles* $l^{-1} s^{-1}$:

```

10     <listOfUnitDefinitions>
11         <unitDefinition id="mmols">
12             <listOfUnits>
13                 <unit kind="mole"    exponent="1"    scale="-3" multiplier="1"/>
14                 <unit kind="litre"   exponent="-1"   scale="0"  multiplier="1"/>
15                 <unit kind="second"  exponent="-1"   scale="0"  multiplier="1"/>
16             </listOfUnits>
17         </unitDefinition>
18     </listOfUnitDefinitions>

```

Comparing and/or converting units in a model

The unit system in SBML allows model quantities to be expressed in units other than the base units of Table 1. For analyses and computations, the consumer of the model (be it a software tool or a human) will want to convert all model quantities to base SI units for purposes such as verifying the consistency of units throughout the model. To understand this process in the context of SBML’s unit system, let us consider how we would convert a quantity y_b , measured in base units, to a new quantity y , measured in alternative units $\{u\}$. The relationship between the two quantities will be given by

$$y_b \cdot \{u_{b_1}\}^{z_1} \cdot \{u_{b_2}\}^{z_2} \cdot \dots \cdot \{u_{b_n}\}^{z_n} = y \cdot \{u\} \quad (5)$$

where the unknown exponents of the base units are denoted with z_i . Substituting Equation 4 for $\{u\}$ into Equation 5 above results in the following expression:

$$y_b \cdot \{u_{b_1}\}^{z_1} \cdot \{u_{b_2}\}^{z_2} \cdot \dots \cdot \{u_{b_n}\}^{z_n} = y \cdot m \cdot 10^s \cdot \{u_{b_1}\}^{x_1} \cdot \{u_{b_2}\}^{x_2} \cdot \dots \cdot \{u_{b_n}\}^{x_n} \quad (6)$$

Since the base units are independent, the equation above can be solved for the unknown exponents on the left-hand side: $z_i = x_i$ and $y_b = y \cdot m \cdot 10^s$.

Some additional points are worth discussing about the unit scheme. First, and most importantly, the equations above are formulated with the assumption that the base units do not require an additive offset as part of their definition. *When temperature values in units other than kelvin are being considered, then a different interpretation must be made*, as discussed in Section 8.2.1.

A second point is that care is needed to avoid seemingly-obvious but incorrect translations into this scheme of units taken from textbooks. The scheme above makes it easy to formulate statements such as “1 foot = 0.3048 metres” in the most natural way. However, the most common expression of the relationship between temperature in Fahrenheit and kelvin, “ $T_{Fahrenheit} = 1.8 \cdot (T_{kelvin} - 273.15) + 32$ ” might lead one to believe that defining Fahrenheit degrees in terms of kelvin degrees involves using **multiplier**=“1.8”. *Not so*, when degree changes are being considered and not temperature values. Converting *temperature values* is different from expressing a relationship between degree measurements. The proper value for the multiplier in the latter case is 5/9, i.e., **multiplier**=“0.555556” (where we picked an arbitrary decimal precision). If, on the other hand, the actual temperature is relevant to a quantity (e.g., if a model uses a quantity that has particular values at particular temperatures), then offsets are required in the unit calculations. While the SBML unit system does not allow a means of specifying offsets directly, it is still possible to encode models using such units. Section 8.2.1 provides suggestions for possible ways of handling such cases.

4.4.4 References to units

An attribute that defines the units of a mathematical entity (e.g., the attribute **units** on **Parameter**) can refer to a defined unit whose identifier is chosen from among the following:

- The base units listed in Table 1 on page 38; and
- A new unit identifier defined by a **UnitDefinition** as described in Section 4.4.

4.5 Compartments

A *compartment* in SBML represents a bounded space in which species are located. Compartments do not necessarily have to correspond to actual structures inside or outside of a biological system, although models are often designed that way. The definition of **Compartment** is shown in Figure 13.

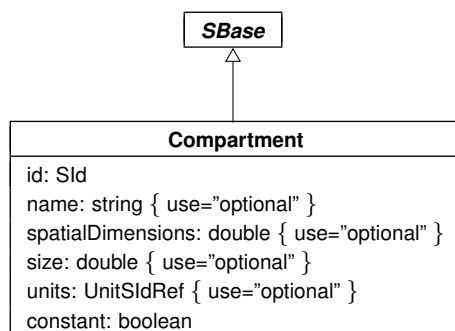


Figure 13: The definition of class **Compartment**. A sequence of one or more instances of **Compartment** objects can be located in an instance of **ListOfCompartments** in **Model**, as shown in Figure 10.

It is important to note that although compartments are optional in the overall definition of **Model** (see Section 4.2), every species in an SBML model must be located in a compartment. This in turn means that if a model defines any species, the model must also define at least one compartment. The reason is simply that species represent physical things, and therefore must exist *somewhere*. Compartments represent the *somewhere*.

4.5.1 The id and name attributes

Compartment has one required attribute, **id**, of type **SId**, to give the compartment a unique identifier by which other parts of an SBML model definition can refer to it. A compartment can also have an optional **name** attribute of type **string**. Identifiers and names must be used according to the guidelines described in Section 3.3.

4.5.2 The spatialDimensions attribute

A **Compartment** object has an optional floating-point attribute named **spatialDimensions** whose value indicates the number of spatial dimensions possessed by the compartment. Most modeling scenarios involve compartments with integer values of **spatialDimensions**="3" (i.e., a three-dimensional compartment, which is to say, a volume), **spatialDimensions**="2" (i.e., a two-dimensional compartment, a surface), or **spatialDimensions**="1" (i.e., a one-dimensional compartment, which is to say, a line). However, SBML Level 3 does not restrict compartments' **spatialDimensions** values, in order to allow for the possibility of representing structures with fractal dimensions.

In SBML Level 3 Version 1 Core, the number of spatial dimensions possessed by a compartment cannot enter into mathematical formulas, and therefore cannot alter the numerical interpretation of a model. However, the value of **spatialDimensions** does affect the interpretation of the units associated with a compartment's **size** value. Specifically, the value of **spatialDimensions** is used to select among the **Model** attributes **volumeUnits**, **areaUnits** and **lengthUnits** when a **Compartment** object does not define a value for its **units** attribute. This is described in more detail below in Section 4.5.4.

4.5.3 The size attribute

The optional **Compartment** attribute **size**, with a data type of **double**, can be used to set the initial size of the compartment. The size may correspond to a volume (if the compartment is a three-dimensional one), or it may be an area (if the compartment is two-dimensional), or a length (if the compartment is one-dimensional).

A compartment's size is set by its **size** attribute exactly once. If the compartment's **constant** attribute value is "true", then the size is fixed and cannot be changed except by an **InitialAssignment** in the model. The approach of using an **InitialAssignment** differs from setting the **size** attribute in that **size** can only be used to set the compartment size to a literal scalar value, whereas **InitialAssignment** allows the value to be set using an arbitrary mathematical expression. If the compartment's **constant** attribute is "false", the size value may be overridden by an **InitialAssignment** or changed by an **AssignmentRule** or **AlgebraicRule**, and in addition, for simulation time $t > 0$, it may also be changed by a **RateRule** or **Events**. (However, some constructs are mutually exclusive; see Sections 4.9 and 4.12.) It is not an error to set the value of **size** on a compartment and also redefine the value using an **InitialAssignment**, but the original **size** value in that case is ignored. Section 3.4.8 provides additional information about the semantics of assignments, rules and values for simulation time $t \leq 0$.

It is important to note that in SBML Level 3, a missing **size** value *does not imply that the compartment size is "1"*. A missing value for **size** for a given compartment signifies that the value either is unknown, or to be obtained from an external source, or determined by an initial assignment (Section 4.8) or other SBML construct elsewhere in the model. Further, due to the fact that alternative methods available for setting the size of a compartment, the **size** attribute must be defined as optional; however, *it is good practice to specify a value the size of every compartment in a model*, no matter what method is used, when compartment size values are available. The reasons for this are explained in Section 8.2.2.

4.5.4 The units attribute

The units associated with the compartment's **size** value may be set using the optional **Compartment** attribute **units**. The attribute's value must have the data type **UnitSIdRef** (Section 3.1.10).

The **units** attribute may be left unspecified for a given compartment in a model; in that case, the compartment inherits the units defined by one of the attributes on the enclosing **Model** object instance. The applicable attribute on **Model** depends on the value of the compartment's **spatialDimensions** attribute; the relationship is shown in Table 2 on the following page. If the **Model** object does not define the relevant attribute (**volumeUnits**, **areaUnits** or **lengthUnits**) for a given **spatialDimensions** value, the units associated with that **Compartment** object's size are undefined. If *both* **spatialDimensions** and **units** are left unset on a given **Compartment** object instance, then no units can be chosen from among the **Model**'s **volumeUnits**, **areaUnits** or **lengthUnits** attributes (even if the **Model** instance provides values for those attributes), because there is no basis to select between them and there is no default value of **spatialDimensions**. Leaving the units of compartment sizes undefined in an SBML model does not render the model invalid; however, as a matter of best practice, we strongly recommend that all models specify the units of measurement for every compartment size. A list of recommended units is given in Section 8.2.1.

The units of the compartment size, as defined by the **units** attribute or (if **units** is not set) the inherited value from **Model** according to Table 2 on the next page, are used in the following ways:

- The value of the **units** attribute is used as the units of the compartment identifier when the identifier appears as a numerical quantity in a mathematical formula expressed in MathML (discussed in Section 3.4.3).
- When a **Species** is to be treated in terms of concentrations or density, the units of the spatial size portion of the concentration value (i.e., the denominator in the units formula *amount/size*) are those indicated by the value of the **units** attribute on the compartment in which the species is located.
- The **math** element of an **AssignmentRule** or **InitialAssignment** referring to this compartment should have identical units (see Sections 4.9.3 and 4.8).

Value of attribute <code>spatialDimensions</code>	Attribute of Model used for inheriting units of size	Allowable kinds of units
"3"	<code>volumeUnits</code>	units of volume, or dimensionless
"2"	<code>areaUnits</code>	units of area, or dimensionless
"1"	<code>lengthUnits</code>	units of length, or dimensionless
other	<i>no units inherited</i>	<i>no restrictions</i>

Table 2: When a **Compartment** object instance does not specify a value for the `units` attribute, the units associated with the compartment's size are inherited from the enclosing **Model** instance according to the rules above. The left-hand column indicates the value of the compartment's `spatialDimensions` attribute, and the middle column indicates the **Model** attribute whose value should be used in that case. The right-hand column lists the types of units allowed in each case; the restrictions allow software tools to perform validation on the most common situations.

- In **RateRule** objects that set the rate of change of the compartment's size (Section 4.9.4), the units of the rule's `math` element should be identical to the compartment's `units` attribute divided by the model-wide *time* units. (In other words, the units for the rate of change of compartment size are *compartment size/time* units.)

4.5.5 The constant attribute

A **Compartment** also has a mandatory boolean attribute called `constant` that indicates whether the compartment's size stays constant or can vary during a simulation. A value of `"false"` indicates the compartment's size can be changed by other constructs in SBML. A value of `"true"` indicates the compartment's size cannot be changed by any construct except **InitialAssignment**. Section 4.5.3 provides more information.

4.5.6 The `sboTerm` attribute

Compartment inherits an optional `sboTerm` attribute of type **SBOTerm** from its parent class **SBase** (see Sections 3.1.11 and 5). When a value is given to this attribute in a **Compartment** instance, it should be an SBO identifier belonging to the branch for type **Compartment** indicated in Table 5. The relationship is of the form "the compartment *is a* X", where X is the SBO term. The term chosen should be the most precise (narrow) one that captures the role of the compartment in the model.

As discussed in Section 5, SBO labels are optional information on a model. Applications are free to ignore `sboTerm` values. A model must be interpretable without the benefit of SBO labels.

4.5.7 Examples

The following example illustrates two compartments in an abbreviated SBML example of a model definition. The compartment definitions do not set their `units` attribute, so both of these compartments inherit units of size (**litre**) from `model`.

```

<model volumeUnits="litre" ...>
  ...
  <listOfCompartments>
    <compartment id="cytosol" size="2.5" constant="true" spatialDimensions="3"/>
    <compartment id="mitochondria" size="0.3" constant="true" spatialDimensions="3"/>
  </listOfCompartments>
  ...
</model>

```

4.6 Species

A *species* in SBML refers to a pool of entities that (a) are considered indistinguishable from each other for the purposes of the model, (b) participate in reactions, and (c) are located in a specific *compartment*. The SBML **Species** object class is intended to represent these pools. Its definition is shown in Figure 14.

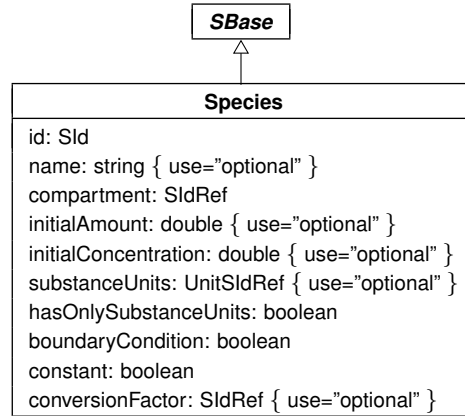


Figure 14: The definition of class **Species**. One or more instances of **Species** objects can be located in an instance of **ListOfSpecies** in **Model**, as shown in Figure 10.

4.6.1 The id and name attributes

As with other major objects in SBML, **Species** has a mandatory attribute, **id**, used to give the species an identifier. The identifier must be a text string conforming to the syntax permitted by the **SId** data type described in Section 3.1.7. **Species** also has an optional **name** attribute, of type **string**. The **name** and **id** attributes must be used as described in Section 3.3.

4.6.2 The sboTerm attribute

Species inherits an optional **sboTerm** attribute of type **SBOTerm** from its parent class **SBase** (see Sections 3.1.11 and 5). When a value is given to this attribute in a **Species** instance, it should be an SBO identifier belonging to the branch for type **Species** indicated in Table 5. The relationship is of the form “the species *is a* X”, where X is the SBO term. The term chosen should be the most precise (narrow) one that captures the role of the species in the model.

As discussed in Section 5, SBO labels are optional information on a model. Applications are free to ignore **sboTerm** values. A model must be interpretable without the benefit of SBO labels.

4.6.3 The compartment attribute

The required attribute **compartment**, of type **SIdRef**, is used to identify the compartment in which the species is located. The attribute’s value must be the identifier of an existing **Compartment** object in the model. Note that SBML does not have a concept of a default compartment—every species in an SBML model must be assigned a compartment *explicitly*, by setting the value of the **compartment** attribute. (This also implies that every model with one or more **Species** objects must define at least one **Compartment** object.)

4.6.4 The initialAmount, initialConcentration, and substanceUnits attributes

The optional attributes **initialAmount** and **initialConcentration**, both having a data type of **double**, can be used to set the initial quantity of the species in the compartment where the species is located. These two attributes are mutually exclusive—either one can be used, but *only one* can have a value on any given instance of a **Species** object. Missing **initialAmount** and **initialConcentration** values implies that their values either are unknown, or to be obtained from an external source, or determined by an initial assignment (Section 4.8) or other SBML construct elsewhere in the model.

A species’ initial quantity is set by the **initialAmount** or **initialConcentration** attributes exactly once. If the **constant** attribute is “**true**”, then the value of the species’ quantity is fixed and cannot be changed except by an **InitialAssignment**. These methods differ in that the **initialAmount** and **initialConcentration**

attributes can only be used to set the species quantity to a literal scalar value, whereas **InitialAssignment** allows the value to be set using an arbitrary mathematical expression. If the species' **constant** attribute is "false", the species' quantity value may be overridden by an **InitialAssignment** or changed by **AssignmentRule** or **AlgebraicRule**, and in addition, for $t > 0$, it may also be changed by a **RateRule**, **Events**, and as a result of being a reactant or product in one or more **Reactions**. (However, some constructs are mutually exclusive; see Sections 4.9 and 4.12.) It is not an error to define **initialAmount** or **initialConcentration** on a species and also redefine the value using an **InitialAssignment**, but the **initialAmount** or **initialConcentration** setting in that case is ignored. Section 3.4.8 provides additional information about the semantics of assignments, rules and values for simulation time $t \leq 0$.

When the **initialAmount** attribute is set, the units associated with its value are specified by the **Species** attribute **substanceUnits**, whose value must have the data type **UnitSIdRef** (Section 3.1.10). When the **initialConcentration** attribute is set, the units associated with its value are *amount/size* units, where the units of *amount* are those specified by the **substanceUnits** attribute, and the *size* units are those associated with the definition of the size of the **Compartment** in which the species is located. Note that in either case, *amount* units are involved and determined by the **substanceUnits** attribute. If the **substanceUnits** attribute is not set on a given **Species** object instance, then the value for that species is inherited from the **substanceUnits** attribute on the enclosing **Model** object instance. If that attribute on **Model** is not set either, then the units of the species' quantity are undefined. Leaving the units of species quantities undefined in an SBML model does not render the model invalid; however, as a matter of best practice, we strongly recommend that all models specify the units of measurement for every species quantity. A list of recommended units is given in Section 8.2.1.

Setting **initialAmount** or **initialConcentration** does *not* determine whether a species identifier represents an amount or a concentration when it appears elsewhere in an SBML model. Instead, that aspect is something controlled by the attribute **hasOnlySubstanceUnits**, discussed in Section 4.6.5 below.

4.6.5 The **hasOnlySubstanceUnits** attribute and the units of the species

When the identifier of a species appears in a mathematical expression or as the subject of SBML rules, initial assignments or event assignments, the identifier may represent either a concentration value or an amount value. The boolean attribute **hasOnlySubstanceUnits** determines which interpretation applies for a given **Species** object instance. The interpretation rule is the following. The *units of the species* are of the form *amount/size* units (i.e., *concentration* units, using a broad definition of concentration) if **hasOnlySubstanceUnits** has the value "false". Otherwise, if **hasOnlySubstanceUnits** has the value "true", then the *units of the species* are of the form *amount*.

Although it may seem as though this intention could be determined based on whether **initialConcentration** or **initialAmount** is set, the fact that these two **Species** attributes are optional means that a separate flag is needed. (Consider the situation where neither is set, and instead the species' quantity is established by an **InitialAssignment** or **AssignmentRule**.)

The *units of the species* are used in the following ways in SBML:

- The species identifier has these units when the identifier appears as a numerical quantity in a mathematical formula expressed in MathML (discussed in Section 3.4.3).
- The **math** element of an **AssignmentRule** or **InitialAssignment** referring to this species should have identical units (see Sections 4.9.3 and 4.8).
- In **RateRule** objects that set the rate of change of the species' quantity (Section 4.9.4), the units of the rule's **math** element should be identical to the *units of the species* divided by the model's *time* units.

4.6.6 The **constant** and **boundaryCondition** attributes

The **Species** object has two other mandatory boolean attributes named **constant** and **boundaryCondition**, used to indicate whether and how the quantity of that species can vary during a simulation. Table 3 shows how to interpret the combined values of the **boundaryCondition** and **constant** attributes.

constant value	boundaryCondition value	can have assignment or rate rule	can be reactant or product	species' quantity can be changed by
true	true	no	yes	(never changes)
false	true	yes	yes	rules and events
true	false	no	no	(never changes)
false	false	yes	yes	reactions <i>or</i> rules (but not both), and events

Table 3: How to interpret the values of the constant and boundaryCondition attributes on *Species*. Note that column four is specifically about reactants and products and not also about species acting as modifiers; the latter are by definition unchanged by reactions.

When a species is a product or reactant of one or more reactions, its quantity is determined by those reactions. In SBML, it is possible to indicate that a given species' quantity is *not* determined by the set of reactions even when that species occurs as a product or reactant; i.e., the species is on the *boundary* of the reaction system, and its quantity is not determined by the reactions. The boolean attribute **boundaryCondition** with value **"true"** can be used to indicate this. A value of **"false"** indicates that the species *is* part of the reaction system.

The **constant** attribute indicates whether the species' quantity can be changed at all, regardless of whether by reactions, rules, or constructs other than **InitialAssignment**. A value of **"false"** indicates that the species' quantity can be changed. This is the most common situation because the purpose of many models is precisely to calculate changes in species quantities over time. Note that the initial quantity of a species can be set by an **InitialAssignment** irrespective of the value of the **constant** attribute.

In practice, a **boundaryCondition** value of **"true"** means a differential equation derived from the reaction definitions should not be generated for the species. However, the species' quantity may still be changed by **AssignmentRule**, **RateRule**, **AlgebraicRule**, **Event**, and **InitialAssignment** constructs if its **constant** attribute is **"false"**. Conversely, if the species' **constant** attribute is **"true"**, then its value cannot be changed by anything except **InitialAssignment**.

A species having **boundaryCondition**=**"false"** and **constant**=**"false"** can appear as a product and/or reactant of one or more reactions in the model. If the species is a reactant or product of a reaction, it must not also appear as the target of any **AssignmentRule** or **RateRule** object in the model. If instead the species has **boundaryCondition**=**"false"** and **constant**=**"true"**, then it cannot appear as a reactant or product, or as the target of any **AssignmentRule**, **RateRule** or **EventAssignment** object in the model.

The example model in section 7.5 contains all four possible combinations of the **boundaryCondition** and **constant** attributes on **species** elements. Section 7.6 gives an example of how one can translate into ODEs a model that uses **boundaryCondition** and **constant** attributes.

4.6.7 The conversionFactor attribute

The attribute **conversionFactor** defines a conversion factor that applies to a particular species. The value of the attribute must have the data type **SidRef** and must be the identifier of a **Parameter** object instance defined in the model. That **Parameter** object must be a constant, meaning its **constant** attribute must be set to **"true"**. If a given **Species** object definition defines a value for its **conversionFactor** attribute, it takes precedence over any factor defined by the **Model** object's **conversionFactor** attribute.

In SBML, the units of species can be different from the units of extent of reactions in the model; thus, when calculating how a reaction affects a species, unit conversions may need to be applied. The use of a conversion factor in computing the effects of reactions on a species' quantity is explained in Section 4.11.7. Because the value of the **conversionFactor** attribute is the identifier of a **Parameter** object in the model, and because parameters can have units attached to them, the transformation from reaction extent units to species units can be completely specified using this approach.

Note that the unit conversion factor is only applied when calculating the effect of a reaction on a species. It is not used in any rules or other SBML constructs that affect the species, and it is also not used when the value of the species is referenced in a mathematical expression.

4.6.8 Additional considerations for interpreting the numerical value of a species

Species are unique in SBML in that they have a kind of duality: a species identifier may stand for either substance amount (meaning, a count of the number of individual entities) or a concentration or density (meaning, amount divided by a compartment size). The previous sections explain the meaning of a species identifier when it is referenced in a mathematical formula or in rules or other SBML constructs; however, it remains to specify what happens to a species when the compartment in which it is located changes in size.

When a species definition has the attribute value `hasOnlySubstanceUnits="false"` and the size of the compartment in which the species is located changes, the default in SBML is to assume that it is the concentration that must be updated to account for the size change. This follows from the principle that, all other things held constant, if a compartment simply changes in size, the size change does not in itself cause an increase or decrease in the number of entities of any species in that compartment. In a sense, the default is that the amount of a species is preserved across compartment size changes. Upon such size changes, the value of the concentration or density must be recalculated from the simple relationship $concentration = amount / size$ if the value of the concentration is needed (for example, if the species identifier appears in a mathematical formula or is otherwise referenced in an SBML construct). There is one exception: if the species' quantity is determined by an [AssignmentRule](#), [RateRule](#), [AlgebraicRule](#), or an [EventAssignment](#) and the species has the attribute value `hasOnlySubstanceUnits="false"`, it means that the *concentration* is assigned by the rule or event; in that case, the *amount* must be calculated when the compartment size changes. (Events also require additional care in this situation, because an event with multiple assignments could conceivably reassign both a species quantity and a compartment size simultaneously. Section 4.12.4 describes the handling of species in event assignments.)

Note that the above only matters if a species has the attribute `hasOnlySubstanceUnits="false"`, meaning that the species identifier refers to a concentration wherever the identifier appears in a mathematical formula. If instead the attribute's value is `"true"`, then the identifier of the species *always* stands for an amount wherever it appears in a mathematical formula or is referenced by an SBML construct. In that case, there is never a question about whether an assignment or event is meant to affect the amount or concentration: it is always the amount.

A particularly confusing situation occurs when the species has attribute `constant="true"`. Suppose that such a species is defined with `hasOnlySubstanceUnits="false"` and is given a value for `initialConcentration`. Does `constant="true"` mean that the concentration is held constant if the compartment size changes? No; it is still the amount that is kept constant across a compartment size change. The fact that the species was initialized using a concentration value is irrelevant.

4.6.9 Example

The following example shows two species definitions within an abbreviated SBML model definition. The example shows that species are listed under the heading `listOfSpecies` in the model:

```
<model ...>
  ...
  <listOfSpecies>
    <species id="Glucose" compartment="cell" initialConcentration="4"
      hasOnlySubstanceUnits="false" boundaryCondition="false"
      constant="false"/>
    <species id="Glucose_6_P" compartment="cell" initialConcentration="0.75"
      hasOnlySubstanceUnits="false" boundaryCondition="false"
      constant="false"/>
  </listOfSpecies>
  ...
</model>
```

4.7 Parameters

A **Parameter** is used in SBML to define a symbol associated with a value; this symbol can then be used in mathematical formulas in a model. The definition of **Parameter** is shown in Figure 15.

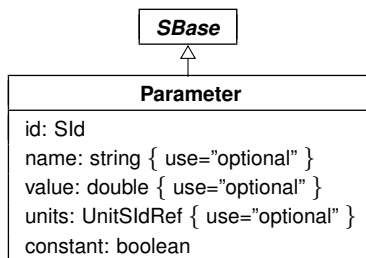


Figure 15: The definition of class **Parameter**. A sequence of one or more instances of **Parameter** objects can be located in an instance of **ListOfParameters** in **Model**, as shown in Figure 10.

The use of the term *parameter* in SBML sometimes leads to confusion among readers who have a particular notion of what something called “parameter” should be. It has been the source of heated debate, but despite this, no one has yet found an adequate replacement term that does not have different connotations to different people and hence leads to confusion among *some* subset of users. Perhaps it would have been better to have two constructs, one called “constants” and the other called “variables”. The current approach in SBML is simply more parsimonious, using a single **Parameter** construct with the boolean flag **constant** indicating which flavor it is. In any case, readers are implored to look past their particular definition of a “parameter” and simply view SBML’s **Parameter** as a single mechanism for defining both constants and (additional) variables in a model. (We write *additional* because the species in a model are usually considered to be the central variables.) After all, software tools are not required to expose to users the actual names of particular SBML constructs, and thus tools can present to their users whatever terms their designers feel best matches their target audience.

4.7.1 The id and name attributes

Parameter has one required attribute, **id**, of type **SId**, to give the parameter a unique identifier by which other parts of an SBML model definition can refer to it. A parameter can also have an optional **name** attribute of type **string**. Identifiers and names must be used according to the guidelines described in Section 3.3.

4.7.2 The value attribute

The optional attribute **value** determines the value (of type **double**) assigned to the identifier. A missing **value** implies that the **value** either is unknown, or to be obtained from an external source, or determined by an initial assignment (Section 4.8) or other SBML construct elsewhere in the model.

A parameter’s value is set by its **value** attribute exactly once. If the parameter’s **constant** attribute (Section 4.7.4) has the value “**true**”, then the value is fixed and cannot be changed except by an **InitialAssignment**. These two methods of setting the parameter’s value differ in that the **value** attribute can only be used to set it to a literal scalar value, whereas **InitialAssignment** allows the value to be set using an arbitrary mathematical expression. If the parameter’s **constant** attribute has the value “**false**”, the parameter’s value may be overridden by an **InitialAssignment** or changed by **AssignmentRule** or **AlgebraicRule**, and in addition, for simulation time $t > 0$, it may also be changed by a **RateRule** or **Events**. (However, some of these constructs are mutually exclusive; see Sections 4.9 and 4.12.) It is not an error to define **value** on a parameter and also redefine the value using an **InitialAssignment**, but the **value** in that case is ignored. Section 3.4.8 provides additional information about the semantics of assignments, rules and values for simulation time $t \leq 0$.

4.7.3 The units attribute

The units associated with the value of the parameter can be specified by the optional attribute **units**. The attribute's value must have the data type **UnitSidRef** (Section 3.1.10). There are no constraints on the units that can be assigned to a parameter; there are also no units to inherit from the enclosing **Model** object (unlike the case for, e.g., **Species** and **Compartment**).

The units of the parameter are used in the following ways:

- When the parameter identifier appears in mathematical formulas expressed in MathML in a model, the units associated with the value are those declared by the parameter's **units** attribute.
- The units of the **math** element of an **AssignmentRule**, **InitialAssignment** or **EventAssignment** setting the value of the parameter should be identical to the units declared by the parameter's **units** attribute.
- The units of the **math** element of a **RateRule** that references the parameter should be identical to *parameter units/time*, where *parameter units* are the units declared for the parameter using the **units** attribute and *time* is the model-wide **time** units.

The fact that the **units** attribute value is optional means that models can define parameters with undeclared units. Leaving the units of parameter values undefined in an SBML model does not render the model invalid; however, as mentioned elsewhere, as a matter of best practice, we strongly recommend that all models specify the units of measurement for every parameter.

4.7.4 The constant attribute

The **Parameter** object has a mandatory boolean attribute named **constant** that indicates whether the parameter's value can vary during a simulation. A value of "true" indicates the parameter's value cannot be changed by any construct except **InitialAssignment**. Conversely, if **constant**="false", other constructs in SBML, such as rules and events, can change the value of the parameter. More information about the effects of **constant** on **value** is presented in Section 4.7.2.

What if a parameter has its **constant** attribute set to "false", but the model does not contain any rules, events or other constructs that ever change its value over time? Although the model may be suspect (why is the parameter in question not flagged as being constant?), this situation is not strictly an error. A value of "false" for **constant** only indicates that a parameter *can* change value, not that it *must*.

4.7.5 The sboTerm attribute

Parameter inherits an optional **sboTerm** attribute of type **SBOTerm** from its parent class **SBase** (see Sections 3.1.11 and 5). When a value is given to this attribute in a **Parameter** instance, it should be an SBO identifier belonging to the branch for type **Parameter** indicated in Table 5. The relationship is of the form "the parameter *is a* X", where X is the SBO term. The term chosen should be the most precise (narrow) one that captures the role of the parameter in the model.

As discussed in Section 5, SBO labels are optional information on a model. Applications are free to ignore **sboTerm** values. A model must be interpretable without the benefit of SBO labels.

4.7.6 Example

The following is an example of parameters defined at the **Model** level:

```
<model ...>
  ...
  <listOfParameters>
    <parameter id="tau2" value="3e-2" units="second" constant="true"/>
    <parameter id="Km1" value="10.7" units="moleperlitre" constant="true"/>
  </listOfParameters>
  ...
</model>
```


4.8 Initial assignments

SBML Level 3 Version 1 Core provides two ways of assigning initial values to entities in a model. The simplest and most basic is to set the values of the appropriate attributes in the relevant components; for example, the initial value of a model parameter (whether it is a constant or a variable) can be assigned by setting its **value** attribute directly in the model definition (Section 4.7). However, this approach is not suitable when the value must be calculated, because the initial value attributes on different components such as species, compartments, and parameters are single values and not mathematical expressions. This is the reason for the introduction of **InitialAssignment**: to permit the calculation of the value of a constant or the initial value of a variable from the values of *other* quantities in a model. The definition of **InitialAssignment** is shown in Figure 16.




Figure 16: The definition of class **InitialAssignment**. The contents of the **Math** class can be any MathML permitted in SBML; see Section 3.4.1. A sequence of one or more instances of **InitialAssignment** objects can be located in an instance of **ListOfInitialAssignments** in **Model**, as shown in Figure 10.

As explained below, the provision of **InitialAssignment** does not mean that models necessarily must use this construct when defining initial values of quantities. If a value can be set using the relevant attribute of a component in a model, then that approach may be more efficient and more portable to other software tools. **InitialAssignment** should be used when the other mechanism is insufficient for the needs of a particular model.

Initial assignments have some similarities to assignment rules (Section 4.9.3). The main differences are (a) an **InitialAssignment** can set the value of a constant whereas an **AssignmentRule** cannot, and (b) unlike **AssignmentRule**, an **InitialAssignment** definition only applies up to and including the beginning of simulation time, i.e., $t \leq 0$, while an **AssignmentRule** applies at all times.

4.8.1 The symbol attribute

InitialAssignment contains the mandatory attribute **symbol**, of type **SIdRef**. The value of this attribute can be the identifier (i.e., the value of the **id** attribute) of a **Compartment**, **Species**, **SpeciesReference** or global **Parameter** elsewhere in the model. The purpose of the **InitialAssignment** is to define the initial value of the constant or variable referred to by the **symbol** attribute. (The attribute's name is **symbol** rather than **variable** because it may assign values to constants as well as variables in a model; see Section 4.8.4 below.)

An initial assignment cannot be made to reaction identifiers, that is, the **symbol** attribute value of an **InitialAssignment** cannot be an identifier that is the **id** attribute value of a **Reaction** object in the model. This is identical to a restriction placed on rules (see Section 4.9.5).

4.8.2 The math element

The **math** element contains a MathML expression used to calculate the value of the entity referenced by **symbol**. The units of the value should match the units associated with the identifier given in the **symbol** attribute. (That is, the units should be the units of the species, species reference, compartment, or parameter, as appropriate for the kind of object identified by the value of **symbol**.)

4.8.3 The sboTerm attribute

InitialAssignment inherits an optional **sboTerm** attribute of type **SBOTerm** from its parent class **SBBase** (see Sections 3.1.11 and 5). When a value is given to this attribute in a **InitialAssignment** instance, it should be an SBO identifier belonging to the branch for type **InitialAssignment** indicated in Table 5. The relationship is of the form “the initial assignment is a X”, where X is the SBO term. The term chosen should be the most precise (narrow) one that captures the role of the initial assignment in the model.

As discussed in Section 5, SBO labels are optional information on a model. Applications are free to ignore `sboTerm` values. A model must be interpretable without the benefit of SBO labels.

4.8.4 Semantics of initial assignments

The value calculated by an `InitialAssignment` object overrides the value assigned to the given symbol by the object defining that symbol. For example, if a `Compartment`'s `size` is set in its definition, and the model also contains an `InitialAssignment` having that compartment's `id` as its `symbol` value, then the interpretation is that the `size` assigned in the `Compartment` definition should be ignored and the value assigned based on the computation defined in the `InitialAssignment`. Initial assignments can take place for `Compartment`, `Species`, `SpeciesReference` and global `Parameter` objects regardless of the value of their `constant` attribute.

This does not mean that a definition of a symbol can be omitted if there is an `InitialAssignment` object for that symbol; the symbols must always be defined even if they are assigned a value separately. For example, there must be a `Parameter` definition for a given parameter if there is an `InitialAssignment` for that parameter.

The actions of all `InitialAssignment` objects are in general terms the same, but differ in the precise details depending on the type of variable being set:

- *In the case of a species*, an `InitialAssignment` sets the referenced species' initial quantity (*concentration* or *amount*) to the value determined by the formula in `math`. (See Section 4.6.5 for an explanation of how the units of the species' quantity are determined.)
- *In the case of a species reference*, an `InitialAssignment` sets the initial stoichiometry of the referenced `SpeciesReference` to the value determined by the formula in `math`. The overall units of the formula in `math` should be **dimensionless**, because reactant and product stoichiometries in reactions are dimensionless quantities.
- *In the case of a compartment*, an `InitialAssignment` sets the referenced compartment's initial size to the size determined by the formula in `math`. The overall units of the formula should be the same as the units specified for the size of the compartment. (See Section 4.5.4 for an explanation of how the units of the compartment's size are determined.)
- *In the case of a parameter*, an `InitialAssignment` sets the referenced parameter's initial value to that determined by the formula in `math`. The overall units of the formula should be the same as the units defined for the parameter. (See Section 4.7.3 for an explanation of how the units of the parameter are determined.)

In the context of a simulation, initial assignments establish values that are in effect prior to and including the start of simulation time, i.e., $t \leq 0$. Section 3.4.8 provides information about the interpretation of assignments, rules, and entity values for simulation time up to and including the start time $t = 0$; this is important for establishing the initial conditions of a simulation if the model involves expressions containing the `delay` `csymbol` (Section 3.4.6).

There cannot be two initial assignments for the same symbol in a model; that is, a model must not contain two or more `InitialAssignment` objects that both have the same identifier as their `symbol` attribute value. A model must also not define initial assignments *and* assignment rules for the same entity. That is, there cannot be *both* an `InitialAssignment` and an `AssignmentRule` for the same symbol in a model, because both kinds of constructs apply prior to and at the start of simulated time—allowing both to exist for a given symbol would result in indeterminism). (See also Section 4.9.5.)

The ordering of `InitialAssignment` objects is not significant. The combined set of `InitialAssignment`, `AssignmentRule` and `KineticLaw` objects form a set of assignment statements that must be considered as a whole. The combined set of assignment statements should not contain algebraic loops: a chain of dependency between these statements should terminate. (More formally, consider the directed graph of assignment statements where nodes are a model's assignment statements and directed arcs exist for each occurrence of a symbol in an assignment statement `math` attribute. The directed arcs in this graph start from the statement assigning the symbol and end at the statement that contains the symbol in their `math` elements. Such a graph must be acyclic.) Examples of valid and invalid set of assignment statements are given in Section 4.9.5.

Finally, it is worth being explicit about the expected behavior in the following situation. Suppose (1) a given symbol has a value x assigned to it in its definition, and (2) there is an initial assignment having the identifier as its **symbol** value and reassigning the value to y , and (3) the identifier is also used in the mathematical formula of a second initial assignment. What value should the second initial assignment use? It is y , the value assigned to the symbol by the first initial assignment, not whatever value was given in the symbol's definition. This follows directly from the behavior at the defined at the beginning of this section and in Section 3.4.8: if an **InitialAssignment** object exists for a given symbol, then the symbol's value is overridden by that initial assignment.

4.8.5 Example

The following example shows how the species “ x ” can assigned the initial value $2 \cdot y$, where “ y ” is an identifier defined elsewhere in the model:

```
<listOfSpecies>
  <species id="x" compartment="C" units="mole"
    hasOnlySubstanceUnits="true" boundaryCondition="false" constant="false"/>
</listOfSpecies>
...
<listOfInitialAssignments>
  <initialAssignment symbol="x">
    <math xmlns="http://www.w3.org/1998/Math/MathML"
      xmlns:sbml="http://www.sbml.org/sbml/level3/version1/core">
      <apply>
        <times/>
        <ci> y </ci>
        <cn sbml:units="dimensionless"> 2 </cn>
      </apply>
    </math>
  </initialAssignment>
</listOfInitialAssignments>
```

The next example illustrates the more complex behavior discussed above, when a symbol has a value assigned in its definition but there also exists an **InitialAssignment** for it and another **InitialAssignment** uses its value in its mathematical formula.

```
<listOfSpecies>
  <species id="x" initialAmount="50" compartment="C" units="item"
    hasOnlySubstanceUnits="true" boundaryCondition="false" constant="false"/>
</listOfSpecies>
...
<listOfInitialAssignments>
  <initialAssignment symbol="x">
    <math xmlns="http://www.w3.org/1998/Math/MathML"
      xmlns:sbml="http://www.sbml.org/sbml/level3/version1/core">
      <cn sbml:units="item"> 10 </cn>
    </math>
  </initialAssignment>
  <initialAssignment symbol="othersymbol">
    <math xmlns="http://www.w3.org/1998/Math/MathML"
      xmlns:sbml="http://www.sbml.org/sbml/level3/version1/core">
      <apply>
        <times/>
        <ci> x </ci>
        <cn sbml:units="dimensionless"> 2 </cn>
      </apply>
    </math>
  </initialAssignment>
</listOfInitialAssignments>
```

The value of “**othersymbol**” in the SBML excerpt above will be “20”. The case illustrates the rule of thumb that if there is an initial assignment for a symbol, the value assigned to the symbol in its definition (here, the value of **initialAmount**) must be ignored and the value created by the initial assignment used instead.

4.9 Rules

In SBML, *Rules* provide additional ways to define the values of variables in a model, their relationships, and the dynamical behaviors of those variables. *Rules* enable the encoding of relationships that cannot be expressed using reactions alone (Section 4.11) nor by the assignment of an initial value to a variable in a model (Section 4.8).

SBML separates rules into three subclasses for the benefit of model analysis software. The three subclasses are based on the following three different possible functional forms (where x is a variable, f is some arbitrary function returning a numerical result, \mathbf{V} is a vector of variables that does not include x , and \mathbf{W} is a vector of variables that may include x):

<i>Algebraic</i>	left-hand side is zero:	$0 = f(\mathbf{W})$
<i>Assignment</i>	left-hand side is a scalar:	$x = f(\mathbf{V})$
<i>Rate</i>	left-hand side is a rate-of-change:	$dx/dt = f(\mathbf{W})$

In their general form given above, there is little to distinguish between *assignment* and *algebraic* rules. They are treated as separate cases for the following reasons:

- *Assignment* rules can simply be evaluated to calculate intermediate values for use in numerical methods;
- SBML needs to place restrictions on assignment rules, for example the restriction that assignment rules cannot contain algebraic loops (discussed further in Section 4.9.5);
- Some simulators do not contain numerical solvers capable of solving unconstrained algebraic equations, and providing more direct forms such as assignment rules may enable those simulators to process models they could not process if the same assignments were put in the form of general algebraic equations;
- Those simulators that *can* solve these algebraic equations make a distinction between the different categories listed above; and
- Some specialized numerical analyses of models may only be applicable to models that do not contain *algebraic* rules.

The approach taken to covering these cases in SBML is to define an abstract *Rule* class containing an element, `math`, to hold the right-hand side expression, then to derive subclasses of *Rule* that add attributes to distinguish the cases of algebraic, assignment and rate rules. Figure 17 gives the definitions of *Rule* and the subclasses derived from it. The figure shows there are three subclasses, *AlgebraicRule*, *AssignmentRule* and *RateRule* derived directly from *Rule*. These correspond to the cases *Algebraic*, *Assignment*, and *Rate* described above, respectively.




Figure 17: The definition of *Rule* and derived types *AlgebraicRule*, *AssignmentRule* and *RateRule*.

4.9.1 Common attributes in Rule

The classes derived from [Rule](#) inherit **math** and the attributes and elements from [SBase](#), including **sboTerm**.

The math element

A [Rule](#) object has a required element called **math**, containing a MathML expression defining the mathematical formula of the rule. This MathML formula must return a numerical value. The formula can be an arbitrary expression referencing the variables and other entities in an SBML model. The interpretation of **math** and the units of the formula are described in more detail in Sections [4.9.2](#), [4.9.3](#) and [4.9.4](#) below.

The sboTerm attribute

[Rule](#) inherits an optional **sboTerm** attribute of type **SBOTerm** from its parent class [SBase](#) (see Sections [3.1.11](#) and [5](#)). When a value is given to this attribute in a [AlgebraicRule](#), [AssignmentRule](#), or [RateRule](#) instance, it should be an SBO identifier belonging to the branch for type [AlgebraicRule](#), [AssignmentRule](#), or [RateRule](#) indicated in Table [5](#). The relationship is of the form “the rule *is a* X”, where X is the SBO term. The term chosen should be the most precise (narrow) one that captures the role of the rule in the model.

As discussed in Section [5](#), SBO labels are optional information on a model. Applications are free to ignore **sboTerm** values. A model must be interpretable without the benefit of SBO labels.

4.9.2 AlgebraicRule

The rule type [AlgebraicRule](#) is used to express equations that are neither assignments of model variables nor rates of change. [AlgebraicRule](#) does not add any attributes to the basic [Rule](#); its role is simply to distinguish this case from the other cases. An example of the use of [AlgebraicRule](#) is given in Section [7.4](#).

In the context of a simulation, algebraic rules are in effect at all times, $t \geq 0$. For purposes of evaluating expressions that involve the *delay* **csymbol** (Section [3.4.6](#)), algebraic rules are considered to apply also at $t \leq 0$. Section [3.4.8](#) provides additional information about the semantics of assignments, rules, and entity values for simulation time $t \leq 0$.

An SBML model must not be overdetermined. The ability to define arbitrary algebraic expressions in an SBML model introduces the possibility that a model is mathematically overdetermined by the overall system of equations constructed from its rules, reactions and events. Therefore, if an algebraic rule is introduced in a model, for at least one of the entities referenced in the rule’s **math** element the value of that entity must not be completely determined by other constructs in the model. This means that at least this entity must not have the attribute **constant**=“true” and there must also not be a rate rule or assignment rule for it. Furthermore, if the entity is a [Species](#) object, its value must not be determined by reactions, which means that it must either have the attribute **boundaryCondition**=“true” or else not be involved in any reaction at all. These restrictions are explained in more detail in Section [4.9.5](#) below.

4.9.3 AssignmentRule

The rule type [AssignmentRule](#) is used to express equations that set the values of variables. The left-hand side (the **variable** attribute) of an assignment rule can refer to the identifier of a [Species](#), [SpeciesReference](#), [Compartment](#), or [Parameter](#) object in the model (but not a reaction). The entity identified must not have its **constant** attribute set to “true”. The effects of an [AssignmentRule](#) are in general terms the same, but differ in the precise details depending on the type of variable being set:

- In the case of a species, an [AssignmentRule](#) sets the referenced species’ quantity (*concentration* or *amount*) to the value determined by the formula in **math**. The units of the formula should match the *units of the species* (Section [4.6.5](#)) for the species identified by the **variable** attribute.

Restrictions: There must not be both an [AssignmentRule](#) **variable** attribute and a [SpeciesReference](#) **species** attribute having the same value, unless that species has its **boundaryCondition** attribute set to “true”. In other words, an assignment rule cannot be defined for a species that is created or destroyed in a reaction unless that species is defined as a boundary condition in the model.

- In the case of a species reference, an **AssignmentRule** sets the stoichiometry of the reactant or product referenced by the **SpeciesReference** object instance to the value determined by the formula in **math**. The overall units of the formula in **math** should be **dimensionless**, because reactant and product stoichiometries in reactions are dimensionless quantities.
- In the case of a compartment, an **AssignmentRule** sets the referenced compartment's size to the value determined by the formula in **math**. The overall units of the formula in **math** should be the same as the units of the size of the compartment (Section 4.5.4).
- In the case of a parameter, an **AssignmentRule** sets the referenced parameter's value to that determined by the formula in **math**. The overall units of the formula in **math** should be the same as the units defined for the parameter (Section 4.7.3).

In the context of a simulation, assignment rules are in effect at all times, $t \geq 0$. For purposes of evaluating expressions that involve the *delay* **csymbol** (Section 3.4.6), assignment rules are considered to apply also at $t \leq 0$. Section 3.4.8 provides additional information about the semantics of assignments, rules, and entity values for simulation time $t \leq 0$.

A model must not contain more than one **AssignmentRule** or **RateRule** object having the same value of **variable**; in other words, in the set of all assignment rules and rate rules in an SBML model, each variable appearing in the left-hand sides can only appear once. This simply follows from the fact that an indeterminate system would result if a model contained more than one assignment rule for the same variable or both an assignment rule and a rate rule for the same variable.

Similarly, a model must also not contain *both* an **AssignmentRule** and an **InitialAssignment** for the same variable, because both kinds of constructs apply prior to and at the start of simulation time, i.e., $t \leq 0$. If a model contained both an initial assignment and an assignment rule for the same variable, an indeterminate system would result. (See also Section 4.8.4.)

The value calculated by an **AssignmentRule** object overrides the value assigned to the given symbol by the object defining that symbol. For example, if a **Compartment's** **size** is set in its definition, and the model also contains an **AssignmentRule** having that compartment's **id** as its **variable** value, then the **size** assigned in the **Compartment** definition is ignored and the value assigned based on the computation defined in the **AssignmentRule**. This does *not* mean that a definition for a given symbol can be omitted if there is an **AssignmentRule** object for it. For example, there must be a **Parameter** definition for a given parameter if there is an **AssignmentRule** for that parameter.

4.9.4 RateRule

The rule type **RateRule** is used to express equations that determine the rates of change of variables. The left-hand side (the **variable** attribute) can refer to the identifier of a species, species reference, compartment, or parameter (but not a reaction). The entity identified must have its **constant** attribute set to "**false**". The effects of a **RateRule** are in general terms the same, but differ in the precise details depending on which variable is being set:

- In the case of a species, a **RateRule** sets the rate of change of the species' quantity (*concentration* or *amount*) to the value determined by the formula in **math**. The overall units of the formula in **math** should be *species quantity* divided by *time*, where the *time* units are the model-wide units of time described in Section 4.2.4 and the *species quantity* units are the *units of the species* as defined in Section 4.6.5.
Restrictions: There must not be both a **RateRule** **variable** attribute and a **SpeciesReference** **species** attribute having the same value, unless that species has its **boundaryCondition** attribute is set to "**true**". This means a rate rule cannot be defined for a species that is created or destroyed in a reaction, unless that species is defined as a boundary condition in the model.
- In the case of a species reference, a **RateRule** sets the rate of change of the stoichiometry of the reactant or product referenced by the **SpeciesReference** object instance to the value determined by the formula in **math**. The overall units of the formula in **math** should be **dimensionless** divided by *time* units,

where the *time* units are the model-wide units of time described in Section 4.2.4, because reactant and product stoichiometries in reactions are dimensionless quantities.

- In the case of a compartment, a **RateRule** sets the rate of change of the compartment's size to the value determined by the formula in **math**. The overall units of the formula should be *size* units divided by *time* units, where the *time* units are the model-wide units of time described in Section 4.2.4 and the *size* units are the units of size on the compartment (Section 4.5.4).
- In the case of a parameter, a **RateRule** sets the rate of change of the parameter's value to that determined by the formula in **math**. The overall units of the formula should be $x/time$, where x are the units of the parameter (Section 4.7.3) and the *time* units are the model-wide units of time described in Section 4.2.4.

In the context of a simulation, rate rules are in effect for simulation time $t > 0$. Other types of rules and initial assignments are in effect at different times; Section 3.4.8 describes these conditions.

As mentioned in Section 4.9.3 for **AssignmentRule**, a model must not contain more than one **RateRule** or **AssignmentRule** object having the same value of **variable**; in other words, in the set of all assignment rules and rate rules in an SBML model, each variable appearing in the left-hand sides can only appear once. This simply follows from the fact that an indeterminate system would result if a model contained more than one assignment rule for the same variable or both an assignment rule and a rate rule for the same variable.

4.9.5 Additional restrictions on rules

An important design goal of SBML rule semantics is to ensure that a model's simulation and analysis results will not be dependent on when or how often rules are evaluated. To achieve this, SBML needs to place two additional restrictions on rule use in addition to the conditions described above regarding the use of **AlgebraicRule**, **AssignmentRule** and **RateRule**. The first concerns algebraic loops in the system of assignments in a model, and the second concerns overdetermined systems.

The model must not contain algebraic loops

The combined set of **InitialAssignment**, **AssignmentRule** and **KineticLaw** objects constitute a set of assignment statements that should be considered as a whole. (A **KineticLaw** object is counted as an assignment because it assigns a value to the symbol contained in the **id** attribute of the **Reaction** object in which it is defined.) This combined set of assignment statements must not contain algebraic loops—dependency chains between these statements must terminate. To put this more formally, consider a directed graph in which nodes are assignment statements and directed arcs exist for each occurrence of an SBML species, species reference, compartment or parameter symbol in an assignment statement's **math** element. Let the directed arcs point from the statement assigning the symbol to the statements that contain the symbol in their **math** element expressions. This graph must be acyclic.

SBML does not specify when or how often rules should be evaluated. Eliminating algebraic loops ensures that assignment statements can be evaluated any number of times without the result of those evaluations changing. As an example, consider the following equations:

$$x = x + 1, \quad y = z + 200, \quad z = y + 100$$

If this set of equations were interpreted as a set of assignment statements, it would be invalid because the rule for x refers to x (exhibiting one type of loop), and the rule for y refers to z while the rule for z refers back to y (exhibiting another type of loop).

Conversely, the following set of equations would constitute a valid set of assignment statements:

$$x = 10, \quad y = z + 200, \quad z = x + 100$$

The model must not be overdetermined

An SBML model must not be overdetermined; that is, a model must not define more equations than there are unknowns in a model. An SBML model without **AlgebraicRule** objects cannot be overdetermined.

Assessing whether a given continuous, deterministic, mathematical model is overdetermined does not require dynamic analysis; it can be done by analyzing the system of equations created from the model. It should be noted that where a model contains both reactions and events there are several sets of equations to consider when determining whether a model is overdetermined. The set of equations derived from the combined set of rules and reactions and, for each event, the set of equations derived from the combined set of rules and event assignments for the particular event.

One approach is to construct a bipartite graph in which one set of vertices represents the variables and the other the set of vertices represents the equations. Place edges between vertices such that variables in the system are linked to the equations that determine them. A mathematical model is overdetermined if the maximal matchings (Chartrand, 1977) of the bipartite graph contain disconnected vertexes representing equations. (If one maximal matching has this property, then all the maximal matchings will have this property; i.e., it is only necessary to find one maximal matching.) Appendix D describes a method of applying this procedure to specific SBML data objects. In some cases it is possible to avoid the use of an AlgebraicRule. This is discussed in more detail in Section 8.2.3.

4.9.6 Example of rule use

This section contains an example set of rules. Consider the following set of equations:

$$k = \frac{k_3}{k_2}, \quad s_2 = \frac{k \cdot x}{1 + k_2}, \quad A = 0.10 \cdot x$$

This can be encoded by the following scalar rule set (where the definitions of x , s , k , k_2 , k_3 and A are assumed to be located elsewhere in the model and not shown in this abbreviated example):

```
<listOfRules>
  <assignmentRule variable="k">
    <math xmlns="http://www.w3.org/1998/Math/MathML">
      <apply> <divide/> <ci> k3 </ci> <ci> k2 </ci> </apply>
    </math>
  </assignmentRule>
  <assignmentRule variable="s2">
    <math xmlns="http://www.w3.org/1998/Math/MathML">
      <apply>
        <divide/>
        <apply> <times/> <ci> k </ci> <ci> x </ci> </apply>
        <apply> <plus/> <cn> 1 </cn> <ci> k2 </ci> </apply>
      </apply>
    </math>
  </assignmentRule>
  <assignmentRule variable="A">
    <math xmlns="http://www.w3.org/1998/Math/MathML">
      <apply> <times/> <cn> 0.10 </cn> <ci> x </ci> </apply>
    </math>
  </assignmentRule>
</listOfRules>
```

4.10 Constraints

The **Constraint** object is a mechanism for stating the assumptions under which a model is designed to operate. The *constraints* are statements about permissible values of different quantities in a model. Figure 18 shows the definition of the **Constraint** object class.

The essential meaning of a constraint is this: if a dynamical analysis of a model (such as a simulation) reaches a state in which a constraint is no longer satisfied, the results of the analysis are deemed invalid beginning with that point in time. The exact behavior of a software tool, upon encountering a constraint violation, is left up to the software; *however*, a software tool must somehow indicate to the user when a model's constraints are no longer satisfied. (Otherwise, a user may not realize that the analysis has reached an invalid state and is potentially producing nonsense results.) If a software tool does not have support for constraints, it should indicate this to the user when encountering a model containing constraints.

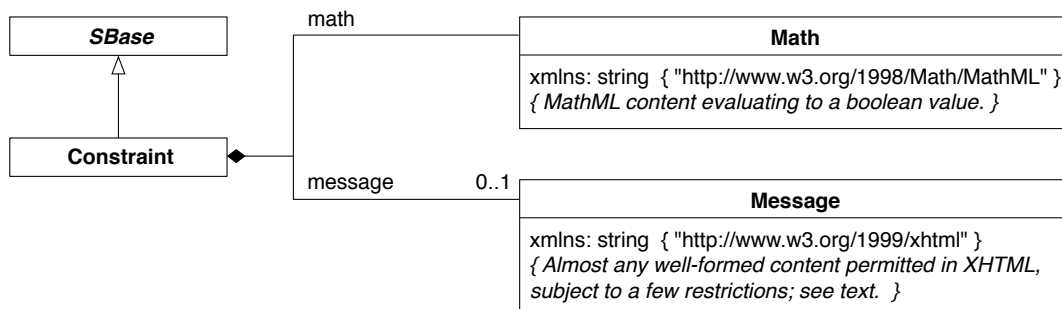


Figure 18: The definition of class **Constraint**. The contents of the **Math** class can be any MathML permitted in SBML, but it must return a boolean value. As shown above, an instance of **Constraint** can also contain zero or one instances of **Message**; this element is simply a wrapper (in the XML form, within `<message> ... </message>` tags) for XHTML content. The same guidelines for XHTML content as explained in Section 3.2.3 for notes on **SBase** also apply to the XHTML within messages in a **Constraint**. A sequence of one or more instances of **Constraint** objects can be located in an instance of **ListOfConstraints** in **Model**, as shown in Figure 10.

4.10.1 The math element

Constraint has one required subelement, **math**, containing a MathML formula defining the condition of the constraint. This formula must return a boolean value of “true” when the model is in a *valid* state. The formula can be an arbitrary expression referencing the variables and other entities in an SBML model. The evaluation of **math** and behavior of constraints are described in more detail in Section 4.10.4 below.

4.10.2 The message element

A **Constraint** object has an optional element called **message**. This can contain a message in XHTML format that may be displayed to the user when the condition of the constraint in **math** evaluates to a value of “false”. Software tools are not required to display the message, but it is recommended that they do so as a matter of best practice.

The XHTML content within a **message** element must follow the same restrictions as for the **notes** element on **SBase** described in Section 3.2.3. For example, **message** must not contain an XML declaration or a DOCTYPE declaration, and the permitted content can only take one of the following general forms: (1) a complete XHTML document beginning with the element `<html>` and ending with `</html>`; (2) the “body” portion of a document beginning with the element `<body>` and ending with `</body>`; or (3) XHTML content that is permitted within a `<body> ... </body>` elements. Appendix F describes one approach to reading the **message** content.

4.10.3 The sboTerm attribute

Constraint inherits an optional **sboTerm** attribute of type **SBOTerm** from its parent class **SBase** (see Sections 3.1.11 and 5). When a value is given to this attribute in a **Constraint** instance, it should be an SBO identifier belonging to the branch for type **Constraint** indicated in Table 5. The relationship is of the form “the constraint *is a* X”, where X is the SBO term. The term chosen should be the most precise (narrow) one that captures the role of the constraint in the model.

As discussed in Section 5, SBO labels are optional information on a model. Applications are free to ignore **sboTerm** values. A model must be interpretable without the benefit of SBO labels.

4.10.4 Semantics of constraints

In the context of a simulation, a **Constraint** has effect at all times $t \geq 0$. Each **Constraint**’s **math** element is first evaluated after any **InitialAssignment** definitions in a model at $t = 0$ and can conceivably trigger at that point. (In other words, a simulation could fail a constraint immediately.)

Constraint definitions *cannot and should not* be used to compute the dynamical behavior of a model as part of, for example, simulation. Constraints may be used as input to non-dynamical analysis, for instance by expressing flux constraints for flux balance analysis.

The results of a simulation of a model containing a constraint are invalid from any simulation time at and after a point when the function given by the **math** returns a value of “false”. Invalid simulation results do not make a prediction of the behavior of the biochemical reaction network represented by the model. The precise behavior of simulation tools is left undefined with respect to constraints. If invalid results are detected with respect to a given constraint, the **message** element (Section 4.10.2) may optionally be displayed to the user. The simulation tool may also halt the simulation or clearly delimit in output data the simulation time point at which the simulation results become invalid.

SBML does not impose restrictions on duplicate **Constraint** definitions or the order of evaluation of **Constraint** objects in a model. It is possible for a model to define multiple constraints all with the same **math** element. Since the failure of any constraint indicates that the model simulation has entered an invalid state, a system is not required to attempt to detect whether other constraints in the model have failed once any one constraint has failed.

4.10.5 Example

As an example, the following SBML fragment demonstrates the constraint that species “S1” should only have values between 1 and 100:

```
<model ...>
  ...
  <listOfConstraints>
    <constraint>
      <math xmlns="http://www.w3.org/1998/Math/MathML"
            xmlns:sbml="http://www.sbml.org/sbml/level3/version1/core">
        <apply>
          <and/>
            <apply>
              <lt/>
              <cn sbml:units="mole"> 1 </cn>
              <ci> S1 </ci>
            </apply>
            <apply>
              <lt/>
              <ci> S1 </ci>
              <cn sbml:units="mole"> 100 </cn>
            </apply>
          </and>
        </apply>
      </math>
      <message>
        <p xmlns="http://www.w3.org/1999/xhtml">
          Species S1 is out of range.
        </p>
      </message>
    </constraint>
  </listOfConstraints>
  ...
</model>
```

4.11 Reactions

A *reaction* in SBML represents any kind of process that can change the quantity of one or more species in a model. Examples of such processes can include transformation, transport, molecular interactions, and more. In SBML, the notion of a reaction is generalized to allow entities that may not be chemical substances; thus, a reaction in SBML does not necessarily have to be a biochemical reaction—a biochemical reaction is just one possible kind of process.

At minimum, to describe a reaction in SBML, it is necessary to define its *structural* properties, specifically the participating reactants and/or products (and their corresponding stoichiometries) and the reversibility

of the process. In addition, an SBML reaction can also contain a *quantitative* description of the rate of the reaction; this aspect consists of a mathematical formula expressing describing the rate at which the reaction process takes place, together with an optional list of modifier species and parameters influencing the reaction rate. The various parts of a reaction are recorded in the SBML **Reaction** object class and other supporting data classes, defined in Figure 19.

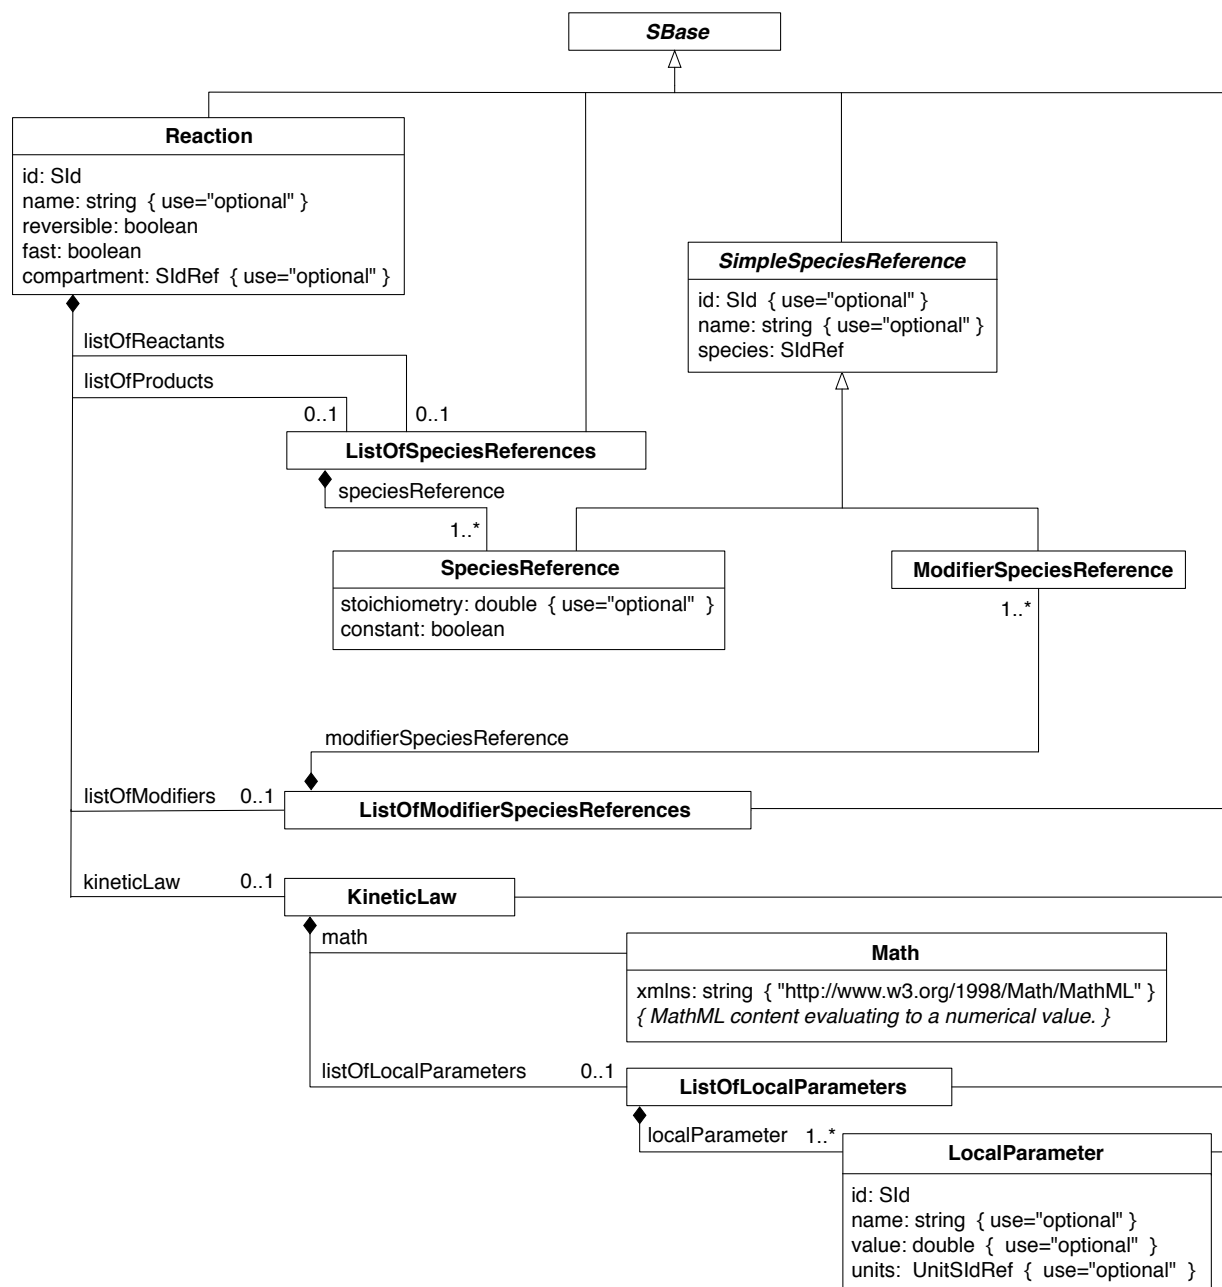


Figure 19: The definitions of **Reaction**, **KineticLaw**, **SpeciesReference**, **ModifierSpeciesReference**, **LocalParameter**, as well as the container classes **ListOfReactants**, **ListOfProducts**, **ListOfModifiers**, and **ListOfLocalParameters**. Note that **SimpleSpeciesReference** is an abstract class used only to provide some common attributes to its derived classes.

4.11.1 Reaction

Each reaction in an SBML model is defined using an instance of a **Reaction** object. As shown in Figure 19 on the preceding page, it contains several scalar attributes and several lists of other objects.

The id and name attributes

As with most other main kinds of objects in SBML, the **Reaction** object class includes a mandatory attribute called **id**, of type **SId**, and an optional attribute **name**, of type **string**. The **id** attribute is used to give the reaction a unique identifier in the model. This identifier can be used in mathematical formulas elsewhere in an SBML model to represent the rate of that reaction; this usage is explained in detail in Section 4.11.8 below. The **name** attribute can be used to give the reaction a more free-form, descriptive name. The **name** and **id** attributes must be used as described in Section 3.3.

The lists of reactants, products and modifiers

The species participating as reactants, products, and/or modifiers in a reaction are declared using lists of **SpeciesReference** and/or **ModifierSpeciesReference** instances stored in **listOfReactants**, **listOfProducts** and **listOfModifiers**. **SpeciesReference** and **ModifierSpeciesReference** are described in more detail in Sections 4.11.3 and 4.11.4 below. Certain restrictions are placed on the appearance of species in reaction definitions:

- The ability of a species to appear as a reactant or product of any reaction in a model is governed by certain combinations of the attributes **constant** and **boundaryCondition** on the **Species** object instance; see Section 4.6.6 for more information.
- Any species appearing in the **math** element of the **kineticLaw** of a **Reaction** instance must be declared in at least one of that **Reaction**'s lists of reactants, products, and/or modifiers. It is an error for a reaction's kinetic law formula to refer to species that have not been declared for that reaction.
- A reaction definition can contain an empty list of reactants *or* an empty list of products, but it must have at least one reactant or product; in other words, a reaction without any reactant or product species is not permitted. (This restriction does not apply to modifier species, which are always optional.)

The kineticLaw element

A reaction can contain up to one **KineticLaw** object in the **kineticLaw** element of the **Reaction**. This “kinetic law” defines the speed at which the process defined by the reaction takes place. A detailed description of **KineticLaw** is left to Section 4.11.6 below.

The inclusion of a **KineticLaw** object in an instance of a **Reaction** is optional. For some modeling purposes, models containing reactions without defined rates is an acceptable alternative (and may even be the only possible option, such as when the kinetics of the reactions are unknown). However, missing kinetic laws preclude the application of many model analysis techniques, including simulation.

The reversible attribute

The mandatory boolean attribute **reversible** indicates whether the reaction is reversible. To say that a reaction is *reversible* is to say it can proceed in either the forward or the reverse direction. This information may be redundant in cases where the reversibility of the reaction can be deduced by inspecting the rate formula given in the kinetic law. However, a reaction is not required to have a kinetic law, and besides, when a rate expression is present, it may not always be possible to deduce the reversibility by inspecting it. Having a separate attribute for **reversible** allows certain kinds of structural analysis, such as elementary mode analysis, even in these cases.

Mathematically, the **reversible** attribute on **Reaction** has no impact on the construction of the equations for change of the species quantities. However, labeling a reaction as irreversible is interpreted as an assertion that the rate expression will not have negative values during a simulation. Software developers may wish to provide their software systems with a means of testing that this condition holds.

The presence of reversibility information in two places (i.e., the rate expression in the kinetic law, and the **reversible** flag) leaves open the possibility that a model could contain contradictory information, but this would be considered to be an error of the encoded model, rather than an invalid SBML encoding.

The fast attribute

The boolean attribute **fast** is another required boolean attribute of **Reaction**. When a model contains a value of “**true**” for **fast** on any of its reactions, it indicates that the creator of the model is distinguishing different time scales of reactions in the system. If a model does not distinguish between time scales, the **fast** attribute should be set to “**false**” for all reactions.

The model’s reaction definitions are divided into two sets by the values of the **fast** attributes. The set of reactions having **fast**=“**true**” (known as *fast reactions*) should be assumed to be operating on a time scale significantly faster than the other reactions (the *slow reactions*). Fast reactions are considered to be instantaneous relative to the slow reactions. Software tools should use a pseudo steady-state approximation for the set of fast reactions when constructing the system of equations for the model. More specifically, the set of reactions that have the **fast** attribute set to “**true**” forms a subsystem that should be described by a pseudo steady-state approximation in relationship to all other reactions in the model. Under this description, relaxation from any initial condition or perturbation from any intermediate state of this subsystem would be infinitely fast. Appendix E provides a technical explanation of an approach to solving systems with fast reactions.

The correctness of the approximation requires a significant separation of time scales between the fast reactions and other processes. It is the responsibility of the modeler or of the software system writing the SBML model to ensure this condition is fulfilled.

Note that the **fast** attribute has a significant effect on the mathematical interpretation of a model and cannot be safely ignored if a software tool does not implement support for the corresponding concept. Software systems should indicate to users when they encounter models with reactions having **fast**=“**true**” and do not have the capacity to analyze the model using a pseudo steady-state approximation.

The compartment attribute on Reaction

The optional attribute **compartment**, of data type **SIdRef**, can be used to indicate the compartment in which the reaction is assumed to take place. If the attribute is present, its value must be the identifier of a **Compartment** object defined in the enclosing **Model** object.

Similar to the **reversible** attribute, the value of the **compartment** attribute has no direct impact on the construction of mathematical equations for the SBML model. When a kinetic law is given for a reaction, the compartment location may already be implicit in the kinetic law (although this cannot always be guaranteed). Nevertheless, software tools may find the **compartment** attribute value useful for such purposes as analyzing the structure of the model, guiding the modeler in constructing correct rate formulas, and visualization purposes.

The sboTerm attribute on Reaction

Reaction inherits an optional **sboTerm** attribute of type **SBOTerm** from its parent class **SBase** (see Sections 3.1.11 and 5). When a value is given to this attribute in a **Reaction** instance, it should be an SBO identifier belonging to the branch for type **Reaction** indicated in Table 5. The relationship is of the form “the reaction *is a* X”, where X is the SBO term. The term chosen should be the most precise (narrow) one that captures the role of the reaction in the model.

As discussed in Section 5, SBO labels are optional information on a model. Applications are free to ignore **sboTerm** values. A model must be interpretable without the benefit of SBO labels.

4.11.2 SimpleSpeciesReference

As mentioned above, every species that enters into a given reaction must appear in that reaction's lists of reactants, products and/or modifiers. In an SBML model, all species that may participate in any reaction are listed in the `listOfSpecies` element of the top-level **Model** instance (see Section 4.2). Lists of products, reactants and modifiers in **Reaction** objects do not introduce new species, but rather, they refer back to those listed in the model's top-level `listOfSpecies`. For reactants and products, the connection is made using a **SpeciesReference** object; for modifiers, it is made using a **ModifierSpeciesReference** object. **SimpleSpeciesReference**, defined in Figure 19 on page 60, is an abstract type that serves as the parent class of both **SpeciesReference** and **ModifierSpeciesReference**. It is used simply to hold the attributes and elements that are common to the latter two objects.

The id and name attributes

SimpleSpeciesReference has optional attributes for an identifier (`id`, of data type `SId`) and name (`name`, of data type `string`). The `id` and `name` attributes must be used as described in Section 3.3.

The `id` value (whether it is in a **SpeciesReference** or **ModifierSpeciesReference** object) exists in the global namespace of the model. In SBML Level 3 Version 1 Core, the only meaning defined for the use of such identifiers concerns the `id` of a **SpeciesReference** object (discussed further in Section 4.11.3); no meaning or value is associated with the identifiers of **ModifierSpeciesReference**. However, the identifiers are present on both object classes for possible use by SBML Level 3 packages.

The species attribute

The **SimpleSpeciesReference** object class has a required attribute, `species`, of data type `SIdRef`. As with the other attributes, it is inherited by **SpeciesReference** and **ModifierSpeciesReference**. The value of `species` must be the identifier of a species defined in the enclosing **Model**; the referenced species is thereby declared as participating in the reaction being defined. The precise role of that species as a reactant, product, or modifier in the reaction is determined by the subtype of **SimpleSpeciesReference** (i.e., either **SpeciesReference** or **ModifierSpeciesReference**) in which the identifier appears and by the specific list of species references in which the **SpeciesReference** appears.

The sboTerm attribute

SimpleSpeciesReference inherits an optional `sboTerm` attribute of type `SBOTerm` from its parent class **SBase** (see Sections 3.1.11 and 5). When a value is given to this attribute in a **SimpleSpeciesReference** instance, it should be an SBO identifier belonging to the branch for type **SimpleSpeciesReference** indicated in Table 5. The relationship is of the form “the species reference *is a* X”, where X is the SBO term. The term chosen should be the most precise (narrow) one that captures the role of the species reference in the model.

As discussed in Section 5, SBO labels are optional information on a model. Applications are free to ignore `sboTerm` values. A model must be interpretable without the benefit of SBO labels.

4.11.3 SpeciesReference

Reaction provides a way to express which species act as reactants and which species act as products in a reaction, and to declare their stoichiometries. This is done using **SpeciesReference** objects. As mentioned above in Section 4.11.2, **SpeciesReference** inherits the mandatory attribute `species` and optional attributes `id`, `name`, and `sboTerm` from the parent type **SimpleSpeciesReference**. It also defines the optional attribute `stoichiometry` and the mandatory attribute `constant`, described below.

The stoichiometry attribute

The *stoichiometry* of a species in a reaction describes how much of the species changes when a reaction event takes place. In SBML, product and reactant stoichiometries are specified using the optional `stoichiometry` on **SpeciesReference** object. The `stoichiometry` attribute is of type `double` and should contain values greater than zero (0). A missing `stoichiometry` implies that the stoichiometry is either unknown, or to

be obtained from an external source, or determined by an initial assignment (Section 4.8) or other SBML construct elsewhere in the model.

A species reference's stoichiometry is set by its **stoichiometry** attribute exactly once. If the **SpeciesReference** object's **constant** attribute (see below) has the value "true", then the stoichiometry is fixed and cannot be changed except by an **InitialAssignment**. These two methods of setting the stoichiometry (i.e., using **stoichiometry** directly, or using an **InitialAssignment**) differ in that the **stoichiometry** attribute can only be set to a literal scalar value, whereas **InitialAssignment** allows the stoichiometry to be set using an arbitrary mathematical expression and to values that are not expressible in the data type **double**. (As an example, the approach could be used to set the stoichiometry to a rational number, something that is occasionally useful in the context of biochemical reaction networks.) If the species reference's **constant** attribute has the value "false", the species reference's value may be overridden by an **InitialAssignment** or changed by a **AssignmentRule** or **AlgebraicRule**, and in addition, for simulation time $t > 0$, it may also be changed by a **RateRule** or **Events**. (However, some of these constructs are mutually exclusive; see Sections 4.9 and 4.12.) It is not an error to define **stoichiometry** on a species reference and also redefine the stoichiometry using an **InitialAssignment**, but the **stoichiometry** attribute in that case is ignored. Section 3.4.8 provides additional information about the semantics of assignments, rules and values for simulation time $t \leq 0$.

An explanation of how exactly the stoichiometry is used in the mathematical interpretation of the model is given in Section 4.11.7.

The constant attribute

The **SpeciesReference** attribute **constant** is a mandatory boolean attribute used to indicate whether the **stoichiometry** value can vary during a simulation. If **constant**="true", the corresponding species' stoichiometry in the reaction cannot be changed by other constructs elsewhere in the model except by an **InitialAssignment**. A value of "false" means the stoichiometry can be changed by other SBML constructs such as rules (see Section 4.9), as described above in the section on the **stoichiometry** attribute.

Use of species reference identifiers in mathematical expressions

The value of the **id** attribute of a **SpeciesReference** can be used as the content of a **ci** element in MathML formulas elsewhere in the model. When the identifier appears in a **ci** element, it represents the stoichiometry of the corresponding species in the reaction where the **SpeciesReference** object instance appears. More specifically, it represents the value of the **stoichiometry** attribute on the **SpeciesReference** object.

*The units of a **SpeciesReference**'s stoichiometry value*

The units associated with the stoichiometry of a species reference are always considered to be **dimensionless**. This has the following implications:

- When a species reference's identifier appears in mathematical formulas elsewhere in the model, the units associated the value are **dimensionless**.
- The units of the **math** element of an **AssignmentRule**, **InitialAssignment** or **EventAssignment** setting the stoichiometry of the species reference should be **dimensionless**.
- The units of the **math** element of a **RateRule** that acts on the species reference's stoichiometry should be identical to **dimensionless/time**, where **time** is the model-wide **time** units (Section 4.2.4).

Examples

The following is a simple example of a species reference for species "X0", with stoichiometry "2", in a list of reactants within a reaction having the identifier "J1":

```
<model ...>
  ...
  <listOfReactions>
    <reaction id="J1" reversible="false" fast="false">
```



```

1         <listOfReactants>
2             <speciesReference species="X0" stoichiometry="2" constant="true"/>
3         </listOfReactants>
4         ...
5     </reaction>
6     ...
7 </listOfReactions>
8 ...
9 </model>

```

The following is a more complex example of a species reference with an id “**sr01**” and an initial assignment that assigns a rational number to the stoichiometry:

```

12 <model ...>
13     ...
14     <listOfInitialAssignments>
15         <initialAssignment symbol="sr01">
16             <math xmlns="http://www.w3.org/1998/Math/MathML"
17                 xmlns:sbml="http://www.sbml.org/sbml/level3/version1/core">
18                 <cn type="rational" sbml:units="dimensionless"> 3 <sep/> 2 </cn>
19             </math>
20         </initialAssignment>
21     </listOfInitialAssignments>
22     ...
23 ...
24     <listOfReactions>
25         <reaction id="J1" reversible="true" fast="false">
26             <listOfReactants>
27                 <speciesReference id="sr01" species="X0" constant="true"/>
28             </listOfReactants>
29             ...
30         </reaction>
31     </listOfReactions>
32     ...
33 </model>
34

```

A species can occur more than once in the lists of reactants and products of a given [Reaction](#) instance. The effective stoichiometry for the species is the sum of the stoichiometry values given in the [SpeciesReference](#) objects in the list of products *minus* the sum of stoichiometry values given in the [SpeciesReference](#) objects in the list of reactants. A positive value indicates the species is effectively a product and a negative value indicates the species is effectively a reactant. SBML places no restrictions on the effective stoichiometry of a species in a reaction; for example, it can be zero. In the following SBML fragment, the two reactions have the same effective stoichiometry for all their species:

```

42 <reaction id="x" reversible="false" fast="false">
43     <listOfReactants>
44         <speciesReference species="a" stoichiometry="1" constant="true"/>
45         <speciesReference species="a" stoichiometry="1" constant="true"/>
46         <speciesReference species="b" stoichiometry="1" constant="true"/>
47     </listOfReactants>
48     <listOfProducts>
49         <speciesReference species="c" stoichiometry="1" constant="true"/>
50         <speciesReference species="b" stoichiometry="1" constant="true"/>
51     </listProducts>
52 </reaction>
53 <reaction id="y" reversible="false" fast="false">
54     <listOfReactants>
55         <speciesReference species="a" stoichiometry="2" constant="true"/>
56     </listOfReactants>
57     <listOfProducts>
58         <speciesReference species="c" stoichiometry="1" constant="true"/>
59     </listProducts>
60 </reaction>

```

4.11.4 *ModifierSpeciesReference*

Sometimes a species appears in the kinetic rate formula of a reaction but is itself neither created nor destroyed in that reaction (for example, because it acts as a catalyst or inhibitor). In SBML, all such species are simply called *modifiers* without regard to the detailed role of those species in the model (though a model could use SBO terms to clarify the roles; see Section 5). The **Reaction** object class provides a way to express which species act as modifiers in a given reaction. This is the purpose of the list of modifiers available in **Reaction**. The list contains instances of **ModifierSpeciesReference** object.

As shown in Figure 19 on page 60, the **ModifierSpeciesReference** class inherits the mandatory attribute **species** and optional attributes **id** and **name** from the parent class **SimpleSpeciesReference**; see Section 4.11.2 for their precise definitions. As already explained in Section 4.11.2, no meaning is assigned to the identifier of **ModifierSpeciesReference** object instances in SBML Level 3 Version 1 Core, but the identifiers are available for possible use by SBML Level 3 packages. Note also that modifiers in reactions also have no stoichiometries and therefore do not possess a **stoichiometry** attribute.

The value of the **species** attribute must be the identifier of a species defined in the enclosing **Model**; this species is designated as a modifier for the current reaction. A reaction may have any number of modifiers. It is permissible for a modifier species to appear simultaneously in the list of reactants and products of the same reaction where it is designated as a modifier, as well as to appear in the list of reactants, products and modifiers of other reactions in the model.

4.11.5 *LocalParameter*

The kinetic law of a **Reaction** instance can contain a list of local parameters that are only accessible by the kinetic law formula of that particular reaction. The list contains **LocalParameter** objects, each of which associates an identifier with a value. This identifier can then be used in the kinetic law. The definition of **LocalParameter** is shown in Figure 19 on page 60.

The id and name attributes

LocalParameter has a required attribute **id**, of data type **SId**, to give the local parameter an identifier by which the kinetic law formula can refer to it. A local parameter can also have an optional **name** attribute of type **string**. The identifier of a local parameter needs to be unique only within the list of local parameters of one reaction. The details about the scope for identifiers are given in Sections 3.3.1 and 4.11.6, and about the use of names in Section 3.3.2.

The value attribute

The optional attribute **value** determines the value (of type **double**) assigned to the identifier. A missing **value** attribute implies that the value either is unknown, or to be obtained from an external source. (Note that, unlike the case with global **Parameter** objects, there is no way in SBML Level 3 Version 1 Core for **InitialAssignment** or other SBML constructs to be used for setting the value of **LocalParameter** objects, because local parameters are local to reactions.)

The units attribute

The units associated with the value of the local parameter can be specified by the optional attribute **units**. The attribute's value must have the data type **UnitSIdRef** (Section 3.1.10). There are no constraints on the units that can be assigned to a parameter; there are also no units to inherit from the enclosing **Model** object (unlike the case for, e.g., **Species** and **Compartment**).

The units of the parameter are used in the following way: when a local parameter's identifier appears in the content of the **math** element of the **KineticLaw** object, the units associated with the value are those declared by the **LocalParameter** object's **units** attribute.

The `sboTerm` attribute

`LocalParameter` inherits an optional `sboTerm` attribute of type `SBOTerm` from its parent class `SBase` (see Sections 3.1.11 and 5). When a value is given to this attribute in a `LocalParameter` instance, it should be an SBO identifier belonging to the branch for type `LocalParameter` indicated in Table 5. The relationship is of the form “the local parameter *is a* X”, where X is the SBO term. The term chosen should be the most precise (narrow) one that captures the role of the local parameter in the model.

As discussed in Section 5, SBO labels are optional information on a model. Applications are free to ignore `sboTerm` values. A model must be interpretable without the benefit of SBO labels.

4.11.6 `KineticLaw`

The `KineticLaw` object class is used to describe the rate at which the process defined by the `Reaction` takes place. As shown in Figure 19 on page 60, `KineticLaw` has elements called `math` and `listOfLocalParameters`, in addition to the attributes and elements it inherits from `SBase`.

The `math` element

As shown in Figure 19 on page 60, `KineticLaw` has an element called `math` for holding a MathML formula defining the rate of the reaction. The expression in `math` may refer to global identifiers defined in the model, or `LocalParameter` object identifiers from the `KineticLaw`’s list of local parameters (see below). However, the only `Species` identifiers that can be used in `math` are those declared in the lists of reactants, products and modifiers in the `Reaction` object (see Sections 4.11.2, 4.11.3 and 4.11.4).

Section 4.11.7 provides important discussions about the meaning and interpretation of SBML “kinetic laws”.

The list of local parameters

An instance of `KineticLaw` can contain a list of one or more `LocalParameter` objects (Section 4.11.5) defining new parameters whose identifiers can be used in the `math` formula. These “local parameters” are optional—a kinetic law can always refer to global `Parameter` objects. The local parameter facility simply provides a way to add additional parameters that may be relevant only to a specific reaction, and that may therefore be better handled by encapsulating their definitions within that kinetic law.

As discussed in Section 3.3.1, reactions introduce local namespaces for local parameter identifiers, and within a `KineticLaw` object, a local parameter whose identifier is identical to a global identifier defined in the model takes precedence over the value associated with the global identifier. Note that this introduces the potential for a local parameter definition to shadow a global identifier. In SBML’s simple symbol system, there is no separation of symbols by class of object; consequently, inside the kinetic law mathematical formula, the value of a local parameter having the same identifier as a species, compartment, parameter or other global model entity will override the global value. Modelers and software developers may wish to take precautions to avoid this happening accidentally.

The `sboTerm` attribute

`KineticLaw` inherits an optional `sboTerm` attribute of type `SBOTerm` from its parent class `SBase` (see Sections 3.1.11 and 5). When a value is given to this attribute in a `KineticLaw` instance, it should be an SBO identifier belonging to the branch for type `KineticLaw` indicated in Table 5. The relationship is of the form “the kinetic law *is a* X”, where X is the SBO term. The term chosen should be the most precise (narrow) one that captures the role of the kinetic law in the model.

As discussed in Section 5, SBO labels are optional information on a model. Applications are free to ignore `sboTerm` values. A model must be interpretable without the benefit of SBO labels.

Example

The following is an example of a `Reaction` object that defines a reaction with identifier J_1 , in which $X_0 \rightarrow S_1$ at a rate given by $k \cdot [X_0] \cdot [S_2]$, where S_2 is a catalyst and k is a parameter, and the square brackets symbolizes

that the species quantities have units of concentration. The reaction is assumed to take place all in one compartment identified as “c1”. The example demonstrates the use of species references and [KineticLaw](#) objects. The units associated with the species identifiers here are *amount/volume* (see Section 4.6), and so the rate expression $k \cdot [X_0] \cdot [S_2]$ needs to be multiplied by the compartment volume (represented by its identifier, “c1”) to produce the desired units of *amount/time* for the rate expression.

```

6      <model timeUnits="second" extentUnits="mole" substanceUnits="mole">
7        <listOfUnitDefinitions>
8          <unitDefinition id="per_concent_per_time">
9            <listOfUnits>
10              <unit kind="litre" exponent="1" scale="0" multiplier="1"/>
11              <unit kind="mole"   exponent="-1" scale="0" multiplier="1"/>
12              <unit kind="second" exponent="-1" scale="0" multiplier="1"/>
13            </listOfUnits>
14          </unitDefinition>
15        </listOfUnitDefinitions>
16        <listOfCompartments>
17          <compartment id="c1" units="litre" size="0.001" spatialDimensions="3" constant="true"/>
18        </listOfCompartments>
19        <listOfSpecies>
20          <species id="S1" compartment="c1" initialConcentration="2.0"
21            hasOnlySubstanceUnits="false" boundaryCondition="false" constant="false"/>
22          <species id="S2" compartment="c1" initialConcentration="0.5"
23            hasOnlySubstanceUnits="false" boundaryCondition="false" constant="false"/>
24          <species id="X0" compartment="c1" initialConcentration="1.0"
25            hasOnlySubstanceUnits="false" boundaryCondition="false" constant="false"/>
26        </listOfSpecies>
27        <listOfReactions>
28          <reaction id="J1" reversible="false" fast="false">
29            <listOfReactants>
30              <speciesReference species="X0" stoichiometry="1" constant="true"/>
31            </listOfReactants>
32            <listOfProducts>
33              <speciesReference species="S1" stoichiometry="1" constant="true"/>
34            </listOfProducts>
35            <listOfModifiers>
36              <modifierSpeciesReference species="S2"/>
37            </listOfModifiers>
38            <kineticLaw>
39              <math xmlns="http://www.w3.org/1998/Math/MathML">
40                <apply>
41                  <times/> <ci> k </ci> <ci> S2 </ci> <ci> X0 </ci> <ci> c1 </ci>
42                </apply>
43              </math>
44              <listOfLocalParameters>
45                <localParameter id="k" value="0.1" units="per_concent_per_time"/>
46              </listOfLocalParameters>
47            </kineticLaw>
48          </reaction>
49        </listOfReactions>
50      </model>

```

4.11.7 Mathematical interpretation of SBML reactions and kinetic laws

In SBML, *reactions* are the central mechanism for describing processes that change the quantities of species in a model. The *kinetic law* of an SBML reaction provides a quantitative description of the speed with which this happens. In this section, we provide an interpretation of SBML kinetic laws in the framework of a system of ordinary differential equations (ODEs). However, the choice of ODEs as the framework is only for exposition purposes here, in order to allow us to present a concrete mathematical expression of the model in terms that many readers will be familiar with; it is equally possible to translate a model into other frameworks, and some formulations, such as discrete stochastic systems, are indeed quite common.

Semantics of rate law and stoichiometry

The *stoichiometry* of a species S in a reaction describes the proportion, relative to other species participating in that reaction, of S involved in each reaction event. For example, in a reaction $S_1 + 2S_2 \rightarrow S_3$, twice as many entities of S_2 as entities of S_1 are involved each time a reaction event is counted. The value of the expression in the [KineticLaw](#)’s `math` element describes the *rate* at which the reaction takes place. The product of the reaction rate (of a given reaction) and the stoichiometry (of a given species in the reaction) describes the reaction’s contribution to the rate of change of the species’s quantity in the overall system.

It is important to make clear that a “kinetic law” in SBML is *not* identical to a traditional rate law. When modeling species as continuous amounts (e.g., concentrations), the rate laws used are traditionally expressed in terms of *concentration per time*. Unfortunately, this approach only works well in cases where certain assumptions hold. Three assumptions in particular are incompatible with generalized multicompartment modeling; they are listed in Table 4 along with the problems they entail.

Assumption	Problem
All species that participate in a given reaction are located in one compartment	SBML must support reaction processes (e.g., transport) that move species between compartments
Compartment are three-dimensional volume containers	SBML must support models where reactions may take place at interfaces (e.g., 2-D membranes) between compartments, thus involving compartments with different dimensions
Compartment volumes are constant over time	SBML must support systems with compartments that can change size over time

Table 4: Assumptions behind “traditional” rate laws, and the problems they imply for general multicompartmental modeling.

A simple example can illustrate the problems that arise when describing reactions between multiple volumes using *concentration/time* units (which is to say, *amount/volume/time*). Suppose we have two species pools S_1 and S_2 , with S_1 located in a compartment having volume V_1 , and S_2 located in a compartment having volume V_2 . Let the volume $V_2 = 3V_1$. Now consider a transport reaction $S_1 \rightarrow S_2$ in which the species S_1 is moved from the first compartment to the second. Assume we only want to model the overall effect, without getting into the molecular details (which might in reality involve such things as pores in a membrane between the compartments). Let us use the simplest type of chemical kinetics, in which the rate of the transport reaction is controlled by the activity of S_1 and this rate is equal to some constant k times the activity of S_1 . For the sake of simplicity, assume S_1 is in a diluted solution and thus that the activity of S_1 can be taken to be equal to its concentration $[S_1]$. The rate expression will therefore be $k \cdot [S_1]$, with the units of k being $1/time$. Then:

$$\frac{d[S_2]}{dt} = -\frac{d[S_1]}{dt} = k \cdot [S_1]$$

So far, this looks normal—until we consider the number of molecules of S_1 that disappear from the compartment of volume V_1 and appear in the compartment of volume V_2 . The number of molecules of S_1 (call this n_{S_1}) is given by $[S_1] \cdot V_1$ and the number of molecules of S_2 (call this n_{S_2}) is given by $[S_2] \cdot V_2$. Since our volumes have the relationship $V_2/V_1 = 3$, the relationship above implies that $n_{S_1} = k \cdot [S_1] \cdot V_1$ molecules disappear from the first compartment per unit of time and $n_{S_2} = 3 \cdot k \cdot [S_1] \cdot V_1$ molecules appear in the second compartment. In other words, we have created matter out of nothing!

The problem lies in the use of concentrations as the measure of what is transferred by the reaction, because concentrations depend on volumes and the scenario involves multiple unequal volumes. The problem is not limited to using concentrations or volumes; the same problem also exists when using density, i.e., *mass/volume*, as well as dependency on other spatial distributions (i.e., areas or lengths). What must be

done instead is to consider the number of “items” being acted upon by a reaction process irrespective of their distribution in space (volume, area or length). An “item” in this context may be a molecule, particle, mass, or other “thing”, as long as the substance measurement is independent of the size of the space in which the items are located and the processes take place.

In multicompartment models, to be able to specify a rate law only once and then use it unmodified in equations for different species, the rate law needs to be expressed in terms of an intensive property, that is, *species quantity/time*, rather than the extensive property typically used, *species quantity/size/time*. As a result, modelers and software tools in general cannot insert traditional textbook rate laws unmodified into the **math** element of a **KineticLaw**. The unusual term “kinetic law” was chosen to alert users to this difference.

Constructing rate-of-change equations for the species

A consequence of approach to “kinetic laws” discussed in the previous section is this: when constructing equations describing the time-rates of change of different species defined by an SBML model, the equations are assumed to be in terms of time-rates of changes to *amounts*, *not concentrations* (or more generally *densities*, i.e., amount per size of compartment). A kinetic law does *not* describe how often a reaction would take place in a compartment of unit size, but rather how often it takes place (per time unit) given the actual size of the compartment. The dimension of the kinetic law is therefore *number of reaction events per time*.

When constructing a system of equations dictating the rates of change of the different species in an SBML model, we only need to consider those species having attribute values **constant**=“false” and **boundaryValue**=“false”, because as discussed in Section 4.6.6, these are the only species affected by the reactions in the model. (Other species not meeting these criteria may be affected by other SBML constructs, but here, we are focusing specifically on the implications of reactions.)

Assume now a model in which N species S_1, S_2, \dots, S_N having attributes **constant**=“false” and **boundaryValue**=“false” participate in M reactions R_1, R_2, \dots, R_M . Let v_{R_j} represent the rate or velocity of reaction R_j as given by the formula in the **math** element of **KineticLaw** object for R_j . The units associated with this rate expression are *extent/time* units, where the extent and time units are specified by the **extentUnits** and **timeUnits** attributes on the **Model** object, respectively. Let stoich_{S_i, R_j} represent the effective stoichiometry of species S_i in reaction R_j . (By “effective stoichiometry”, we mean the number that results from taking the sum of the stoichiometry values of all references to S_i in R_j ’s **listOfReactants** and subtracting the sum of the stoichiometric values of all references to S_i in R_j ’s **listOfProducts**.) If S_i is neither a reactant nor product in some reaction R_x , then $\text{stoich}_{S_i, R_x} = 0$. Finally, let n_{S_i} represent the amount of species S_i in the model (and note that this value is *not* a concentration or density).

There are three possible cases to consider when constructing rate-of-change equations for the species:

1. *No conversion factors defined.* If the **Species** object instance for S_i does not define a value for its **conversionFactor** attribute, and the **Model** object also does not define a value for its **conversionFactor** attribute, then the implication is that the units of the reaction extent are identical to the units of amount of species S_i . In that case, the formula for the rate of change of the species amount is as follows:

$$\frac{dn_{S_i}}{dt} = \sum_{j=1}^M \text{stoich}_{S_i, R_j} \cdot v_{R_j}$$

2. *Global conversion factor defined.* If the **Model** object instance defines a value for its **conversionFactor** attribute, and the **Species** object for S_i does not define a value for its **conversionFactor**, then the global conversion factor is used to convert between the units of reaction *extent* in the model and the units of amount of species S_i . Let c_{model} represent the value of the parameter object identified by the **conversionFactor** attribute value on **Model** (see Section 4.2.7). The formula for the rate of change of S_i ’s amount then becomes the following:

$$\frac{dn_{S_i}}{dt} = c_{\text{model}} \cdot \sum_{j=1}^M \text{stoich}_{S_i, R_j} \cdot v_{R_j}$$

3. *Conversion factor defined for the species.* If the **Species** object instance for S_i defines a value for its **conversionFactor** attribute, then this factor is used to convert between the units of reaction *extent* in the model and the units of amount of species S_i . (The factor defined by the species is used regardless of whether a value also exists for the **Model** object's **conversionFactor**.) Let c_{S_i} represent the value of the parameter object identified by S_i 's **conversionFactor** attribute value (see Section 4.6.7). The formula for the rate of change of S_i 's amount then becomes the following:

$$\frac{dn_{S_i}}{dt} = c_{S_i} \cdot \sum_{j=1}^M \text{stoich}_{S_i, R_j} \cdot v_{R_j}$$

In Section 8.2.4, we present some recommendations for how to encode rate laws and models in SBML.

4.11.8 Use of reaction identifiers in mathematical expressions

The value of the **id** attribute of a **Reaction** can be used as the content of a **ci** element in MathML formulas elsewhere in the model. Such a **ci** element or symbol represents the rate of the given reaction as given by the reaction's **KineticLaw** object. A **KineticLaw** object in effect forms an assignment statement assigning the evaluated value of the **math** element to the symbol value contained in the **Reaction id** attribute. No other object can assign a value to such a reaction symbol; i.e., the **variable** attributes of **InitialAssignment**, **RateRule**, **AssignmentRule** and **EventAssignment** objects cannot contain the value of a **Reaction id** attribute.

The combined set of **InitialAssignment**, **AssignmentRule** and **KineticLaw** objects form a set of assignment statements that should be considered as a whole. The combined set of assignment rules should not contain algebraic loops: a chain of dependency between these statements should terminate. (More formally, consider the directed graph of assignment statements where nodes are statements and directed arcs exist for each occurrence of a symbol in an assignment statement **math** element. The directed arcs start from the statement defining the symbol to the statements that contain the symbol in their **math** elements. Such a graph must be acyclic.) Examples of valid and invalid set of assignment statements are given in Section 4.9.5.

4.12 Events

Model has an optional list of **Event** objects that describe the time and form of explicit instantaneous discontinuous state changes in the model. For example, an event may describe that one species quantity is halved when another species' quantity exceeds a given threshold value.

An **Event** object defines when the event can occur, the variables that are affected by the event, and how the variables are affected. The effect of the event can optionally be delayed after the occurrence of the condition which invokes it. Conceptually, the operation of an event is divided into two phases (even when the event is not delayed): one when the event is *fired* and the other when the event is *executed*. The **Event** type is defined in Figure 20 on the following page. The object classes **Event**, **Trigger**, **Delay** and **EventAssignment** are derived from **SBase** (see Section 3.2). An example of a model which uses events is given below.

4.12.1 Event

An **Event** definition has two required parts: a trigger condition and at least one **EventAssignment**. In addition, an event can include an optional delay. These features of **Event** are described below.

The id and name attributes

As with most components in SBML, an **Event** has **id** and **name** attributes, but in the case of **Event**, both are optional. These attributes operate in the manner described in Section 3.3.

The optional sboTerm attribute on Event

Event inherits an optional **sboTerm** attribute of type **SBOTerm** from its parent class **SBase** (see Sections 3.1.11 and 5). When a value is given to this attribute in a **Event** instance, it should be an SBO identifier belonging

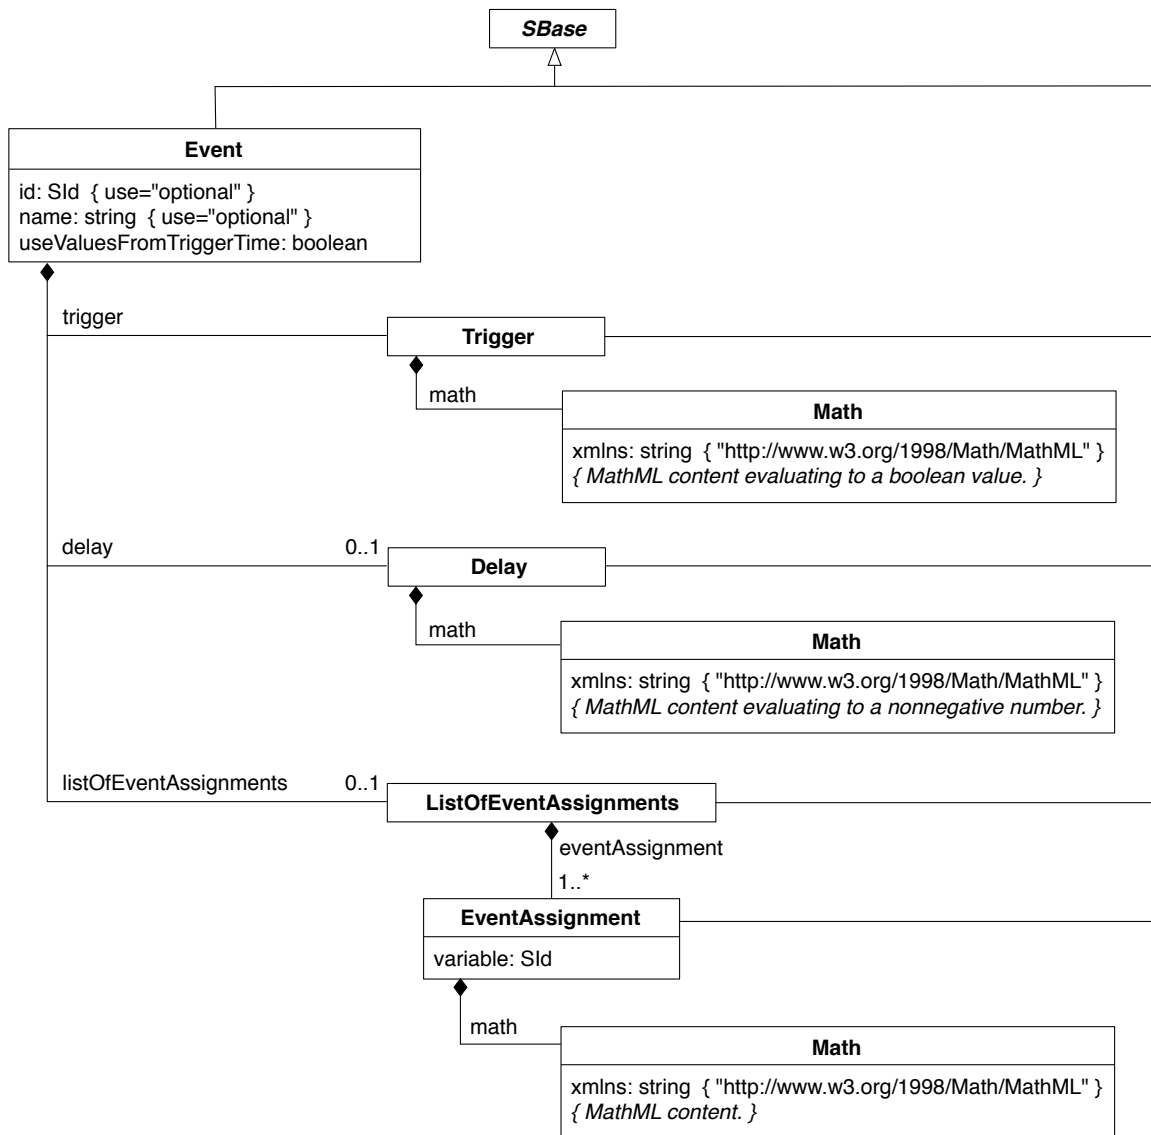


Figure 20: The definitions of *Event*, *Trigger*, *Delay*, *EventAssignment*, and *ListOfEventAssignment*.

to the branch for type *Event* indicated in Table 5. The relationship is of the form “the event *is a* X”, where X is the SBO term. The term chosen should be the most precise (narrow) one that captures the role of the event in the model.

As discussed in Section 5, SBO labels are optional information on a model. Applications are free to ignore *sboTerm* values. A model must be interpretable without the benefit of SBO labels.

The *useValuesFromTriggerTime* attribute

The optional *Delay* on *Event* means there are two times to consider when computing the results of an event: the time at which the event *fires*, and the time at which assignments are *executed*. It is also possible to distinguish between the time at which the *EventAssignment*’s expression is calculated, and the time at which the assignment is made: the expression could be evaluated at the same time the assignments are performed, i.e., when the event is *executed*, but it could also be defined to be evaluated at the time the event *fires*.

The mandatory attribute `useValuesFromTriggerTime` on **Event** allows a model to indicate the time at which the event's assignments are intended to be evaluated. A value of `"true"` indicates that values of the assignment formulas are to be computed at the moment the event fired, not after the delay. If `useValuesFromTriggerTime="false"`, it means that the formulas in the event's assignments are to be computed after the delay, at the time the event is executed.

4.12.2 Trigger

As shown in Figure 20, the `trigger` element of an **Event** must contain exactly one object of class **Trigger**. This object contains one `math` element containing a MathML expression. The expression must evaluate to a value of type `boolean`. The exact moment at which the expression evaluates to `"true"` is the time point when the **Event** is *fired*.

An event only fires when its **Trigger** expression makes the transition in value from `"false"` to `"true"`. The event will also fire at any future time points when the `trigger` make this transition; in other words, an event can fire multiple times during a simulation if its trigger condition makes the transition from `"false"` to `"true"` more than once.

An important question is whether an event can fire prior to, or at, initial simulation time, i.e., $t \leq 0$. The answer is no: an event can only be triggered immediately after initial simulation time i.e., $t > 0$.

The optional `sboTerm` attribute on **Trigger**

Trigger inherits an optional `sboTerm` attribute of type `SBOTerm` from its parent class **SBase** (see Sections 3.1.11 and 5). When a value is given to this attribute in a **Trigger** instance, it should be an SBO identifier belonging to the branch for type **Trigger** indicated in Table 5. The relationship is of the form "the trigger is a X", where X is the SBO term. The term chosen should be the most precise (narrow) one that captures the role of the trigger in the model.

As discussed in Section 5, SBO labels are optional information on a model. Applications are free to ignore `sboTerm` values. A model must be interpretable without the benefit of SBO labels.

4.12.3 Delay

As shown in Figure 20, an **Event** object can contain an optional `delay` element of class **Delay**. The **Delay** is derived from **SBase** and contains a mathematical formula stored in `math`. The formula is used to compute the length of time between when the event has *fired* and when the event's assignments (see below) are actually *executed*. If no delay is present on a given **Event**, no delay is defined for that event.

The expression in the **Delay** object's `math` element must be evaluated at the time the event is *fired*. The expression must always evaluate to a nonnegative number (otherwise, a nonsensical situation could arise where an event is defined to fire before it is triggered!).

Units of delay expressions

The units of the numerical value computed by a **Delay** instance's `math` expression should match the model's units of *time* (meaning the definition of the `"timeUnits"` in the model; see Section 4.2.4). Note that, as in other cases of MathML expressions in SBML, units are *not* automatically predefined or assumed. As discussed in Section 3.4.10, expressions containing only literal numbers and/or **Parameter** objects without declared units, are considered to have unspecified units. In such cases, the correspondence between the needed units and the (unknown) units of the **Delay** `math` expression cannot be proven, and while such expressions are not considered inconsistent, all that can be assumed by model interpreters (whether software or human) is that the units *may* be consistent.

The following **Event** example fragment helps illustrate this:

```
<model timeUnits="second" ...>
  ...
  <listOfEvents>
    <event useValuesFromTriggerTime="true">
```

```

1      ...
2      <delay>
3        <math xmlns="http://www.w3.org/1998/Math/MathML">
4          <cn> 10 </cn>
5        </math>
6      </delay>
7      ...
8    </event>
9  </listOfEvents>
10   ...
11 </model>

```

Note that the “<cn> 10 </cn>” within the mathematical formula has no specified units attached to it. The model is not invalid because of this, but a recipient of the model may justifiably be concerned about what “10” really means. (Ten seconds? What if the global units of time on the model were changed from seconds to milliseconds? Would the modeler remember to change “10” to “10 000”?) A better approach would be to declare the units explicitly, as in the following example:

```

17 <model timeUnits="second" ...>
18   ...
19   <listOfEvents>
20     <event useValuesFromTriggerTime="true">
21       ...
22       <delay>
23         <math xmlns="http://www.w3.org/1998/Math/MathML"
24               xmlns:sbml="http://www.sbml.org/sbml/level3/version1/core">
25           <cn sbml:units="second"> 10 </cn>
26         </math>
27       </delay>
28       ...
29     </event>
30   </listOfEvents>
31   ...
32 </model>

```

Another approach would be to define a global **Parameter** object for the desired time delay, assign units to that parameter, and replace the **cn** element in the example above with a **ci** element referencing the parameter’s identifier.

The optional **sboTerm** attribute on **Delay**

Delay inherits an optional **sboTerm** attribute of type **SBOTerm** from its parent class **SBase** (see Sections 3.1.11 and 5). When a value is given to this attribute in a **Delay** instance, it should be an SBO identifier belonging to the branch for type **Delay** indicated in Table 5. The relationship is of the form “the delay *is a* X”, where X is the SBO term. The term chosen should be the most precise (narrow) one that captures the role of the delay in the model.

As discussed in Section 5, SBO labels are optional information on a model. Applications are free to ignore **sboTerm** values. A model must be interpretable without the benefit of SBO labels.

4.12.4 **EventAssignment**

Event contains an optional element called **listOfEventAssignments**, of class **ListOfEventAssignment**. In every instance of an event definition in a model, the object’s **listOfEventAssignments** element must have a non-empty list of one or more **eventAssignment** elements of class **EventAssignment**. The object class **EventAssignment** has one required attribute, **variable**, and a required element, **math**. Being derived from **SBase**, it also has all the usual attributes and elements of its parent class.

An “event assignment” has effect when the event is *executed*; that is, at the end of any given delay period (if given) following the moment that the **Event** is triggered. See Section 4.12.6 below for more information about events and event assignments in SBML.

The variable attribute

The **variable** attribute is of type **SidRef** and can contain the identifier of a **Compartment**, **Species**, **Species-Reference**, or **Parameter** instance defined in the model. When the event is executed, the value of the model component identified by **variable** is changed by the **EventAssignment** to the value computed by the **math** element; that is, a species' quantity, species reference's stoichiometry, compartment's size or parameter's value are reset to the value computed by **math**.

Certain restrictions are placed on what can appear in **variable**:

- The object identified by the value of the **variable** attribute must not have its **constant** attribute set to **"true"**. (Constants cannot be affected by events.)
- The **variable** attribute must not contain the identifier of a reaction; only species, species references, compartment and parameter values may be set by an **Event**.
- The value of every **variable** attribute must be unique among the set of **EventAssignment** objects within a given **Event** instance. In other words, a single event cannot have multiple **EventAssignments** assigning the same variable. (All of them would be performed at the same time, when that particular **Event** triggers, resulting in indeterminacy.) Separate **Event** instances can refer to the same variable.
- A variable cannot be assigned a value in an **EventAssignment** object instance and also be assigned a value by an **AssignmentRule**, i.e., the value of the **variable** attribute in an **EventAssignment** instance cannot be the same as the value of a **variable** attribute in a **AssignmentRule** instance. (Assignment rules hold at all times, therefore it would be inconsistent to also define an event that reassigns the value of the same variable.)

Note that the time of assignment of the object identified by the value of **variable** is always the time at which the **Event** is *executed*, not when it is *fired*. The timing is controlled by the optional **Delay** in an **Event**. The time of assignment is not affected by the **useValuesFromTriggerTime** attribute on **Event**—that attribute affects the time at which the **EventAssignment**'s **math** expression is evaluated. In other words, SBML allows decoupling the time at which the **variable** is assigned from the time at which its value expression is calculated.

The optional sboTerm attribute on EventAssignment

EventAssignment inherits an optional **sboTerm** attribute of type **SBOTerm** from its parent class **SBase** (see Sections 3.1.11 and 5). When a value is given to this attribute in a **EventAssignment** instance, it should be an SBO identifier belonging to the branch for type **EventAssignment** indicated in Table 5. The relationship is of the form “the event assignment *is a* X”, where X is the SBO term. The term chosen should be the most precise (narrow) one that captures the role of the event assignment in the model.

As discussed in Section 5, SBO labels are optional information on a model. Applications are free to ignore **sboTerm** values. A model must be interpretable without the benefit of SBO labels.

EventAssignment's math

The **math** element contains a MathML expression that defines the new value of the object identified by the **variable**.

The time at which this expression is evaluated is determined by **Event**'s **useValuesFromTriggerTime** attribute. If the attribute value is **"true"**, the expression must be evaluated when the event is *fired*; more precisely, the values of identifiers occurring in MathML **ci** attributes in the **EventAssignment**'s **math** expression are the values they have at the point when the event *fired*. If, instead, **useValuesFromTriggerTime**'s value is **"false"**, it means the values at *execution* time should be used; that is, the values of identifiers occurring in MathML **ci** attributes in the **EventAssignment**'s **math** expression are the values they have at the point when the event *executed*.

Units of the `math` formula in **EventAssignment**

In all cases, as would be expected, the units of the formula contained in the `math` element of **EventAssignment** should be consistent with the units of the object identified by the `variable` attribute. More precisely:

- In the case of a species, an **EventAssignment** sets the referenced species' quantity (*concentration* or *amount*) to the value determined by the formula in `math`. The overall units of the `math` formula should be identical to the *units of the species* as defined in Section 4.6.5.
- In the case of a species reference, an **EventAssignment** sets the value of the `stoichiometry` attribute of the referenced **SpeciesReference** object to the value determined by the formula in `math`. The overall units of the formula in `math` should be **dimensionless**, because reactant and product stoichiometries in reactions are dimensionless quantities.
- In the case of a compartment, an **EventAssignment** sets the referenced compartment's size to the size determined by the formula in `math`. The overall units of the formula should be identical to the units specified for the size of the compartment identified by the value of the **EventAssignment**'s `variable` attribute. (See Section 4.5.4 for an explanation of how the units of the compartment's size are determined.)
- In the case of a parameter, an **EventAssignment** sets the referenced parameter's value to that determined by the formula in `math`. The overall units of the formula should be identical to the units defined for the parameter identified by the value of the **EventAssignment**'s `variable` attribute. (See Section 4.7.3 for an explanation of how the units of the parameter are determined.)

Note that the formula placed in the `math` element has no assumed units. The consistency of the units of the formula, and the units of the entity which the assignment affects, should be explicitly established just as in the case of the value of `delay`.

4.12.5 Example Event definitions

A example of an **Event** object follows. This structure makes the assignment $k_2 = 0$ at the point when $P_1 \leq P_2$:

```
<event>
  ...
  <listOfUnitDefinitions>
    <unitDefinition id="per_second">
      <listOfUnits>
        <unit kind="second" exponent="-1" multiplier="1" scale="0"/>
      </listOfUnits>
    </unitDefinition>
  </listOfUnitDefinitions>
  ...
  <listOfParameters>
    ...
    <parameter id="k2" value="0.05" units="per_second" constant="false"/>
    <parameter id="k2reset" value="0.0" units="per_second" constant="true"/>
    ...
  </listOfParameters>
  ...
  <listOfEvents>
    <event useValuesfromTriggerTime="true">
      <trigger>
        <math xmlns="http://www.w3.org/1998/Math/MathML">
          <apply>
            <leq/>
            <ci> P_1 </ci>
            <ci> P_2 </ci>
          </apply>
        </math>
      </trigger>
    </event>
  </listOfEvents>
</event>
```

```

1      <listOfEventAssignments>
2          <eventAssignment variable="k2">
3              <math xmlns="http://www.w3.org/1998/Math/MathML">
4                  <ci> k2reset </ci>
5              </math>
6          </eventAssignment>
7      </listOfEventAssignments>
8  </event>
9  </listOfEvents>
10  ...
11 </model>

```

A complete example of a model using events is given in Section 7.9.

4.12.6 Detailed semantics of events

The description of events above describes the action of events in isolation from each other. This section describes how events interact.

Events whose **trigger** expression is true at the start of a simulation do not *fire* at the start of the simulation ($t = 0$). Events *fire* only when the trigger *becomes* true, i.e., the trigger expression transitions from “false” to “true”, which cannot happen at $t = 0$ but can happen at $t > 0$.

Any transition of a **trigger** expression from “false” to “true” will cause an event to *fire*. Consider an **Event** object definition E with delay d where the **trigger** expression makes a transition from “false” to “true” at times t_1 and t_2 . The **EventAssignment** within the **Event** object will have effect at $t_1 + d$ and $t_2 + d$ irrespective of the relative times of t_1 and t_2 . For example events can “overlap” so that $t_1 < t_2 < t_1 + d$ still causes an event assignments to occur at $t_1 + d$ and $t_2 + d$.

It is entirely possible for two events to be *executed* simultaneously in simulated time, and it is possible for events to *fire* other events (i.e., an event assignment can cause an event to *fire*). This leads to several points:

- A software package should retest all event triggers after executing an event assignment in order to account for the possibility that the assignment causes another event trigger to transition from “false” to “true”.
- It is assumed that, although the precise time at which these events are executed is not resolved beyond the given point in simulated time, the order in which the events occur *is* resolved. This order can be significant in determining the overall outcome of a given simulation. When an event X *fires* another event Y and event Y has zero delay, then event Y is added to the existing set of simultaneous events that are pending *execution*. Events such as Y do not have a special priority or ordering. Events X and Y form a cascade of events at the same point in simulation time.
- All events in a model are open to being in a cascade. The position of an event in the event list does not affect whether it can be in the cascade: Y can be triggered whether it is before or after X in the list of events. A cascade of events can be potentially infinite (never terminate); when this occurs a simulator should indicate this has occurred—it is incorrect for a simulator to break a cascade arbitrarily and continue the simulation without at least indicating the infinite cascade occurred.
- A variable can change more than once when processing simultaneous events at simulation time t . The model behavior (output) for such a variable is the value of the variable at the end of processing all the simultaneous events at time t .

5 The Systems Biology Ontology and the **sboTerm** attribute

The values of **id** attributes on SBML components allow the components to be cross-referenced within a model. The values of **name** attributes on SBML components provide the opportunity to assign them meaningful labels suitable for display to humans (Section 3.3). The specific identifiers and labels used in a model necessarily must be unrestricted by SBML, so that software and users are free to pick whatever they need. However, this freedom makes it more difficult for software tools to determine, without additional human intervention, the semantics of models more precisely than the semantics provided by the SBML object classes defined in other sections of this document. For example, there is nothing inherent in a parameter with identifier “**k**” that would indicate to a software tool it is a first-order rate constant (if that’s what “**k**” happened to be in some given model). However, one may need to convert a model between different representations (e.g., Henri-Michaelis-Menten vs. elementary steps), or to use it with different modelling approaches (discrete or continuous). One may also need to relate the model components with other description formats such as SBGN (<http://www.sbgn.org/>) using deeper semantics. Although an advanced software tool *might* be able to deduce the semantics of some model components through detailed analysis of the kinetic rate expressions and other parts of the model, this quickly becomes infeasible for any but the simplest of models.

An approach to solving this problem is to associate model components with terms from carefully curated controlled vocabularies (CVs). This is the purpose of the optional **sboTerm** attribute provided on the SBML class **SBase**. The **sboTerm** attribute always refers to terms belonging to the Systems Biology Ontology (SBO, <http://biomodels.net/SBO/>). In this section, we discuss the **sboTerm** attribute, SBO, the motivations and theory behind their introduction, and guidelines for their use.

SBO is not part of SBML; it is being developed separately, to allow the modeling community to evolve the ontology independently of SBML. However, the terms in the ontology are being designed keeping SBML components in mind, and are classified into subsets that can be directly related with SBML components such as reaction rate expressions, parameters, and a few others, see below. The use of **sboTerm** attributes is optional, and the presence of **sboTerm** on an element does not change the way the model is *interpreted*. Annotating SBML elements with SBO terms adds additional semantic information that may be used to *convert* the model into another model, or another format. Although SBO support provides an important source of information to understand the meaning of a model, software does not need to support **sboTerm** to be considered SBML-compliant.

5.1 Principles

Labeling model components with terms from shared controlled vocabularies allows a software tool to identify each component using identifiers that are not tool-specific. An example of where this is useful is the desire by many software developers to provide users with meaningful names for reaction rate equations. Software tools with editing interfaces frequently provide these names in menus or lists of choices for users. However, without a standardized set of names or identifiers shared between developers, a given software package cannot reliably interpret the names or identifiers of reactions used in models written by other tools.

The first solution that might come to mind is to stipulate that certain common reactions always have the same name (e.g., “Michaelis-Menten”), but this is simply impossible to do: not only do humans often disagree on the names themselves, but it would not allow for correction of errors or updates to the list of predefined names except by issuing new releases of the SBML specification—to say nothing of many other limitations with this approach. Moreover, the parameters and variables that appear in rate expressions also need to be identified in a way that software tools can interpret mechanically, implying that the names of these entities would also need to be regulated.

The Systems Biology Ontology (SBO) provides terms for identifying most elements of SBML. The relationship implied by an **sboTerm** on an SBML model component is “is a” between the characteristic of the component meant to be described by SBO on this element and the SBO term identified by the value of the **sboTerm**. By adding SBO term references on the components of a model, a software tool can provide additional details using independent, shared vocabularies that can enable *other* software tools to recognize precisely what the component is meant to be. Those tools can then act on that information. For example,

if the SBO identifier **SBO:0000049** is assigned to the concept of “first-order irreversible mass-action kinetics, continuous framework”, and a given **KineticLaw** object in a model has an **sboTerm** attribute with this value, then regardless of the identifier and name given to the reaction itself, a software tool could use this to inform users that the reaction is a first-order irreversible mass-action reaction. This kind of reverse engineering of the meaning of reactions in a model would be difficult to do otherwise, especially for more complex reaction types.

The presence of an SBO label on a compartment, species, or reaction, can help map SBML elements to equivalent concepts in other standards, such as (but not limited to) BioPAX (<http://www.biopax.org/>), PSI-MI (<http://www.psidev.info/index.php?q=node/60>), or the Systems Biology Graphical Notation (SBGN, <http://www.sbgn.org/>). Such mappings can be used in conversion procedures, or to build interfaces, with SBO becoming a kind of “glue” between standards of representation.

The presence of the label on a kinetic expression can also allow software tools to make more intelligent decisions about reaction rate expressions. For example, an application could recognize certain types of reaction formulas as being ones it knows how to solve with optimized procedures. The application could then use internal, optimized code implementing the rate formula indexed by identifiers such as **SBO:0000049** appearing in SBML models.

Finally, SBO labels may be a very valuable tool when it comes to model integration, by helping identify interfaces, convert mathematical expressions and parameters etc.

Although the use of SBO can be beneficial, it is critical to keep in mind that the presence of an **sboTerm** value on an object *must not change the fundamental mathematical meaning* of the model. An SBML model must be defined such that it stands on its own and does not depend on additional information added by SBO terms for a correct mathematical interpretation. SBO term definitions will not imply any alternative mathematical semantics for any SBML object labeled with that term. Two important reasons motivate this principle. First, it would be too limiting to require all software tools to be able to understand the SBO vocabularies in addition to understanding SBML. Supporting SBO is not only additional work for the software developer; for some kinds of applications, it may not make sense. If SBO terms on a model are optional, it follows that the SBML model *must* remain unambiguous and fully interpretable without them, because an application reading the model may ignore the terms. Second, we believe allowing the use of **sboTerm** to alter the mathematical meaning of a model would allow too much leeway to shoehorn inconsistent concepts into SBML objects, ultimately reducing the interoperability of the models.

5.2 Using SBO and sboTerm

The **sboTerm** attribute data type is always **SBOTerm**, defined in Section 3.1.11. When present in a given model object instance, the attribute’s value must be an identifier taken from the Systems Biology Ontology (SBO; <http://biomodels.net/SBO/>). This identifier must refer to a single SBO term that best defines the entity encoded by the SBML object in question. An example of the type of relationship intended is: *the KineticLaw in reaction R1 is a first-order irreversible mass action rate law*.

Note the careful use of the words “defines” and “entity encoded by the SBML object” in the paragraph above. As mentioned, the relationship between the SBML object and the URI is:

The “thing” encoded by this SBML object has a characteristic that *is an* instance of the “thing” represented by the referenced SBO term.

The characteristic relevant for each SBML object is described in the second column of Table 5.

5.2.1 The structure of the Systems Biology Ontology

The goal of SBO labeling for SBML is to clarify to the fullest extent possible the nature of each element in a model. The approach taken in SBO begins with a hierarchically-structured set of controlled vocabularies with six main divisions: (1) entity, (2) participant role, (3) quantitative parameter, (4) modeling framework, (5) mathematical expression, and (6) interaction. Figure 21 on the next page illustrates the highest level of SBO.

Each of the six branches of Figure 21 have a hierarchy of terms underneath them. At this time, we can only begin to list some initial concepts and terms in SBO; what follows is not meant to be complete, comprehensive or even necessarily consistent with future versions of SBO. The web site for SBO (<http://biomodels.net/SBO/>) should be consulted for the current version of the ontology. Section 5.4.1 describes how the impact of SBO changes on software applications is minimized.




Figure 21: The six controlled vocabularies (CVs) that make up the main branches of SBO.

Figure 22 shows the structure for the *entity* branch, which reflects the hierarchical groupings of the types of entities that can be represented by a **compartment** or a **species**. Note that the values taken by the `sboTerm` attribute on those elements should refer to SBO terms belonging to the *material entity* branch, so as to distinguish whether the element represents a macromolecule, a simple chemical, etc. Indeed, this information remains valid for the whole model. The term should not belong to the *functional entity* branch, representing the function of the entity within a certain functional context. If one wants to use this information, one should refer to the SBO terms using a controlled RDF annotation instead (Section 6), carefully choosing the qualifiers (Section 6.5) to reflect the fact that a given **species**, for instance, can fulfill different functions within a given model (e.g., EGF receptor is a receptor and an enzyme).

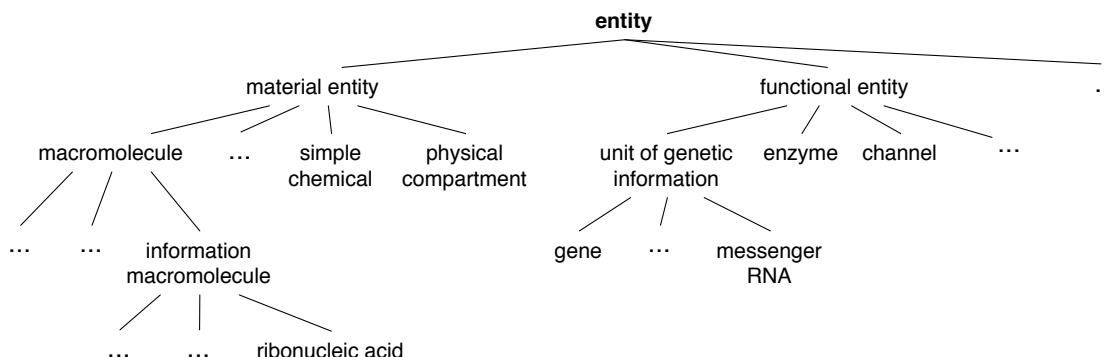


Figure 22: Partial expansion of some of the terms in the entity branch of SBO.

Figure 23 shows the structure for the *participant role* branch, also grouping the concepts in a hierarchical manner. For example, in reaction rate expressions, there are a variety of possible modifiers. Some classes of modifiers can be further subdivided and grouped. All of this is easy to capture in the ontology. As more agreement is reached in the modeling community about how to define and name modifiers for different cases, the ontology can grow to accommodate it.




Figure 23: Partial expansion of some of the terms in the participant role branch of SBO.

The controlled vocabulary for quantitative parameters is illustrated in Figure 24. Note the separation of *kinetic constant* into separate terms for unimolecular, bimolecular, etc. reactions, as well as for forward and reverse reactions. The need to have separate terms for forward and reverse rate constants arises in reversible mass-action reactions. This distinction is not always necessary for all quantitative parameters; for example, there is no comparable concept for the Michaelis constant. Another distinction for some quantitative parameters is a decomposition into different versions based on the modeling framework being assumed. For example, different terms for continuous and discrete formulations of kinetic constants represent specializations of the constants for particular simulation frameworks. Not all quantitative parameters will need to be distinguished along this dimension.

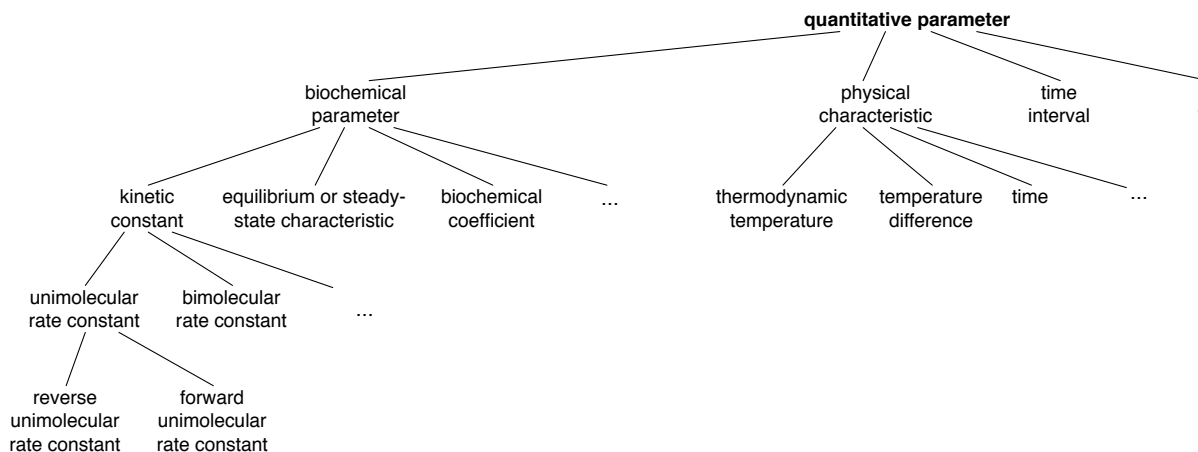


Figure 24: Partial expansion of some of the terms in the quantitative parameter branch.

The terms of the SBO quantitative parameter branch contain mathematical formulas encoded using MathML 2.0 expressing the parameter using other SBO parameters. The main use of that approach is to avoid listing all the variants of a mathematical expression, escaping a combinatorial explosion.

The *modeling framework* controlled vocabulary is needed to precise how to simulate a mathematical expression used in models. Figure 25 illustrates the structure of this branch, which is at this point extremely simple, but we expect that more terms will evolve in the future.




Figure 25: Partial expansion of some of the terms in the modeling framework branch.

The *mathematical expression* vocabulary encompasses the various mathematical expressions that constitute a model. Figure 26 on the following page illustrates a portion of the hierarchy. Rate law or conservation law formulas are part of the mathematical expression hierarchy, and subdivided by successively more refined distinctions until the leaf terms represent precise statements of common reaction or rule types. Other types of mathematical expressions may be included in the future in order to be able to further characterize mathematical components of a model, such as initial assignments, assignment rules, rate rules, algebraic rules, constraints, and event triggers and assignments.

The leaf terms of the mathematical expression branch contain the mathematical formulas encoded using MathML 2.0. There are many potential uses for this. One is to allow a software application to obtain the

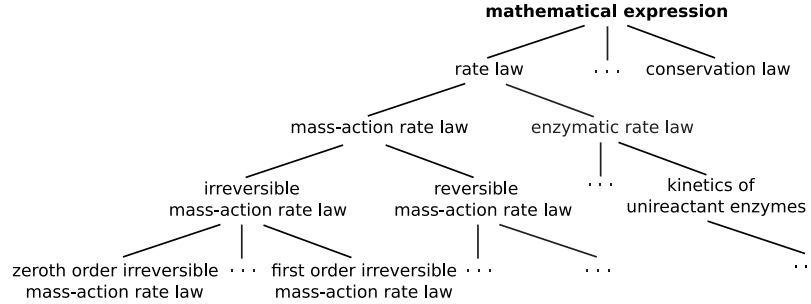


Figure 26: Partial expansion of some of the terms in the mathematical expression branch.

formula corresponding to a term and use it as the basis of an expression to insert into a model. In effect, the formulas given in the CV act as templates for what to put into an SBML construct such as [KineticLaw](#) or [Rule](#). The MathML definition also acts as a precise statement about the rate law in question. In particular, it carries information about the modeling framework to use in order to interpret the formula. Some of the non-leaf terms also contain formulas encoded using MathML 2.0. In that case, the formulas contained in the children terms are specific versions of the formula contained in the parent term. Those formulas may be generic, containing MathML constructs not yet supported by SBML, and need to be expanded into the MathML subset allowed in SBML before they can be used in conjunction with SBML models.

To make this discussion concrete, here is an example definition of an entry in the SBO rate law hierarchy at the time of this writing. This term represents second-order, irreversible, mass-action rate laws with one reactant, formulated for use in a continuous modeling framework:

ID: SBO:0000052

Name: second-order irreversible mass action rate law, one reactant, continuous scheme.

Definition: Reaction scheme where the products are created from the reactants and the change of a product quantity is proportional to the product of reactant activities. The reaction scheme does not include any reverse process that creates the reactants from the products. The change of a product quantity is proportional to the square of one reactant quantity. It is to be used in a reaction modelled using a continuous framework.

Parent(s):

- SBO:0000050 second order irreversible mass action rate law, one reactant (is a).
- SBO:0000163 irreversible mass action rate law, continuous scheme (is a).

MathML:

```

<math xmlns="http://www.w3.org/1998/Math/MathML">
  <semantics definitionURL="http://biomodels.net/SBO/#SBO:0000062">
    <lambda>
      <bvar><ci definitionURL="http://biomodels.net/SBO/#SBO:0000036">k</ci></bvar>
      <bvar><ci definitionURL="http://biomodels.net/SBO/#SBO:0000010">R</ci></bvar>
      <apply>
        <times/>
        <ci>k</ci>
        <ci>R</ci>
        <ci>R</ci>
      </apply>
    </lambda>
  </semantics>
</math>

```

In the MathML definition of the term shown above, the bound variables in the `lambda` expression are tagged with references to terms in the SBO quantitative parameter (for `k`) and SBO participant role (for `R`) branches. This makes it possible for software applications to interpret the intended meanings of the parameters in the expression. This also permits to convert an expression into another, by using the MathML 2.0 formula contained in the SBO terms associated with the parameters.

The *interaction* branch of SBO defines types of biological processes, events or relationship involving entities. It lists the types of biochemical reactions, such as binding, conformational transition, or cleavage, and also the different controls that modify a biochemical reaction, such as inhibition, catalysis, etc.

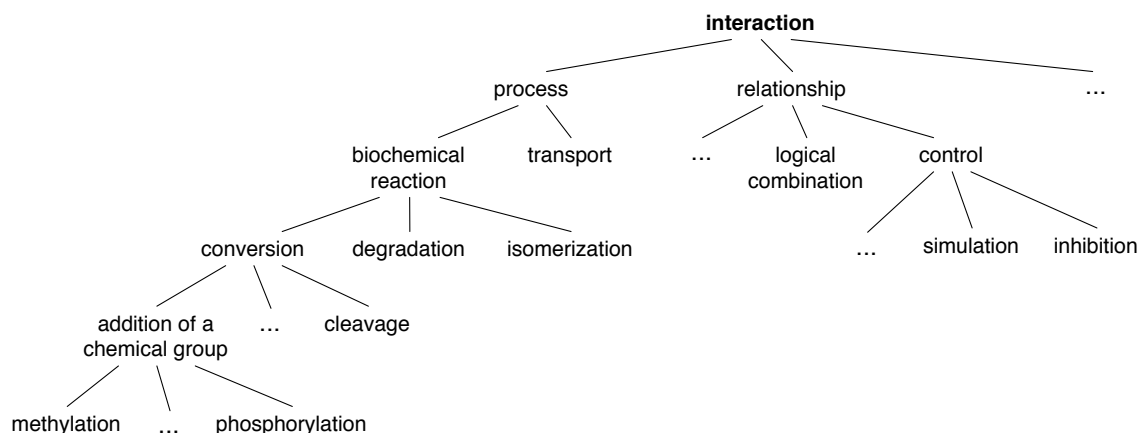


Figure 27: Partial expansion of some of the terms in the *interaction* branch.

One of the goals of SBO is to permit a tool to traverse up and down the hierarchy in order to find equivalent terms in different frameworks. The hope is that when a software tool encounters a given rate formula in a model, the formula will be a specific form (say, “mass-action rate law, second order, one reactant, for discrete simulation”), but by virtue of the consistent organization of the reaction rate CV into framework-specific definitions, and the declaration of every parameters involved in each expression, the tool should in principle be able to determine the definitions for other frameworks (say, “mass-action rate law, second order, one reactant for *continuous* simulation”). If the software tool is designed for continuous simulation and it encounters an SBML model with rate laws formulated for discrete simulation, it could in principle look up the rate laws’ identifiers in the CV and search for alternative definitions intended for discrete simulation. And of course, the converse is true, for when a tool designed for discrete simulation encounters a model with rate laws formulated for continuous simulation.

5.2.2 Relationships between individual SBML components and SBO terms

The `sboTerm` attribute is defined on the abstract class `SBase` and can be used in all derived elements. However, not all SBO terms should be used to annotate all SBML elements. Table 5 summarizes the relationships between SBML components and the branches within SBO that apply to that component. (There are currently no specific SBO term that correspond to the `Sbml`, `UnitDefinition`, `Unit`, and various `ListOf_____` list classes.)

The parent identifiers shown in Table 5 are provided for reference. They are the highest-level terms in their respective branch; however, these are *not* the terms that would be used to annotate an element in SBML, because there are more specific terms underneath the parents shown here. A software tool should use the most specific SBO term available for a given concept rather than using the top-level identifier acting as the root of that particular vocabulary.

5.2.3 Tradeoffs in using SBO terms

The SBO-based approach to annotating SBML components with controlled terms has the following strengths:

1. The syntax is minimally intrusive and maximally simple, requiring only one string-valued attribute.
2. It supports a significant fraction of what SBML users have wanted to do with controlled vocabularies.
3. It does not interfere with any other scheme. The more general annotation-based approach described in Section 6 can still be used simultaneously in the same model.

SBML Component	SBO Branch	Branch Identifier
Model	interaction	SBO:0000231
FunctionDefinition	mathematical expression	SBO:0000064
Compartment	material entity	SBO:0000240
Species	material entity	SBO:0000240
Reaction	interaction	SBO:0000231
Parameter	quantitative parameter	SBO:0000002
SpeciesReference	participant role	SBO:0000003
ModifierSpeciesReference	participant role	SBO:0000003
KineticLaw	rate law	SBO:0000001
LocalParameter	quantitative parameter	SBO:0000002
InitialAssignment	mathematical expression	SBO:0000064
AlgebraicRule	mathematical expression	SBO:0000064
AssignmentRule	mathematical expression	SBO:0000064
RateRule	mathematical expression	SBO:0000064
Constraint	mathematical expression	SBO:0000064
Event	interaction	SBO:0000231
Trigger	mathematical expression	SBO:0000064
Delay	mathematical expression	SBO:0000064
EventAssignment	mathematical expression	SBO:0000064

Table 5: SBML components and the main types of SBO terms that may be assigned to them. The identifiers of the highest-level SBO terms in each branch are provided for guidance, but actual values used for `sboTerm` attributes should be more specific child terms within these branches. Note that the important aspect here is the set of specific SBO identifiers, not the SBO term names, because the names may change as SBO continues to evolve. See text for further explanations.

The scheme has the following weaknesses:

1. An object can only have one `sboTerm` attribute, therefore, it can only be related to a single term in SBO. (This also impacts the design of SBO: it must be structured such that a class of SBML elements can logically only be associated with one class of terms in the ontology.)
2. The only relationship that can be expressed by `sboTerm` is “is a”. It is not possible to represent different relationships (known as *verbs* in ontology-speak). This limits what can be expressed using SBO.

The weaknesses are not shared by the annotation scheme described in Section 6.

5.3 Relationships to the SBML annotation element

Another way to provide this information would be to place SBO terms inside the **SBase annotation** element (Sections 3.2 and 6). However, in the interest of making the use of SBO in SBML as interoperable as possible between software tools, the best-practice recommendation is to place SBO references in the `sboTerm` attribute rather than inside the **annotation** element of an object. If instead the approach of using **annotation** is taken, the qualifiers (Section 6.5) linking the SBML element and SBO term should be chosen extremely carefully, since it will no longer be possible to assume an “instance to class” relationship.

Although `sboTerm` is just another kind of optional annotation in SBML, SBO references are separated into their own attribute on SBML components, both to simplify their use for software tools and because doing so asserts a stronger and more focused connection in a more regimented fashion. SBO references are intended to allow a modeler to make a statement of the form “this object is identical in meaning and intention to the object defined in the term X of SBO”, and do so in a way that a *software tool can interpret unambiguously*.

Some software applications may have their own vocabulary of terms similar in purpose to SBO. For maximal software interoperability, the best-practice recommendation in SBML is nonetheless to use SBO terms in preference to using application-specific annotation schemes. Software applications should therefore attempt to translate their private terms to and from SBO terms when writing and reading SBML, respectively.

5.4 Discussion

Here we discuss some additional points about the SBO-based approach.

5.4.1 *Frequency of change in the ontology*

The SBO development approach follows conventional ontology development approaches in bioinformatics. One of the principles being followed is that identifiers and meanings of terms in the CVs never change and the terms are never deleted. Where some terms are deemed obsolete, the introduction of new terms refine or supersede existing terms, but the existing identifiers are left in the CV. Thus, references never end up pointing to nonexistent entries. In the case where synonymous terms are merged after agreement that multiple terms are identical, the term identifiers are again left in the CV and they still refer to the same concept as before. Out-of-date terms cached or hard-coded by an application remain usable in all cases. (Moreover, machine-readable CV encodings and appropriate software design should render possible the development of API libraries that automatically map older terms to newer terms as the CVs evolve.)

Therefore, a model is never in danger of ending up with SBO identifiers that cannot be dereferenced. If an application finds an old model with a term **SBO:0000065**, it can be assured that it will be able to find this term in SBO, even if it has been superseded by other, more preferred terms.

5.4.2 *Consistency of information*

If you have a means of linking (say) a reaction rate formula to a term in a CV, it is possible to have an inconsistency between the formula in the SBML model and the one defined for the CV term. However, this is not a new problem; it arises in other situations involving SBML models already. The guideline for these situations is that the model must be self-contained and stand on its own. Therefore, in cases where they differ, the definitions in the SBML model take precedence over the definitions referenced by the CV. In other words, the model (and its MathML) is authoritative.

5.4.3 *Implications for network access*

A software tool does not need to have the ability to access the network or read the CV every time it encounters a model or otherwise works with SBML. Since the SBO will likely stabilize and change infrequently once a core set of terms is defined, applications can cache the controlled vocabulary, and not make network accesses to the master SBO copy unless something forces them to (e.g., detecting a reference in a model to an SBO term that the application does not recognize). Applications could have user preference settings indicating how often the CV definitions should be refreshed (similar to how modern applications provide a setting dictating how often they should check for new versions of themselves). Simple applications may go further and hard code references to terms in SBO that have reached stability and community consensus. SBO is available for download under different formats (<http://biomodels.net/SBO/>). Web services are also available to provide programmatic access to the ontology.

5.4.4 *Implications for software tools*

If a software tool does not pay attention to the SBO annotations described here, one is faced with exactly the situation that exists today: the SBML model must be interpreted as-is, without benefit of the information added by the SBO terms. The purpose of introducing an ontology scheme and guidelines for its use is to give tools enough information that they *could* perform added processing, if they were designed to take advantage of that information.

6 A standard format for the annotation element

This section describes the recommended non-proprietary format for the content of **annotation** elements in SBML when (a) referring to controlled vocabulary terms and database identifiers which define and describe biological and biochemical entities, and (b) describing the creator of a model and its modification history. Such a structured format should facilitate the generation of models compliant with the MIRIAM guidelines for model curation (Le Novère et al., 2005).

The format described in this section is intended to be the form of one of the top-level elements that could reside in an **annotation** element attached to an SBML object derived from **SBase**. The element is named **rdf:RDF**. The format described here is compliant with the constraints placed on the form of annotation elements described in Section 3.2.4. We refer readers to Section 3.2.4 for important information on the structure and organization of application-specific annotations; these are not described here.

The annotations described in this section are optional on a model, but if present, they must conform to the details specified here in order to be considered valid annotations in this format. If they do not conform to the format described here, it does not render the overall SBML model invalid, but the annotations are then considered to be in a proprietary format rather than being *SBML MIRIAM annotations*.

6.1 Motivation

The SBML structures described elsewhere in this document do not have any biochemical or biological semantics. This section provides a scheme for linking SBML structures to external resources so that those structures can be given semantics. The motivation for the introduction of this scheme is similar to that given for the introduction of **sboTerm**; however, the general annotation scheme here is more flexible.

It is generally not recommended that this format be used to refer to SBO terms. In most cases, the SBO terms should be referred to using the attribute **sboTerm** part of **SBase** (Section 5). However in certain situations, for instance to be able to add further information about the functional role of a species, it is necessary to add this additional information using the annotation format described here.

6.2 XML namespaces in the standard annotation

This format uses a restricted form of Dublin Core (Dublin Core Metadata Initiative, 2005) and BioModels qualifier elements (see <http://sbml.org/miriam/qualifiers/>) embedded in the XML form of RDF (W3C, 2004b). The scheme defined here uses a number of external XML standards and associated XML namespaces. Table 6 lists these namespaces and relevant documentation on those namespaces. The format constrains the order of elements in these namespaces beyond the constraints defined in the standard definitions for those namespaces. For each standard listed, the format only uses a subset of the possible syntax defined by the given standard. Thus, it is possible for an **annotation** element to include XML that is compliant with those external standards but is not compliant with the format described here.

Prefix used in examples here	Namespace URI	Reference/description
dc	http://purl.org/dc/elements/1.1/	Powell and Johnston (2003)
rdf	http://www.w3.org/1999/02/22-rdf-syntax-ns#	W3C (2004a)
dcterms	http://purl.org/dc/terms/	Kokkelink and Schwänzl (2002), DCMI Usage Board (2005)
vcard	http://www.w3.org/2001/vcard-rdf/3.0#	Iannella (2001)
bqbiol	http://biomodels.net/biology-qualifiers/	http://sbml.org/miriam/qualifiers/
bqmodel	http://biomodels.net/model-qualifiers/	http://sbml.org/miriam/qualifiers/

Table 6: The XML standards used in the SBML standard format for annotation. The namespace prefix are shown to indicate only the prefix used in the main text.

6.3 General syntax for the standard annotation

An outline of the format syntax is shown below.

```
<SBML_ELEMENT +++ metaid="SBML_META_ID" +++ >
+++
<annotation>
+++
  <rdf:RDF xmlns:rdf="http://www.w3.org/1999/02/22-rdf-syntax-ns#"
            xmlns:dc="http://purl.org/dc/elements/1.1/"
            xmlns:dcterms="http://purl.org/dc/terms/"
            xmlns:vcard="http://www.w3.org/2001/vcard-rdf/3.0#"
            xmlns:bqbiol="http://biomodels.net/biology-qualifiers/"
            xmlns:bqmodel="http://biomodels.net/model-qualifiers/"
  >
    <rdf:Description rdf:about="#SBML_META_ID">
      [HISTORY]
      <RELATION_ELEMENT>
        <rdf:Bag>
          <rdf:li rdf:resource="URI" />
          ...
        </rdf:Bag>
      </RELATION_ELEMENT>
      ...
    </rdf:Description>
  </rdf:RDF>
</annotation>
+++
</SBML_ELEMENT>
```

The above outline shows the order of the elements. The capitalized identifiers refer to generic strings of a particular type: **SBML_ELEMENT** refers to any SBML element name that can contain an **annotation** element; **SBML_META_ID** is a XML ID string; **RELATION_ELEMENT** refers to element names in either the namespace <http://biomodels.net/biology-qualifiers/> or <http://biomodels.net/model-qualifiers/>; and **URI** is a URI (See Section 6.4). **[HISTORY]** refers to an optional section described in Section 6.6. ‘+++’ is a placeholder for either no content or valid XML syntax that is not defined by the standard annotation scheme but is consistent with the relevant standards for the enclosing elements. ‘...’ is a placeholder for zero or more elements of the same form as the immediately preceding element. The precise form of whitespace and the XML namespace prefix definitions is not constrained; however, the elements and attributes must be in the namespaces shown. The rest of this section describes the format formally in English.

In this format, the annotation of an element is located in a single **rdf:RDF** element contained within an SBML **annotation** element. The annotation element can contain other elements in any order as described in Section 3.2.4. The format described in this section only defines the form of the **rdf:RDF** element. The containing SBML **SBase** element must have a **metaid** attribute value. (As this attribute is of the type **ID** its value must be unique to the entire SBML document.)

The first element of the **rdf:RDF** element must be an **rdf:Description** element with an **rdf:about** attribute. The value of the **rdf:about** attribute must be of the form **#<string>** where the string component is equal to the value of the **metaid** attribute of the containing SBML element. This format doesn’t define the form of subsequent subelements of the **rdf:RDF** element. In particular, the unique **rdf:RDF** element contained in the annotation can contain other **rdf:Description**, which content can be any valid RDF.

The **rdf:Description** element can contain only an optional history section (see Section 6.6) followed by a sequence of zero or more BioModels relation elements. The specific type of the relation elements will vary depending on the relationship between the SBML component and referenced information or resource.

Although Section 6.5 describes the detailed semantics of each of the relation element types, the content of these elements follows the same form. The BioModels qualifiers relation elements must only contain a single **rdf:Bag** element which in turn must only contain one or more **rdf:li** elements. The **rdf:li** elements must only have a **rdf:resource** attribute containing a URI referring to an information resource (See Section 6.4).

6.4 Use of URIs

The format represents a set of relationships between SBML elements and resources referred to by **rdf:resource** attribute values. The BioModels relation elements simply define the relationship.

For example, a **Species** element representing a protein could be annotated with a reference to the database UniProt by the **urn:miriam:uniprot:P12999** resource identifier, identifying exactly the protein described by the **Species** element. This identifier maps to a unique entry in UniProt which is never deleted from the database. In the case of UniProt, this is the “accession” of the entry. When the entry is merged with another one, both “accession” are conserved. Similarly in a controlled vocabulary resource, each term is associated with a perennial identifier. The UniProt entry also possess an “entry name” (the Swiss-Prot “identifier”), a “protein name”, “synonyms” etc. Only the “accession” is perennial and should be used.

The value of a **rdf:resource** attribute is a URI that both uniquely identifies the resource, and the data in the resource. The value of the **rdf:resource** attribute is a URI, not a URL; as such, a URI does not have to reference a physical web object but simply identifies a controlled vocabulary term or database object (a URI is a label). For instance, a true URL for an Internet resource such as <http://www.uniprot.org/entry/P12999> might correspond to the URI **urn:miriam:uniprot:P12999**.

SBML does not specify how a parser is to interpret a URI. In the case of a transformation into a physical URL, there could be several solutions. For example, the URI **urn:miriam:obo.go:GO%3A0007268** can be translated into:

```
http://www.ebi.ac.uk/ego/DisplayGoTerm?selected=GO:0007268  
http://www.godatabase.org/cgi-bin/amigo/go.cgi?view=details&query=GO:0007268  
http://www.informatics.jax.org/searches/GO.cgi?id=GO:0007268
```

Similarly the URI **urn:miriam:ec-code:3.5.4.4** can refer to:

```
http://www.ebi.ac.uk/intenz/query?cmd=SearchEC&ec=3.5.4.4  
http://www.expasy.org/cgi-bin/nicezyme.pl?3.5.4.4  
http://www.chem.qmul.ac.uk/iubmb/enzyme/EC3/5/4/4.html  
http://www.genome.jp/dbget-bin/www.bget?ec:3.5.4.4
```

etc.

To enable interoperability, the community has agreed on an initial set of standardized valid URI syntax rules which may be used within the standard annotation format. This set of rules is not part of the SBML standard but will grow independently from specific SBML Levels and Versions. As the set changes, a given URI syntax rule will not be modified, although the physical resources associated with the rule may change. These URIs will always be composed as **resource:id**. An up-to-date list and explanation of the URIs is always available online at <http://sbml.org/miriam/qualifiers>. Each entry contains a list of SBML and relation elements in which the given URI can be appropriately embedded. To enable consistent and thus useful links to external resources, the URI syntax rule set must have a consistent view of the concepts represented by the different SBML elements for the purposes of this format. For example, as the rule set is designed to link SBML biological and biochemical resources the rule set assumes that a **Species** element represents the concept of a biochemical entity type rather than mathematical symbol. The URI rule list will evolve with the evolution of databases and resources.

Note that this means that all **rdf:resource** *must* be MIRIAM URIs and thus cannot refer to for example other elements in the model. While it would be possible to place such information in other RDF content elsewhere (for example after the first **rdf:Description** element), this would place this information outside the simple annotation scheme described here, and as such there would be no guarantee that other software programs could read this information.

6.5 Relation elements

Different BioModels qualifier elements encode different types of relationships. When appearing in an annotation, each qualifier element encloses a set of **rdf:li** elements. Its appearance in a relation element implies

a specific relationship between the enclosing SBML object and the resources referenced by the `rdf:li` elements. The detailed semantics of each qualifier is defined online at <http://sbml.org/miriam/qualifiers/>. Generally, each qualifier is defined with the assumption that the biological entity represented by a given SBML element is the concept linked to the reference resource.

Several relation elements with a given tag, enclosed in the same SBML element, each represent an alternative annotation to the SBML element. For example two `bqbiol:hasPart` elements within a `Species` SBML element represent two different sets of references to the parts making up the the chemical entity represented by the species. (The species is not made up of all the entities represented by all the references combined).

The complete list of the qualifier elements in the BioModels qualifier namespaces is documented online at <http://sbml.org/miriam/qualifiers/>. The list is divided into two symbol namespaces, one for model qualifiers; this has the URI <http://biomodels.net/model-qualifiers/> (and we use the prefix `bqmodel` in examples shown in this section). The other is for biological qualifiers; this has the URI <http://biomodels.net/biology-qualifiers/> (and we use the prefix `bqbiol`). The list will only grow; i.e., no element will be removed from the list. The following is the list of elements at the time of writing:

- `bqmodel:is` The modeling object encoded by the SBML component is the subject of the referenced resource. For instance, this qualifier might be used to link the model to a model database.
- `bqmodel:isDescribedBy` The modeling object encoded by the SBML component is described by the referenced resource. This relation might be used to link SBML components to the literature that describes this model or this kinetic law.
- `bqbiol:is` The biological entity represented by the SBML component is the subject of the referenced resource. This relation might be used to link a reaction to its exact counterpart in (e.g.) KEGG or Reactome.
- `bqbiol:hasPart` The biological entity represented by the SBML component includes the subject of the referenced resource, either physically or logically. This relation might be used to link a complex to the description of its components.
- `bqbiol:isPartOf` The biological entity represented by the SBML component is a physical or logical part of the subject of the referenced resource. This relation might be used to link a component to the description of the complex it belongs to.
- `bqbiol:isVersionOf` The biological entity represented by the SBML component is a version or an instance of the subject of the referenced resource.
- `bqbiol:hasVersion` The subject of the referenced resource is a version or an instance of the biological entity represented by the SBML component.
- `bqbiol:isHomologTo` The biological entity represented by the SBML component is homolog, to the subject of the referenced resource, i.e. they share a common ancestor.
- `bqbiol:isDescribedBy` The biological entity represented by the SBML component is described by the referenced resource. This relation should be used, for example, to link a species or a parameter to the literature that describes the quantity of the species or the value of the parameter.
- `bqbiol:isEncodedBy` The biological entity represented by the model component is encoded, either directly or by virtue of transitivity, by the subject of the referenced resource.
- `bqbiol:encodes` The biological entity represented by the model component encodes, either directly or by virtue of transitivity, the subject of the referenced resource.
- `bqbiol:occursIn` The biological entity represented by the model component takes place in the subject of the reference resource.

6.6 History

The format described in previous sections can include additional elements to describe the history of the model or a portion of the model. If this history data is present, it must occur immediately before the first BioModels relation elements. These additional elements encode information on the model creator and a sequence of dates recording changes to the model. The syntax for this section is outlined below.

```
<dc:creator>
  <rdf:Bag>
    <rdf:li rdf:parseType="Resource">
      [[
        +++
        <vCard:N rdf:parseType="Resource">
          <vCard:Family>FAMILY_NAME</vCard:Family>
          <vCard:Given>GIVEN_NAME</vCard:Given>
        </vCard:N>
        +++
        [<vCard:EMAIL>EMAIL_ADDRESS</vCard:EMAIL>]
        +++
        [<vCard:ORG rdf:parseType="Resource" >
          <vCard:Orgname>ORGANIZATION_NAME</vCard:Orgname>
        </vCard:ORG>]
        +++
      ]]
    </rdf:li>
    ...
  </rdf:Bag>
</dc:creator>
<dcterms:created rdf:parseType="Resource">
  <dcterms:W3CDTF>DATE</dcterms:W3CDTF>
</dcterms:created>
<dcterms:modified rdf:parseType="Resource">
  <dcterms:W3CDTF>DATE</dcterms:W3CDTF>
</dcterms:modified>
...
```

The order of elements is as shown above, except that elements of the format contained between `[[` and `]]` can occur in any order. The capitalized identifiers refer to generic strings of a particular type: **FAMILY_NAME** is the family name of a person who created the model; **GIVEN_NAME** is the first name of the same person who created the model; **EMAIL_ADDRESS** is the email address of the same person who created the model; and **ORGANIZATION_NAME** is the name of the organization with which the same person who created the model is affiliated. **DATE** is a date in W3C date format (Wolf and Wicksteed, 1998). **W3CDTF**, **N**, **ORG** and **EMAIL** are literal strings. The elements of the format contained between `[` and `]` are optional. `+++` is a placeholder for either no content or valid XML syntax that is not defined by this scheme but is consistent with the relevant standards for the enclosing elements. Ellipses (`...`) are placeholders for zero or more elements of the same form as the immediately preceding element. The precise form of whitespace and the XML namespace prefix definitions is not constrained. The remaining text in this section describes the syntax formally in English.

The additional elements of the history sub-format consist in sequence of a **dc:creator** element, a **dcterms:created** element and zero or more **dcterms:modified** elements. The last two elements must have the attribute **rdf:parseType** set to **Resource**.

The **dc:creator** element describes the person who created the SBML encoding of the model and contains a single **rdf:Bag** element. The **rdf:Bag** element can contain any number of elements; however, the first element must be a **rdf:li** element. The **rdf:li** element can contain any number of elements in any order. The set of elements contained with the **rdf:li** element can include the following informative elements: **vCard:N**, **vCard:EMAIL** and **vCard:ORG**. The **vCard:N** contains the name of the creator and must consist of a sequence of two elements: **vCard:Family** and the **vCard:Given** whose content is the family (surname) and given (first) names of the creator respectively. The **vCard:N** must have the attribute **rdf:parseType** set to **Resource**. The content of the **vCard:EMAIL** element must be the email address of the creator. The content of the **vCard:ORG** element must contain a single **vCard:Orgname** element. The **vCard:Orgname** element must contain the name of an organization to which the creator is affiliated.

The `dcterms:created` and `dcterms:modified` elements must each contain a single `dcterms:W3CDTF` element whose content is a date in W3C date format (Wolf and Wicksteed, 1998) which is a a profile of (restricted form of) ISO 8601.

6.7 Examples

The following shows the annotation of a model with model creation data and links to external resources:

```
<model metaid="_180340" id="GMO" name="Goldbeter1991_MinMitOscil">
  <annotation>
    <rdf:RDF
      xmlns:rdf="http://www.w3.org/1999/02/22-rdf-syntax-ns#"
      xmlns:dc="http://purl.org/dc/elements/1.1/"
      xmlns:dcterms="http://purl.org/dc/terms/"
      xmlns:vCard="http://www.w3.org/2001/vcard-rdf/3.0#"
      xmlns:bqbiol="http://biomodels.net/biology-qualifiers/"
      xmlns:bqmodel="http://biomodels.net/model-qualifiers/"
    >
      <rdf:Description rdf:about="#_180340">
        <dc:creator>
          <rdf:Bag>
            <rdf:li rdf:parseType="Resource">
              <vCard:N rdf:parseType="Resource">
                <vCard:Family>Shapiro</vCard:Family>
                <vCard:Given>Bruce</vCard:Given>
              </vCard:N>
              <vCard:EMAIL>bshapiro@jpl.nasa.gov</vCard:EMAIL>
              <vCard:ORG rdf:parseType="Resource">
                <vCard:Orgname>NASA Jet Propulsion Laboratory</vCard:Orgname>
              </vCard:ORG>
            </rdf:li>
          </rdf:Bag>
        </dc:creator>
        <dcterms:created rdf:parseType="Resource">
          <dcterms:W3CDTF>2005-02-06T23:39:40+00:00</dcterms:W3CDTF>
        </dcterms:created>
        <dcterms:modified rdf:parseType="Resource">
          <dcterms:W3CDTF>2005-09-13T13:24:56+00:00</dcterms:W3CDTF>
        </dcterms:modified>
        <bqmodel:is>
          <rdf:Bag>
            <rdf:li rdf:resource="urn:miriam:biomodels.db:BIOMD0000000003"/>
          </rdf:Bag>
        </bqmodel:is>
        <bqmodel:isDescribedBy>
          <rdf:Bag>
            <rdf:li rdf:resource="urn:miriam:pubmed:1833774"/>
          </rdf:Bag>
        </bqmodel:isDescribedBy>
        <bqbiol:isVersionOf>
          <rdf:Bag>
            <rdf:li rdf:resource="urn:miriam:kegg.pathway:hsa04110"/>
            <rdf:li rdf:resource="urn:miriam:reactome:REACT_152"/>
          </rdf:Bag>
        </bqbiol:isVersionOf>
      </rdf:Description>
    </rdf:RDF>
  </annotation>
</model>
```

The following example shows a [Reaction](#) structure annotated with a reference to its exact Reactome counterpart.

```
<reaction id="cdc2Phospho" metaid="jb007" reversible="true" fast="false">
  <annotation>
    <rdf:RDF
      xmlns:bqbiol="http://biomodels.net/biology-qualifiers/"
      xmlns:rdf="http://www.w3.org/1999/02/22-rdf-syntax-ns#"
    >
```

```

1      >
2      <rdf:Description rdf:about="#jb007">
3          <bqbiol:is>
4              <rdf:Bag>
5                  <rdf:li rdf:resource="urn:miriam:reactome:REACT_6327"/>
6              </rdf:Bag>
7          </bqbiol:is>
8      </rdf:Description>
9  </rdf:RDF>
10 </annotation>
11 <listOfReactants>
12   <speciesReference species="cdc2" stoichiometry="1"/>
13 </listOfReactants>
14 <listOfProducts>
15   <speciesReference species="cdc2-Y15P" stoichiometry="1"/>
16 </listOfProducts>
17 <listOfModifiers>
18   <modifierSpeciesReference species="wee1"/>
19 </listOfModifiers>
20 </reaction>

```

The following example describes a species that represents a complex between the protein calmodulin and calcium ions:

```

23 <species id="Ca_calmodulin" metaid="cacam" compartment="C"
24   hasOnlySubstanceUnits="false" boundaryCondition="false"
25   constant="false">
26   <annotation>
27     <rdf:RDF
28       xmlns:rdf="http://www.w3.org/1999/02/22-rdf-syntax-ns#"
29       xmlns:bqbiol="http://biomodels.net/biology-qualifiers/"
30     >
31       <rdf:Description rdf:about="#cacam">
32         <bqbiol:hasPart>
33           <rdf:Bag>
34             <rdf:li rdf:resource="urn:miriam:uniprot:P62158"/>
35             <rdf:li rdf:resource="urn:miriam:kegg.compound:C00076"/>
36           </rdf:Bag>
37         </bqbiol:hasPart>
38       </rdf:Description>
39     </rdf:RDF>
40   </annotation>
41 </species>

```

The following example describes a species that represents either “Calcium/calmodulin-dependent protein kinase type II alpha chain” or “Calcium/calmodulin-dependent protein kinase type II beta chain”. This is the case, for example, in the somatic cytoplasm of striatal medium-size spiny neurons, where both are present but they cannot be functionally differentiated.

```

46 <species id="calcium_calmodulin" metaid="cacam" compartment="C"
47   hasOnlySubstanceUnits="false" boundaryCondition="false"
48   constant="false">
49   <annotation>
50     <rdf:RDF
51       xmlns:rdf="http://www.w3.org/1999/02/22-rdf-syntax-ns#"
52       xmlns:bqbiol="http://biomodels.net/biology-qualifiers/"
53     >
54       <rdf:Description rdf:about="#cacam">
55         <bqbiol:hasVersion>
56           <rdf:Bag>
57             <rdf:li rdf:resource="urn:miriam:uniprot:Q9UQM7"/>
58             <rdf:li rdf:resource="urn:miriam:uniprot:Q13554"/>
59           </rdf:Bag>
60         </bqbiol:hasVersion>
61       </rdf:Description>
62     </rdf:RDF>
63   </annotation>
64 </species>

```


The above approach should not be used to describe “any Calcium/calmodulin-dependent protein kinase type II chain”, because such an annotation requires referencing the products of other genes such as gamma or delta. All the known proteins could be enumerated, but such an approach would almost surely lead to inaccuracies because biological knowledge continues to evolve. Instead, the annotation should refer to generic information such as Ensembl family ENSF00000000194 “CALCIUM/CALMODULIN DEPENDENT KINASE TYPE II CHAIN” or PIR superfamily PIRSF000594 “Calcium/calmodulin-dependent protein kinase type II”.

The following two examples show how to use the qualifier `isVersionOf`. The first example is the relationship between a reaction and an EC code. An EC code describes an enzymatic activity and an enzymatic reaction involving a particular enzyme can be seen as an instance of this activity. For example, the following reaction represents the phosphorylation of a glutamate receptor by a complex calcium/calmodulin kinase II.

```
<reaction id="NMDAR_phosphorylation" metaid="thx1138"
  reversible="true" fast="false">
  <annotation>
    <rdf:RDF
      xmlns:bqbiol="http://biomodels.net/biology-qualifiers/"
      xmlns:rdf="http://www.w3.org/1999/02/22-rdf-syntax-ns#"
    >
      <rdf:Description rdf:about="#thx1138">
        <bqbiol:isVersionOf>
          <rdf:Bag>
            <rdf:li rdf:resource="urn:miriam:ec-code:2.7.1.17"/>
          </rdf:Bag>
        </bqbiol:isVersionOf>
      </rdf:Description>
    </rdf:RDF>
  </annotation>
  <listOfReactants>
    <speciesReference species="NMDAR" stoichiometry="1"/>
  </listOfReactants>
  <listOfProducts>
    <speciesReference species="P-NMDAR" stoichiometry="1"/>
  </listOfProducts>
  <listOfModifiers>
    <modifierSpeciesReference species="CaMKII"/>
  </listOfModifiers>
  <kineticLaw>
    <math xmlns="http://www.w3.org/1998/Math/MathML">
      <apply>
        <times/>
        <ci>CaMKII</ci>
        <ci>kcat</ci>
      </apply>
      <divide/>
      <ci>NMDAR</ci>
      <apply> </times> <ci>NMDAR</ci> <ci>Km</ci> </apply>
    </apply>
  </math>
  <listOfLocalParameters>
    <localParameter id="kcat" value="1"/>
    <localParameter id="Km" value="5e-10"/>
  </listOfLocalParameters>
</kineticLaw>
</reaction>
```

The second example of the use of `isVersionOf` is the complex between Calcium/calmodulin-dependent protein kinase type II alpha chain and Calcium/calmodulin, that is only one of the “calcium- and calmodulin-dependent protein kinase complexes” described by the Gene Ontology term GO:0005954. (Note also how the GO identifier is written—we return to this point below.)

```
<species id="CaCaMKII" metaid="C8H10N4O2" compartment="C"
  hasOnlySubstanceUnits="false" boundaryCondition="false"
  constant="false">
  <annotation>
```

```

1      <rdf:RDF
2        xmlns:rdf="http://www.w3.org/1999/02/22-rdf-syntax-ns#"
3        xmlns:bqbiol="http://biomodels.net/biology-qualifiers/"
4      >
5        <rdf:Description rdf:about="#C8H10N4O2">
6          <bqbiol:isVersionOf>
7            <rdf:Bag>
8              <rdf:li rdf:resource="urn:miriam:obo.go:GO%3A0005954"/>
9            </rdf:Bag>
10          </bqbiol:isVersionOf>
11        </rdf:Description>
12      </rdf:RDF>
13    </annotation>
14  </species>

```

In the example above, the URN for the GO term is written as `GO%3A0005954`, yet in reality, the actual GO identifier is `GO:0005954`. The reason for this rests in the definition of RDF/XML and URNs. The essential point is that the colon character (“:”) is a reserved character representing the component separator in URNs. Thus, when an identifier contains a colon character as part of it (as GO, ChEBI, and certain other identifiers do), the colon characters must be percent-encoded. The sequence “%3A” is the percent-encoded form of “:”. Applications must percent-encode “:” characters that appear in entity identifiers (whether from ECO, ChEBI, GO, or other) when writing them in MIRIAM URIs, and percent-decode the identifiers when reading the URIs. More examples of this appear throughout the rest of this section.

The previous case is different from the following one, although they could seem similar at first sight. The “Calcium/calmodulin-dependent protein kinase type II alpha chain” is a part of the above mentioned “calcium- and calmodulin-dependent protein kinase complex”.

```

26    <species id="CaMKIIalpha" metaid="C10H14N2" compartment="C"
27      hasOnlySubstanceUnits="false" boundaryCondition="false"
28      constant="false">>
29    <annotation>
30      <rdf:RDF
31        xmlns:rdf="http://www.w3.org/1999/02/22-rdf-syntax-ns#"
32        xmlns:bqbiol="http://biomodels.net/biology-qualifiers/"
33      >
34        <rdf:Description rdf:about="#C10H14N2">
35          <bqbiol:isPartOf>
36            <rdf:Bag>
37              <rdf:li rdf:resource="urn:miriam:obo.go:GO%3A0005954"/>
38            </rdf:Bag>
39          </bqbiol:isPartOf>
40        </rdf:Description>
41      </rdf:RDF>
42    </annotation>
43  </species>

```

It is possible to describe a component with several alternative sets of qualified annotations. For example, the following species represents a pool of GMP, GDP and GTP. We annotate it with the three corresponding KEGG compound identifiers but also with the three corresponding ChEBI identifiers. The two alternative annotations are encoded in separate `hasVersion` qualifier elements.

```

48    <species id="GXP" metaid="GXP" compartment="C"
49      hasOnlySubstanceUnits="false" boundaryCondition="false"
50      constant="false">>
51    <annotation>
52      <rdf:RDF
53        xmlns:rdf="http://www.w3.org/1999/02/22-rdf-syntax-ns#"
54        xmlns:bqbiol="http://biomodels.net/biology-qualifiers/"
55      >
56        <rdf:Description rdf:about="#GXP">
57          <bqbiol:hasVersion>
58            <rdf:Bag>
59              <rdf:li rdf:resource="urn:miriam:obo.chebi:CHEBI%3A17345"/>
60              <rdf:li rdf:resource="urn:miriam:obo.chebi:CHEBI%3A17552"/>
61              <rdf:li rdf:resource="urn:miriam:obo.chebi:CHEBI%3A17627"/>

```

```

1         </rdf:Bag>
2     </bqbiol:hasVersion>
3     <bqbiol:hasVersion>
4         <rdf:Bag>
5             <rdf:li rdf:resource="urn:miriam:kegg.compound:C00035"/>
6             <rdf:li rdf:resource="urn:miriam:kegg.compound:C00044"/>
7             <rdf:li rdf:resource="urn:miriam:kegg.compound:C00144"/>
8         </rdf:Bag>
9     </bqbiol:hasVersion>
10 </rdf:Description>
11 </rdf:RDF>
12 </annotation>
13 </species>

```

The following example presents a reaction being actually the combination of three different elementary molecular reactions. We annotate it with the three corresponding KEGG reactions, but also with the three corresponding enzymatic activities. Again the two **hasPart** elements represent two alternative annotations. The process represented by the **Reaction** structure is composed of three parts, and not six parts.

```

18 <reaction id="adenineProd" metaid="adeprod" reversible="true" fast="false">
19   <annotation>
20     <rdf:RDF
21       xmlns:bqbiol="http://biomodels.net/biology-qualifiers/"
22       xmlns:rdf="http://www.w3.org/1999/02/22-rdf-syntax-ns#"
23     >
24       <rdf:Description rdf:about="#adeprod">
25         <bqbiol:hasPart>
26           <rdf:Bag>
27             <rdf:li rdf:resource="urn:miriam:ec-code:2.5.1.22"/>
28             <rdf:li rdf:resource="urn:miriam:ec-code:3.2.2.16"/>
29             <rdf:li rdf:resource="urn:miriam:ec-code:4.1.1.50"/>
30           </rdf:Bag>
31         </bqbiol:hasPart>
32         <bqbiol:hasPart>
33           <rdf:Bag>
34             <rdf:li rdf:resource="urn:miriam:kegg.reaction:R00178"/>
35             <rdf:li rdf:resource="urn:miriam:kegg.reaction:R01401"/>
36             <rdf:li rdf:resource="urn:miriam:kegg.reaction:R02869"/>
37           </rdf:Bag>
38         </bqbiol:hasPart>
39       </rdf:Description>
40     </rdf:RDF>
41   </annotation>
42 </reaction>

```

It is possible to mix different URIs in a given set. The following example presents two alternative annotations of the human hemoglobin, the first with ChEBI heme and the second with KEGG heme.

```

45 <species id="heme" metaid="heme" compartment="C"
46   hasOnlySubstanceUnits="false" boundaryCondition="false"
47   constant="false">
48   <annotation>
49     <rdf:RDF
50       xmlns:rdf="http://www.w3.org/1999/02/22-rdf-syntax-ns#"
51       xmlns:bqbiol="http://biomodels.net/biology-qualifiers/"
52     >
53       <rdf:Description rdf:about="#heme">
54         <bqbiol:hasPart>
55           <rdf:Bag>
56             <rdf:li rdf:resource="urn:miriam:uniprot:P69905"/>
57             <rdf:li rdf:resource="urn:miriam:uniprot:P68871"/>
58             <rdf:li rdf:resource="urn:miriam:obo.chebi:CHEBI%3A17627">
59           </rdf:Bag>
60         </bqbiol:hasPart>
61         <bqbiol:hasPart>
62           <rdf:Bag>
63             <rdf:li rdf:resource="urn:miriam:uniprot:P69905"/>
64             <rdf:li rdf:resource="urn:miriam:uniprot:P68871"/>

```

```

1         <rdf:li rdf:resource="urn:miriam:kegg.compound:C00032"/>
2     </rdf:Bag>
3 </bqbiol:hasPart>
4 </rdf:Description>
5 </rdf:RDF>
6 </annotation>
7 </species>

```

As formally defined above it is possible to use different qualifiers in the same annotation element. The following phosphorylation is annotated by its exact KEGG counterpart and by the generic GO term “phosphorylation”.

```

11 <reaction id="phosphorylation" metaid="phosphorylation"
12     reversible="true" fast="false">
13     <annotation>
14         <rdf:RDF
15             xmlns:bqbiol="http://biomodels.net/biology-qualifiers/"
16             xmlns:rdf="http://www.w3.org/1999/02/22-rdf-syntax-ns#"
17         >
18             <rdf:Description rdf:about="#phosphorylation">
19                 <bqbiol:is>
20                     <rdf:Bag>
21                         <rdf:li rdf:resource="urn:miriam:kegg.reaction:R03313" />
22                     </rdf:Bag>
23                 </bqbiol:is>
24                 <bqbiol:isVersionOf>
25                     <rdf:Bag>
26                         <rdf:li rdf:resource="urn:miriam:obo.go:GO%3A0016310" />
27                     </rdf:Bag>
28                 </bqbiol:isVersionOf>
29             </rdf:Description>
30         </rdf:RDF>
31     </annotation>
32 </reaction>


```

7 Example models expressed in XML using SBML

In this section, we present several examples of complete models encoded in XML using SBML Level 3.

7.1 A simple example application of SBML

Consider the following representation of an enzymatic reaction:



Suppose we have the following initial species concentrations and parameter values,

$$\begin{aligned}[E] &= 5 \cdot 10^{-21} \text{ mole litre}^{-1} \\ [S] &= 10^{-20} \text{ mole litre}^{-1} \\ [P] &= 0 \text{ mole litre}^{-1} \\ [ES] &= 0 \text{ mole litre}^{-1} \\ k_{\text{on}} &= 1\,000\,000 \text{ litre mole}^{-1} \text{ second}^{-1} \\ k_{\text{off}} &= 0.2 \text{ second}^{-1} \\ k_{\text{cat}} &= 0.1 \text{ second}^{-1}\end{aligned}$$

and we place everything in a single compartment called **comp** whose volume is 10^{-14} liters. The following is a minimal but complete SBML document encoding this model:

```
<?xml version="1.0" encoding="UTF-8"?>
<sbml level="3" version="1" xmlns="http://www.sbml.org/sbml/level3/version1/core">
  <model extentUnits="mole" timeUnits="second">
    <listOfUnitDefinitions>
      <unitDefinition id="per_second">
        <listOfUnits>
          <unit kind="second" exponent="-1" scale="0" multiplier="1"/>
        </listOfUnits>
      </unitDefinition>
      <unitDefinition id="litre_per_mole_second">
        <listOfUnits>
          <unit kind="mole" exponent="-1" scale="0" multiplier="1"/>
          <unit kind="litre" exponent="1" scale="0" multiplier="1"/>
          <unit kind="second" exponent="-1" scale="0" multiplier="1"/>
        </listOfUnits>
      </unitDefinition>
    </listOfUnitDefinitions>
    <listOfCompartments>
      <compartment id="comp" size="1e-14" spatialDimensions="3" units="litre" constant="true"/>
    </listOfCompartments>
    <listOfSpecies>
      <species compartment="comp" id="ES" initialAmount="0" boundaryCondition="false"
        hasOnlySubstanceUnits="false" substanceUnits="mole" constant="false"/>
      <species compartment="comp" id="P" initialAmount="0" boundaryCondition="false"
        hasOnlySubstanceUnits="false" substanceUnits="mole" constant="false"/>
      <species compartment="comp" id="S" initialAmount="1e-20" boundaryCondition="false"
        hasOnlySubstanceUnits="false" substanceUnits="mole" constant="false"/>
      <species compartment="comp" id="E" initialAmount="5e-21" boundaryCondition="false"
        hasOnlySubstanceUnits="false" substanceUnits="mole" constant="false"/>
    </listOfSpecies>
    <listOfReactions>
      <reaction id="veq" reversible="true" fast="false">
        <listOfReactants>
          <speciesReference species="E" stoichiometry="1" constant="true"/>
          <speciesReference species="S" stoichiometry="1" constant="true"/>
        </listOfReactants>
        <listOfProducts>
          <speciesReference species="ES" stoichiometry="1" constant="true"/>
        </listOfProducts>
        <kineticLaw>
          <math xmlns="http://www.w3.org/1998/Math/MathML">
```

```

1      <apply>
2        <times/>
3        <ci>comp</ci>
4        <apply>
5          <minus/>
6          <apply>
7            <times/>
8            <ci>kon</ci>
9            <ci>E</ci>
10           <ci>S</ci>
11          </apply>
12         <apply>
13           <times/>
14           <ci>koff</ci>
15           <ci>ES</ci>
16         </apply>
17       </apply>
18     </math>
19   <listOfLocalParameters>
20     <localParameter id="kon" value="1000000" units="litre_per_mole_second"/>
21     <localParameter id="koff" value="0.2" units="per_second"/>
22   </listOfLocalParameters>
23 </kineticLaw>
24 </reaction>
25 <reaction id="vcat" reversible="false" fast="false">
26   <listOfReactants>
27     <speciesReference species="ES" stoichiometry="1" constant="true"/>
28   </listOfReactants>
29   <listOfProducts>
30     <speciesReference species="E" stoichiometry="1" constant="true"/>
31     <speciesReference species="P" stoichiometry="1" constant="true"/>
32   </listOfProducts>
33   <kineticLaw>
34     <math xmlns="http://www.w3.org/1998/Math/MathML">
35       <apply>
36         <times/>
37         <ci>comp</ci>
38         <ci>kcat</ci>
39         <ci>ES</ci>
40       </apply>
41     </math>
42     <listOfLocalParameters>
43       <localParameter id="kcat" value="0.1" units="per_second"/>
44     </listOfLocalParameters>
45   </kineticLaw>
46 </reaction>
47 </listOfReactions>
48 </model>
49 </sbml>

```

In this example, the model contains one compartment (with identifier “**comp**”), four species (with identifiers “**ES**”, “**P**”, “**S**”, and “**E**”), and two reactions (“**veq**” and “**vcat**”). The elements in the **listOfReactants** and **listOfProducts** in each reaction refer to elements listed in the **listOfSpecies**. The correspondences between the various elements is explicitly stated by the **speciesReference** elements.

The model also features local parameter definitions in each reaction. In this case, the three parameters (“**kon**”, “**koff**”, “**kcat**”) all have unique identifiers and they could also have just as easily been declared global parameters in the model. Local parameters frequently become more useful in larger models, where it may become tedious to assign unique identifiers for all the different parameters.


Finally, the example above also demonstrates the use of unit specifications throughout the model. The **model** component defines the units of kinetic laws as being *mole/second* by virtue of the values of the attributes **extentUnits** and **timeUnits**. Elsewhere in the model, species, parameters and compartments are defined with appropriate units so that the mathematical formulas inside the **kineticLaw** elements work out to be *mole/second*.

7.2 Example of a discrete version of a simple dimerization reaction

(SBO annotations for this model contributed by Lukas Endler, EMBL-EBI, Cambridge, UK.)

This example illustrates subtle differences between models formulated for use in a continuous simulation framework (e.g., using differential equations) and those intended for a discrete simulation framework. The model shown here is suitable for use with a discrete stochastic simulation algorithm of the sort developed by Gillespie (1977). In such an approach, species are described in terms of molecular counts and simulation proceeds by computing the probability of the time and identity of the next reaction, then updating the species amounts appropriately.

The model involves a simple dimerization reaction for a protein named “P”:



The SBML representation is shown below. There are several important points to note. First, the species “P” and “P2” declare they are always in discrete amounts by using the flag `hasOnlySubstanceUnits=“true”`. This indicates that when the species identifiers appear in mathematical formulas, the units are *substance*, not *substance/size*. A second point is that, as a result, the corresponding “kinetic law” formulas do not need volume corrections. In Gillespie’s approach, the constants in the rate expressions (here, “c1” and “c2”) contain a contribution from the kinetic constants of the reaction and the size of the compartment in which the reactions take place. This is a convention commonly adopted by stochastic modellers, but is in no way essential—it is perfectly reasonable to factor volume out of the rate constants, and in certain situations it may be desirable to do so (e.g., for models having time-varying compartment volume), but due to the use of substance units, it must be done differently compared to the deterministic case. Thirdly, although the reaction is reversible, it is encoded as two separate irreversible reactions, one each for the forward and reverse directions, as averaging over the reactions will affect the stochasticity. Finally, it is worth noting the rate expression for the forward reaction is a second-order mass-action reaction, but it is the *discrete* formulation of such a reaction rate (Gillespie, 1977).

```
<?xml version="1.0" encoding="UTF-8"?>
<sbml xmlns="http://www.sbml.org/sbml/level3/version1/core" level="3" version="1">
  <model id="dimerization" substanceUnits="item" timeUnits="second"
    volumeUnits="litre" extentUnits="item">
    <listOfUnitDefinitions>
      <unitDefinition id="per_second">
        <listOfUnits>
          <unit kind="second" exponent="-1" scale="0" multiplier="1"/>
        </listOfUnits>
      </unitDefinition>
      <unitDefinition id="per_item_per_second">
        <listOfUnits>
          <unit kind="item" exponent="-1" scale="0" multiplier="1"/>
          <unit kind="second" exponent="-1" scale="0" multiplier="1"/>
        </listOfUnits>
      </unitDefinition>
    </listOfUnitDefinitions>
    <listOfCompartments>
      <compartment id="Cell" size="1e-15" spatialDimensions="3"
        constant="true" sboTerm="SBO:0000290"/>
    </listOfCompartments>
    <listOfSpecies>
      <species id="P" compartment="Cell" initialAmount="301"
        hasOnlySubstanceUnits="true" boundaryCondition="false"
        constant="false" sboTerm="SBO:0000252"/>
      <species id="P2" compartment="Cell" initialAmount="0"
        hasOnlySubstanceUnits="true" boundaryCondition="false"
        constant="false" sboTerm="SBO:0000420"/>
    </listOfSpecies>
    <listOfReactions>
      <reaction id="Dimerization" reversible="false" fast="false" sboTerm="SBO:0000177">
        <listOfReactants>
          <speciesReference species="P" stoichiometry="2" constant="true"
            sboTerm="SBO:0000010"/>
        </listOfReactants>
      </reaction>
    </listOfReactions>
  </model>
</sbml>
```



```

1      </listOfReactants>
2      <listOfProducts>
3          <speciesReference species="P2" stoichiometry="1" constant="true"
4              sboTerm="SBO:0000011"/>
5      </listOfProducts>
6      <kineticLaw sboTerm="SBO:0000142">
7          <math xmlns="http://www.w3.org/1998/Math/MathML"
8              xmlns:sbml="http://www.sbml.org/sbml/level3/version1/core">
9              <apply>
10                  <divide/>
11                  <apply>
12                      <times/>
13                      <ci> c1 </ci>
14                      <ci> P </ci>
15                      <apply>
16                          <minus/>
17                          <ci> P </ci>
18                          <cn type="integer" sbml:unit="item"> 1 </cn>
19                      </apply>
20                  </apply>
21                  <cn type="integer" sbml:unit="dimensionless"> 2 </cn>
22              </apply>
23          </math>
24          <listOfLocalParameters>
25              <localParameter id="c1" value="0.00166" units="per_item_per_second"
26                  sboTerm="SBO:0000067"/>
27          </listOfLocalParameters>
28      </kineticLaw>
29  </reaction>
30  <reaction id="Dissociation" reversible="false" fast="false" sboTerm="SBO:0000180">
31      <listOfReactants>
32          <speciesReference species="P2" stoichiometry="1" constant="true"
33              sboTerm="SBO:0000010"/>
34      </listOfReactants>
35      <listOfProducts>
36          <speciesReference species="P" stoichiometry="2" constant="true"
37              sboTerm="SBO:0000011"/>
38      </listOfProducts>
39      <kineticLaw sboTerm="SBO:0000141">
40          <math xmlns="http://www.w3.org/1998/Math/MathML">
41              <apply>
42                  <times/>
43                  <ci> c2 </ci>
44                  <ci> P </ci>
45              </apply>
46          </math>
47          <listOfLocalParameters>
48              <localParameter id="c2" value="0.2" units="per_second"
49                  sboTerm="SBO:0000066"/>
50          </listOfLocalParameters>
51      </kineticLaw>
52  </reaction>
53  </listOfReactions>
54  </model>
55  </sbml>

```

This example also illustrates the need to provide additional information in a model so that software tools using different mathematical frameworks can properly interpret it. In this case, a simulation tool designed for continuous ODE-based simulation would likely misinterpret the model (in particular the reaction rate formulas), unless it deduced that a discrete stochastic simulation was intended. One of the purposes of SBO annotations (Section 5) is to enable such interpretation without the need for deduction. However, the interpretation of the model is essentially the same irrespective of whether the model is to be simulated in a deterministic or stochastic manner, and a properly SBML-compliant deterministic simulator will in most cases correctly simulate the continuous deterministic approximation of the stochastic model even if it has no stochastic simulation capability.

The interpretation of rate laws for stochastic models is similar to, yet different from, that of deterministic models. Taking the first reaction as an example, the rate law is $c_1 P(P-1)/2$ reaction events per second. In the continuous deterministic case, the interpretation of this is that the extent of the reaction in time dt is $[c_1 P(P-1)/2]dt$ (and this leads naturally to the usual ODE formulation of the model). In the stochastic case, the interpretation is that the *propensity* (or *rate*, or *hazard*) of the reaction is $c_1 P(P-1)/2$. That is, the *probability* of a single reaction event occurring in time dt is $[c_1 P(P-1)/2]dt$ (and note that the *expected* extent of the reaction will be $[c_1 P(P-1)/2]dt$). This interpretation leads to a Markov jump process for the system dynamics, where the inter-event times are exponentially distributed. Such dynamics can be simulated using a discrete event simulation algorithm such as the *Gillespie algorithm*. In this case, the algorithm for simulating the model can be described as follows:

1. Initialise $t := 0$, $c_1 := 0.00166$, $c_2 := 0.2$, $P := 301$, $P_2 := 0$
2. Compute $h_1 := c_1 P(P-1)/2$, $h_2 := c_2 P_2$
3. Compute $h_0 = h_1 + h_2$
4. Simulate $t' \sim \text{Exp}(h_0)$ and set $t := t + t'$
5. With probability h_1/h_0 set $P := P - 2$, $P_2 := P_2 + 1$, otherwise set $P := P + 2$, $P_2 := P_2 - 1$.
6. Output t , P , P_2
7. If $t < T_{max}$, return to step 2, otherwise stop.

Although this is a simulation algorithm is a very practical way of describing how to construct exact realisations of the Markov jump process corresponding to the discrete stochastic kinetic model, it is not a concise mathematical description. Such a description can be provided by writing the model as a time change of a pair of independent unit Poisson processes. Let $N_1(t)$ and $N_2(t)$ be the counting functions of these processes, so that for each $i = 1, 2$, $t > 0$, $N_i(t) \sim \text{Poisson}(t)$. Then, writing $P(t)$ and $P_2(t)$ for the numbers of molecules of P and P_2 at time t , respectively, we have that the stochastic process $\{P(t), P_2(t) | t > 0\}$ satisfies the stochastic integral equation

$$P_2(t) = N_1 \left(\int_0^t c_1 \frac{P(\tau)[P(\tau) - 1]}{2} d\tau \right) - N_2 \left(\int_0^t c_2 P_2(\tau) d\tau \right)$$

$$P(t) = 301 - 2P_2(t).$$

The above representation is arguably the most useful for mathematical analysis of the stochastic model; see [Ball et al. \(2006\)](#) for details. Another popular representation is the so-called chemical Master equation (CME) for the probability distribution of the possible states at all times ([Gillespie, 1992](#)). In this case, since there are 151 possible states of the system (corresponding to the 151 possible values of P_2), the CME consists of 151 coupled ODEs,

$$\frac{d}{dt} p(P, P_2, t) = \begin{cases} -\frac{c_1}{2} \times 301 \times 299 p(301, 0, t) + c_2 p(299, 1, t), & P = 301, P_2 = 0, \\ \frac{c_1}{2} (P+2)(P+1) p(P+2, P_2-1, t) - \frac{c_1}{2} P(P-1) p(P, P_2, t) & P = 301-x, P_2 = x, \\ + c_2 (P_2+1) p(P-2, P_2+1, t) - c_2 P_2 p(P, P_2, t), & x = 1, 2, \dots, 149, \\ \frac{c_1}{2} \times 2 \times 3 p(3, 149, t) - c_2 \times 150 p(1, 150, t), & P = 1, P_2 = 150, \end{cases}$$


where $p(P, P_2, t)$ denotes the probability that there are P molecules of P and P_2 molecules of P_2 at time t , and the ODEs are subject to the initial conditions

$$p(301, 0, 0) = 1, p(301-2x, x, 0) = 0, x = 1, 2, \dots, 150.$$

See [Evans et al. \(2008\)](#) for further examples of discrete stochastic kinetic models encoded in SBML and [Wilkinson \(2006\)](#) for an introduction to discrete stochastic modelling using SBML.


7.3 Example involving assignment rules

This section contains a model that simulates a system containing a fast reaction. This model uses rules to express the mathematics of the fast reaction explicitly rather than using the **fast** attribute on a reaction element. The system modeled is



$$k_1 = 0.1, \quad k_2 = 0.15, \quad k_f = K_{eq}10000, \quad k_r = 10000, \quad K_{eq} = 2.5.$$

where $[X_0]$, $[S_1]$, $[S_2]$, and $[X_1]$ are species in concentration units, and k_1 , k_2 , k_f , k_r , and K_{eq} are parameters. This system of reactions can be approximated with the following new system:



$$[S_1] = \frac{[T]}{1 + K_{eq}}$$

$$[S_2] = K_{eq}[S_1]$$

where T is a new species. The following example SBML model encodes the second system.

```
<?xml version="1.0" encoding="UTF-8"?>
<sbml xmlns="http://www.sbml.org/sbml/level3/version1/core" level="3" version="1"
  xmlns:math="http://www.w3.org/1998/Math/MathML">
  <model volumeUnits="litre" substanceUnits="mole" timeUnits="second" extentUnits="mole">
    <listOfUnitDefinitions>
      <unitDefinition id="per_second">
        <listOfUnits>
          <unit kind="second" exponent="-1" scale="0" multiplier="1"/>
        </listOfUnits>
      </unitDefinition>
    </listOfUnitDefinitions>
    <listOfCompartments>
      <compartment id="cell" size="1" spatialDimensions="3" constant="true"/>
    </listOfCompartments>
    <listOfSpecies>
      <species id="X0" compartment="cell" initialConcentration="1" constant="false"
        hasOnlySubstanceUnits="false" boundaryCondition="false"/>
      <species id="X1" compartment="cell" initialConcentration="0" constant="false"
        hasOnlySubstanceUnits="false" boundaryCondition="false"/>
      <species id="T" compartment="cell" initialConcentration="0" constant="false"
        hasOnlySubstanceUnits="false" boundaryCondition="false"/>
      <species id="S1" compartment="cell" initialConcentration="0" constant="false"
        hasOnlySubstanceUnits="false" boundaryCondition="false"/>
      <species id="S2" compartment="cell" initialConcentration="0" constant="false"
        hasOnlySubstanceUnits="false" boundaryCondition="false"/>
    </listOfSpecies>
    <listOfParameters>
      <parameter id="Keq" value="2.5" units="dimensionless" constant="true"/>
    </listOfParameters>
    <listOfRules>
      <assignmentRule variable="S1">
        <math xmlns="http://www.w3.org/1998/Math/MathML"
          xmlns:sbml="http://www.sbml.org/sbml/level3/version1/core">
          <apply>
            <divide/>
            <ci> T </ci>
          </apply>
        </math>
      </assignmentRule>
    </listOfRules>
  </model>
</sbml>
```


```

1          <plus/>
2          <cn sbml:units="dimensionless"> 1 </cn>
3          <ci> Keq </ci>
4        </apply>
5      </math>
6    </assignmentRule>
7    <assignmentRule variable="S2">
8      <math xmlns="http://www.w3.org/1998/Math/MathML">
9        <apply>
10          <times/>
11          <ci> Keq </ci>
12          <ci> S1 </ci>
13        </apply>
14      </math>
15    </assignmentRule>
16  </listOfRules>
17  <listOfReactions>
18    <reaction id="in" reversible="false" fast="false">
19      <listOfReactants>
20        <speciesReference species="X0" stoichiometry="1" constant="true"/>
21      </listOfReactants>
22      <listOfProducts>
23        <speciesReference species="T" stoichiometry="1" constant="true"/>
24      </listOfProducts>
25      <kineticLaw>
26        <math xmlns="http://www.w3.org/1998/Math/MathML">
27          <apply>
28            <times/>
29            <ci> k1 </ci>
30            <ci> X0 </ci>
31            <ci> cell </ci>
32          </apply>
33        </math>
34        <listOfLocalParameters>
35          <localParameter id="k1" value="0.1" units="per_second"/>
36        </listOfLocalParameters>
37      </kineticLaw>
38    </reaction>
39    <reaction id="out" reversible="false" fast="false">
40      <listOfReactants>
41        <speciesReference species="T" stoichiometry="1" constant="true"/>
42      </listOfReactants>
43      <listOfProducts>
44        <speciesReference species="X1" stoichiometry="1" constant="true"/>
45      </listOfProducts>
46      <listOfModifiers>
47        <modifierSpeciesReference species="S2"/>
48      </listOfModifiers>
49      <kineticLaw>
50        <math xmlns="http://www.w3.org/1998/Math/MathML">
51          <apply>
52            <times/>
53            <ci> k2 </ci>
54            <ci> S2 </ci>
55            <ci> cell </ci>
56          </apply>
57        </math>
58        <listOfLocalParameters>
59          <localParameter id="k2" value="0.15" units="per_second"/>
60        </listOfLocalParameters>
61      </kineticLaw>
62    </reaction>
63  </listOfReactions>
64 </model>
65 </sbml>

```

7.4 Example involving algebraic rules

This section contains an example model that contains two **AlgebraicRule** objects that are necessary to determine the values of two variables within the model. In this particular case, the rules cannot be rewritten in terms of **AssignmentRule**. This example illustrates a more rigorous analysis of the enzymatic reaction given in the example of Section 7.1.



In this example, we describe a quasi-steady-state approximation of the enzymatic reaction equation shown above. It is based on two assumptions. First, the rate at which the concentration of the substrate bound enzyme ($[ES]$) changes is assumed to be slow compared to the rate of change of concentration of both the substrate ($[S]$) and product ($[P]$). Second, the total concentration of the enzyme is assumed to stay constant over time. This means we can assume the concentration of $[ES]$ and $[E]$ are not governed by the reactions, and so some other equations must be used to determine the values of these concentrations in order to be able to simulate the model.

Applying the first assumption means that the rate of change of $[ES]$ should be set to zero:

$$\frac{d[ES]}{dt} = k_{1\text{on}} \cdot [E] \cdot [S] - (k_{1\text{off}} + k_2) \cdot [ES] = 0$$

The second assumption can be written as

$$[E_{\text{total}}] = [E] + [ES]$$

which, after rearranging, becomes

$$[E_{\text{total}}] - ([E] + [ES]) = 0$$

Thus, we have two algebraic rules that must be applied to determine the values of $[E]$ and $[ES]$. The SBML encoding of this model is given below. Note that the species E and ES have their **boundaryCondition** attribute set to “true”. This means that a simulation tool should not construct equations for them based on the reactions in the system. Their values are instead set using the rules in the model. Also, the model uses a dummy species E_{total} with its **constant** attribute set to “true”; its role is to assign the total concentration of the enzyme in the model. This could just as easily have been done using a parameter instead of a constant dummy species, but we use the latter approach as an illustration.

```
<?xml version="1.0" encoding="UTF-8" ?>
<sbml xmlns="http://www.sbml.org/sbml/level3/version1/core" level="3" version="1">
  <model substanceUnits="mole" volumeUnits="litre" timeUnits="second" extentUnits="mole">
    <listOfUnitDefinitions>
      <unitDefinition id="per_second">
        <listOfUnits>
          <unit kind="second" exponent="-1" scale="0" multiplier="1"/>
        </listOfUnits>
      </unitDefinition>
      <unitDefinition id="mole_per_litre">
        <listOfUnits>
          <unit kind="mole" exponent="1" scale="0" multiplier="1"/>
          <unit kind="litre" exponent="-1" scale="0" multiplier="1"/>
        </listOfUnits>
      </unitDefinition>
      <unitDefinition id="litre_per_mole_per_second">
        <listOfUnits>
          <unit kind="litre" exponent="1" scale="0" multiplier="1"/>
          <unit kind="mole" exponent="-1" scale="0" multiplier="1"/>
          <unit kind="second" exponent="-1" scale="0" multiplier="1"/>
        </listOfUnits>
      </unitDefinition>
    </listOfUnitDefinitions>
    <listOfCompartments>
```

```

1      <compartment id="cell" size="1" spatialDimensions="3" constant="true"/>
2  </listOfCompartments>
3  <listOfSpecies>
4      <species id="E" compartment="cell" initialConcentration="0.5" constant="false"
5          hasOnlySubstanceUnits="false" boundaryCondition="true"/>
6      <species id="S" compartment="cell" initialConcentration="1.0" constant="false"
7          hasOnlySubstanceUnits="false" boundaryCondition="false"/>
8      <species id="ES" compartment="cell" initialConcentration="0.5" constant="false"
9          hasOnlySubstanceUnits="false" boundaryCondition="true"/>
10     <species id="P" compartment="cell" initialConcentration="0" constant="false"
11         hasOnlySubstanceUnits="false" boundaryCondition="false"/>
12     <species id="E_total" compartment="cell" initialConcentration="1.0" constant="true"
13         hasOnlySubstanceUnits="false" boundaryCondition="true"/>
14 </listOfSpecies>
15 <listOfParameters>
16     <parameter id="k1_on" value="1" units="litre_per_mole_per_second" constant="true"/>
17     <parameter id="k1_off" value="0.5" units="per_second" constant="true"/>
18     <parameter id="k2" value="0.5" units="per_second" constant="true"/>
19 </listOfParameters>
20 <listOfRules>
21     <algebraicRule>
22         <math xmlns="http://www.w3.org/1998/Math/MathML">
23             <apply>
24                 <minus/>
25                 <apply> <times/> <ci> k1_on </ci> <ci> E </ci> <ci> S </ci> </apply>
26                 <apply>
27                     <times/>
28                     <apply> <plus/> <ci> k1_off </ci> <ci> k2 </ci> </apply>
29                     <ci> ES </ci>
30                 </apply>
31             </math>
32         </algebraicRule>
33     <algebraicRule>
34         <math xmlns="http://www.w3.org/1998/Math/MathML">
35             <apply>
36                 <minus/>
37                 <apply> <plus/> <ci> E </ci> <ci> ES </ci> </apply>
38                 <ci> E_total </ci>
39             </apply>
40         </math>
41     </algebraicRule>
42 </listOfRules>
43 <listOfReactions>
44     <reaction id="r1" reversible="true" fast="false">
45         <listOfReactants>
46             <speciesReference species="E" stoichiometry="1" constant="true"/>
47             <speciesReference species="S" stoichiometry="1" constant="true"/>
48         </listOfReactants>
49         <listOfProducts>
50             <speciesReference species="ES" stoichiometry="1" constant="true"/>
51         </listOfProducts>
52         <kineticLaw>
53             <math xmlns="http://www.w3.org/1998/Math/MathML">
54                 <apply>
55                     <times/>
56                     <ci> cell </ci>
57                     <apply>
58                         <minus/>
59                         <apply> <times/> <ci> k1_on </ci> <ci> E </ci> <ci> S </ci> </apply>
60                         <apply> <times/> <ci> k1_off </ci> <ci> ES </ci> </apply>
61                     </apply>
62                 </math>
63             </kineticLaw>
64         </reaction>
65     <reaction id="r2" reversible="false" fast="false">
66         <listOfReactants>
67             <speciesReference species="ES" stoichiometry="1" constant="true"/>

```


```

1      </listOfReactants>
2      <listOfProducts>
3        <speciesReference species="E" stoichiometry="1" constant="true"/>
4        <speciesReference species="P" stoichiometry="1" constant="true"/>
5      </listOfProducts>
6      <kineticLaw>
7        <math xmlns="http://www.w3.org/1998/Math/MathML">
8          <apply> <times/> <ci> cell </ci> <ci> k2 </ci> <ci> ES </ci> </apply>
9        </math>
10     </kineticLaw>
11 </reaction>
12 </listOfReactions>
13 </model>
14 </sbml>

```

7.5 Example with combinations of boundaryCondition and constant values on Species with RateRule objects

In this section, we discuss a model that includes four species, each with a different combination of values for their **boundaryCondition** and **constant** attributes. The model represents a hypothetical system containing one reaction,



where S_3 is a species that catalyzes the conversion of species S_1 and S_2 into S_4 . Species S_1 and S_2 are on the boundary of the system (i.e., S_1 and S_2 are reactants but their values are not determined by kinetic laws). The value of S_1 in the system is determined over time by the rate rule:

$$\frac{d[S_1]}{dt} = k_2$$

The species S_2 and S_3 are not affected by either the reaction or the rate rule, and have the following initial concentration values:

$$[S_2] = 1, \quad [S_3] = 2$$

The values of constant parameters in the system are:

$$k_1 = 0.5, \quad k_2 = 0.1$$

and the initial values of varying species are:

$$[S_1] = 0, \quad [S_4] = 0$$

The value of $[S_1]$ varies over time and it is on the boundary, so in the SBML representation, **S1** has a **constant** attribute with a value of “false” and a **boundaryCondition** attribute with a value of “true”. The value of $[S_2]$ is fixed and it is also on the boundary, so **S2** has a **constant** attribute value of “false” and a **boundaryCondition** attribute value of “true”. $[S_3]$ is fixed but not on the boundary, so the **constant** attribute is “true” and the **boundaryCondition** attribute is “false”. Finally, $[S_4]$ is a product whose value is determined by a kinetic law and therefore in the SBML representation has “false” for both its **boundaryCondition** and **constant** attributes.

The following is the SBML rendition of the model shown above:

```

40 <?xml version="1.0" encoding="UTF-8"?>
41 <sbml xmlns="http://www.sbml.org/sbml/level3/version1/core" level="3" version="1">
42   <model id="BoundaryCondExampleModel"
43     volumeUnits="litre" substanceUnits="mole" timeUnits="second" extentUnits="mole">
44     <listOfUnitDefinitions>
45       <unitDefinition id="mole_per_litre_per_second">
46         <listOfUnits>
47           <unit kind="mole" exponent="1" scale="0" multiplier="1"/>
48           <unit kind="litre" exponent="-1" scale="0" multiplier="1"/>

```



```

1         <unit kind="second" exponent="-1" scale="0" multiplier="1"/>
2     </listOfUnits>
3 </unitDefinition>
4 <unitDefinition id="litre_sq_per_mole_sq_per_second">
5     <listOfUnits>
6         <unit kind="mole" exponent="-2" scale="0" multiplier="1"/>
7         <unit kind="litre" exponent="2" scale="0" multiplier="1"/>
8         <unit kind="second" exponent="-1" scale="0" multiplier="1"/>
9     </listOfUnits>
10 </unitDefinition>
11 </listOfUnitDefinitions>
12 <listOfCompartments>
13     <compartment id="compartmentOne" size="1" spatialDimensions="3" constant="true"/>
14 </listOfCompartments>
15 <listOfSpecies>
16     <species id="S1" initialConcentration="0" compartment="compartmentOne"
17         hasOnlySubstanceUnits="false" boundaryCondition="true"
18         constant="false"/>
19     <species id="S2" initialConcentration="1" compartment="compartmentOne"
20         hasOnlySubstanceUnits="false" boundaryCondition="true"
21         constant="true"/>
22     <species id="S3" initialConcentration="3" compartment="compartmentOne"
23         hasOnlySubstanceUnits="false" boundaryCondition="false"
24         constant="true"/>
25     <species id="S4" initialConcentration="0" compartment="compartmentOne"
26         hasOnlySubstanceUnits="false" boundaryCondition="false"
27         constant="false"/>
28 </listOfSpecies>
29 <listOfParameters>
30     <parameter id="k1" value="0.5" units="litre_sq_per_mole_sq_per_second"
31         constant="true"/>
32     <parameter id="k2" value="0.1" units="mole_per_litre_per_second"
33         constant="true"/>
34 </listOfParameters>
35 <listOfRules>
36     <rateRule variable="S1">
37         <math xmlns="http://www.w3.org/1998/Math/MathML">
38             <ci> k2 </ci>
39         </math>
40     </rateRule>
41 </listOfRules>
42 <listOfReactions>
43     <reaction id="reaction_1" reversible="false" fast="false">
44         <listOfReactants>
45             <speciesReference species="S1" stoichiometry="1" constant="true"/>
46             <speciesReference species="S2" stoichiometry="1" constant="true"/>
47         </listOfReactants>
48         <listOfProducts>
49             <speciesReference species="S4" stoichiometry="1" constant="true"/>
50         </listOfProducts>
51         <listOfModifiers>
52             <modifierSpeciesReference species="S3"/>
53         </listOfModifiers>
54         <kineticLaw>
55             <math xmlns="http://www.w3.org/1998/Math/MathML">
56                 <apply>
57                     <times/>
58                     <ci> k1 </ci>
59                     <ci> S1 </ci>
60                     <ci> S2 </ci>
61                     <ci> S3 </ci>
62                     <ci> compartmentOne </ci>
63                 </apply>
64             </math>
65         </kineticLaw>
66     </reaction>
67 </listOfReactions>
68 </model>
69 </sbml>

```

7.6 Example of translation from a multi-compartmental model to ODEs

This section contains a model with two compartments and four reactions. The model is derived from Lotka-Volterra, with the addition of a reversible transport step. When observed in a time-course simulation, three of this model's species display damped oscillations.




Figure 28: A example multi-compartmental model.

Figure 28 illustrates the arrangement of compartments and reactions in the model `LotkaVolterra_transport`. The reaction between the compartments called `cytosol` and `nucleus` is a transport reaction whose mechanisms are not modeled here; in particular, the reaction does not take place on the membrane between the compartments, and is modeled here simply as a process that spans the two three-dimensional compartments.

The text of the SBML representation of the model is shown below, and it is followed by its complete translation into ordinary differential equations. As usual, in this SBML model, the reaction rate equations in the kinetic laws are in substance per time units. The reactions have also been simplified to reduce common stoichiometric factors in the original system depicted in Figure 28. The species variables in this SBML representation are in concentration units; their initial quantities are declared using the attribute `initialAmount` on the `species` definitions, but since the attribute `hasOnlySubstanceUnits` is *not* set to true, the identifiers of the species represent their concentrations when those identifiers appear in mathematical expressions elsewhere in the model. Note that the species whose identifier is “X” is a boundary condition, as indicated by the attribute `boundaryCondition=“true”` in its definition.

```
<?xml version="1.0" encoding="UTF-8"?>
<sbml xmlns="http://www.sbml.org/sbml/level3/version1/core" level="3" version="1">
  <model name="LotkaVolterra_transport" substanceUnits="mole" volumeUnits="litre"
    extentUnits="mole" timeUnits="second">
    <listOfUnitDefinitions>
      <unitDefinition id="per_second">
        <listOfUnits>
          <unit kind="second" exponent="-1" scale="0" multiplier="1"/>
        </listOfUnits>
      </unitDefinition>
      <unitDefinition id="litre_per_mole_per_second">
        <listOfUnits>
          <unit kind="mole" exponent="-1" scale="0" multiplier="1"/>
          <unit kind="litre" exponent="1" scale="0" multiplier="1"/>
          <unit kind="second" exponent="-1" scale="0" multiplier="1"/>
        </listOfUnits>
      </unitDefinition>
    </listOfUnitDefinitions>
    <listOfCompartments>
      <compartment id="cytoplasm" size="5" constant="true"/>
      <compartment id="nucleus" size="1" constant="true"/>
    </listOfCompartments>
    <listOfSpecies>
      <species id="X" compartment="nucleus" initialAmount="1" constant="true"
        boundaryCondition="true" hasOnlySubstanceUnits="false"/>
      <species id="Y1n" compartment="nucleus" initialAmount="1" constant="true"
        boundaryCondition="false" hasOnlySubstanceUnits="false"/>
      <species id="Y1c" compartment="cytoplasm" initialAmount="0" constant="true"
        boundaryCondition="false" hasOnlySubstanceUnits="false"/>
      <species id="Y2" compartment="cytoplasm" initialAmount="1" constant="true"
        boundaryCondition="false" hasOnlySubstanceUnits="false"/>
    </listOfSpecies>
    <listOfParameters>
      <parameter id="k1" value="2500" units="litre_per_mole_per_second" constant="true"/>
    </listOfParameters>
  </model>
</sbml>
```

```

1      <parameter id="k2" value="2500" units="litre_per_mole_per_second" constant="true"/>
2      <parameter id="KT" value="25000" units="per_second" constant="true"/>
3      <parameter id="k3" value="2500" units="per_second" constant="true"/>
4  </listOfParameters>
5  <listOfReactions>
6      <reaction id="production" reversible="false" fast="false">
7          <listOfReactants>
8              <speciesReference species="X" stoichiometry="1" constant="true"/>
9              <speciesReference species="Y1n" stoichiometry="1" constant="true"/>
10         </listOfReactants>
11         <listOfProducts>
12             <speciesReference species="Y1n" stoichiometry="1" constant="true"/>
13             <speciesReference species="Y1n" stoichiometry="1" constant="true"/>
14         </listOfProducts>
15         <kineticLaw>
16             <math xmlns="http://www.w3.org/1998/Math/MathML">
17                 <apply>
18                     <times/>
19                     <ci>nucleus</ci>
20                     <ci>k1</ci>
21                     <ci>X</ci>
22                     <ci>Y1n</ci>
23                 </apply>
24             </math>
25         </kineticLaw>
26     </reaction>
27     <reaction id="transport" reversible="true" fast="false">
28         <listOfReactants>
29             <speciesReference species="Y1n" stoichiometry="1" constant="true"/>
30         </listOfReactants>
31         <listOfProducts>
32             <speciesReference species="Y1c" stoichiometry="1" constant="true"/>
33         </listOfProducts>
34         <kineticLaw>
35             <math xmlns="http://www.w3.org/1998/Math/MathML">
36                 <apply>
37                     <times/>
38                     <ci>cytoplasm</ci>
39                     <ci>KT</ci>
40                     <apply>
41                         <minus/>
42                         <ci>Y1n</ci>
43                         <ci>Y1c</ci>
44                     </apply>
45                 </apply>
46             </math>
47         </kineticLaw>
48     </reaction>
49     <reaction id="transformation" reversible="false" fast="false">
50         <listOfReactants>
51             <speciesReference species="Y1c" stoichiometry="1" constant="true"/>
52             <speciesReference species="Y2" stoichiometry="1" constant="true"/>
53         </listOfReactants>
54         <listOfProducts>
55             <speciesReference species="Y2" stoichiometry="2" constant="true"/>
56         </listOfProducts>
57         <kineticLaw>
58             <math xmlns="http://www.w3.org/1998/Math/MathML">
59                 <apply>
60                     <times/>
61                     <ci>cytoplasm</ci>
62                     <ci>k2</ci>
63                     <ci>Y1c</ci>
64                     <ci>Y2</ci>
65                 </apply>
66             </math>
67         </kineticLaw>
68     </reaction>
69     <reaction id="degradation" reversible="false" fast="false">

```

```

1      <listOfReactants>
2        <speciesReference species="Y2" stoichiometry="1" constant="true"/>
3      </listOfReactants>
4      <kineticLaw>
5        <math xmlns="http://www.w3.org/1998/Math/MathML">
6          <apply>
7            <times/>
8            <ci>cytoplasm</ci>
9            <ci>k3</ci>
10           <ci>Y2</ci>
11          </apply>
12        </math>
13      </kineticLaw>
14    </reaction>
15  </listOfReactions>
16 </model>
17 </sbml>

```

The ODE translation of this model is as follows. First, we give the values of the constant parameters:

$$\begin{aligned}
 k_1 &= 2500 \text{ litre mole}^{-1} \text{ second}^{-1} \\
 k_2 &= 2500 \text{ litre mole}^{-1} \text{ second}^{-1} \\
 K_3 &= 25000 \text{ second}^{-1} \\
 K_T &= 25000 \text{ second}^{-1}
 \end{aligned}$$

Now on to the initial conditions of the variables. In the following, the terms $[X]$, $[Y_{1n}]$, $[Y_{1c}]$, and $[Y_2]$ refer to the species' concentrations. Note that the corresponding species identifiers **X**, **Y_{1n}**, **Y_{1c}** and **Y₂** in the model are in concentration units, even though all the values in the model are initialized in terms of amounts. (The reason the species identifiers in the model are still in concentration units goes back to the meaning of the **hasOnlySubstanceUnits** attribute on a **Species**; if the attribute is set to a value of “**false**”, a species' symbol in a model is interpreted as a concentration or density regardless of whether its initial value is set using **initialAmount** or **initialConcentration**.) We use V_n to represent the size of compartment “**nucleus**” and V_c the size of compartment “**cytoplasm**”:

$$\begin{aligned}
 V_n &= 1 \text{ litre} \\
 V_c &= 5 \text{ litre} \\
 X &= 1 \text{ mole} \\
 Y_{1n} &= 1 \text{ mole} \\
 Y_{1c} &= 0 \text{ mole} \\
 Y_2 &= 1/5 \text{ mole}
 \end{aligned}$$


And finally, here are the differential equations:

$$\begin{aligned}
 \frac{d[X]}{dt} &= 0 \\
 V_n \frac{d[Y_{1n}]}{dt} &= k_1[X][Y_{1n}][V_n] - K_T([Y_{1n}] - [Y_{1c}])V_c && \text{reactions production and transport} \\
 V_c \frac{d[Y_{1c}]}{dt} &= K_T([Y_{1n}] - [Y_{1c}])V_c - k_2[Y_{1c}][Y_2]V_c && \text{reactions transport and transformation} \\
 V_c \frac{d[Y_2]}{dt} &= k_2[Y_{1c}][Y_2]V_c - k_3[Y_2]V_c && \text{reactions transformation and degradation}
 \end{aligned}$$

As formulated here, this example assumes constant volumes. If the sizes of the compartments “**cytoplasm**” or “**nucleus**” could change during simulation, then it would be preferable to use a different approach to constructing the differential equations. In this alternative approach, the ODEs would compute substance change rather than concentration change, and the concentration values would be computed using separate equations. This approach is used in Section 4.11.7.

7.7 Example involving function definitions

This section contains a model that uses the function definition feature of SBML. Consider the following hypothetical system:



where

$$f(x) = 2x$$

The following is the XML document that encodes the model shown above:

```
<?xml version="1.0" encoding="UTF-8"?>
<sbml xmlns="http://www.sbml.org/sbml/level3/version1/core" level="3" version="1">
  <model id="Example" substanceUnits="mole" volumeUnits="litre"
    timeUnits="second" extentUnits="mole"
    <listOfUnitDefinitions>
      <unitDefinition id="conc">
        <listOfUnits>
          <unit kind="mole" multiplier="1" scale="0" exponent="1"/>
          <unit kind="litre" multiplier="1" scale="0" exponent="-1"/>
        </listOfUnits>
      </unitDefinition>
    </listOfUnitDefinitions>
    <listOfFunctionDefinitions>
      <functionDefinition id="f">
        <math xmlns="http://www.w3.org/1998/Math/MathML"
          xmlns:sbml="http://www.sbml.org/sbml/level3/version1/core">
          <lambda>
            <bvar>
              <ci> x </ci>
            </bvar>
            <apply>
              <times/>
              <ci> x </ci>
              <cn sbml:units="conc"> 2 </cn>
            </apply>
          </lambda>
        </math>
      </functionDefinition>
    </listOfFunctionDefinitions>
    <listOfCompartments>
      <compartment id="compartmentOne" size="1" spatialDimensions="3" constant="true"/>
    </listOfCompartments>
    <listOfSpecies>
      <species id="S1" initialConcentration="1" compartment="compartmentOne"
        hasOnlySubstanceUnits="false" boundaryCondition="false"
        constant="false"/>
      <species id="S2" initialConcentration="0" compartment="compartmentOne"
        hasOnlySubstanceUnits="false" boundaryCondition="false"
        constant="false"/>
    </listOfSpecies>
    <listOfParameters>
      <parameter id="t" value = "1" constant="true"/>
    </listOfParameters>
    <listOfReactions>
      <reaction id="reaction_1" reversible="false" fast="false">
        <listOfReactants>
          <speciesReference species="S1" stoichiometry="1" constant="true"/>
        </listOfReactants>
        <listOfProducts>
          <speciesReference species="S2" stoichiometry="1" constant="true"/>
        </listOfProducts>
        <kineticLaw>
          <math xmlns="http://www.w3.org/1998/Math/MathML">
            <apply>
              <divide/>
            </apply>
          </math>
        </kineticLaw>
      </reaction>
    </listOfReactions>
  </model>
</sbml>
```

```

1          <times/>
2          <apply>
3              <ci> f </ci>
4              <ci> S1 </ci>
5          </apply>
6          <ci> compartmentOne </ci>
7      </apply>
8      <ci> t</ci>
9  </apply>
10 </math>
11 </kineticLaw>
12 </reaction>
13 </listOfReactions>
14 </model>
15 </sbml>

```

7.8 Example involving *delay* functions

The following is a simple model illustrating the use of *delay* to represent a gene that suppresses its own expression. The model can be expressed in a single rule:

$$\frac{d[P]}{dt} = \frac{1}{1 + m[P_{\text{delayed}}]^q} - [P]$$

where

$[P_{\text{delayed}}]$ is *delay*([P], Δ_t) or [P] at $t - \Delta_t$
 $[P]$ is protein concentration
 τ is the response time
 m is a multiplier or equilibrium constant
 q is the Hill coefficient

and the species quantities are in concentration units. The text of an SBML encoding of this model is given below:

```

24 <?xml version="1.0" encoding="UTF-8"?>
25 <sbml xmlns="http://www.sbml.org/sbml/level3/version1/core" level="3" version="1">
26   <model substanceUnits="mole" volumeUnits="litre"
27     extentUnits="mole" timeUnits="second">
28     <listOfUnitDefinitions>
29       <unitDefinition id="conc">
30         <listOfUnits>
31           <unit kind="mole" multiplier="1" scale="0" exponent="1"/>
32           <unit kind="litre" multiplier="1" scale="0" exponent="-1"/>
33         </listOfUnits>
34       </unitDefinition>
35       <unitDefinition id="conc_sq">
36         <listOfUnits>
37           <unit kind="mole" multiplier="1" scale="0" exponent="2"/>
38           <unit kind="litre" multiplier="1" scale="0" exponent="-2"/>
39         </listOfUnits>
40       </unitDefinition>
41     </listOfUnitDefinitions>
42     <listOfCompartments>
43       <compartment id="cell" size="1" spatialDimensions="3" constant="true"/>
44     </listOfCompartments>
45     <listOfSpecies>
46       <species id="P" compartment="cell" initialConcentration="0"
47         hasOnlySubstanceUnits="false" boundaryCondition="false"
48         constant="false"/>
49     </listOfSpecies>
50     <listOfParameters>
51       <parameter id="tau" value="1" units="second" constant="true"/>
52       <parameter id="m" value="0.5" units="dimensionless" constant="true"/>
53       <parameter id="q" value="1" units="dimensionless" constant="true"/>
54       <parameter id="delta_t" value="1" units="second" constant="true"/>

```

```

1      </listOfParameters>
2      <listOfRules>
3          <rateRule variable="P">
4              <math xmlns="http://www.w3.org/1998/Math/MathML"
5                  xmlns:sbml="http://www.sbml.org/sbml/level3/version1/core">
6                  <apply>
7                      <divide/>
8                      <apply>
9                          <minus/>
10                         <apply>
11                             <divide/>
12                             <cn sbml:units="conc_sq"> 1 </cn>
13                             <apply>
14                                 <plus/>
15                                 <cn sbml:units="conc"> 1 </cn>
16                                 <apply>
17                                     <times/>
18                                     <ci> m </ci>
19                                     <apply>
20                                         <power/>
21                                         <apply>
22                                             <csymbol
23                                                 encoding="text"
24                                                 definitionURL="http://www.sbml.org/sbml/symbols/delay">
25                                                 delay
26                                             </csymbol>
27                                             <ci> P </ci>
28                                             <ci> delta_t </ci>
29                                         </apply>
30                                     <ci> q </ci>
31                                 </apply>
32                             </apply>
33                         </apply>
34                     <ci> P </ci>
35                 </apply>
36                 <ci> tau </ci>
37             </apply>
38         </math>
39     </rateRule>
40 </listOfRules>
41 </model>
42 </sbml>

```

7.9 Example involving events

This section presents a simple model system that demonstrates the use of events in SBML. Consider a system with two genes, G_1 and G_2 . G_1 is initially on and G_2 is initially off. When turned on, the two genes lead to the production of two products, P_1 and P_2 , respectively, at a fixed rate. When P_1 reaches a given concentration, G_2 switches on. This system can be represented mathematically as follows:

$$\begin{aligned}
 \frac{d[P_1]}{dt} &= k_1([G_1] - [P_1]) \\
 \frac{d[P_2]}{dt} &= k_2([G_2] - [P_2]) \\
 [G_2] &= \begin{cases} 0 & \text{when } [P_1] \leq \tau, \\ 1 & \text{when } [P_1] > \tau. \end{cases}
 \end{aligned}$$

The initial values are:

$$[G_1] = 1, \quad [G_2] = 0, \quad \tau = 0.25, \quad P_1 = 0, \quad P_2 = 0, \quad k_1 = k_2 = 1.$$

The SBML Level 3 representation of this is as follows:


```

1 <?xml version="1.0" encoding="UTF-8"?>
2 <sbml xmlns="http://www.sbml.org/sbml/level3/version1/core" level="3" version="1"
3   xmlns:math="http://www.w3.org/1998/Math/MathML"
4   xmlns:sbml="http://www.sbml.org/sbml/level3/version1/core">
5   <model substanceUnits="mole" volumeUnits="litre" timeUnits="second"
6     extentUnits="mole">
7     <listOfUnitDefinitions>
8       <unitDefinition id="per_second">
9         <listOfUnits>
10           <unit kind="second" exponent="-1" scale="0" multiplier="1"/>
11         </listOfUnits>
12       </unitDefinition>
13       <unitDefinition id="concentration">
14         <listOfUnits>
15           <unit kind="mole" exponent="1" scale="0" multiplier="1"/>
16           <unit kind="litre" exponent="-1" scale="0" multiplier="1"/>
17         </listOfUnits>
18       </unitDefinition>
19     </listOfUnitDefinitions>
20     <listOfCompartments>
21       <compartment id="cell" size="1" spatialDimensions="3" constant="true"/>
22     </listOfCompartments>
23     <listOfSpecies>
24       <species id="P1" compartment="cell" initialConcentration="0"
25         hasOnlySubstanceUnits="false" boundaryCondition="false"
26         constant="false"/>
27       <species id="P2" compartment="cell" initialConcentration="0"
28         hasOnlySubstanceUnits="false" boundaryCondition="false"
29         constant="false"/>
30     </listOfSpecies>
31     <listOfParameters>
32       <parameter id="k1" value="1" units="per_second" constant="true"/>
33       <parameter id="k2" value="1" units="per_second" constant="true"/>
34       <parameter id="tau" value="0.25" units="concentration" constant="true"/>
35       <parameter id="G1" value="1" units="concentration" constant="false"/>
36       <parameter id="G2" value="0" units="concentration" constant="false"/>
37     </listOfParameters>
38     <listOfRules>
39       <rateRule variable="P1">
40         <math:math>
41           <math:apply>
42             <math:times/>
43             <math:ci> k1 </math:ci>
44             <math:apply>
45               <math:minus/>
46               <math:ci> G1 </math:ci>
47               <math:ci> P1 </math:ci>
48             </math:apply>
49           </math:apply>
50         </math:math>
51       </rateRule>
52       <rateRule variable="P2">
53         <math:math>
54           <math:apply>
55             <math:times/>
56             <math:ci> k2 </math:ci>
57             <math:apply>
58               <math:minus/>
59               <math:ci> G2 </math:ci>
60               <math:ci> P2 </math:ci>
61             </math:apply>
62           </math:apply>
63         </math:math>
64       </rateRule>
65     </listOfRules>
66     <listOfEvents>
67       <event useValuesFromTriggerTime="true">
68         <trigger>

```

```

1         <math:math>
2             <math:apply>
3                 <math:gt/>
4                 <math:ci> P1 </math:ci>
5                 <math:ci> tau </math:ci>
6             </math:apply>
7         </math:math>
8     </trigger>
9     <listOfEventAssignments>
10         <eventAssignment variable="G2">
11             <math:math>
12                 <math:cn sbml:units="concentration"> 1 </math:cn>
13             </math:math>
14         </eventAssignment>
15     </listOfEventAssignments>
16 </event>
17 <event useValuesFromTriggerTime="true">
18     <trigger>
19         <math:math>
20             <math:apply>
21                 <math:leq/>
22                 <math:ci> P1 </math:ci>
23                 <math:ci> tau </math:ci>
24             </math:apply>
25         </math:math>
26     </trigger>
27     <listOfEventAssignments>
28         <eventAssignment variable="G2">
29             <math:math>
30                 <math:cn sbml:units="concentration"> 0 </math:cn>
31             </math:math>
32         </eventAssignment>
33     </listOfEventAssignments>
34 </event>
35 </listOfEvents>
36 </model>
37 </sbml>

```

7.10 Example involving two-dimensional compartments

The following example is a model that uses a two-dimensional compartment. It is a fragment of a larger model of calcium regulation across the plasma membrane of a cell. The model includes a calcium influx channel, “Ca_channel”, and a calcium-extruding PMCA pump, “Ca_Pump”. It also includes two cytosolic proteins that buffer calcium via the “CalciumCalbindin_gt_BoundCytosol” and “CalciumBuffer_gt_BoundCytosol” reactions. Finally, the rate expressions in this model do not include explicit factors of the compartment volumes; instead, the various rate constants are assume to include any necessary corrections for volume.

```

45 <?xml version="1.0" encoding="UTF-8"?>
46 <sbml xmlns="http://www.sbml.org/sbml/level3/version1/core" level="3" version="1">
47     <model id="facilitated_ca_diffusion" substanceUnits="substance"
48         areaUnits="area" volumeUnits="litre" timeUnits="second" extentUnits="substance">
49         <listOfUnitDefinitions>
50             <unitDefinition id="substance">
51                 <listOfUnits>
52                     <unit kind="mole" exponent="1" scale="-6" multiplier="1"/>
53                 </listOfUnits>
54             </unitDefinition>
55             <unitDefinition id="area">
56                 <listOfUnits>
57                     <unit kind="metre" exponent="2" scale="-6" multiplier="1"/>
58                 </listOfUnits>
59             </unitDefinition>
60             <unitDefinition id="per_second">
61                 <listOfUnits>
62                     <unit kind="second" exponent="-1" scale="0" multiplier="1"/>
63                 </listOfUnits>
64             </unitDefinition>

```

```

1      <unitDefinition id="litre_per_mole_per_second">
2          <listOfUnits>
3              <unit kind="mole"    exponent="-1" scale="-6" multiplier="1"/>
4              <unit kind="litre"   exponent="1"  scale="0"  multiplier="1"/>
5              <unit kind="second"  exponent="-1" scale="0"  multiplier="1"/>
6          </listOfUnits>
7      </unitDefinition>
8      <unitDefinition id="subs_per_vol">
9          <listOfUnits>
10             <unit kind="mole"    exponent="1"  scale="-6" multiplier="1"/>
11             <unit kind="litre"   exponent="-1" scale="0"  multiplier="1"/>
12         </listOfUnits>
13     </unitDefinition>
14 </listOfUnitDefinitions>
15 <listOfCompartments>
16     <compartment id="Extracellular"
17         spatialDimensions="3" size="1" constant="true"/>
18     <compartment id="PlasmaMembrane"
19         spatialDimensions="2" size="1"
20         constant="true"/>
21     <compartment id="Cytosol"
22         spatialDimensions="3" size="1"
23         constant="true"/>
24 </listOfCompartments>
25 <listOfSpecies>
26     <species id="CaBPB_C" compartment="Cytosol" initialConcentration="47.17"
27         hasOnlySubstanceUnits="false" boundaryCondition="false"
28         constant="false"/>
29     <species id="B_C" compartment="Cytosol" initialConcentration="396.04"
30         hasOnlySubstanceUnits="false" boundaryCondition="false"
31         constant="false"/>
32     <species id="CaB_C" compartment="Cytosol" initialConcentration="3.96"
33         hasOnlySubstanceUnits="false" boundaryCondition="false"
34         constant="false"/>
35     <species id="Ca_C" name="Ca" compartment="Cytosol" initialConcentration="0.1"
36         hasOnlySubstanceUnits="false" boundaryCondition="false"
37         constant="false"/>
38     <species id="Ca_EC" name="Ca" compartment="Extracellular"
39         initialConcentration="1000"
40         hasOnlySubstanceUnits="false" boundaryCondition="false"
41         constant="false"/>
42     <species id="CaCh_PM" compartment="PlasmaMembrane" initialConcentration="1"
43         hasOnlySubstanceUnits="false" boundaryCondition="false"
44         constant="false"/>
45     <species id="CaPump_PM" compartment="PlasmaMembrane" initialConcentration="1"
46         hasOnlySubstanceUnits="false" boundaryCondition="false"
47         constant="false"/>
48     <species id="CaBP_C" compartment="Cytosol" initialConcentration="202.83"
49         hasOnlySubstanceUnits="false" boundaryCondition="false"
50         constant="false"/>
51 </listOfSpecies>
52 <listOfReactions>
53     <reaction id="CalciumCalbindin_gt_BoundCytosol" reversible="true" fast="true">
54         <listOfReactants>
55             <speciesReference species="CaBP_C" stoichiometry="1" constant="true"/>
56             <speciesReference species="Ca_C" stoichiometry="1" constant="true"/>
57         </listOfReactants>
58         <listOfProducts>
59             <speciesReference species="CaBPB_C" stoichiometry="1" constant="true"/>
60         </listOfProducts>
61         <kineticLaw>
62             <notes>
63                 <p xmlns="http://www.w3.org/1999/xhtml">
64                     (((Kf_CalciumCalbindin_BoundCytosol * CaBP_C) * Ca_C) -
65                     (Kr_CalciumCalbindin_BoundCytosol * CaBPB_C))
66                 </p>
67             </notes>
68             <math xmlns="http://www.w3.org/1998/Math/MathML">
69                 <apply>

```

```

1      <times/>
2      <ci> Cytosol </ci>
3      <apply>
4          <minus/>
5          <apply>
6              <times/>
7              <ci> Kf_CalciumCalbindin_BoundCytosol </ci>
8              <ci> CaBP_C </ci>
9              <ci> Ca_C </ci>
10             </apply>
11             <apply>
12                 <times/>
13                 <ci> Kr_CalciumCalbindin_BoundCytosol </ci>
14                 <ci> CaBPB_C </ci>
15             </apply>
16         </apply>
17     </math>
18     <listOfLocalParameters>
19         <localParameter id="Kf_CalciumCalbindin_BoundCytosol" value="20.0"
20             units="litre_per_mole_per_second"/>
21         <localParameter id="Kr_CalciumCalbindin_BoundCytosol" value="8.6"
22             units="per_second"/>
23     </listOfLocalParameters>
24 </kineticLaw>
25 </reaction>
26 <reaction id="CalciumBuffer_gt_BoundCytosol" reversible="true" fast="true">
27     <listOfReactants>
28         <speciesReference species="Ca_C" stoichiometry="1" constant="true"/>
29         <speciesReference species="B_C" stoichiometry="1" constant="true"/>
30     </listOfReactants>
31     <listOfProducts>
32         <speciesReference species="CaB_C" stoichiometry="1" constant="true"/>
33     </listOfProducts>
34     <kineticLaw>
35         <notes>
36             <p xmlns="http://www.w3.org/1999/xhtml">
37                 (((Kf_CalciumBuffer_BoundCytosol * Ca_C) * B_C) -
38                 (Kr_CalciumBuffer_BoundCytosol * CaB_C))
39             </p>
40         </notes>
41         <math xmlns="http://www.w3.org/1998/Math/MathML">
42             <apply>
43                 <times/>
44                 <ci> Cytosol</ci>
45                 <apply>
46                     <minus/>
47                     <apply>
48                         <times/>
49                         <ci> Kf_CalciumBuffer_BoundCytosol </ci>
50                         <ci> Ca_C </ci>
51                         <ci> B_C </ci>
52                     </apply>
53                     <apply>
54                         <times/>
55                         <ci> Kr_CalciumBuffer_BoundCytosol </ci>
56                         <ci> CaB_C </ci>
57                     </apply>
58                 </apply>
59             </math>
60         </math>
61         <listOfLocalParameters>
62             <localParameter id="Kf_CalciumBuffer_BoundCytosol" value="0.1"
63                 units="litre_per_mole_per_second"/>
64             <localParameter id="Kr_CalciumBuffer_BoundCytosol" value="1.0"
65                 units="per_second"/>
66         </listOfLocalParameters>
67     </kineticLaw>
68 </reaction>
69

```

```

1      <reaction id="Ca_Pump" reversible="true" fast="false">
2          <listOfReactants>
3              <speciesReference species="Ca_C" stoichiometry="1" constant="true"/>
4          </listOfReactants>
5          <listOfProducts>
6              <speciesReference species="Ca_EC" stoichiometry="1" constant="true"/>
7          </listOfProducts>
8          <listOfModifiers>
9              <modifierSpeciesReference species="CaPump_PM"/>
10         </listOfModifiers>
11         <kineticLaw>
12             <notes>
13                 <p xmlns="http://www.w3.org/1999/xhtml">
14                     ((Vmax * kP * ((Ca_C - Ca_Rest) / (Ca_C + kP)) /
15                      (Ca_Rest + kP)) * CaPump_PM)
16                 </p>
17             </notes>
18             <math xmlns="http://www.w3.org/1998/Math/MathML">
19                 <apply>
20                     <times/>
21                     <ci> PlasmaMembrane</ci>
22                     <apply>
23                         <divide/>
24                         <apply>
25                             <times/>
26                             <ci> Vmax </ci>
27                             <ci> kP </ci>
28                             <ci> CaPump_PM </ci>
29                         </apply>
30                         <minus/>
31                         <ci> Ca_C </ci>
32                         <ci> Ca_Rest </ci>
33                     </apply>
34                 </apply>
35                 <apply>
36                     <times/>
37                     <apply>
38                         <plus/>
39                         <ci> Ca_C </ci>
40                         <ci> kP </ci>
41                     </apply>
42                     <apply>
43                         <plus/>
44                         <ci> Ca_Rest </ci>
45                         <ci> kP </ci>
46                     </apply>
47                 </apply>
48             </math>
49             <listOfLocalParameters>
50                 <localParameter id="Vmax" value="4000" units="per_second"/>
51                 <localParameter id="kP" value="0.25" units="subs_per_vol"/>
52                 <localParameter id="Ca_Rest" value="0.1" units="subs_per_vol"/>
53             </listOfLocalParameters>
54         </kineticLaw>
55     </reaction>
56     <reaction id="Ca_channel" reversible="true" fast="false">
57         <listOfReactants>
58             <speciesReference species="Ca_EC" stoichiometry="1" constant="true"/>
59         </listOfReactants>
60         <listOfProducts>
61             <speciesReference species="Ca_C" stoichiometry="1" constant="true"/>
62         </listOfProducts>
63         <listOfModifiers>
64             <modifierSpeciesReference species="CaCh_PM"/>
65         </listOfModifiers>
66         <kineticLaw>
67             <notes>
68

```


```

1      <p xmlns="http://www.w3.org/1999/xhtml">
2          (J0 * Kc * (Ca_EC - Ca_C) / (Kc + Ca_C) * CaCh_PM)
3      </p>
4  </notes>
5  <math xmlns="http://www.w3.org/1998/Math/MathML">
6  <apply>
7      <times/>
8      <ci> PlasmaMembrane </ci>
9      <apply>
10         <divide/>
11         <apply>
12             <times/>
13             <ci> CaCh_PM </ci>
14             <ci> J0 </ci>
15             <ci> Kc </ci>
16             <apply>
17                 <minus/>
18                 <ci> Ca_EC </ci>
19                 <ci> Ca_C </ci>
20             </apply>
21         </apply>
22         <plus/>
23         <ci> Kc </ci>
24         <ci> Ca_C </ci>
25     </apply>
26 </apply>
27 </math>
28 </listOfLocalParameters>
29 <localParameter id="J0" value="0.014" units="litre_per_mole_per_second"/>
30 <localParameter id="Kc" value="0.5" units="subs_per_vol"/>
31 </listOfLocalParameters>
32 </kineticLaw>
33 </reaction>
34 </listOfReactions>
35 </model>
36 </sbml>

```

7.11 Example of a reaction located at a membrane

This section describes a model containing one single enzymatic reaction where substrate and product are located in the same compartment but the enzyme is localized at the membrane surrounding the compartment.



The model contains two compartments, a three dimensional one called “cytosol” and a two dimensional one called “membrane” that is assumed to be the boundary of the cell. The reaction R has a substrate S and a product P that are both located in the cytosol. The enzyme E that catalyzes the reactions is located at the membrane. The kinetic law of reaction R is

$$v = A \cdot \frac{k_{cat} \cdot [E] \cdot [S]}{K_M + [S]}$$

where A is the area of the membrane (measured in μm^2), $[E]$ is the *density* of the enzyme on the membrane (in $\mu mol \mu m^{-2}$), $[S]$ is the *concentration* of the substrate (in $\mu mol l^{-1}$), K_M the Michaelis-Menten constant (also in $\mu mol l^{-1}$), and k_{cat} the rate constant (in min^{-1}). The units of the result of the kinetic law are in $\mu mol min^{-1}$. Since the units for the amounts of all species (S , P , and E) and for the reaction extent are the same (μmol) the model does not require unit conversion factors.

The kinetic law as it is given here scales correctly for changes in cytosol volume, membrane area, or enzyme density. This means that if one of these values is changed (even if it is variable during a simulation) the rate law is still valid.

The following is the text of the model’s SBML representation.

```

1 <?xml version="1.0" encoding="UTF-8"?>
2 <sbml xmlns="http://www.sbml.org/sbml/level3/version1/core" level="3" version="1">
3   <model id="Model_1" name="Reaction on membrane" substanceUnits="micromole"
4     timeUnits="minute" extentUnits="micromole">
5     <listOfFunctionDefinitions>
6       <functionDefinition id="MM_enzyme" name="MM_enzyme">
7         <math xmlns="http://www.w3.org/1998/Math/MathML">
8           <lambda>
9             <bvar> <ci> size </ci> </bvar>
10            <bvar> <ci> k </ci> </bvar>
11            <bvar> <ci> enz </ci> </bvar>
12            <bvar> <ci> subs </ci> </bvar>
13            <bvar> <ci> Km </ci> </bvar>
14            <apply>
15              <divide/>
16              <apply>
17                <times/>
18                <ci> size </ci>
19                <ci> k </ci>
20                <ci> enz </ci>
21                <ci> subs </ci>
22              </apply>
23            <apply>
24              <plus/>
25              <ci> Km </ci>
26              <ci> subs </ci>
27            </apply>
28          </apply>
29        </lambda>
30      </math>
31    </functionDefinition>
32  </listOfFunctionDefinitions>
33  <listOfUnitDefinitions>
34    <unitDefinition id="minute">
35      <listOfUnits>
36        <unit kind="second" exponent="1" scale="0" multiplier="60"/>
37      </listOfUnits>
38    </unitDefinition>
39    <unitDefinition id="per_minute">
40      <listOfUnits>
41        <unit kind="second" exponent="-1" scale="0" multiplier="60"/>
42      </listOfUnits>
43    </unitDefinition>
44    <unitDefinition id="micromole">
45      <listOfUnits>
46        <unit kind="mole" exponent="1" scale="-6" multiplier="1"/>
47      </listOfUnits>
48    </unitDefinition>
49    <unitDefinition id="micromole_per_l">
50      <listOfUnits>
51        <unit kind="mole" exponent="1" scale="-6" multiplier="1"/>
52        <unit kind="litre" exponent="-1" scale="0" multiplier="1"/>
53      </listOfUnits>
54    </unitDefinition>
55    <unitDefinition id="sqrmicrometre">
56      <listOfUnits>
57        <unit kind="metre" exponent="2" scale="-6" multiplier="1"/>
58      </listOfUnits>
59    </unitDefinition>
60  </listOfUnitDefinitions>
61  <listOfCompartments>
62    <compartment id="cyt" name="Cytosol"
63      spatialDimensions="3" units="litre"
64      size="1e-15" constant="true"/>
65    <compartment id="mem" name="Membrane"
66      spatialDimensions="2" units="sqrmicrometre"
67      size="1" constant="true"/>
68  </listOfCompartments>

```



```

1      <listOfSpecies>
2        <species id="species_1" name="substrate" compartment="cyt"
3          hasOnlySubstanceUnits="false" boundaryCondition="false"
4          initialConcentration="1" constant="false"/>
5        <species id="species_2" name="product" compartment="cyt"
6          hasOnlySubstanceUnits="false" boundaryCondition="false"
7          initialConcentration="1" constant="false"/>
8        <species id="species_3" name="enzyme" compartment="mem"
9          hasOnlySubstanceUnits="false" boundaryCondition="false"
10         initialConcentration="1" constant="false"/>
11      </listOfSpecies>
12      <listOfReactions>
13        <reaction id="reaction_1" name="Reaction" reversible="false"
14          fast="false" compartment="mem">
15          <listOfReactants>
16            <speciesReference species="species_1" stoichiometry="1" constant="true"/>
17          </listOfReactants>
18          <listOfProducts>
19            <speciesReference species="species_2" stoichiometry="1" constant="true"/>
20          </listOfProducts>
21          <listOfModifiers>
22            <modifierSpeciesReference species="species_3"/>
23          </listOfModifiers>
24          <kineticLaw>
25            <math xmlns="http://www.w3.org/1998/Math/MathML">
26              <apply>
27                <ci> MM_enzyme </ci>
28                <ci> mem </ci>
29                <ci> k </ci>
30                <ci> species_3 </ci>
31                <ci> species_1 </ci>
32                <ci> Km </ci>
33              </apply>
34            </math>
35            <listOfParameters>
36              <localParameter id="k" value="0.1" units="per_minute"/>
37              <localParameter id="Km" value="0.1" units="micromole_per_l"/>
38            </listOfParameters>
39          </kineticLaw>
40        </reaction>
41      </listOfReactions>
42    </model>
43  </sbml>

```

8 Recommended practices

In this section, we recommend a number of practices for using and interpreting various SBML constructs. These recommendations are non-normative, but we advocate them strongly; ignoring them will not render a model invalid, but may hinder interoperability between different software systems exchanging SBML content.

8.1 Recommended practices concerning common SBML attributes and objects

Many SBML components share some or all of the following attributes and objects. We describe recommendations concerning them here, separately from discussing the specific SBML components. In Section 8.2, we turn to the specific SBML components, but the recommendations described here also apply to them.

8.1.1 Identifiers and names

The **id** attribute is available on most (but not all) objects in SBML, and all objects that have **id** attributes also have an optional **name** attribute. How should models treat identifiers and names?

The following is the recommended practice for handling **name**. If a software tool has the capability for displaying the content of **name** attributes, it should display this content to the user as a component's label instead of the component's **id**. If the user interface does not have this capability (e.g., because it cannot display or use special characters in symbol names), or if the **name** attribute is missing on a given component, then the user interface should display the value of the **id** attribute instead. (Script language interpreters are especially likely to display **id** instead of **name**.)

As a consequence of the above, authors of software systems that automatically generate values for **id** attributes should be aware some other systems may display the **id**'s to the user. Authors therefore may wish to take some care to have their software create **id** values that are: (a) reasonably easy for humans to type and read, and (b) likely to be meaningful (e.g., by making the **id** attribute is an abbreviated form of the **name** attribute value).

8.1.2 Initial Values

SBML allows for the creation of **Compartment**, **Species**, **Parameter**, **LocalParameter** and **SpeciesReference** objects without declaring their initial values directly on the object instances. That is, a **Compartment** object can be created without defining a value for its **size** attribute; a **Species** object can be created without defining a value for either its **initialConcentration** or **initialAmount** attribute; **Parameter** and **LocalParameter** objects can be created without giving a value to their **value** attributes; and a **SpeciesReference** object can be created without assigning a value to its **stoichiometry** attribute. A missing value in the case of **Compartment**, **Species**, **Parameter**, and **SpeciesReference** objects implies that the value is either set via an **InitialAssignment** object elsewhere in the model, or is meant to be obtained from an external source (e.g., by querying the user of a software system), or is unknown. In the case of **LocalParameter** objects, a missing value implies that the value is either unknown or meant to be obtained from an external source.

Where initial values are available and are scalar numbers that *can* be set using the appropriate attribute on an object, the best practice recommendation is to do that in preference to using an **InitialAssignment** construct if there is no particular reason to use **InitialAssignment**. Setting the relevant attribute directly on the **Compartment**, **Species**, and **Parameter** and **SpeciesReference** object is simpler and may be more interoperable with different software systems. This is especially true of **stoichiometry** on **SpeciesReference**, which in the vast majority of models, is never more than a scalar constant value anyway.

8.1.3 The constant flag

There is a mandatory boolean attribute called **constant** on the **Compartment**, **Species**, **SpeciesReference** and **Parameter** components. A value of “**true**” means that the SBML object in question will not be changed by other constructs in SBML (except possibly an **InitialAssignment**). A value of “**false**” indicates an intention to change the element's value by an **AssignmentRule**, **RateRule**, **AlgebraicRule**, **Reaction** or **Event** in the model.

A **constant** attribute value of “false” does not *require* that the object in question is changed; strictly speaking, an SBML model is valid even if it sets all **constant** attributes to “false” but never actually modify any of the values. However, the best practice recommendation is to communicate intentions by setting **constant** to “true” unless an entity in a model really is intended to be changed. The exception to this is **Species**, which are usually part of the reaction system and thus usually need to have **constant**=“false”.

8.1.4 Annotations

Appropriate uses of annotations

In the description of **SBase**’s **annotation** element, we already made the point that it is critical not to put data essential to understanding a model into an **annotation**. This raises a question: what kind of data may be appropriately put into annotations?

Here are examples of the kinds of data that may be appropriately stored in **annotation** elements:

- Application-specific processing instructions that do not change the essential meaning of a model, but help a particular application with tasks such as managing the model, maintaining state data across editing sessions, etc.
- Identification information for cross-referencing components in a model with items in a data resource such as a database. This is the purpose of the annotation scheme described in Section 6.
- Evidence codes for annotating a model with controlled vocabulary terms that indicate (e.g.) the quality of biological evidence supporting the inclusion of each component in the model. The annotation scheme of Section 6 can be used in this capacity.
- Information about the model that cannot be readily encoded in existing SBML elements, but that does not alter the mathematical meaning of the model.

Specificity of annotations

The annotation data (Section 3.2.4) attached to a specific SBML object in a model is assumed by other applications to be directly associated with that SBML object. Therefore, it is important to decompose and locate annotation data appropriately in an SBML document. Applications are advised to avoid encoding all their annotations in a single top-level attribute. The data associated with, for example, an individual **Species** object in a model should be encoded in the **<annotation>** element enclosed within the SBML **<species>** element representing that species in the SBML file.

Syntax of annotations

The annotation scheme described in Section 6 is useful for many, but not all, situations. It is tempting to develop new annotation syntaxes for situations that fall outside the scope of the SBML MIRIAM annotation scheme. However, a proliferation of proprietary annotation schemes will hinder software interoperability in the long run.

We recommend the following approach when faced with a need to use alternate annotation syntaxes:

1. The modular nature of SBML Level 3 Version 1 Core means that data that in SBML Level 2 could only be stored in annotations may now be supported using a full SBML Level 3 package. Therefore, software developers and modelers should first check if there already exists a package that may serve their needs. A list of SBML Level 3 packages is always maintained at the SBML website, <http://sbml.org>.
2. If no package exists, developers and modelers may wish to check if someone else has already developed a similar annotation syntax for use with another software system. A list of known SBML annotation schemes is maintained online at http://sbml.org/Community/Wiki/Known_SBML_annotations.
3. If none of the above alternatives provide a satisfactory result, developers and modelers should query the SBML discussion list (sbml-discuss@caltech.edu) to see if anyone else has been faced with similar problems. Other SBML users may have insights or even partial solutions already available.

8.2 Recommended practices concerning specific SBML components

In this section, we describe expectations and recommendations concerning specific SBML components. We do not reiterate the recommendations presented in Section 8.1, but they apply to many of the SBML components discussed here and should be kept in mind. The order of the components discussed here follows the order of their presentation in Section 4, but we only include here those components for which we have specific recommendations.

8.2.1 Unit definitions

We advise modelers and software tools to declare the units of all quantities in a model, insofar as this is possible, using the various mechanisms provided for this in SBML. Fully declared units can allow software tools to perform dimensional analysis on the units of mathematical expressions, and such analysis can be valuable in helping modelers produce correct models. In addition, it can allow model-wide operations such as conversion or rescaling of units.

Recommendations for choices of units

Table 7 lists the units recommended for different SBML components.

Component or attribute	Unit recommendations	
Model substanceUnits	mole, item, avogadro, dimensionless, kilogram, gram, or units derived from these	
Model timeUnits	second, dimensionless, or units derived from these	
Model volumeUnits	litre, metre ³ , dimensionless, or units derived from these	
Model areaUnits	metre ² , dimensionless, or units derived from these	
Model lengthUnits	metre, dimensionless, or units derived from these	
Model extentUnits	mole, item, avogadro, dimensionless, or units derived from these	
Compartment units	Value of attribute spatialDimensions	Recommended units
	"3"	litre, metre ³ , dimensionless, or units derived from these
	"2"	metre ² , dimensionless, or units derived from these
	"1"	metre, dimensionless, or units derived from these
	other	<i>no specific recommendations</i>
Species substanceUnits	mole, item, avogadro, dimensionless, kilogram, gram, or units derived from these	
Parameter units	<i>no specific recommendations</i>	

Table 7: Units recommended for use on different SBML model components.

Handling units requiring the use of offsets

As mentioned in Section 4.4.2 Unit definitions and conversions requiring offsets cannot be done using the simple approach described in 4.4.3. In fact, SBML does not included Celsius as a predefined unit because its definition requires the use of an offset. Unit definitions involving Celsius, Fahrenheit or other units with offsets require a different approach.

We first address the question of how to handle units that *do* require offsets:

- *Handling Celsius.* A model in which certain quantities are temperatures measured in degrees Celsius can be converted straightforwardly to a model in which those temperatures are in kelvin. A software tool could do this by performing a substitution using the following relationship:

$$T_{kelvin} = T_{Celsius} + 273.15$$

In every mathematical formula of the model where a quantity (call it x) in degrees Celsius appears, replace x with $x_k + 273.15$ where x_k is now in kelvin. An alternative approach would be to use a [FunctionDefinition](#) to define a function encapsulating this relationship above and then using that in the rest of the model as needed. Since Celsius is a commonly-used unit, software tools could help users by providing users with the ability to express temperatures in Celsius in the tools' interfaces, and making substitutions automatically when writing out the SBML.

- *Handling other units requiring offsets.* The only other units requiring offsets in SBML's domain of common applications are other temperature units such as Fahrenheit. Few modern scientists employ Fahrenheit degrees; therefore, this is an unusual situation. The complication inherent in converting between degrees Fahrenheit and kelvin is that both a multiplier and an offset are required:

$$T_{kelvin} = \frac{T_F + 459.67}{1.8}$$

One approach to handling this is to use a [FunctionDefinition](#) to define a function encapsulating the relationship above, then to substitute a call to this function wherever the original temperature in Fahrenheit appears in the model's mathematical formulas. Here is a candidate definition as an example:

```
<functionDefinition id="Fahrenheit_to_kelvin">
  <math xmlns="http://www.w3.org/1998/Math/MathML"
    xmlns:sbml="http://www.sbml.org/sbml/level3/version1/core">
    <lambda>
      <bvar><ci> temp_in_fahrenheit </ci></bvar>
      <apply>
        <divide/>
        <apply>
          <plus/>
          <ci> temp_in_fahrenheit </ci>
          <cn> 459.67 </cn>
        </apply>
        <cn> 1.8 </cn>
      </apply>
    </lambda>
  </math>
</functionDefinition>
```

An alternative approach not requiring the use of function definitions is to use an [AssignmentRule](#) for each variable in Fahrenheit units. The [AssignmentRule](#) could compute the conversion from Fahrenheit to (say) kelvin, assign its value to a variable (in Kelvin units), and then that variable could be used elsewhere in the model. Still another approach is to rewrite the mathematical formulas of a model to directly incorporate the conversion above wherever the quantity appears.

All of these approaches provide general solutions to the problem of supporting any units requiring offsets in the unit system of SBML Level 3. It can be used for other temperature units requiring an offset (e.g., degrees Rankine, degrees Réaumur), although the likelihood of a real-life model requiring such other temperature units seems exceedingly small.

In summary, the fact that SBML units do not support specifying an offset does not impede the creation of models using alternative units. If conversions are needed, then converting between temperature in degrees Celsius and thermodynamic temperature can be handled rather easily by the simple substitution described above. For the rarer case of Fahrenheit and other units requiring combinations of multipliers and offsets, users are encouraged to employ the power of [FunctionDefinition](#), [AssignmentRule](#), or other constructs in SBML.

8.2.2 Compartments

Setting the size attribute on a compartment

As mentioned in Section 4.5.3, we highly recommend that every **Compartment** object in a model has its **size** set. There are three major technical reasons for this. First, if the model contains any species whose initial amounts are given in terms of concentrations, and there is at least one reaction in the model referencing such a species, then the model will be numerically incomplete if it lacks a value for the size of the compartment in which the species is located. The reason is that SBML reactions are expected to be in terms of intensive properties such as *amount/time* (or more generally, *extent units/time units*; see Section 4.11.7), and converting from concentration to amount requires the compartment size. Second, models ideally should be instantiable in a variety of simulation frameworks. A commonly-used one is the discrete stochastic framework (Gillespie, 1977; Wilkinson, 2006) in which species are represented as item counts (e.g., molecule counts). If species' initial quantities are given in terms of concentrations or densities, it is impossible to convert the values to item counts without knowing compartment sizes. Third, if a model contains multiple compartments whose sizes are not all identical to each other, it is impossible to quantify the reaction rate expressions without knowing the compartment volumes. The reason for the latter is again that reaction rates in SBML are defined in terms *extent/time*, and when species quantities are given in terms of concentrations or densities, the compartment sizes usually become factors in the reaction rate expressions.

Indicating a default compartment

Some types of models do not use compartments, for example because they factor out volumes completely. Since SBML requires at least one compartment to be defined if any species exists in a model, the representation of models where no compartments are needed sometimes leaves model creators wishing they could indicate that a compartment is only a “default” in some sense. The recommended approach to handling this situation is to annotate the **Compartment** object by setting its **sboTerm** attribute to an appropriate SBO term, specifically “SBO:0000410”.

8.2.3 Rules

Section 4.9.5 establishes the fact that when **AlgebraicRule** objects are used, it is possible to produce a model that is overdetermined. When a model includes both **Event** and **Reaction** objects, it is necessary to analyze the set of equations produced from the rules and reactions and the set of equations produces from rules and the event assignments of each event. Each set of equations must not be overdetermined. In addition, each set of equations must be fully determined if accurate simulation is to be performed.

The following example illustrates a case where the set of equations is fully determined. First, we present the SBML expression of the model:

```
<?xml version="1.0" encoding="UTF-8"?>
<sbml xmlns="http://www.sbml.org/sbml/level3/version1/core" level="3" version="1">
  <model id="example" substanceUnits="mole" volumeUnits="litre"
    timeUnits="second" extentUnits="model">
    <listOfUnitDefinitions>
      <unitDefinition id="conc">
        <listOfUnits>
          <unit kind="mole" multiplier="1" scale="0" exponent="1"/>
          <unit kind="litre" multiplier="1" scale="0" exponent="-1"/>
        </listOfUnits>
      </unitDefinition>
      <unitDefinition id="per_second">
        <listOfUnits>
          <unit kind="second" exponent="-1" scale="0" multiplier="1"/>
        </listOfUnits>
      </unitDefinition>
    </listOfUnitDefinitions>
    <listOfCompartments>
      <compartment id="C" size="1" spatialDimensions="3" constant="true"/>
    </listOfCompartments>
    <listOfSpecies>
```

```

1      <species id="S1" compartment="C" initialConcentration="1"
2          boundaryCondition="false" hasOnlySubstanceUnits="false"
3          constant="false"/>
4      <species id="S2" compartment="C" initialConcentration="0"
5          boundaryCondition="false" hasOnlySubstanceUnits="false"
6          constant="false"/>
7      <species id="S3" compartment="C" initialConcentration="0"
8          boundaryCondition="false" hasOnlySubstanceUnits="false"
9          constant="false"/>
10     </listOfSpecies>
11     <listOfParameters>
12         <parameter id="p1" value="1" constant="true" units="conc"/>
13         <parameter id="p2" value="1.5" constant="true" units="conc"/>
14     </listOfParameters>
15     <listOfRules>
16         <algebraicRule>
17             <math xmlns="http://www.w3.org/1998/Math/MathML">
18                 <apply> <minus/> <ci> S1 </ci> <ci> S3 </ci> </apply>
19             </math>
20         </algebraicRule>
21     </listOfRules>
22     <listOfReactions>
23         <reaction id="R" reversible="true" fast="false">
24             <listOfReactants>
25                 <speciesReference species="S1" stoichiometry="1" constant="true"/>
26             </listOfReactants>
27             <listOfProducts>
28                 <speciesReference species="S2" stoichiometry="1" constant="true"/>
29             </listOfProducts>
30             <kineticLaw>
31                 <math xmlns="http://www.w3.org/1998/Math/MathML">
32                     <apply> <times/> <ci> C </ci> <ci> k1 </ci> <ci> S1 </ci> </apply>
33                 </math>
34                 <listOfLocalParameters>
35                     <localParameter id="k1" value="0.1" units="per_second"/>
36                 </listOfLocalParameters>
37             </kineticLaw>
38         </reaction>
39     </listOfReactions>
40     <listOfEvents>
41         <event useValuesFromTriggerTime="true">
42             <trigger>
43                 <math xmlns="http://www.w3.org/1998/Math/MathML">
44                     <apply> <gt/> <ci> S2 </ci> <ci> p1 </ci> </apply>
45                 </math>
46             </trigger>
47             <listOfEventAssignments>
48                 <eventAssignment variable="S1">
49                     <math xmlns="http://www.w3.org/1998/Math/MathML">
50                         <ci> p1 </ci>
51                     </math>
52                 </eventAssignment>
53                 <eventAssignment variable="S2">
54                     <math xmlns="http://www.w3.org/1998/Math/MathML">
55                         <ci> p2 </ci>
56                     </math>
57                 </eventAssignment>
58             </listOfEventAssignments>
59         </event>
60     </listOfEvents>
61 </model>
62 </sbml>

```

There are three species in the model above whose values may vary. The first set of equations to consider is the set produced by the [Reaction](#) and the [AlgebraicRule](#) objects:

$$\frac{d[S_1]}{dt} = -C \cdot k_1 \cdot [S_1]$$

$$\begin{aligned}\frac{d[S_2]}{dt} &= C \cdot k_1 \cdot [S_1] \\ [S_1] - [S_3] &= 0\end{aligned}$$

This set of equations is fully determined, i.e., each of the three variables S_1 , S_2 and S_3 are derived from one equation. The second set of equations to consider is the set produced by the **Event** and the **AlgebraicRule** objects:

$$\begin{aligned}[S_1] &= 1 \\ [S_2] &= 1.5 \\ [S_1] - [S_3] &= 0\end{aligned}$$

Again the set of equations is fully determined, but had the event assignment for species S_1 been absent, the algebraic rule would produce an ambiguity regarding which variable should be adjusted.


In this example, as is often the case when an **AlgebraicRule** has been used, the **AlgebraicRule** could be replaced by an **AssignmentRule**.

```
<assignmentRule variable="S3">
  <math xmlns="http://www.w3.org/1998/Math/MathML">
    <apply>
      <ci> S1 </ci>
    </apply>
  </math>
</assignmentRule>
```

Replacing **AlgebraicRule** objects with **AssignmentRule** objects, particularly in models that use events, reduces the possibilities for creating either overdetermined or ambiguous models and produces models that can be exchanged with greater ease.

8.2.4 Reactions

Consider a very simple model consisting of a single enzymatic reaction R that converts S_1 to S_2 for which a traditional kinetic law v_R is given:



where

$$v_R = \frac{v_{\max} \cdot [S_1]}{K_M + [S_1]}$$

with v_R and v_{\max} given in units of concentration per time.

As mentioned above, when a rate law is presented in the traditional way, it usually embodies (implicitly or explicitly) several assumptions: that all species are located in the same compartment, that the compartment size does not change, and that the reaction takes place uniformly throughout the volume of the compartment, i.e. the enzyme is not localized in any special way. Under these circumstances it is possible to construct rate equations for the *concentration* of the species:

$$\frac{d[S_2]}{dt} = -\frac{d[S_1]}{dt} = v_R$$

In SBML, however, the rate equations are constructed for the rate of change of the *amount* of the species:

$$\frac{dn_{S_2}}{dt} = -\frac{dn_{S_1}}{dt} = \hat{v}_R = V \cdot v_R$$

where \hat{v}_R is the modified SBML kinetic law and V is the volume of the compartment. Since the traditional kinetic law v_R describes how fast the amount of the species changes *per volume*, the SBML kinetic law \hat{v}_R simply equals the product of v_R and the compartment volume V . This means that the actual rate of change

1 of the amounts of the species is proportional to the compartment size, which will only be true if the reaction
2 takes place uniformly throughout the compartment. (See Section 7.11 for an example of a reaction that
3 is located at the boundary of a compartment.) The concentrations of the species (that are needed in the
4 definition of v_R) can easily be recovered through the relation $[S_i] = n_{S_i}/V$.

5 An important property of the amount rate equation is that it is still valid if the volume V changes during a
6 simulation. This is not true for the concentration rate equations.

A XML Schema for SBML

The following is an XML Schema definition for SBML Level 3 Version 1 Core, using the W3C Recommendation for XML Schema version 1.0 of 2 May 2001 ([Biron and Malhotra, 2000](#); [Fallside, 2000](#); [Thompson et al., 2000](#)). This Schema does not define all aspects of SBML Level 3: an SBML document validated by this schema is not necessarily a valid SBML Level 3 Version 1 Core document. Appendix B contains a schema for the SBML MathML subset. Appendix C contains a list of the remaining checks required to validate a model in addition to making it consistent with these two schemas.

Note to implementors: the following schema is self-contained and makes reference to the official XML Schema for MathML hosted at the W3. However, for use in software systems, it is more efficient to store the MathML subset Schema of Appendix B in a file on a user's local disk, and change the `schemaLocation` value (text line 25 below) to refer to this local copy of the MathML subset Schema. Doing so will avoid requiring a network access every time this SBML Schema is used.

(Forthcoming.)

B XML Schema for MathML subset

The following XML schema defines the syntax of the MathML syntax that is used in SBML Level 3.

(Forthcoming.)

C Validation and Consistency checks for SBML

The SBML specification is designed to allow modelers a high degree of flexibility. An SBML document that adheres to the specification will be considered valid, but if the modeler has chosen to only partially declare optional attributes there may be inconsistencies within the model, for example with units. This section outlines a set of validation rules that must be followed to produce valid SBML and sets of consistency checks that can be applied to eliminate inconsistencies where this is desirable.

Note that since validation rule numbers remain unique, missing numbers here merely indicate rules that applied in previous levels and versions of SBML that no longer apply.

(Forthcoming.)

D A method for assessing whether an SBML model is overdetermined

As explained in Section 4.9.5, an SBML model must not be overdetermined. It is possible to use purely static analysis to assess this condition for the system of equations implied by a model, by constructing a bipartite graph of the model's variables and equations and then searching for a maximal matching (Chartrand, 1977). A efficient algorithm for finding a maximal matching is described by Hopcroft and Karp (1973). In this appendix, we provide a concrete application to SBML of the general approach described in Section 4.9.5. The approach is defined in terms of the ordinary differential equations (ODEs) implied by an SBML model; despite our use of a differential equation framework for this explanation, it should be understood that this use of ODEs has no implication about the framework actually used to simulate the model.

Definition of the method

First, we treat both assignment rules and formulas in **KineticLaw** objects as SBML-style algebraic equations. (Both are essentially assignment statements, and assignment statements can be trivially converted to algebraic equations with zero on the left-hand side.) We also assume that an ODE is constructed for each species determined by one or more **Reaction**'s **KineticLaw** math expressions. Finally, we assume that the model has already been determined to be valid in all other respects (e.g., there are no undefined variables in the equations), and what remains is to evaluate whether it is overdetermined.

We construct the bipartite graph for a given SBML model as follows:

1. For each of the following in the model, create one vertex representing an equation:
 - (a) Every **Species** object having **boundaryCondition**="false", **constant**="false", and which is referenced as a reactant or product in one or more **Reaction** objects containing **KineticLaw** objects
 - (b) Every **AssignmentRule** object
 - (c) Every **RateRule** object
 - (d) Every **AlgebraicRule** object
 - (e) Every **KineticLaw** object
2. For each of the following in the model, create one vertex representing a variable:
 - (a) Every **Species** object having **constant**="false"
 - (b) Every **Compartment** object having **constant**="false"
 - (c) Every global **Parameter** having **constant**="false"
 - (d) Every **Reaction** object
3. For each of the following, create one edge:
 - (a) Every vertex created in step 2(a) to that species' equation vertex created in step 1(a)
 - (b) Every vertex created in step 1(b) to the particular vertex created in steps 2(a)–2(d) that represents the variable referenced by the **variable** attribute of the rule
 - (c) Every vertex created in step 1(c) to the particular vertex created in steps 2(a)–2(d) that represents the variable referenced by the **variable** attribute of the rule
 - (d) Every vertex created in step 1(e) to the particular vertex created in step 2(d) that is the **Reaction** object containing that particular **KineticLaw** object
 - (e) Every vertex created in steps 2(a)–2(d) representing an identifier appearing as the content of a MathML **ci** element within an expression of an **AssignmentRule**, to the vertex for that particular **AssignmentRule** created in step 1(b)
 - (f) Every vertex created in steps 2(a)–2(d) representing an identifier appearing as the content of a MathML **ci** element within an expression of an **AlgebraicRule**, to the vertex for that particular **AlgebraicRule** created in step 1(d)
 - (g) Every vertex created in steps 2(a)–2(c) representing an identifier appearing as the content of a MathML **ci** element within an expression of a **KineticLaw**, to the vertex for that particular **KineticLaw** created in step 1(e)

Example application of the method

What follows is an example of applying the method above to the SBML model shown below:

```
<?xml version="1.0" encoding="UTF-8"?>
<sbml xmlns="http://www.sbml.org/sbml/level3/version1/core" level="3" version="1">
  <model id="example">
    <listOfCompartments>
      <compartment id="C" size="1" spatialDimensions="3" constant="true"/>
    </listOfCompartments>
    <listOfSpecies>
      <species id="S1" compartment="C" initialConcentration="1"
        boundaryCondition="false" hasOnlySubstanceUnits="false"
        constant="false"/>
      <species id="S2" compartment="C" initialConcentration="0"
        boundaryCondition="false" hasOnlySubstanceUnits="false"
        constant="false"/>
    </listOfSpecies>
    <listOfRules>
      <algebraicRule>
        <math xmlns="http://www.w3.org/1998/Math/MathML">
          <apply>
            <plus/>
            <ci> S1 </ci>
            <ci> S2 </ci>
          </apply>
        </math>
      </algebraicRule>
    </listOfRules>
    <listOfReactions>
      <reaction id="R" reversible="true" fast="false">
        <listOfReactants>
          <speciesReference species="S1" stoichiometry="1" constant="true"/>
        </listOfReactants>
        <listOfProducts>
          <speciesReference species="S2" stoichiometry="1" constant="true"/>
        </listOfProducts>
        <kineticLaw>
          <math xmlns="http://www.w3.org/1998/Math/MathML">
            <apply>
              <times/>
              <ci> C </ci>
              <ci> k1 </ci>
              <ci> S1 </ci>
            </apply>
          </math>
          <listOfLocalParameters>
            <localParameter id="k1" value="0.1"/>
          </listOfLocalParameters>
        </kineticLaw>
      </reaction>
    </listOfReactions>
  </model>
</sbml>
```

For the model above, we create *equation* vertices as follows:

1. [Corresponding to step 1(a) in Section D.] Every **Species** object having **boundaryCondition**="false", **constant**="false", and which is referenced as a reactant or product in one or more **Reaction** objects containing **KineticLaw** objects. This generates two vertices, for "S1" and "S2".
2. [Corresponding to step 1(b) in Section D.] Every **AlgebraicRule** object. This generates one vertex, for the model's lone algebraic rule (call it "rule").
3. [Corresponding to step 1(e) in Section D.] Every **KineticLaw** object. This generates one vertex, for the lone kinetic law in the model (call it "law").

We create *variable* vertices for the following:

1. [Corresponding to step 2(a) in Section D.] Every **Species** object having `constant="false"`. This generates two vertices, for "S1" and "S2".

2. [Corresponding to step 2(b) in Section D.] Every **Compartment** object having `constant="false"`. This generates one vertex, for "C".

3. [Corresponding to step 2(d) in Section D.] Every **Reaction** object. This generates one vertex, for "R".

Note that it is not necessary to include parameters declared within **KineticLaw** objects because they are local to a particular reaction and cannot be affected by rules. With the steps above, we have the following set of graph nodes:


Vertices for equations



Vertices for variables

Next, we create edges following the procedure described above. Doing so results in the following graph:


Vertices for equations



Vertices for variables

The algorithm of [Hopcroft and Karp \(1973\)](#) can now be applied to search for a maximal matching of the bipartite graph. A maximal matching is a graph in which each vertex is connected to at most one other vertex and the maximum possible number of connections have been made. Doing so here results in the following:

Vertices for equations



Vertices for variables

If the maximal matching of the bipartite graph leaves any equation vertex unconnected, then the model is considered overdetermined. That is the case for the example shown here, because the equation vertex for "rule" is unconnected in the maximal matching.

E Mathematical consequences of the fast attribute on Reaction

(Appendix contributed by James C. Schaff, University of Connecticut Health Center, Connecticut, U.S.A.)

Section 4.11.1 described the **fast** flag available on **Reaction**. In this appendix, we discuss the principles involved in interpreting this attribute in the context of a simple biochemical reaction model. The derivation presented here is not fully rigorous and this section is not considered normative; achieving a higher level of rigor would require considerably more background exposition and a much longer appendix. Nevertheless, we hope this section is sufficient to answer unambiguously the question “How should a system of reactions be treated if some of the reactions have **fast=true**?”

Identification of “fast” reactions

First, it is worth noting that the identification of so-called *fast* reactions is actually a modeling issue, not an SBML representation issue. The notion of fast reactions is the following. A system may be decomposable into two sets of reactions, where one set may have characteristic times that are much faster than the other time scales in the system. An approximation that is sometimes useful is to assume that the fast reactions have kinetics that settle infinitely fast compared to the other reactions in the system. In other words, the fast reactions are assumed to be always in equilibrium. This is called a pseudo-steady state approximation (PSSA), and is also known as a quasi-steady state approximation (QSSA). Given a case where the time-scale separation between fast and other reactions in the system is large, an accurate and efficient approach for computing the time-course of the system behavior is to treat the fast reactions as being always in equilibrium.

The key to successful application of a PSSA is that there should be a significant separation of time scales between these fast reactions and other reactions in the system. The determination of which reactions qualify as fast is up to the creator of the model, because there is currently no known general algorithm for doing so.

Simple one-compartment biochemical system model

To explain how to solve a system containing fast reactions, we use a simple model of a biochemical reaction network located in a single compartment. Let \mathbf{x}^* represent a vector of all the species in the system, \mathbf{v}^* a vector of the reaction rates, and \mathbf{A}^* the stoichiometry matrix, with the vector dimension being \mathbf{n}^* . Then the system can be described using the following matrix equation:

$$\frac{d\mathbf{x}^*}{dt} = \mathbf{A}^* \mathbf{v}^*(\mathbf{x}^*)$$

This system can be optionally reduced by noting that mass conservation usually implies there are linear combinations of species quantities in the system and the value of these combinations do not change over time. Identifying these combinations is the topic of structural analysis and is described in the literature (Reder, 1988; Sauro and Ingalls, 2003). Briefly, let \mathbf{N} be defined as the left null space of \mathbf{A}^* :

$$\mathbf{N}\mathbf{A}^* = \mathbf{0}$$

Now, premultiply the previous equation by \mathbf{N} to get

$$\mathbf{N} \frac{d\mathbf{x}^*}{dt} = \mathbf{N}\mathbf{A}^* \mathbf{v}^*(\mathbf{x}^*) = \mathbf{0}$$

Thus, \mathbf{N} captures the space of solutions to the equation

$$\mathbf{m}^T \left(\frac{d\mathbf{x}^*}{dt} \right) = 0$$

where \mathbf{m} is a vector representing the coefficients in a mass conservation relationship, that is, combinations of species that are time-invariant. Now, let

$$r = \text{rank}(\mathbf{A}^*)$$

$$n = \text{dim}(\mathbf{x}^*)$$

Then the system has $n - r$ mass conservation relationships, each of which is a linear equation. We can use these $n - r$ linear equations to solve for $n - r$ dependent variables in terms of r independent variables and the initial masses of all species. Doing that allows us to decompose \mathbf{x}^* into $n - r$ dependent variables \mathbf{x}_d and r independent variables \mathbf{x}_i where \mathbf{L} is an $(n - r) \times r$ matrix that is derived from \mathbf{N} and represents \mathbf{x}_d in terms of \mathbf{x}_i , \mathbf{I} is the $r \times r$ identity matrix, and \mathbf{T} is an $n \times r$ matrix:

$$\mathbf{x}^* \equiv \begin{bmatrix} \mathbf{x}_i \\ \mathbf{x}_d \end{bmatrix} = \begin{bmatrix} \mathbf{I} \\ \mathbf{L} \end{bmatrix} \mathbf{x}_i = \mathbf{T} \mathbf{x}_i$$

Using this equation, we can define a new vector of reaction velocities \mathbf{v} in terms of \mathbf{x}_i only:

$$\mathbf{v}(\mathbf{x}_i) \equiv \mathbf{v}^*(\mathbf{T} \mathbf{x}_i)$$

With this \mathbf{v} , we can now write a reduced system by substituting terms. First we define \mathbf{A} as the r independent rows of \mathbf{A}^* corresponding to \mathbf{x}_i . Then:

$$\frac{d\mathbf{x}_i}{dt} = \mathbf{A} \mathbf{v}(\mathbf{x}_i)$$

This is a set of r independent differential equations in r unknowns (i.e., an r -dimensional system). To simplify the notation slightly, let

$$\mathbf{x} \equiv \mathbf{x}_i$$

and, thus,

$$\frac{d\mathbf{x}}{dt} = \mathbf{A} \mathbf{v}(\mathbf{x})$$

Application of a PSSA to biochemical systems

Assume that we have eliminated redundant variables and equations using the mass conservation analysis above. Further assume that we have some external means of classifying some reactions in a given system as being *fast* as discussed earlier. We now need to apply this to the system under study. We begin by decomposing the vector of reaction velocities \mathbf{v} according to fast and slow reactions:

$$\frac{d\mathbf{x}}{dt} = \mathbf{A}_1 \mathbf{v}_f(\mathbf{x}) + \mathbf{A}_2 \mathbf{v}_s(\mathbf{x})$$

In the expression above, \mathbf{A}_1 represents the stoichiometry of the set of reactions operating on the fast time scale, and \mathbf{A}_2 the stoichiometry of the set of reactions operating on a slower time scale. We find the left null space of \mathbf{A}_1 (i.e., the space of solutions to $\mathbf{m}^T [d\mathbf{x}/dt] = 0$ on a fast time scale), and call this matrix \mathbf{B} :

$$\mathbf{B} \mathbf{A}_1 = \mathbf{0}$$

The matrix \mathbf{B} represents the linear combination of species that do not change on a fast time scale, i.e., the slow species in the system. Now, we premultiply the equation for $d\mathbf{x}/dt$ by \mathbf{B} :

$$\begin{aligned} \mathbf{B} \frac{d\mathbf{x}}{dt} &= \mathbf{B} \mathbf{A}_1 \mathbf{v}_f(\mathbf{x}) + \mathbf{B} \mathbf{A}_2 \mathbf{v}_s(\mathbf{x}) \\ &= \mathbf{B} \mathbf{A}_2 \mathbf{v}_s(\mathbf{x}) \end{aligned}$$

where the second line follows from the fact that $\mathbf{B} \mathbf{A}_1 = \mathbf{0}$. The above is an ordinary differential equation in terms of only the slow dynamics. The remaining fast dynamics are handled by applying the pseudo-steady state approximation, with fast transients assumed to have settled with respect to the slow time scale. This produces a system of nonlinear algebraic equations:

$$\mathbf{A}_1 \mathbf{v}_f = \mathbf{0}$$

The last two equations form the system of equations resulting from the application of the PSSA. If $r_1 = \text{rank}(\mathbf{A}_1)$ and $r = \text{rank}(\mathbf{A})$, then there will be r_1 degrees of freedom that will be determined by solving an algebraic system (the equation $\mathbf{A}_1 \mathbf{v}_f = \mathbf{0}$ above), and there will be $r - r_1$ degrees of freedom that will be determined by ordinary differential equations (the equation for $\mathbf{B} d\mathbf{x}/dt$).

F Processing and validating SBase notes and Constraint message content

In Section 3.2.3 and Section 4.10.2, we discussed the **notes** element on **SBase** and the **message** element on **Constraint**, respectively. These elements can contain a number of possible forms of XHTML content. In this appendix, we describe a general procedure for how application software can process such content. We concentrate on the common case of an SBML-reading application that needs to take the contents and pass it to an XHTML display and/or editing function obtained from a third-party API library. The content of the **notes** and **message** may not be a complete XHTML document, so the application will have to perform some processing before handing it to the XHTML editor or validator. How should this be done?

Based on the three forms of **SBase notes** content described in Section 3.2.3 and the identical forms for **Constraint message** described in Section 4.10.2, there are only three cases possible. Here we give an example approach for handling them, although the actual implementation details will differ depending on various factors such as the requirements of the software libraries being used. This example approach would be performed for each **notes** and **message** to be viewed or edited:

Step 1. If the XHTML viewing/editing function requires a fully compliant XML document, the SBML application could create a temporary data object containing an appropriate XML declaration and a DOCTYPE declaration; otherwise, the XML data object can be initially blank.

Step 2. The application should look at the first element inside the **notes** or **message** (or rather, the first element that is not an XML comment), and take action based on the following three possibilities:

- If the first element begins with `<html xhtml="http://www.w3.org/1999/xhtml">`, the application could assume that the content is a complete XHTML document and insert this into the temporary data object.
- Else, if the first element is `<body>`, the application should insert the following into the temporary data object,

```
<html xhtml="http://www.w3.org/1999/xhtml">
  <head><title></title></head>
```

then insert the content of the **notes** or **message**, and finally insert a closing `</html>`.

- Else, if the content begins with neither of the above elements, the application should insert the following into the temporary data object,

```
<html xhtml="http://www.w3.org/1999/xhtml">
  <head><title></title></head>
  <body>
```

then insert the content of the **notes** or **message**, and finally insert `</body></html>` to close the XHTML document.

The result of the above would be a temporary XML data object that the application could then pass to the XHTML viewing/editing API function.

Acknowledgments

The development of SBML was originally funded by the Japan Science and Technology Agency (JST) under the ERATO Kitano Symbiotic Systems Project during the years 2000–2003. From 2003, general support for development of SBML and associated software such as libSBML and the SBML Test Suite has been provided by the National Institute of General Medical Sciences (USA) via grant numbers GM070923 and GM077671.

We gratefully acknowledge additional sponsorship from the following funding agencies: the National Human Genome Research Institute (USA); the International Joint Research Program of NEDO (Japan); the JST ERATO-SORST Program (Japan); the Japanese Ministry of Agriculture; the Japanese Ministry of Education, Culture, Sports, Science and Technology; the BBSRC e-Science Initiative (UK); the DARPA IPTO Bio-Computation Program (USA); and the Air Force Office of Scientific Research (USA).

Additional support has been or continues to be provided by the following institutions: the California Institute of Technology (USA), EML Research gGmbH (Germany), the European Molecular Biology Laboratory's European Bioinformatics Institute (UK), the Molecular Sciences Institute (USA), the University of Heidelberg (Germany), the University of Hertfordshire (UK), the University of Newcastle (UK), the Systems Biology Institute (Japan), and the Virginia Bioinformatics Institute (USA).

The following individuals served as SBML Editors in past years. Their efforts shaped what SBML is today:

- Hamid Bolouri (California Institute of Technology, CA, USA)
- Andrew M. Finney
- Nicolas Le Novère (EMBL-EBI, UK)
- Herbert M. Sauro (University of Washington, WA, USA)

SBML was first conceived at the JST/ERATO-sponsored *First Workshop on Software Platforms for Systems Biology*, held in April, 2000, at the California Institute of Technology in Pasadena, California, USA. The participants collectively decided to begin developing a common XML-based declarative language for representing models. A draft version of SBML was developed by the Caltech ERATO team and delivered to collaborators in August, 2000. This version underwent extensive discussion over mailing lists and then again during the *Second Workshop on Software Platforms for Systems Biology* held in Tokyo, Japan, November 2000. The Caltech ERATO team issued a revised version in December, 2000, and after further discussions over mailing lists and in meetings, they produced a specification for SBML Level 1 (Hucka et al., 2001).

SBML Level 2 was conceived at the *5th Workshop on Software Platforms for Systems Biology*, held in July 2002, at the University of Hertfordshire, UK. The participants collectively decided to revise the form of SBML in SBML Level 2. The first draft of the Level 2 Version 1 document was released in August 2002. The final set of features in SBML Level 2 Version 1 was finalized in May 2003 at the *7th Workshop on Software Platforms for Systems Biology* in Ft. Lauderdale, Florida.

SBML Level 2 Version 2 was largely finalized after the 2005 SBML Forum meeting in Boston and a final document was issued in September 2006. SBML Level 2 Version 3 was finalized after the 2006 SBML Forum meeting in Yokohama, Japan, and the 2007 SBML Hackathon in Newcastle, UK. SBML Level 2 Version 4 was finalized after the 2008 SBML Forum in Göteborg, Sweden. They were developed with contributions from so many people constituting the worldwide *SBML Forum* that we regret it has become infeasible to list individuals by name. For discussions and help developing SBML, and for feedback about this specification, we are grateful to everyone on the sbml-discuss@caltech.edu and sbml-interoperability@caltech.edu mailing lists, and many other groups and developers at large, notably the creators of CellML (Hedley et al., 2001), the members of the DARPA Bio-SPICE project, and the authors of all of the software systems that support SBML. A guide to software known to support SBML is provided on the SBML.org website at the following URL: http://sbml.org/SBML_Software_Guide.

The idea of SBML Level 3 was first publicly discussed at the *6th Workshop on Software Platforms for Systems Biology*, held in December, 2002, at the Karolinska Institute, Stockholm, Sweden. Development proceeded on and off ever since then.

References

- Abbott, A. (1999). Alliance of US labs plans to build map of cell signalling pathways. *Nature*, 402:219–200.
- Abramowitz, M. and Stegun, I. A., editors (1977). *Mathematical Functions: With Formulas, Graphs, and Mathematical Tables*. Dover Publications Inc.
- Ausbrooks, R., Buswell, S., Carlisle, D., Dalmas, S., Devitt, S., Diaz, A., Froumentin, M., Hunter, R., Ion, P., Kohlhase, M., Miner, R., Poppelier, N., Smith, B., Soiffer, N., Sutor, R., and Watt, S. (2003). Mathematical Markup Language (MathML) Version 2.0 (second edition): W3C Recommendation 21 October 2003. Available via the World Wide Web at <http://www.w3.org/TR/2003/REC-MathML2-20031021/>.
- Ball, K., Kurtz, T. G., Popovic, L., and Rempala, G. (2006). Asymptotic analysis of multiscale approximations to reaction networks. *Annals of Applied Probability*, 16(4):1925–1961.
- Biron, P. V. and Malhotra, A. (2000). XML Schema part 2: Datatypes (W3C candidate recommendation 24 October 2000). Available via the World Wide Web at <http://www.w3.org/TR/xmlschema-2/>.
- Bray, T., D. Hollander, D., and Layman, A. (1999). Namespaces in XML. World Wide Web Consortium 14-January-1999. Available via the World Wide Web at <http://www.w3.org/TR/1999/REC-xml-names-19990114/>.
- Bray, T., Paoli, J., Sperberg-McQueen, C. M., and Maler, E. (2000). Extensible markup language (XML) 1.0 (second edition), W3C recommendation 6-October-2000. Available via the World Wide Web at <http://www.w3.org/TR/1998/REC-xml-19980210/>.
- Bray, T., Paoli, J., Sperberg-McQueen, C. M., Maler, E., and Yergeau, F. (2004). Extensible markup language (XML) 1.0 (third edition), W3C recommendation 4-February-2004. Available via the World Wide Web at <http://www.w3.org/TR/2004/REC-xml-20040204/>.
- Bureau International des Poids et Mesures (2006). The International System of Units (SI) 8th edition (2006). Available via the World Wide Web at http://www.bipm.org/utis/common/pdf/si-brochure_8.pdf.
- Chartrand, G. (1977). *Introductory Graph Theory*. Dover Publishing, Inc., New York.
- DCMI Usage Board (2005). DCMI Metadata Terms. Available via the World Wide Web at <http://www.dublincore.org/documents/dcmi-terms/>.
- Dublin Core Metadata Initiative (2005). Dublin core metadata initiative. Available via the World Wide Web at <http://dublincore.org/>.
- Eriksson, H.-E. and Penker, M. (1998). *UML Toolkit*. John Wiley & Sons, New York.
- Evans, T. W., Gillespie, C. S., and Wilkinson, D. J. (2008). The SBML discrete stochastic models test suite. *Bioinformatics*, 24:285–286.
- Fallside, D. C. (2000). XML Schema part 0: Primer (W3C candidate recommendation 24 October 2000). Available via the World Wide Web at <http://www.w3.org/TR/xmlschema-0/>.
- Gillespie, D. (1977). Exact stochastic simulation of coupled chemical reactions. *J. Phys. Chem.*, 81:2340–2361.
- Gillespie, D. (1992). A rigorous derivation of the chemical master equation. *Physica A*, 188:404–425.
- Gilman, A. (2000). A letter to the signaling community. Alliance for Cellular Signaling, The University of Texas Southwestern Medical Center. Available via the World Wide Web at http://afcs.swmed.edu/afcs/Letter_to_community.htm.
- Harold, E. R. and Means, E. S. (2001). *XML in a Nutshell*. O’Reilly & Associates.

- 1 Hedley, W. J., Nelson, M. R., Bullivant, D., Cuellar, A., Ge, Y., Grehlinger, M., Jim, K., Lett, S., Nickerson,
2 D., Nielsen, P., and Yu, H. (2001). CellML specification. Available via the World Wide Web at [http:](http://www.cellml.org/public/specification/20010810/cellml.specification.html)
3 [//www.cellml.org/public/specification/20010810/cellml.specification.html](http://www.cellml.org/public/specification/20010810/cellml.specification.html).
- 4 Hopcroft, J. E. and Karp, R. M. (1973). An $n^{5/2}$ algorithm for maximum matchings in bipartite graphs.
5 *SIAM Journal on Computing*, 2(4):225–231.
- 6 Hucka, M., Finney, A., Sauro, H. M., and Bolouri, H. (2001). Systems Biology Markup Language (SBML)
7 Level 1: Structures and facilities for basic model definitions. Available via the World Wide Web at
8 <http://www.sbml.org>.
- 9 Hucka, M., Finney, A., Sauro, H. M., Bolouri, H., Doyle, J. C., Kitano, H., Arkin, A. P., Bornstein, B. J.,
10 Bray, D., Cornish-Bowden, A., Cuellar, A. A., Dronov, S., Gilles, E. D., Ginkel, M., Gor, V., Goryanin, I. I.,
11 Hedley, W. J., Hodgman, T. C., Hofmeyr, J.-H., Hunter, P. J., Juty, N. S., Kasberger, J. L., Kremling, A.,
12 Kummer, U., Le Novère, N., Loew, L. M., Lucio, D., Mendes, P., Minch, E., Mjolsness, E. D., Nakayama,
13 Y., Nelson, M. R., Nielsen, P. F., Sakurada, T., Schaff, J. C., Shapiro, B. E., Shimizu, T. S., Spence,
14 H. D., Stelling, J., Takahashi, K., Tomita, M., Wagner, J., and Wang, J. (2003). The Systems Biology
15 Markup Language (SBML): A medium for representation and exchange of biochemical network models.
16 *Bioinformatics*, 19(4):524–531.
- 17 Iannella, R. (2001). Representing vCard objects in RDF/XML. Available via the World Wide Web at
18 <http://www.w3.org/TR/vcard-rdf>.
- 19 Jacobs, I. (2004). World wide web consortium process document. Available via the World Wide Web at
20 <http://www.w3.org/2004/02/Process-20040205/>.
- 21 Kokkeliink, S. and Schwänzl, R. (2002). Expressing qualified Dublin Core in RDF/XML. Available via the
22 World Wide Web at <http://dublincore.org/documents/dcq-rdf-xml/index.shtml>.
- 23 Lassila, O. and Swick, R. (1999). Resource description framework (RDF) model and syntax specification.
24 Available via the World Wide Web at <http://www.w3.org/TR/REC-rdf-syntax/>.
- 25 Le Novère, N., Finney, A., Hucka, M., Bhalla, U., Campagne, F., Collado-Vides, J., Crampin, E. J., Halstead,
26 M., Klipp, E., Mendes, P., Nielsen, P., Sauro, H., Shapiro, B., Snoep, J. L., Spence, H. D., and Wanner,
27 B. L. (2005). Minimum information requested in the annotation of biochemical models (MIRIAM). *Nature*
28 *Biotechnology*, 23:1509–1515.
- 29 Mohr, P. J., Taylor, B. N., and Newell, D. B. (2008). CODATA Recommended Values of the Fundamental
30 Physical Constants: 2006. *Reviews of Modern Physics*, 80:633–731.
- 31 Oestereich, B. (1999). *Developing Software with UML: Object-Oriented Analysis and Design in Practice*.
32 Addison-Wesley Publishing Company.
- 33 Pemberton, S., Austin, D., Axelsson, J., Celik, T., Dominiak, D., Elenbaas, H., Epperson, B., Ishikawa, M.,
34 Matsui, S., McCarron, S., Navarro, Peruvemba, S., Relyea, R., Schnitzenbaumer, S., and Stark, P. (2002).
35 XHTMLTM 1.0 the Extensible HyperText Markup Language (second edition): W3C Recommendation 26
36 January 2000, revised 1 August 2002. Available via the World Wide Web at [http://www.w3.org/TR/](http://www.w3.org/TR/xhtml1/)
37 [xhtml1/](http://www.w3.org/TR/xhtml1/).
- 38 Popel, A. and Winslow, R. L. (1998). A letter from the directors... Center for Computational Medicine
39 & Biology, Johns Hopkins School of Medicine, Johns Hopkins University. Available via the World Wide
40 Web at <http://www.bme.jhu.edu/ccmb/ccmbletter.html>.
- 41 Powell, A. and Johnston, P. (2003). Guidelines for implementing Dublin Core in XML. Available via the
42 World Wide Web at <http://dublincore.org/documents/dc-xml-guidelines/index.shtml>.
- 43 Reder, C. (1988). Metabolic Control Theory: a structural approach. *Journal of Theoretical Biology*, 135:175–
44 201.

- 1 Sauro, H. M. and Ingalls, B. (2003). Conservation analysis in biochemical networks: Computational issues
2 for software writers. Available via the World Wide Web at [http://www.math.uwaterloo.ca/~bingalls/](http://www.math.uwaterloo.ca/~bingalls/Pubs/conservation.pdf)
3 [Pubs/conservation.pdf](http://www.math.uwaterloo.ca/~bingalls/Pubs/conservation.pdf).
- 4 Smaglik, P. (2000). For my next trick... *Nature*, 407:828–829.
- 5 Thompson, H. S., Beech, D., Maloney, M., and Mendelsohn, N. (2000). XML Schema part 1: Structures
6 (W3C candidate recommendation 24 October 2000). Available via the World Wide Web at [http://www.](http://www.w3.org/TR/xmlschema-1/)
7 [w3.org/TR/xmlschema-1/](http://www.w3.org/TR/xmlschema-1/).
- 8 Unicode Consortium (1996). *The Unicode Standard, Version 2.0*. Addison-Wesley Developers Press, Reading,
9 Massachusetts.
- 10 W3C (2000a). Naming and addressing: URIs, URLs, Available via the World Wide Web at [http:](http://www.w3.org/Addressing/)
11 [//www.w3.org/Addressing/](http://www.w3.org/Addressing/).
- 12 W3C (2000b). W3C's math home page. Available via the World Wide Web at <http://www.w3.org/Math/>.
- 13 W3C (2004a). Rdf/xml syntax specification (revised). Available via the World Wide Web at [http://www.](http://www.w3.org/TR/rdf-syntax-grammar/)
14 [w3.org/TR/rdf-syntax-grammar/](http://www.w3.org/TR/rdf-syntax-grammar/).
- 15 W3C (2004b). Resource description framework (RDF). Available via the World Wide Web at [http://www.](http://www.w3.org/RDF/)
16 [w3.org/RDF/](http://www.w3.org/RDF/).
- 17 Wilkinson, D. J. (2006). *Stochastic Modelling for Systems Biology*. Chapman & Hall/CRC.
- 18 Wolf, M. and Wicksteed, C. (1998). Date and time formats. Available via the World Wide Web at [http:](http://www.w3.org/TR/NOTE-datetime)
19 [//www.w3.org/TR/NOTE-datetime](http://www.w3.org/TR/NOTE-datetime).
- 20 Zwillinger, D., editor (1996). *Standard Mathematical Tables and Formulae*. CRC Press LLC, 30th edition.